

Energy Berger

energyberger.github.io / CompSCI-630

Systems Lunch @ Fridays - Fall

Random number

TA Abhinav Jangda

Github!

Hot CRP

energy@cs.umass.edu

- Scribes!
 - course load :
 - 1 core / semester for MS
 - ..
 - 1st semester MS / PhD

> 1 core = too much!
 - schedule - yes, we have a midterm & final!
 - reviews
-
- ## FORTRAN (in COBOL)
- 1954
- | | |
|--|--|
| military <ul style="list-style-type: none"> - calculating - artillery trajectories - decryption | first computers
— people! |
| general-purpose computing machine | automatic computers |
| Turing machine
von Neumann machine |] data & code together
Harvard architecture |
| | fetch-decode-execute cycle |

assembly language

- lowlevel
- direct access to memory
- operations not abstract

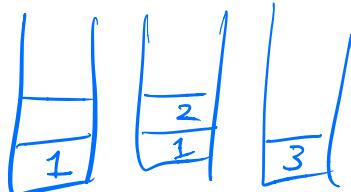
- stack architecture
- register architecture

A, B/C/D
AX BX CX DX

- portability

ASM ≠
~~portable~~

MOV 1, A
MOU 2, B
ADD A, B, C



PUSH 1
PUSH 2
ADD
POP

"Intel" "ARM" (subset of instr.)
JVM, .NET, Python bytecodes,

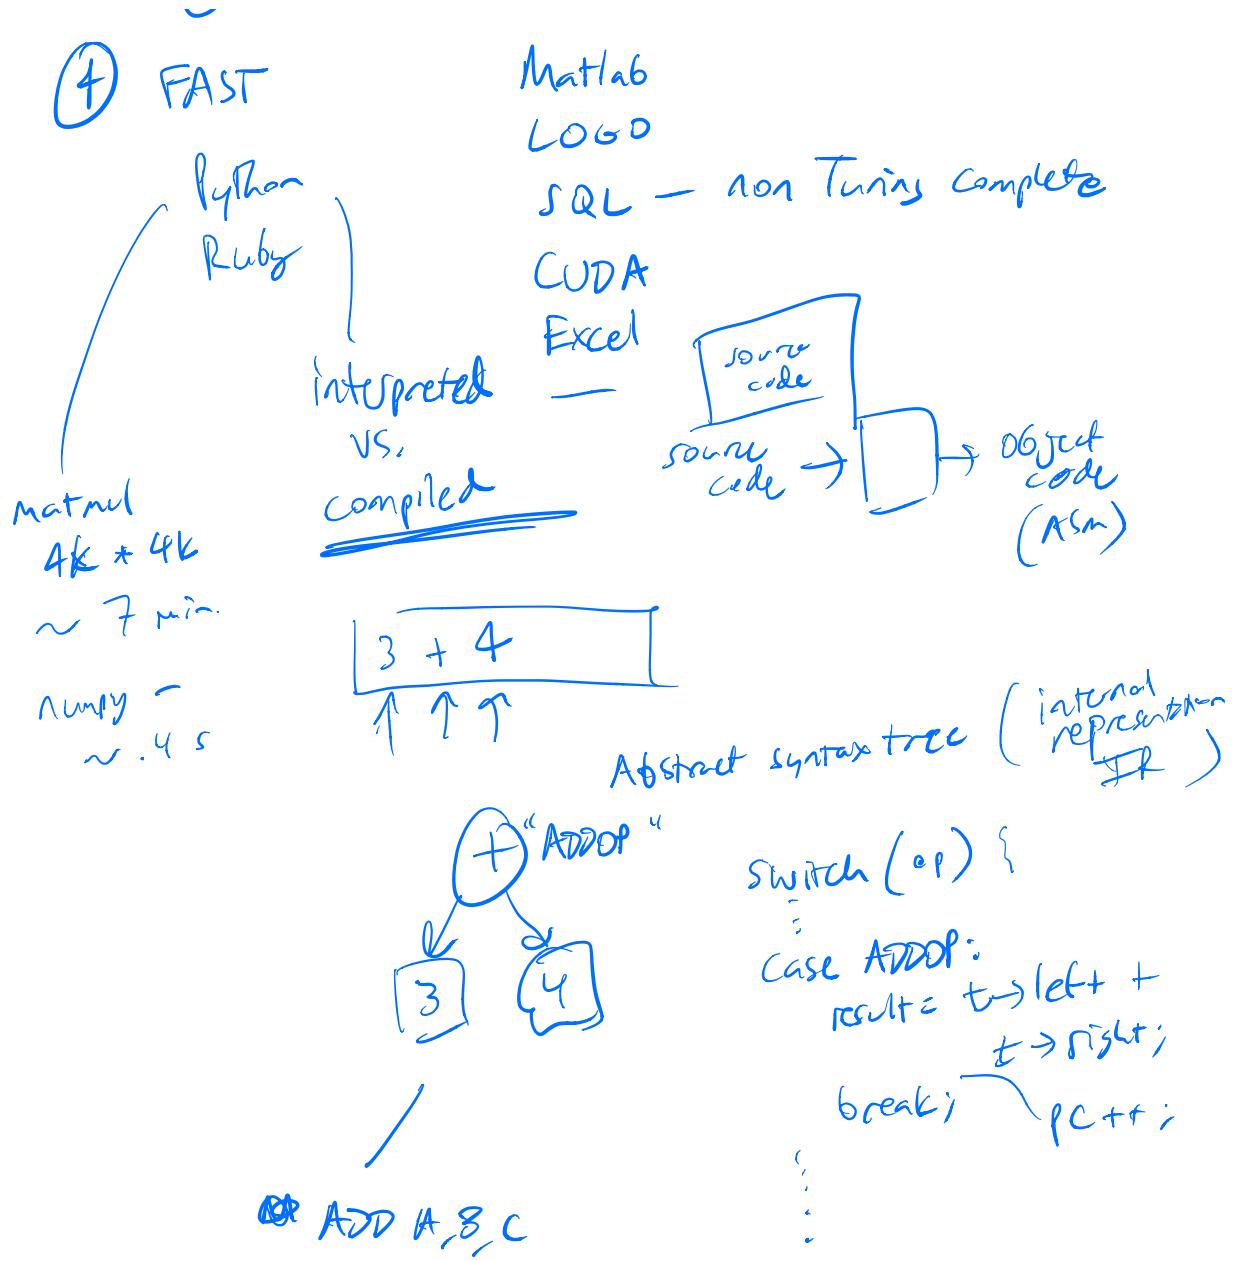
WASM

backward compatibility

- hard to read, hard to debug

- ① PORTABLE
- ② READABLE
- ③ WRITABLE

- COBOL - manager
FORTRAN - scientist } domain-specific languages



programmer time MORE USEFUL than cycles

— modern POU

- productivity
- debugging
- ~ 100x faster!

interpreter → just-in-time compiler (JIT)

Python PyPy
JavaScript V8
 Chrome

WASM

Predicates -
- Parsing
- Compiling
- program analysis
- type coercion
- scope
- object orientation
- modularity
- structured prog.

FORTRAN
1st optimizing compiler

I, J, K integer
X, Y, Z float

DO 10 I = 10 10 *

(was comma) line number

10 CONTINUE

BASIC



$\text{DO } \cup^1 \phi \cup^2 \psi = \cup^1 \phi$

control flow

$\text{DO } \phi \psi = \phi \psi$

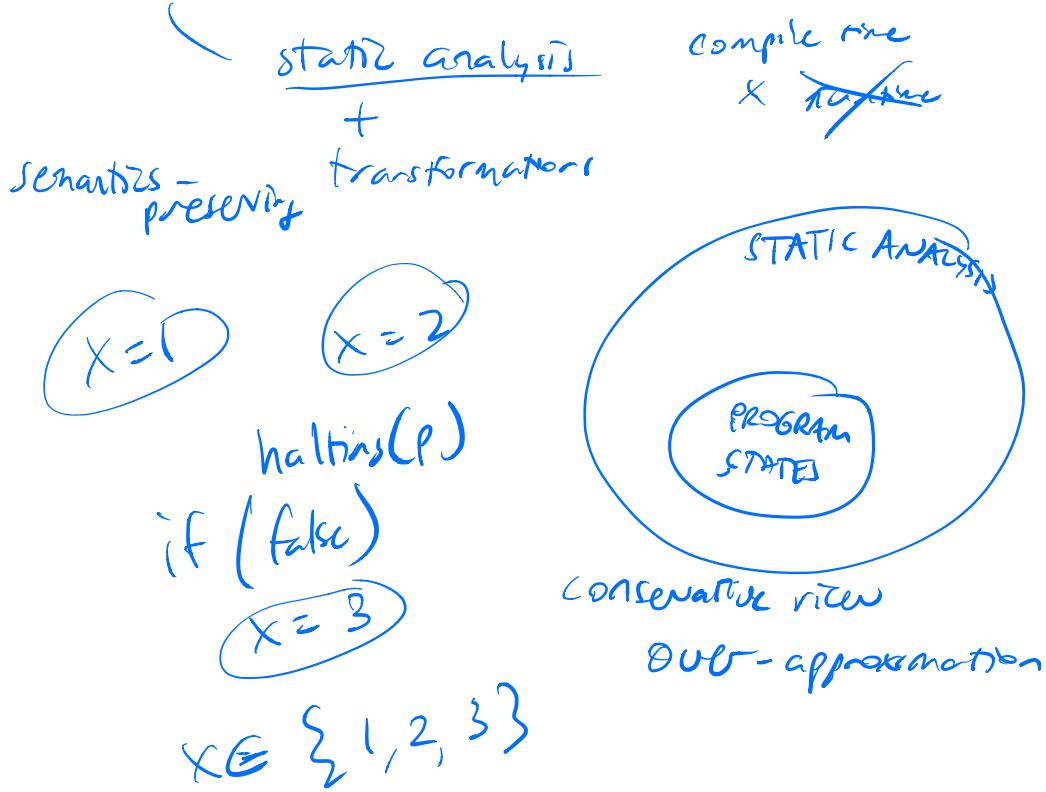
FORTRAN 77, FORTRAN 90 ~

writing

LEGACY
FAST

The Fastest PC today! 2019 !!

optimization



Halting problem

halts (program-text)
→ true if program terminates
 ^{always}
→ false otherwise

p' if halts(p)
 run forever
else terminate

① Constant folding

$$x = 3 + 4$$

↓

$$x = 7$$

$* PI * PI$

② Constant propagation

$$x = 7 \quad z = 12$$
$$y := x + 3 + z$$

↓

$$x = 7 \quad z = 12$$
$$y = 22$$

③ ``strength' reduction

$$x = x \times x$$

↓

$$x = x * x$$

/ %

expensive

cheap

$$x = x \% 8$$

↓

$$x = x \& 7$$

_____ 111

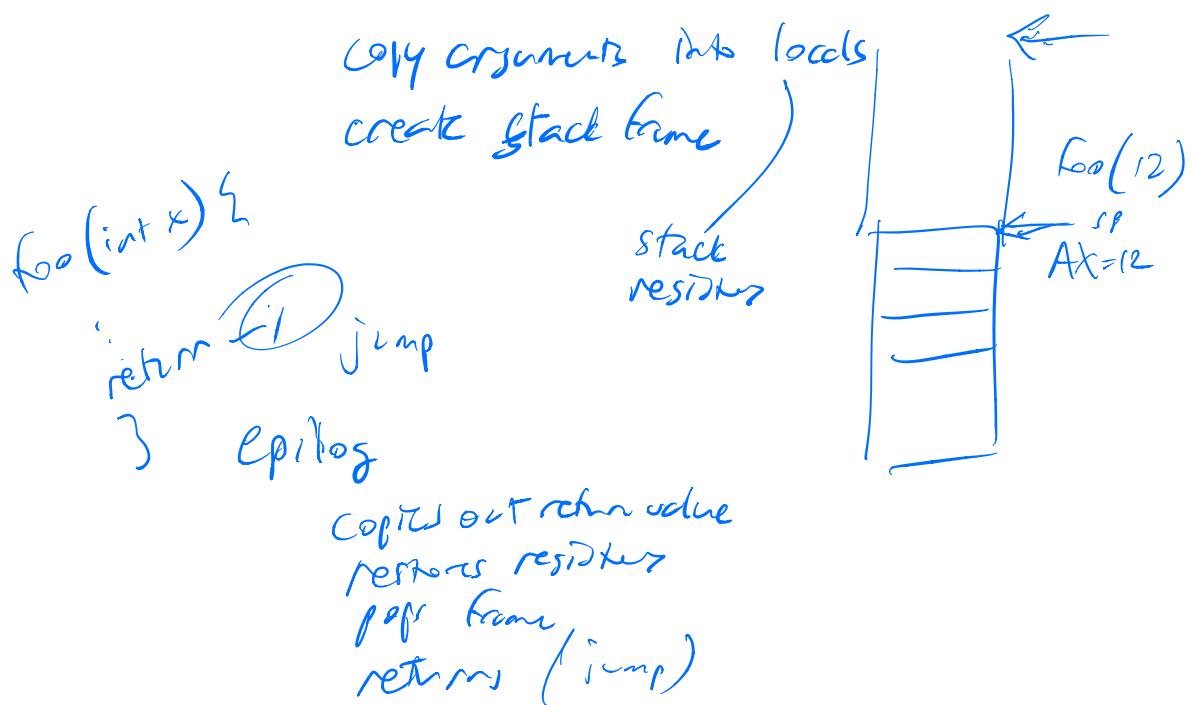
④ vectorization

8	12	13	45
↓			
16	24	26	90

```
for (i=0; i<16; i++) {
    a[i] *= 2;
}
||_
vector ops (4x speedup)
```

⑤ inlining

function prolog

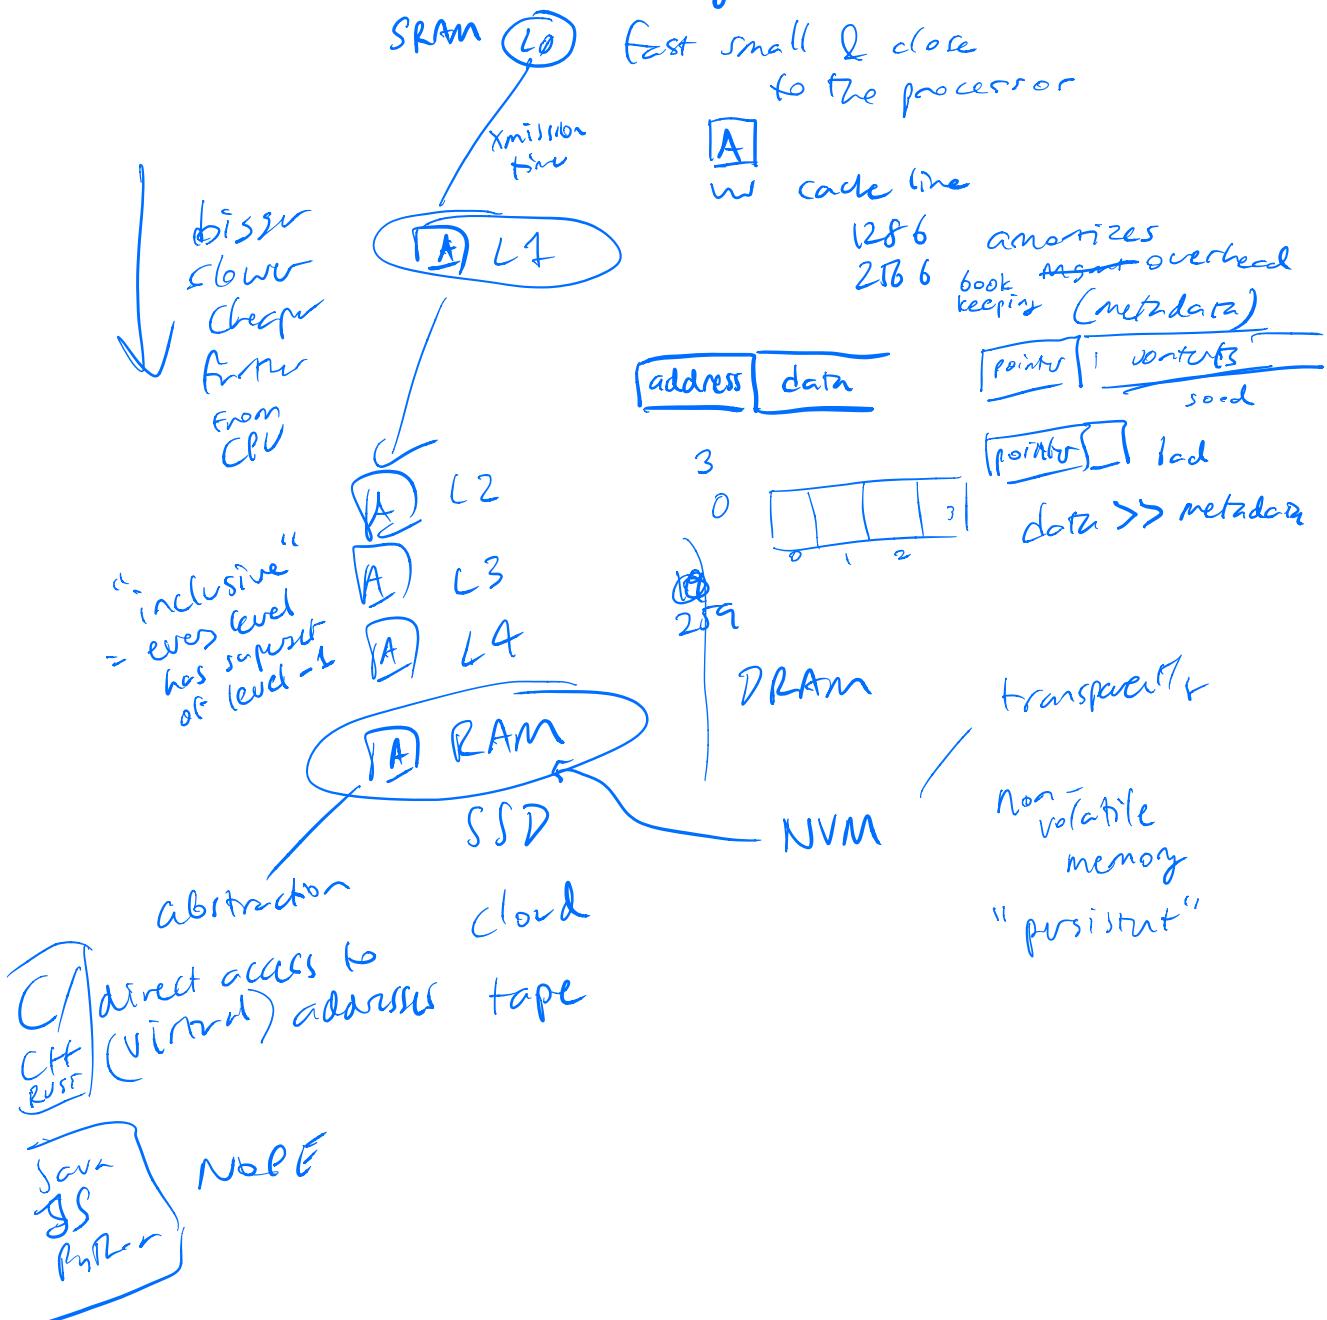


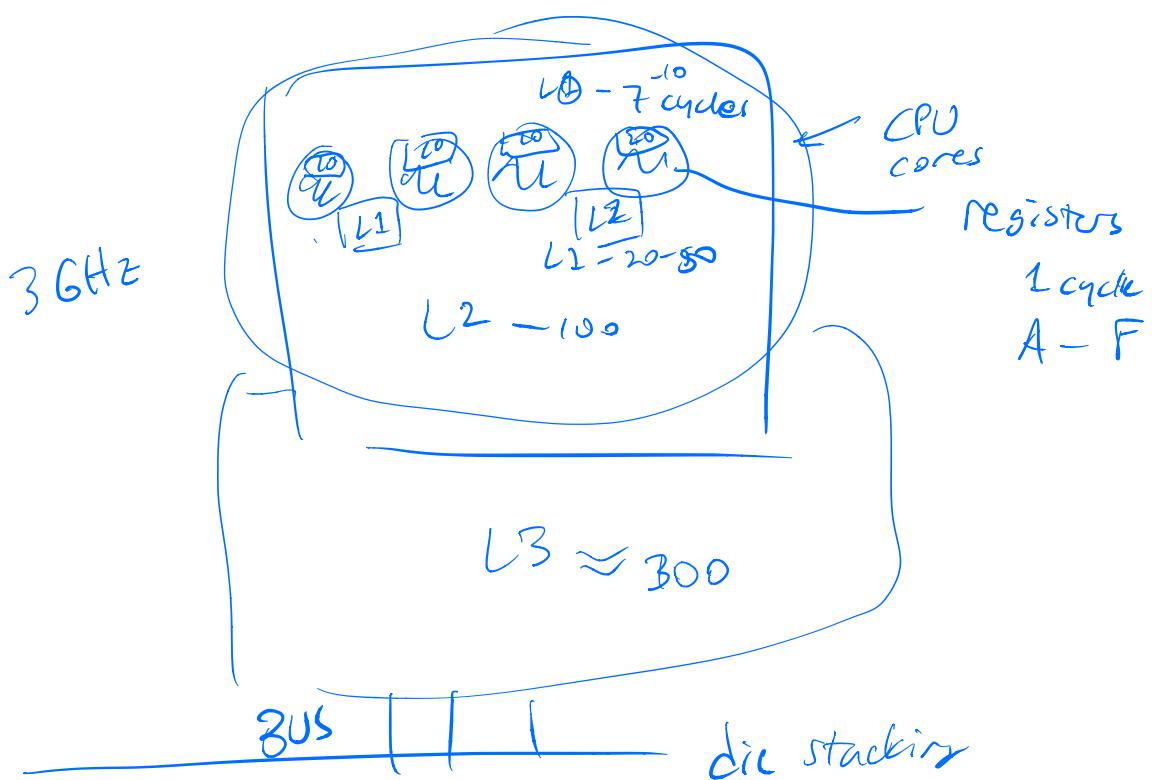
```
int add(int a, int b) {  
    return a + b;  
}
```

foo() {
 z = add(a, b);
}
↓
foo() { z = a + b; }

- Code size
 - I\$ instruction cache

memory hierarchy





RAM "1Gbs"

~~1000~~

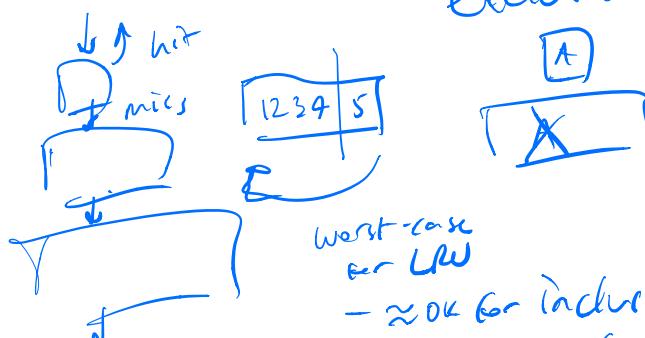
RAM bandwidth

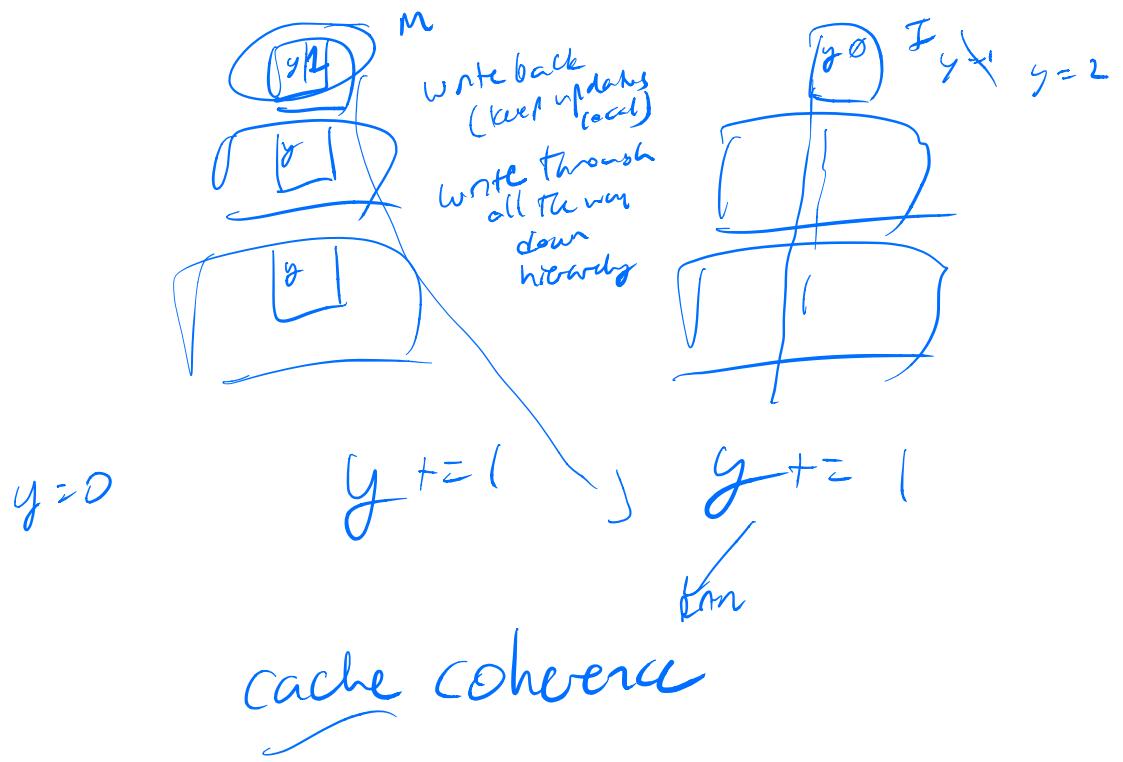
1000
10000

Caches — "replacement policy"

eviction "LFU"

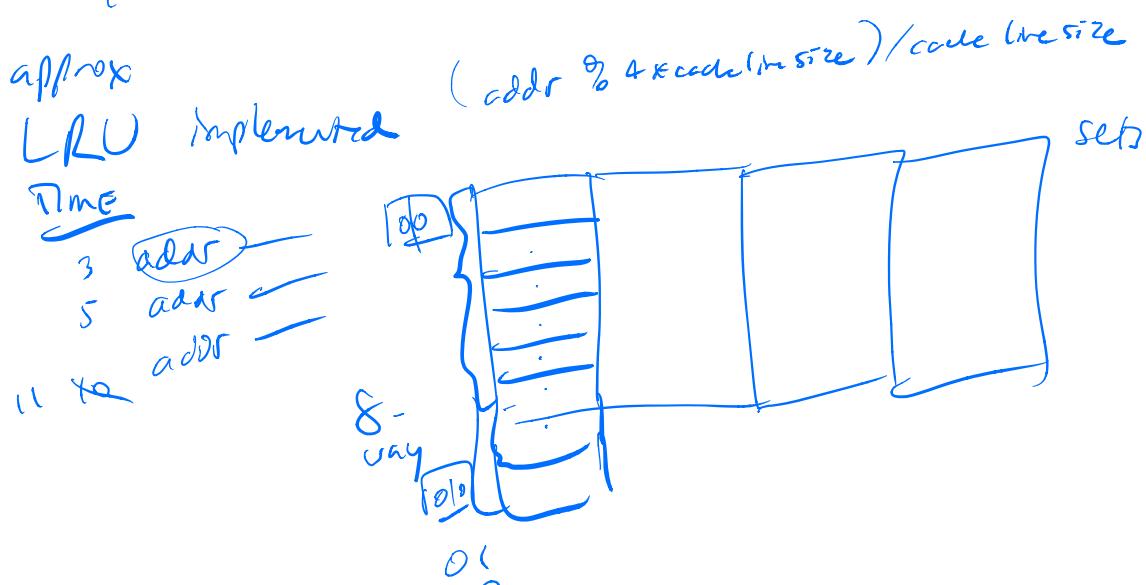
latency
wait time
caches wide latency





MESI protocol

- { Modified
 - { Exclusive
 - { Shared
 - { Invalid
- I changed it just me several CPUs have it (modification)
trash



Cache misses

Can be different
types -

Capacity - "hash collision"
Conflict -
Coherence - 3C's model

direct mapped - 1-way

2-way

4-way

8-way

16-way

fully associative - ∞ way

TLB

translation
lookaside
buffer

Cache of
virtual \rightarrow physical addresses

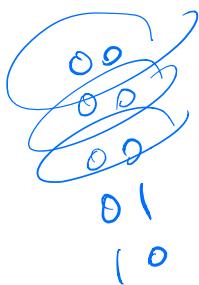
#include < stdio.h >

```
int main() {
    int a;
    printf("a = %p\n", a); 0x10000000
    return 0;
}
```

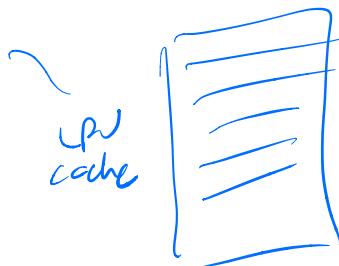
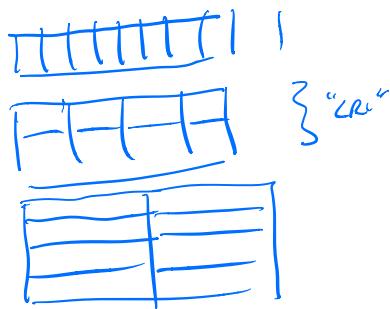
address

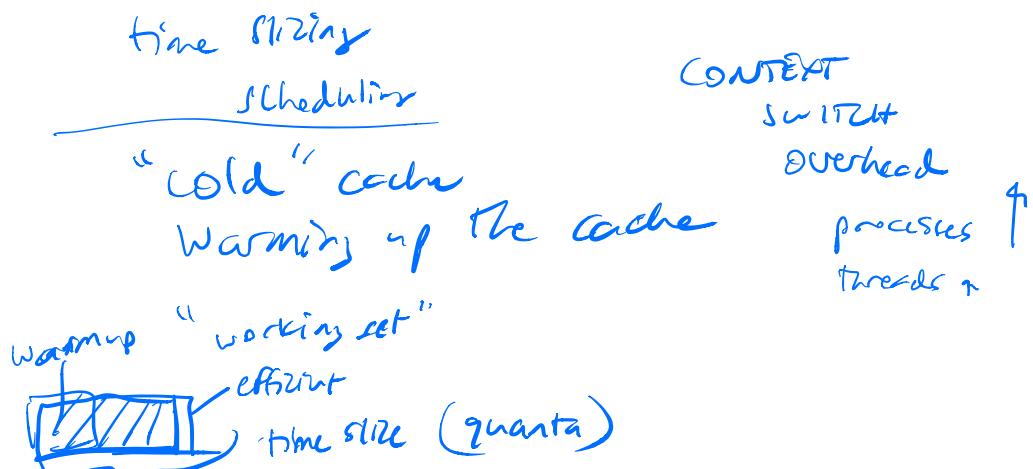
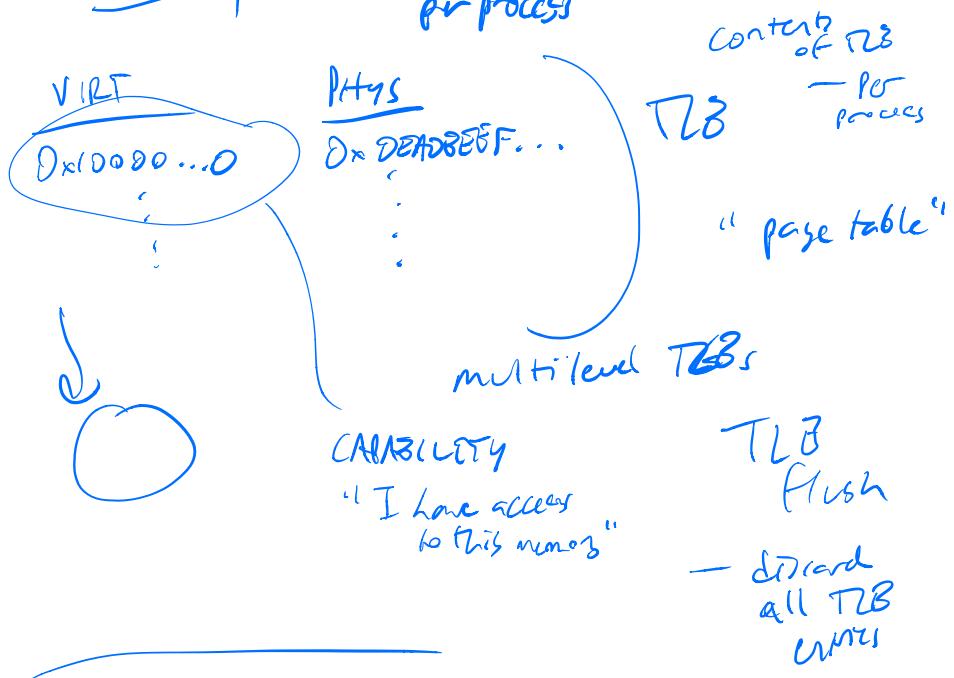
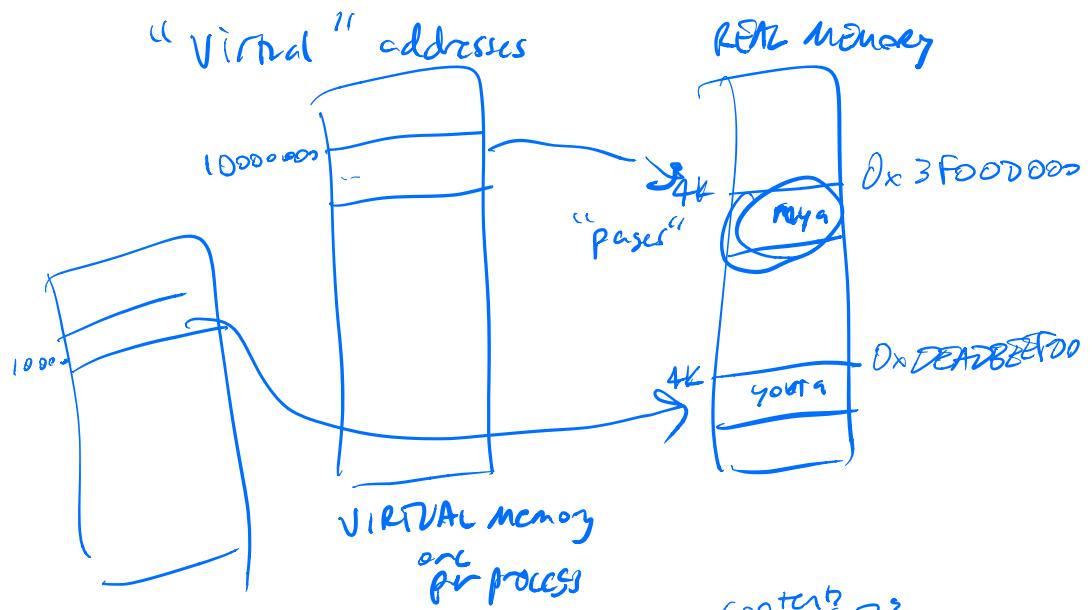
process isolation

10
11



11
11
saturated
counter







anomalous cleanup
across the quantum



REGS < QUANTUM

register allocation

"spill" - copy
contents
of register
to mem "eviction"!

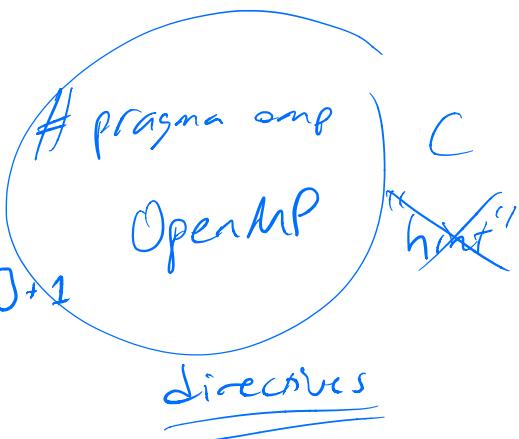
straight-line code
- predictable!

"graph coloring" NP complete!

Why is FORTRAN fast?

~~OMP PARALLEL~~

```
DO I = 1, 100
  DO J = 1, 100
    A(I,J) = A(I,J-1) + 1
  ENDDO
ENDDO
```



automatic
parallelization

"holy grail"

main thread
 S
 ↓
 for ($i=0; i < N; i++$) {
 foo();
 }
 ↗
 are there
 loop carried
 dependencies?

P
 }
 for ($i=0; i < N; i++$) {
 t[i] new thread (foo);
 }
 ↓
 foo₀ foo₁ ...
 [t[i].join()]

$\text{result}(S) \equiv \text{result}(P)$

parallel prog
semantically equivalent
to seq prog.

does `foo()` have side effects?

foo()
 {
 x++
 (_{global})
 }

read x into a register
 increment x
 write the result back

$x = 0$

A
 R₀ ← x
 R₀ ← R₀ + 1
 x ← R₀

$A ; B \Rightarrow z$

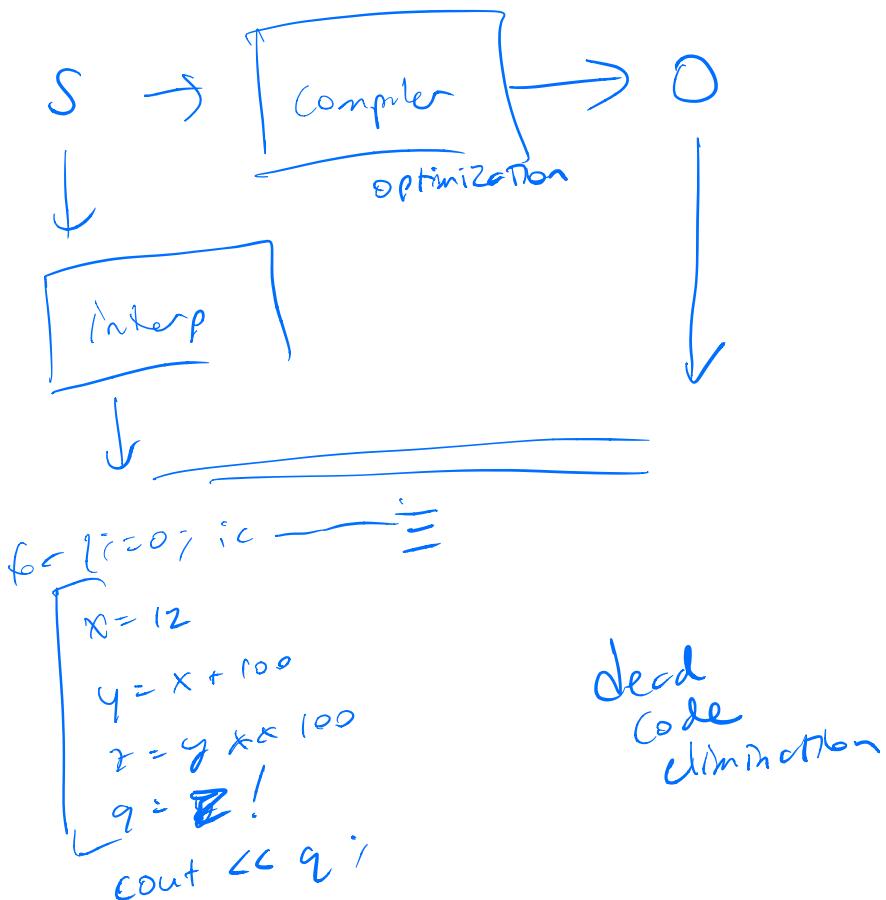
B
 R₀ ← x
 R₀ ← R₀ + 1
 x ← R₀

$$X \in \{1, 2\}$$

non-deterministic!

Race condition

Concurrency errors



```

foo()
{
    ~x ~
}

```

```

foo (int i)
{
    ~i ~
    return ack(i, i+1);
}

```

Ackermann

```
foo (int *int i a,) {  
    a[i] = ack(i, i+1);  
}
```

```
foo (int * a, int *int i b) {  
    a[i] = -;  
    b[i] = a[i] + 2;  
}
```

```
int x[200]; int *y = &x[0];  
v = foo(x, y, i)
```

aliasing

$p \{ \quad \}$ alias analysis
 $q \{ \quad \}$ pointer analysis

$p \cap q = \emptyset$ $\text{int } *p; \quad p \in \{ \}$

if ($__$) {
 $p = q;$ $p \in \{q\}$

} else {
 $p = r;$ $p \in \{r\}$

symbolic execution

abstract interpretation \rightarrow
 $\wp \in \{q, r\}$

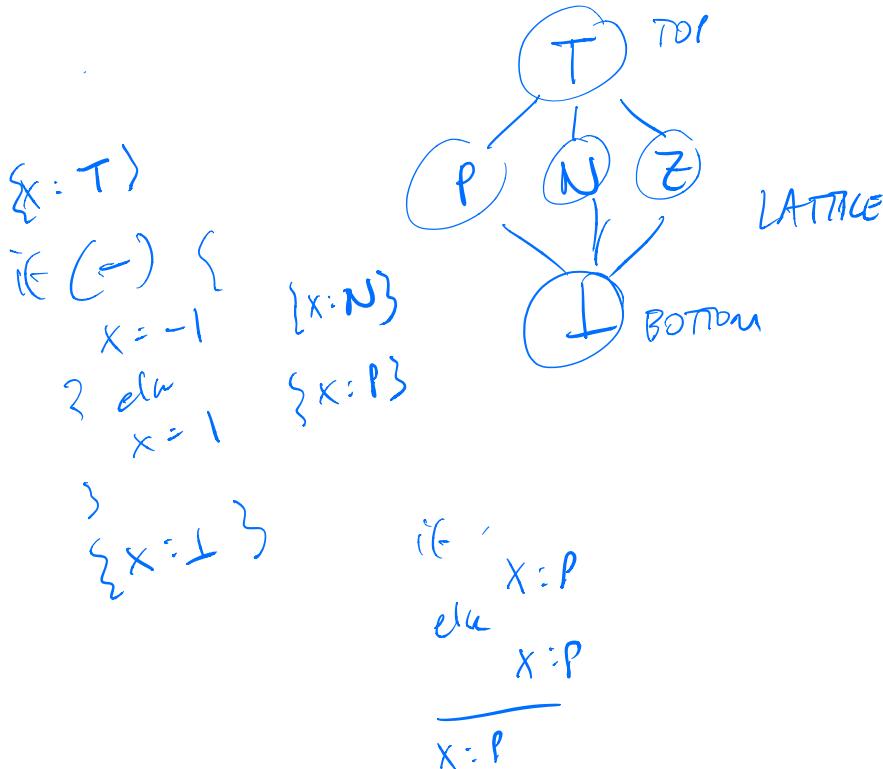
$$\begin{aligned} x &= 1^2 \\ x &= x - 100 \\ x &= \dots \end{aligned}$$

if ($x > 0$) {
 |
 }

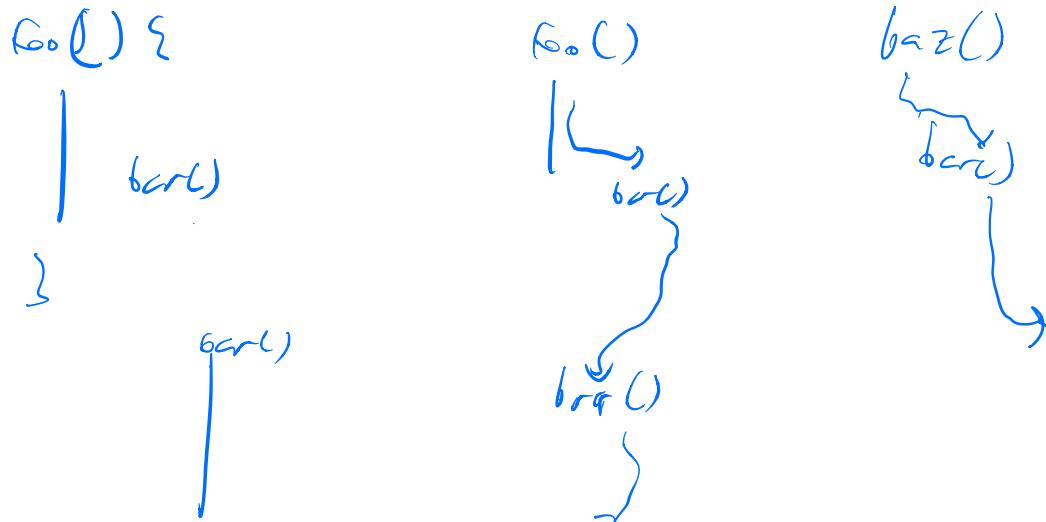
$x \in O, A, P$

$x = _ \quad \{O, N, P\}$
 $x = x - 1 \quad \{\}$
 does this run over?
 $O \Rightarrow O$
 $N \Rightarrow P$
 $P \Rightarrow N$

CONSERVATIVE APPROX
 PROG BEHAVIOR



intraprocedural vs. inter procedural



flow sensitivity



Steensgaard's algorithm

Context sensitivity

PCR sensitivity

false positives

precision -
recall

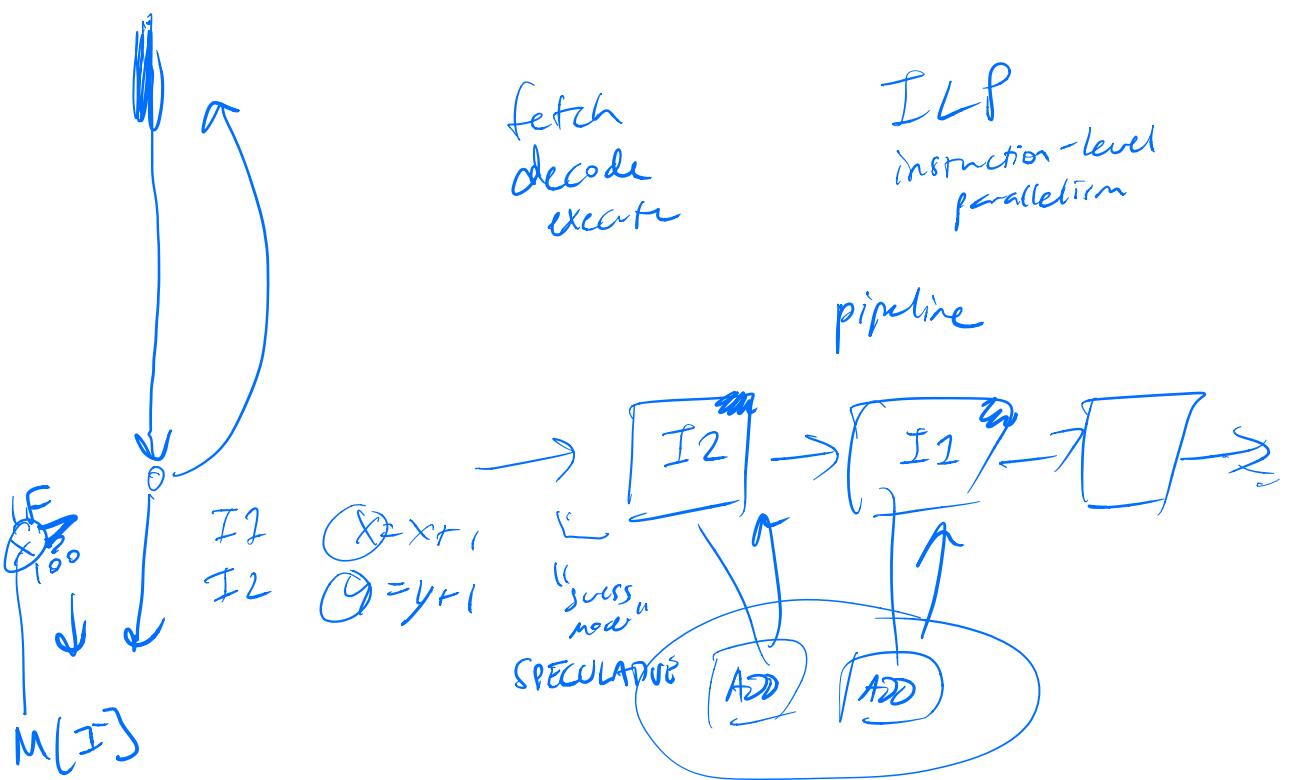
false negatives

tradeoff

$$V \sim \{\emptyset\}$$

-W

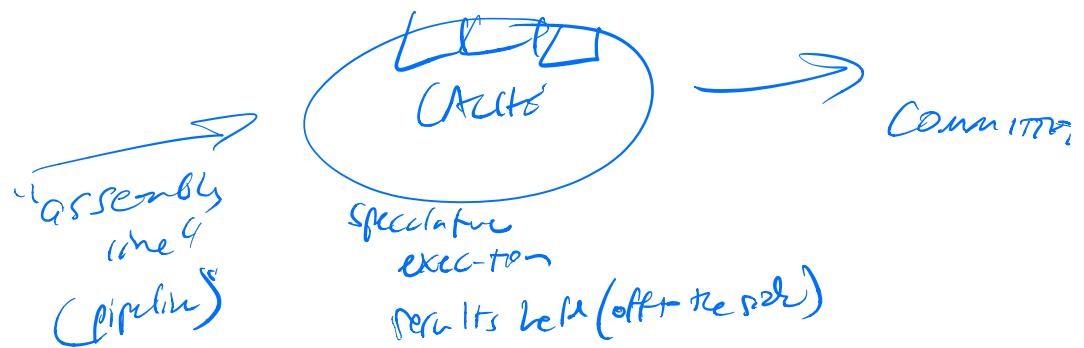
if ($x == \underline{\quad}$)
 $y = x / @ V$



Pipeline stall

"bubble"

"BACK EDGE DATA"



Speculative execution
hide memory latency

SPECTRE / Meltdown

timing channel

Covert channel

side channels

last branch taken

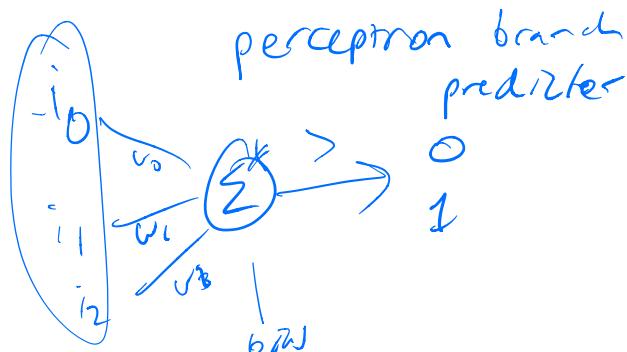
85-90% effective

PHAT

Modern

branch predictors

99.9% effective



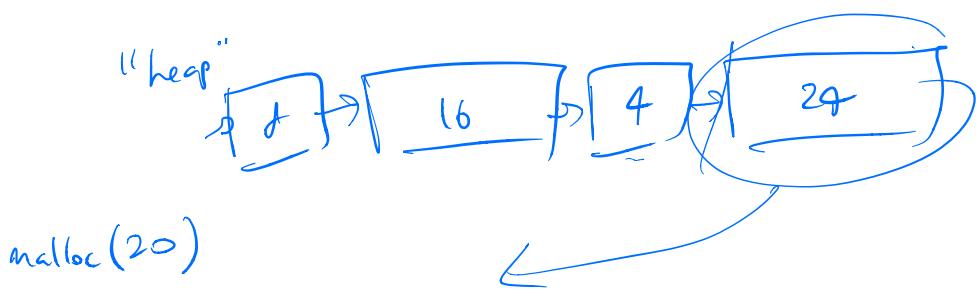
dynam. memory allocation

globally

stack variables

dynamic memory

"the
heap"



Object 24

First-fit → SPACE INEFF

best-fit

SPACE
EFFICIENT

slow ...

$$n \cdot O\left(\log \frac{M}{m}\right) * n$$

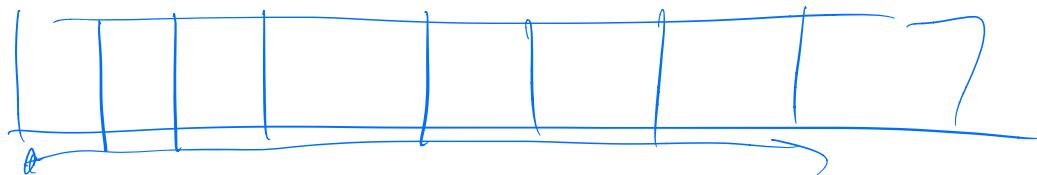
M - biggest
m - smallest

$$O(M \times m)$$

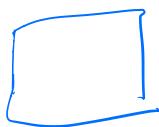
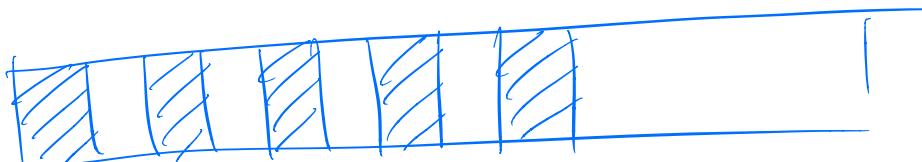
best fit

worst-case
"fragmentation"

$$= \frac{\text{mem consumed}}{\text{mem requested}}$$



D → D → D



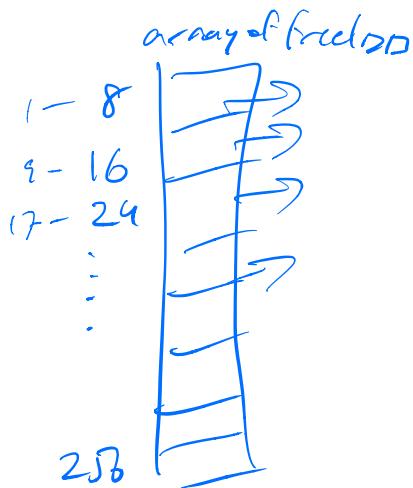
Compaction
NOT IN C/C++

(L1E')

Python, Ruby

"custom allocators"
"ad hoc allocators"

malloc - "too slow"
Reconsidering Custom Mem. Allocation



Py-object
- malloc

if ($sz \leq 28$) {

use small obj.

} else

use malloc

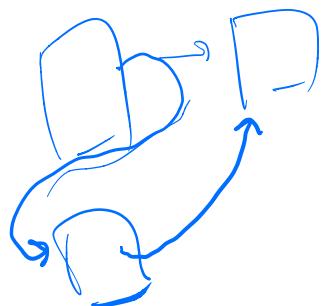
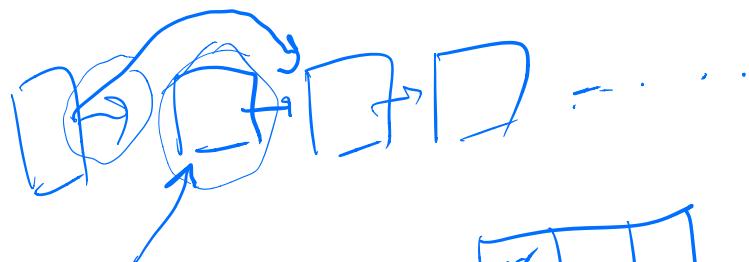
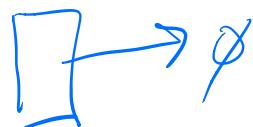
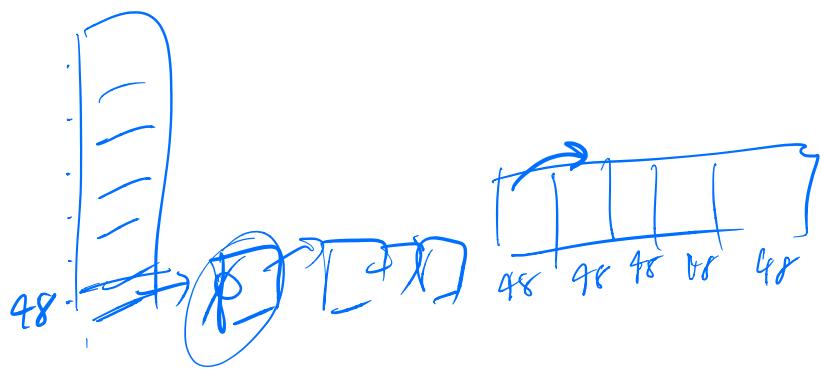
int ≈ 32 bits C Python
long ≈ 64 bits 46 24/28

"a"
1/16 50

{ "a": }
~~2~~

240

60 20 75 100 140

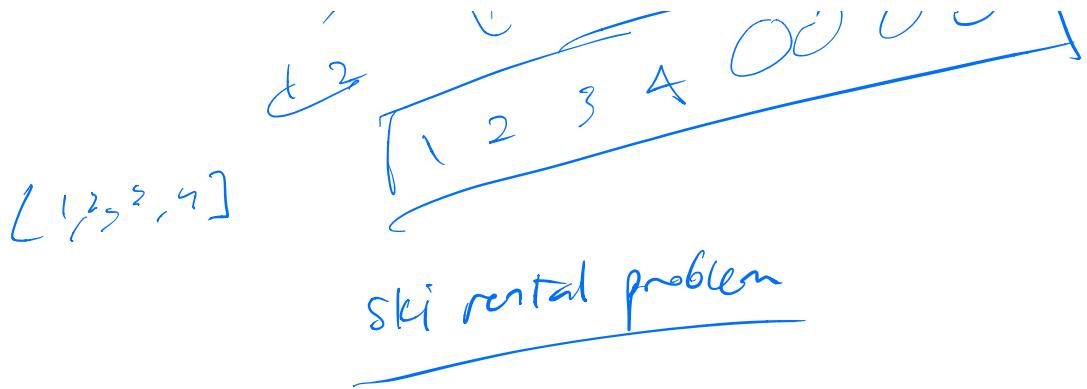


48		
REF		

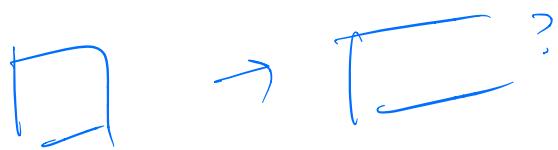
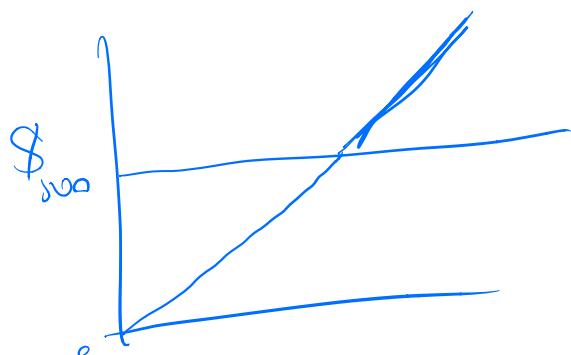
ref count

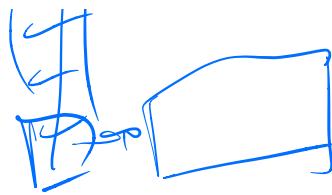
$$x = [1, 2, 3]$$





\$50 rent
\$500 buy
rent until
you reach
cost of buying

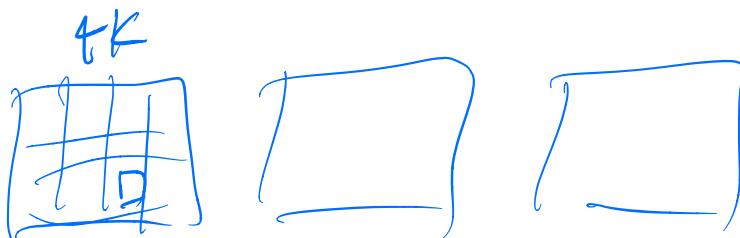
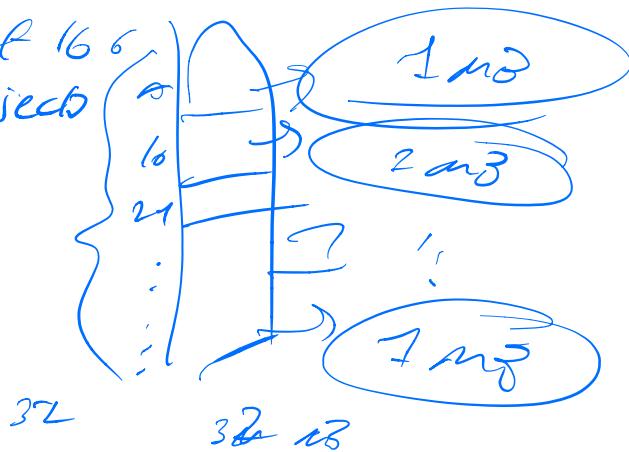




allocate 1MB of 8 64B objects
free them all

allocate 1MB of 16 64B objects
free them all

actual
MMT 1MB
32x

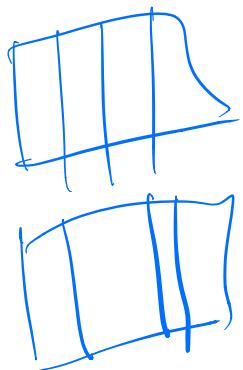


Space-time tradeoff

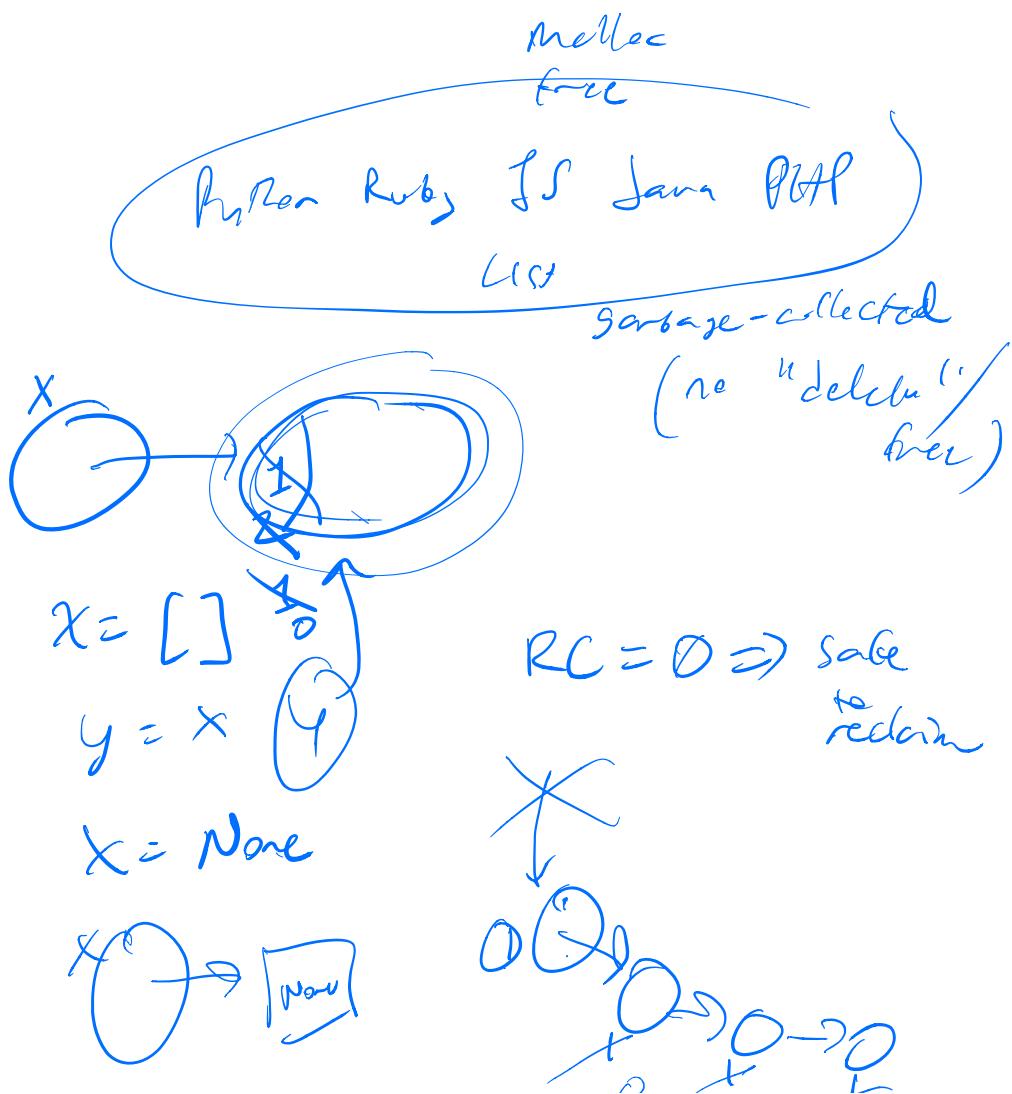
$$f(a) =$$

$$\begin{aligned} f(0) &= a \\ f(1) &= b \\ f(2) &= c \end{aligned}$$

0	a
1	b
2	c

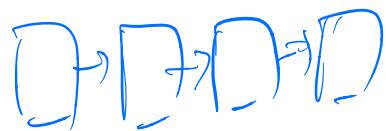


Reference counting



$y = \text{None}$

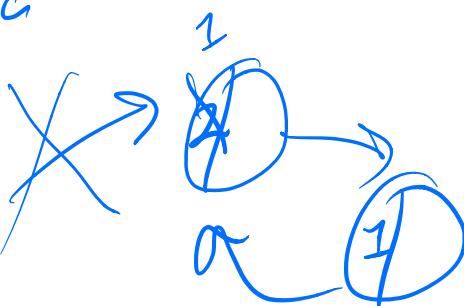
$\leftarrow \circ \rightarrow$



doubly-linked (2+)

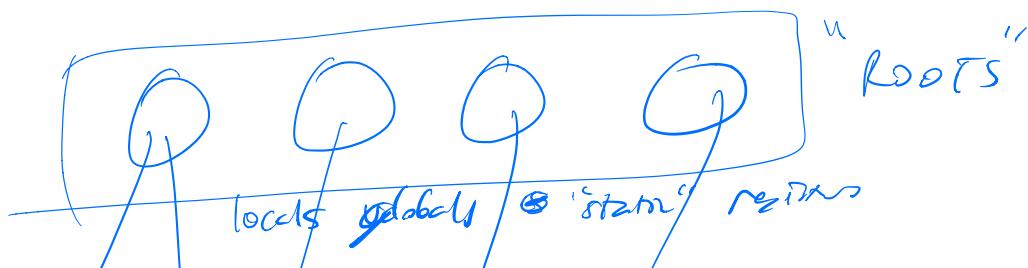
$a = \{\text{'key'}: 0\}$ CYCLE

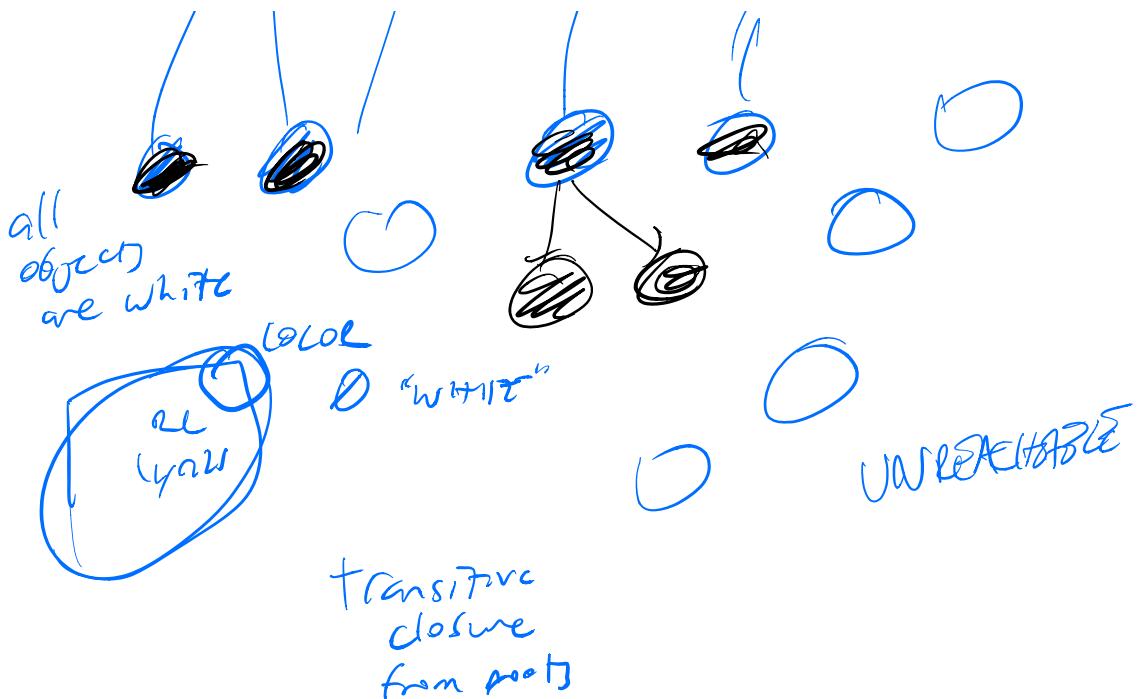
$a[\text{'key'}] = a$



HC
(1956)
Incomplete
GC algorithm

Mark-sweep (1978)
"generational" GC (1983?)





David Unger

smalltalk self

Jg Holzle
Gerd Hanberg

Ole Ascan posthum created

dynamiz

Lco
Rak

Galgorithms

JIT compilation

dynamic
languages

polymorphism

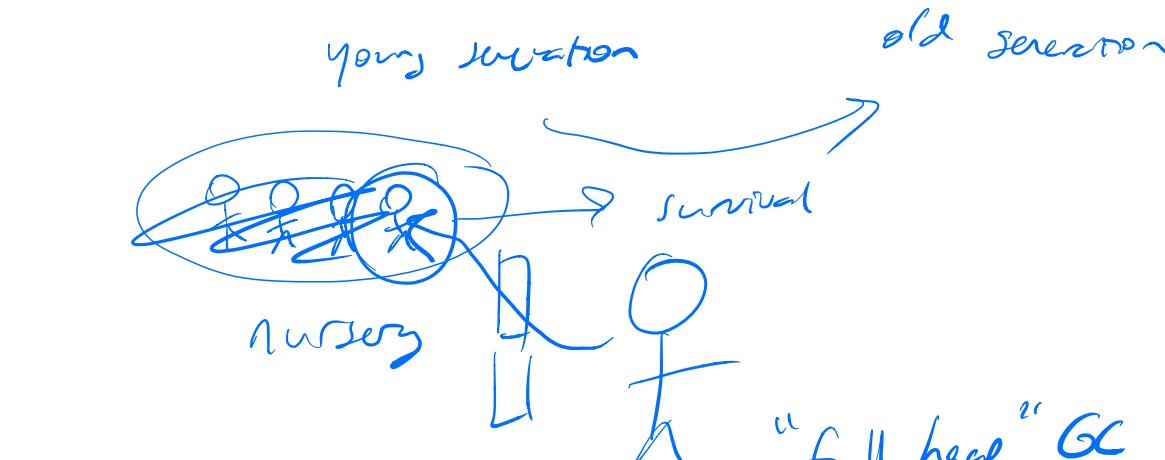
mine
cache

"Pic"

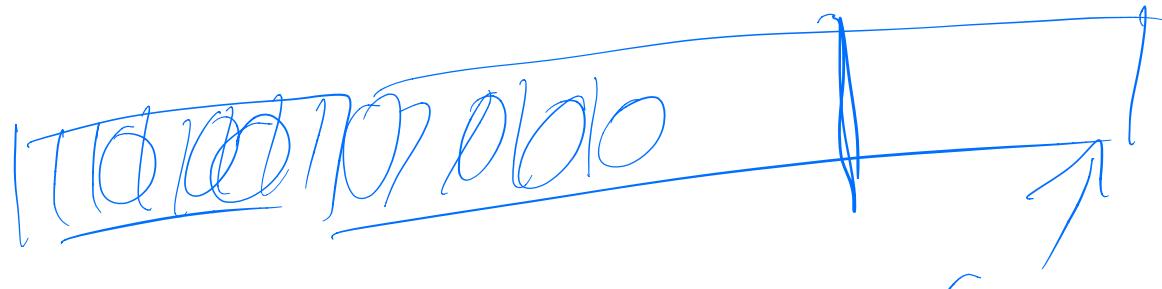
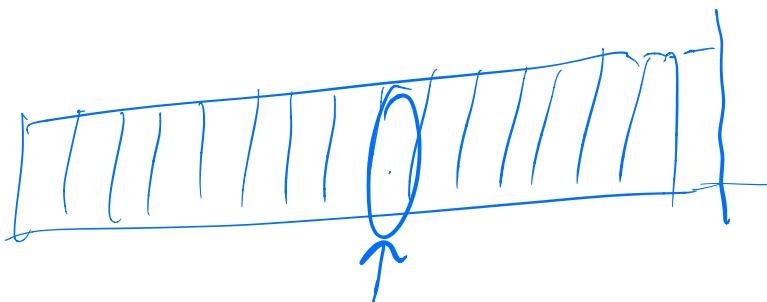
600sl

→ v8

most objects die young



gen GC →
takes pressure
off the full
heap GC

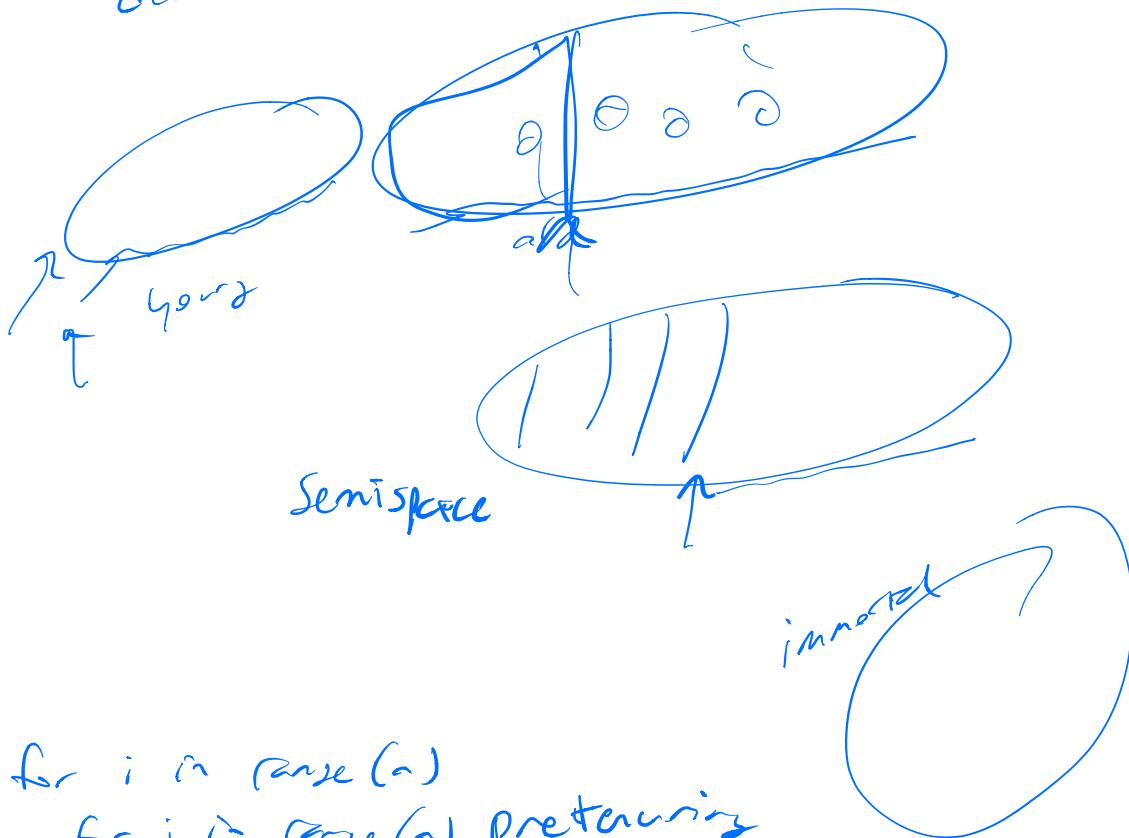


GC

$\sim 2 \rightarrow 5x$

"live
objects"
"dead"

Quantizing the cost of
GC vs. Explicit deallocation



for i in range(a)

 for j in range(n) pretenuring

 for k in range(n)

$$C[i][j] += \delta A[i][k] + \delta[k][j]$$

• pyc

PyPy

Ontogeny recapitulates phylogeny

SCRIBERS

Lillian
& Sven

OS: single user, single process, no protection

multiple users, multiple processes

"time-sharing"

time slices

memory protection

file system protection

:

Mac 1984

1 user

1 process

no protection

VM

"MMU"

memory
management
unit

PC

Apple
IBM PC

:= 1981

Motorola 6802
Intel 8080...
68000 Macs

) no MMU

MS-DOS

Microsoft

Microsoft BASIC

Altair

Gates & Allen

CP/M

8080

Z-80

8088

Gary Kildall

latencies

"security
through
obscenity"

AT&T

Unix

BSD

OS X

iPhone

stripped down Unix!

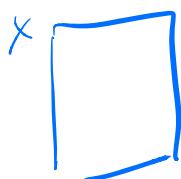
→ multiple processes
mem protection
security model!

PL?

→ dynamically typed lang — Values have types
(statically typed lang) — Variables have types

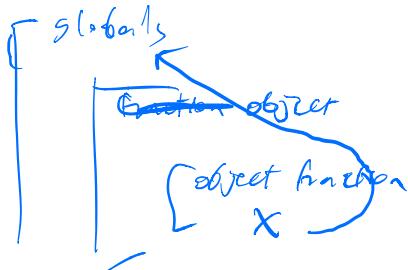
$x = []$
 $x = {}$
 $x = ""$

→ dynamically scoped



static / lexically scoped

"weird
scope
mics"



Match the innermost scope

foo() {

~ x ~
} $x = x + 1$

walking call stack

bar() {

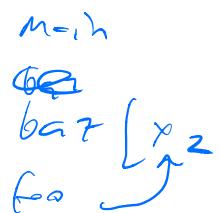
int ~~x₁~~
foo();

}

bar() {

int ~~x₂~~
foo();

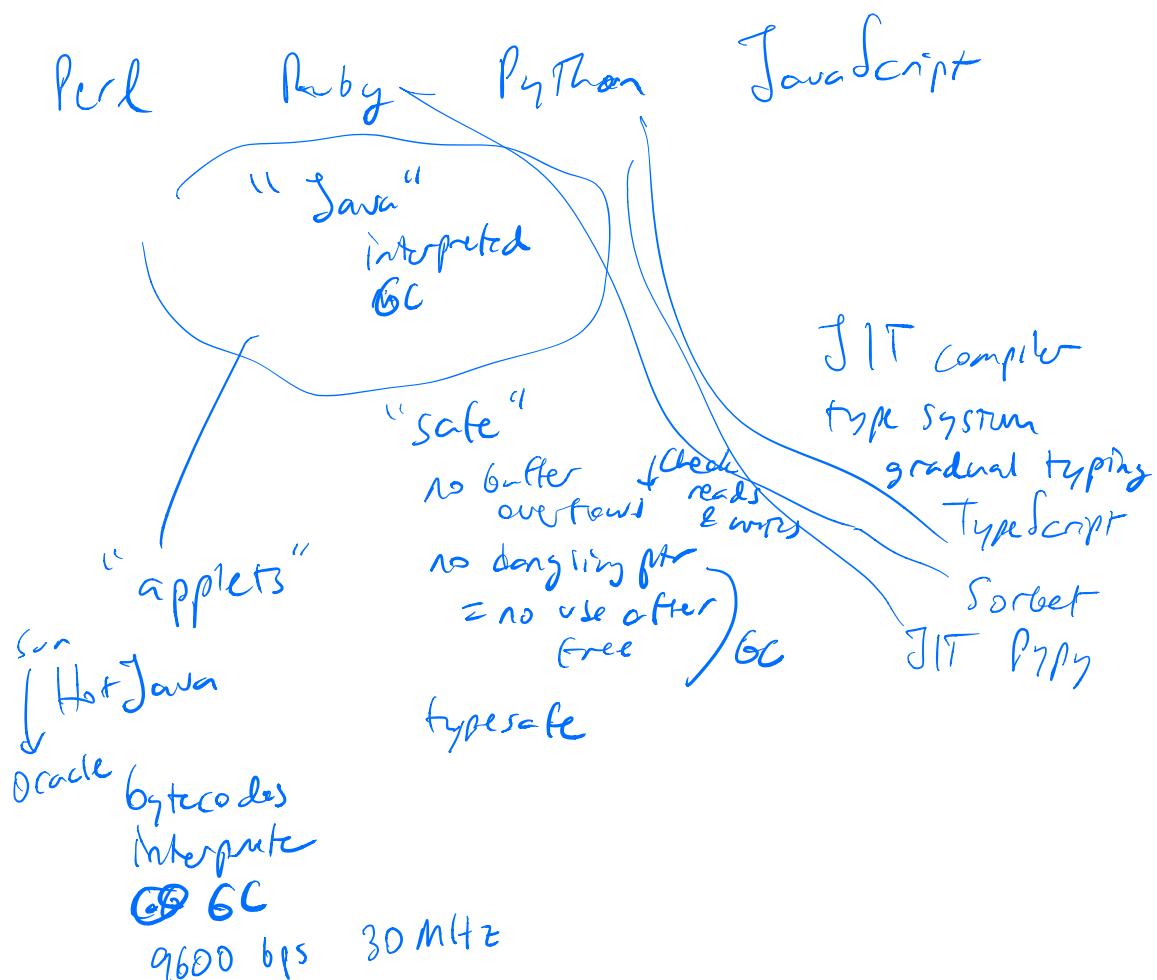
}



→ Interpreter eval

hard to compile
into efficient code

— reference country GC



Netscape → Mozilla
↓ Navigator

1995 Brendan Eich

Mozilla { }
"curly-brace languages"

LISP Scheme
((()))

LiveScript 10 days

Browsix

Doppio

asm.js

subset that is easy to compile

asm.js
 $x = (v \ 1 \ \emptyset);$
 $x = a + b;$

Web Assembly

- safe
- efficient
- fast to parse!!



-- loading form compile connection

- ISA instruction set architecture
 - arbitrary memory (module mem protection)
 - arbitrary structure jmp GOTD
 - hard to analyze

Register-based
stack-based
architecture

A8
CD

NaCl
p NaCl

Compiler
Smart
"safe"
subset
of x86
ARM

|||
PUSH
POP
ADD

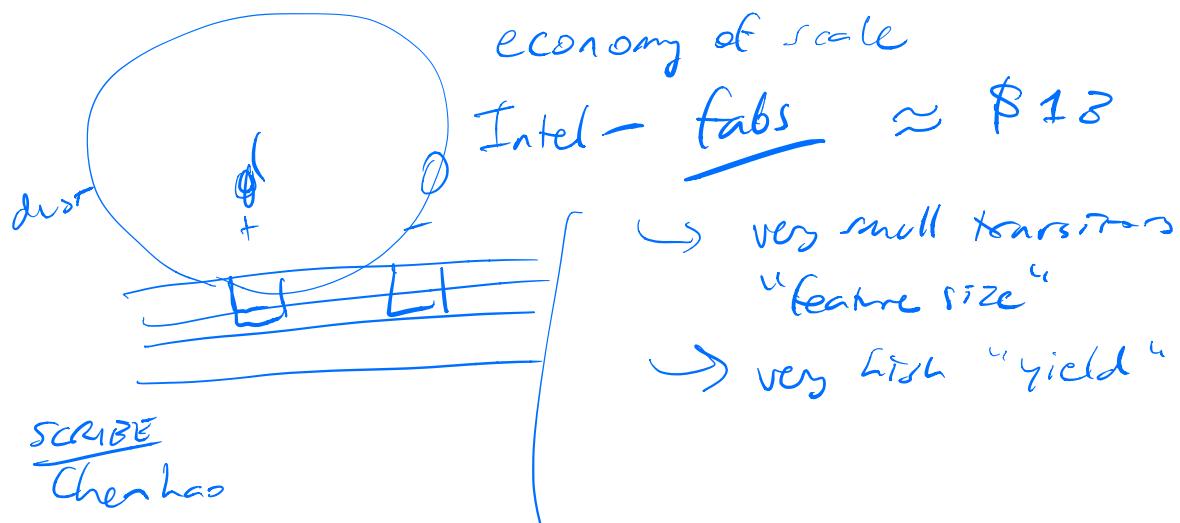
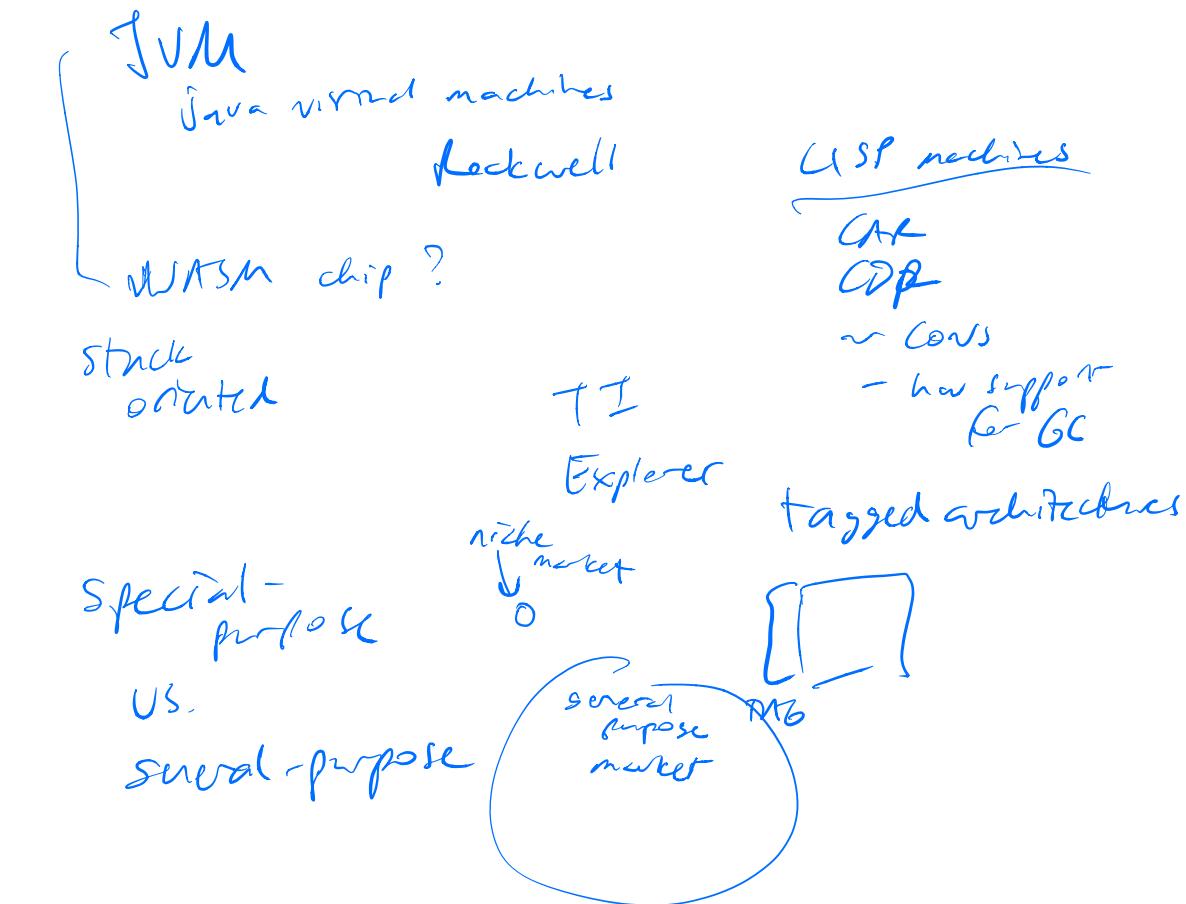
ROP
Return-oriented
programming
Return-to-libc



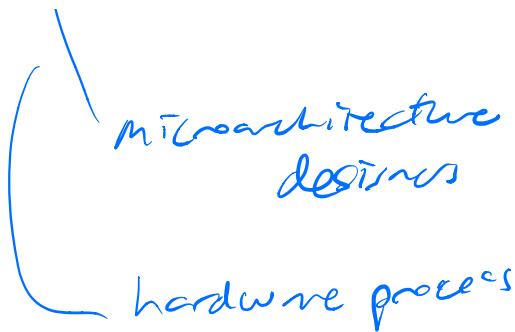
CFI
Control-
Flow
Integrity

WASM - CISC or RISC?

simple small instruction set



Shahrya

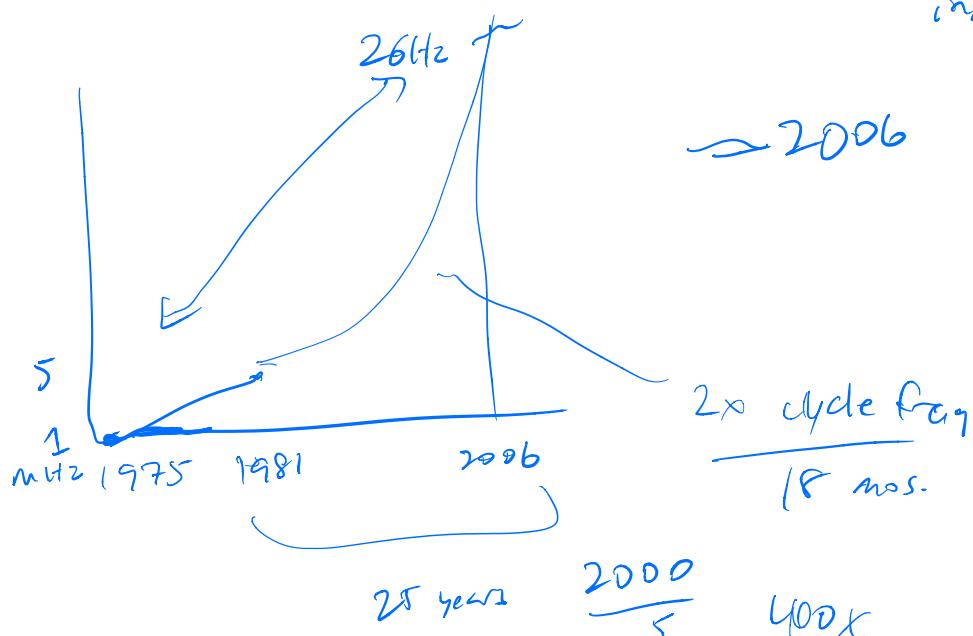


AMD

- fabs - not as good
- ⇒ low yield
- lower profits

→ 32-bit to 64-bit

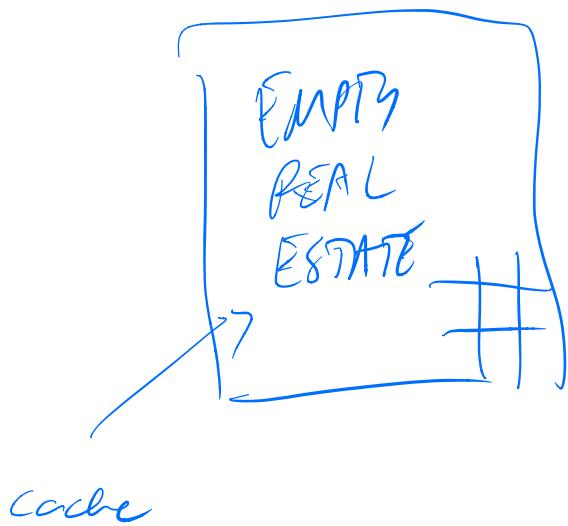
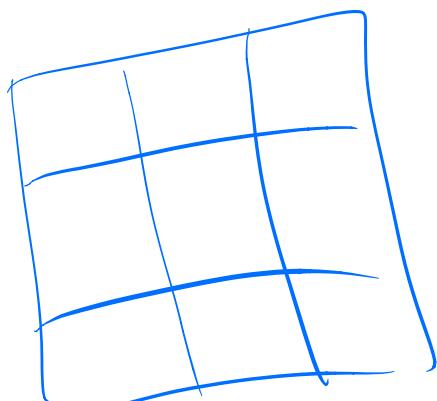
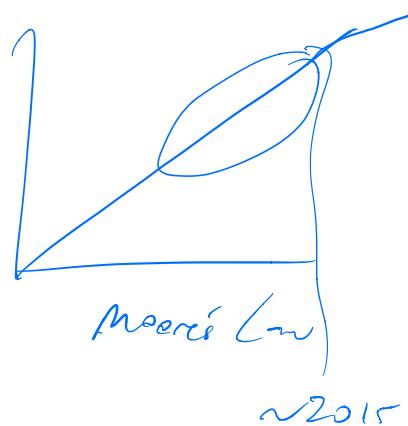
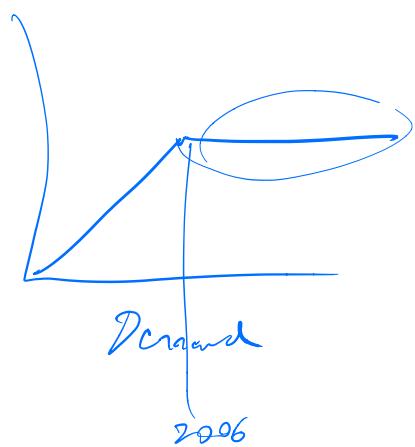
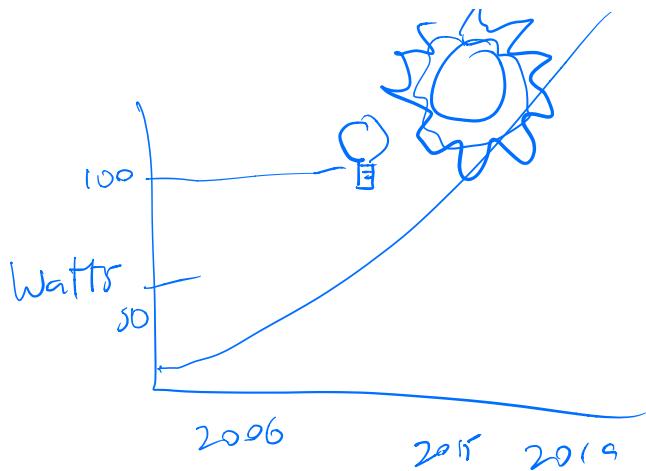
≈ multizone

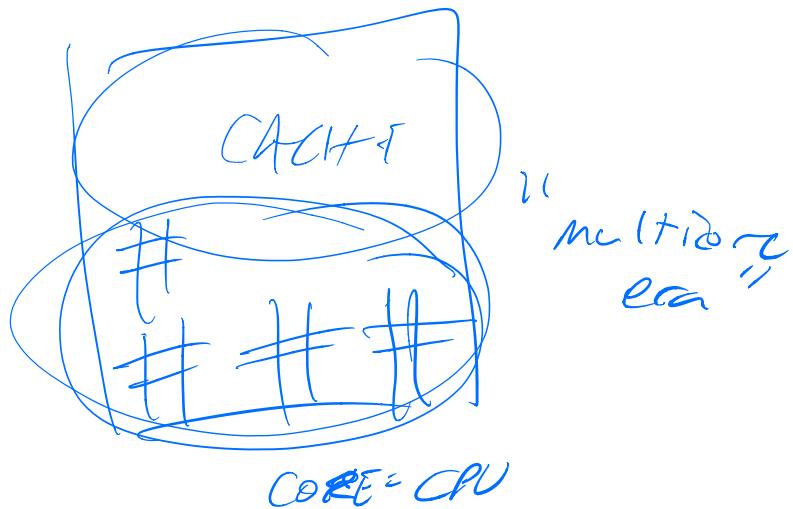


Demand scaling

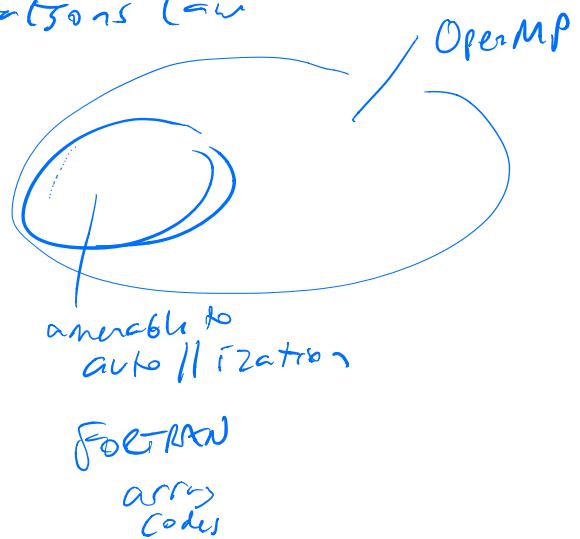
increase clock speed
w/o increasing voltage
→ w/o
increasing
heat.

≈ 2006

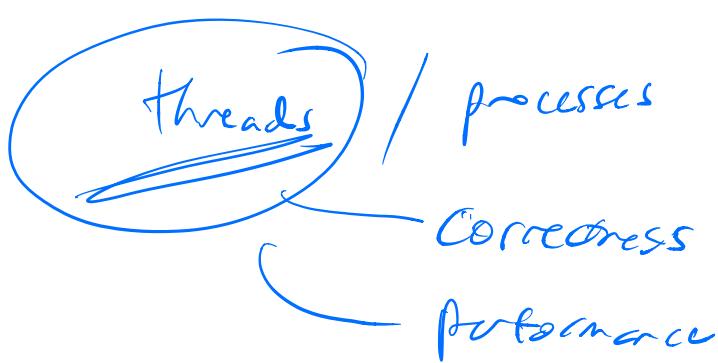


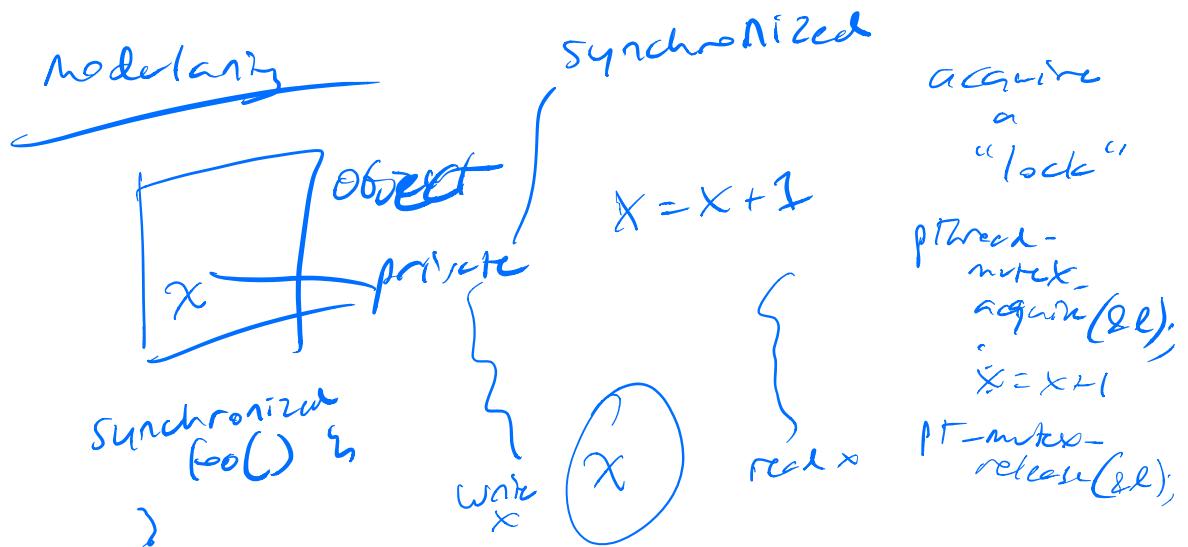


Amdahl's Law
Gustavson's Law



FORTRAN
array
Codes





"atomicity"

transactions: all or nothing

atoms - indivisible

atomic

Demostrates

transaction

- commit
- abort
- rollsback

"conflict"

retry

spin lock

```

lock
while (l == 1) // while locked
    ;
    ; // wait — sleep
(atomic) if (l == 0) { l = 1; } // acquire lock
                                // atomically
TST & Lock —

```

~~~

(unlock)  $l=0;$

"thundering herd"

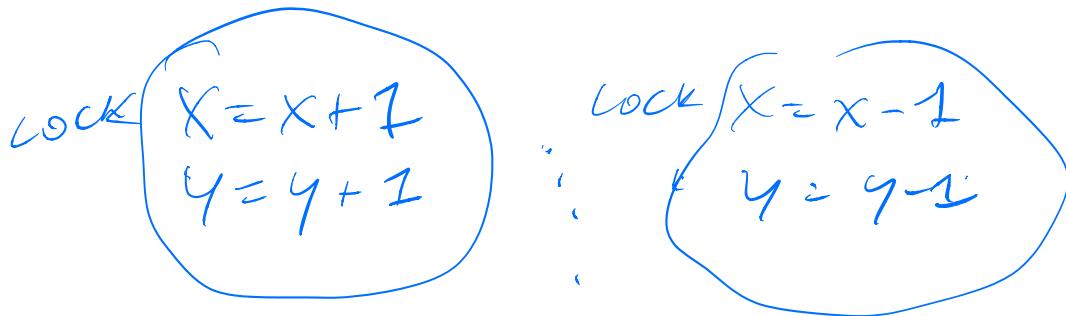
fair locks  
queuing locks

---

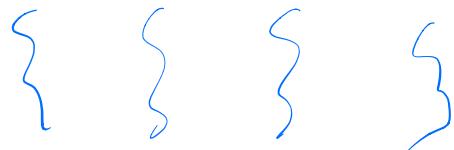
adding locks everywhere

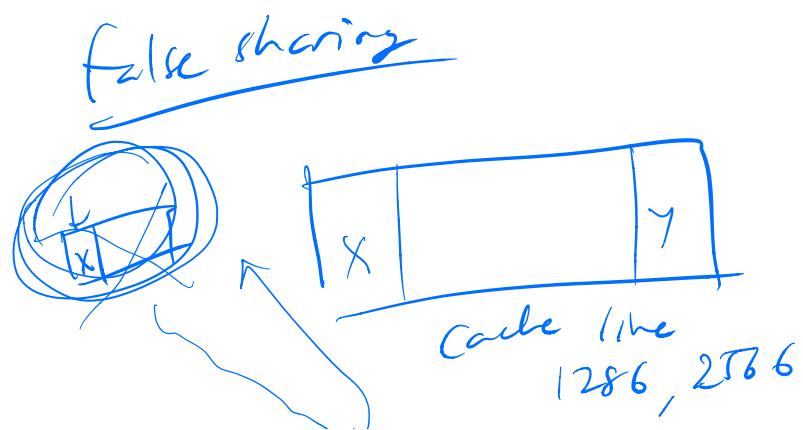
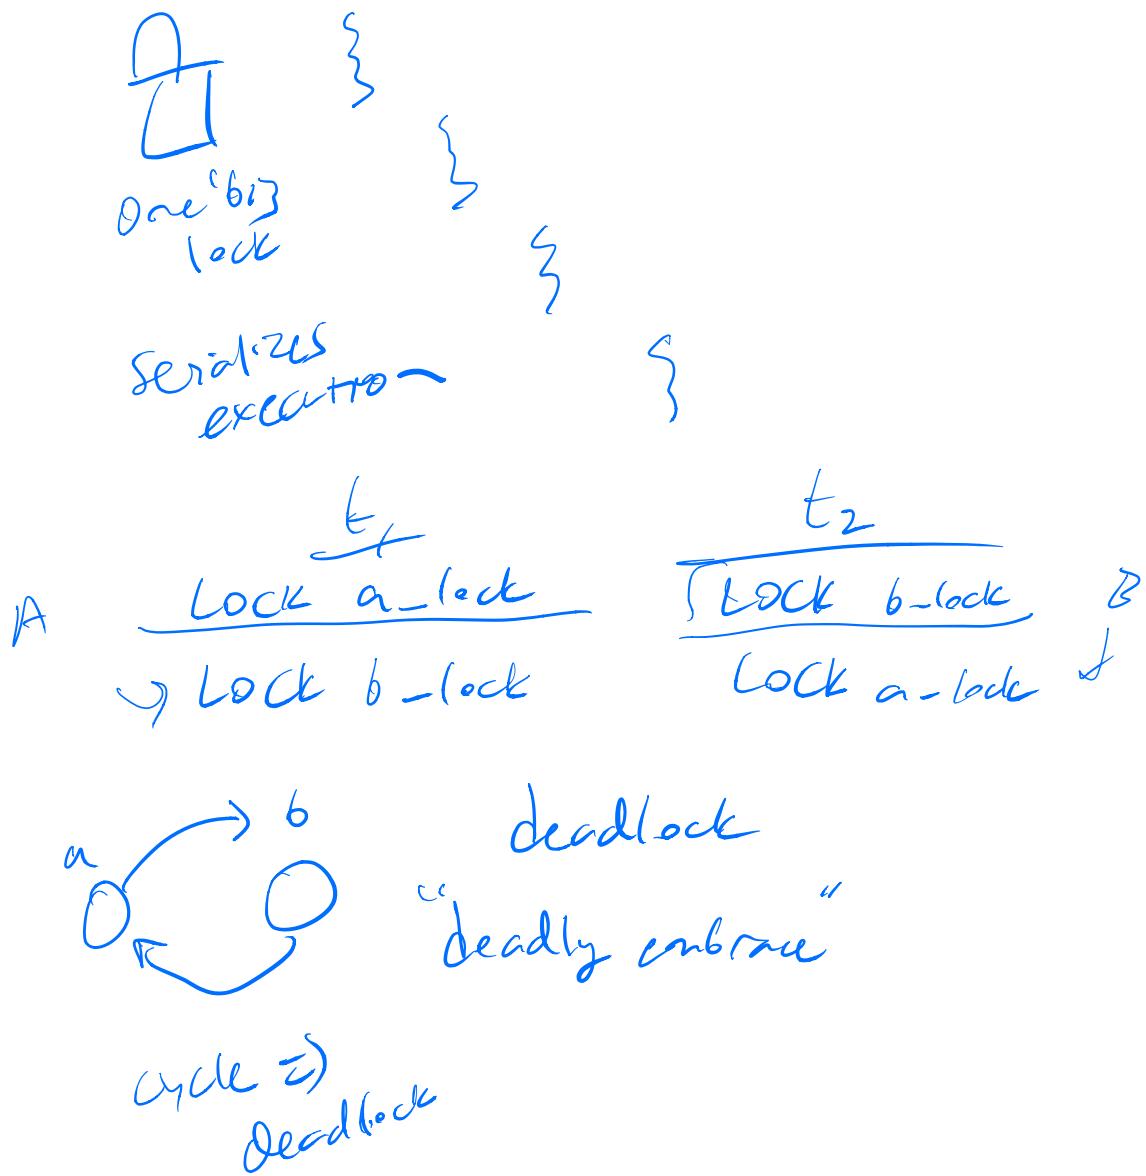
neither necessary (for correctness)  
nor sufficient (for performance)

$$x=0, y=0 \quad \left\{ \begin{array}{l} x=y \\ \end{array} \right.$$



fine-grained locks  $\rightarrow$  coarser  
for atomicity



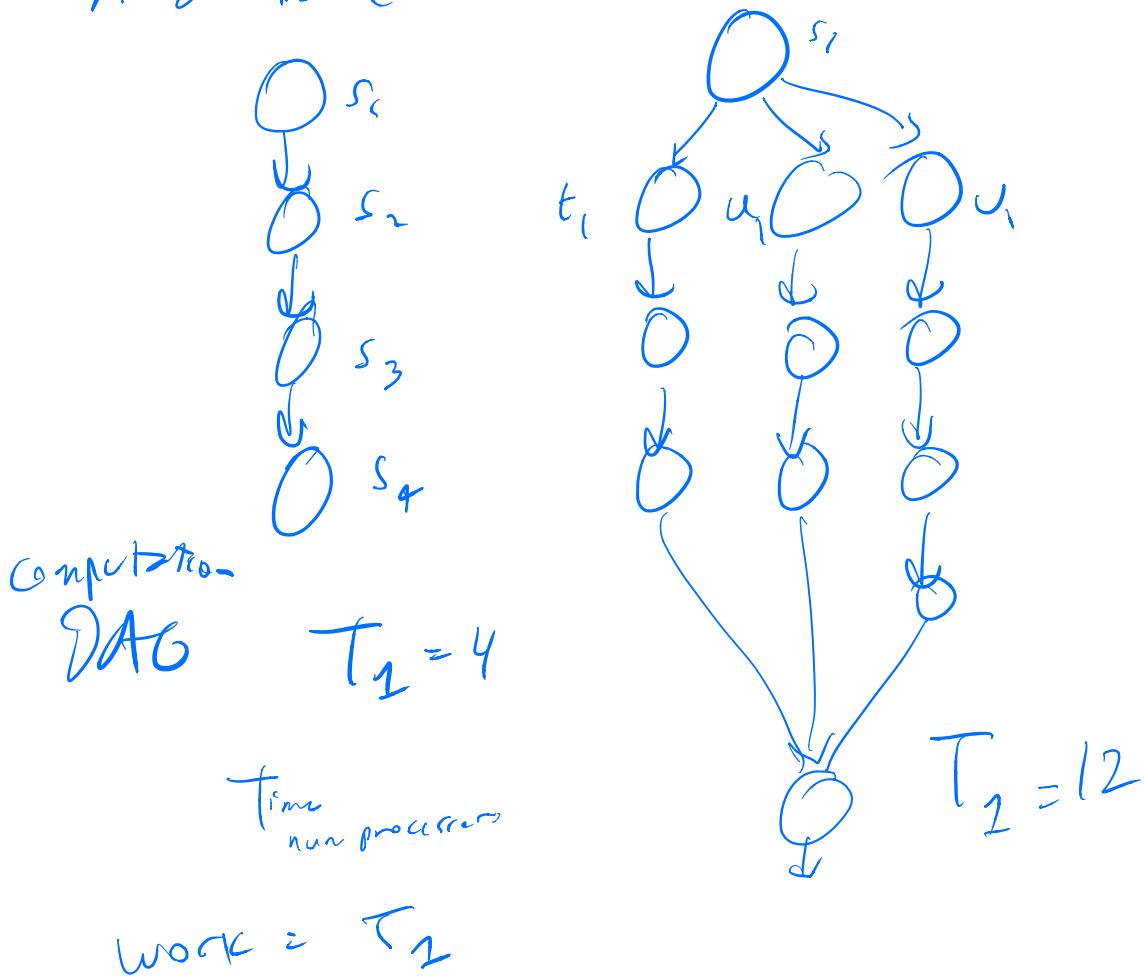


$\sim$

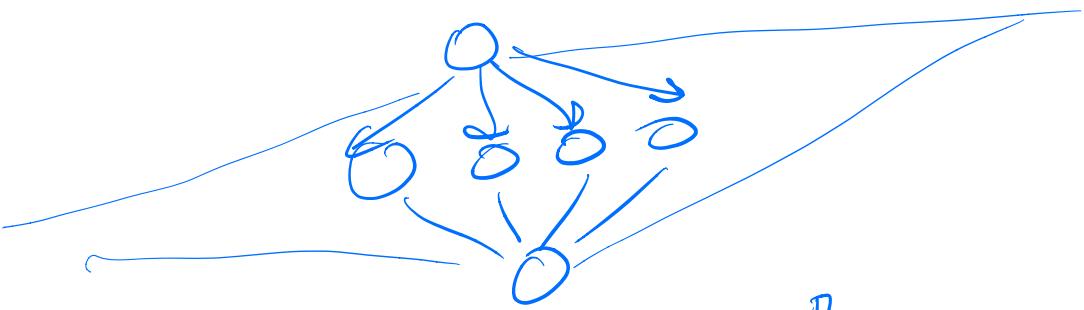
Cache line  
pins parsing

$\approx f\text{-}Ox$  slowdowns

Amdahl's Law



" $T_n \approx \frac{T_1}{n}$ " scale - perfect scaling



$T_{oo}$  = control path  
 max from start to leaf from  
 "header of Data"

"serial part"

$$T_p \approx T_1 + T_{oo}$$

$T_1$

parallel work

P

serial part

parallel work

90% of program — perfectly parallelizable

embarrassingly parallel

$$T_{pp} = \frac{90\%}{PP} + 10\%$$

0

100s → 50s

max speedup  $\rightarrow$  DX

VM hosting  $\approx$  embarrassingly II  
scales w/ # of instances

Amdahl's Law  
→ 1 workload fixed size

today's world  
Netflix → indep. customer moves  
emb.-II  
 $\rightarrow$  # customers

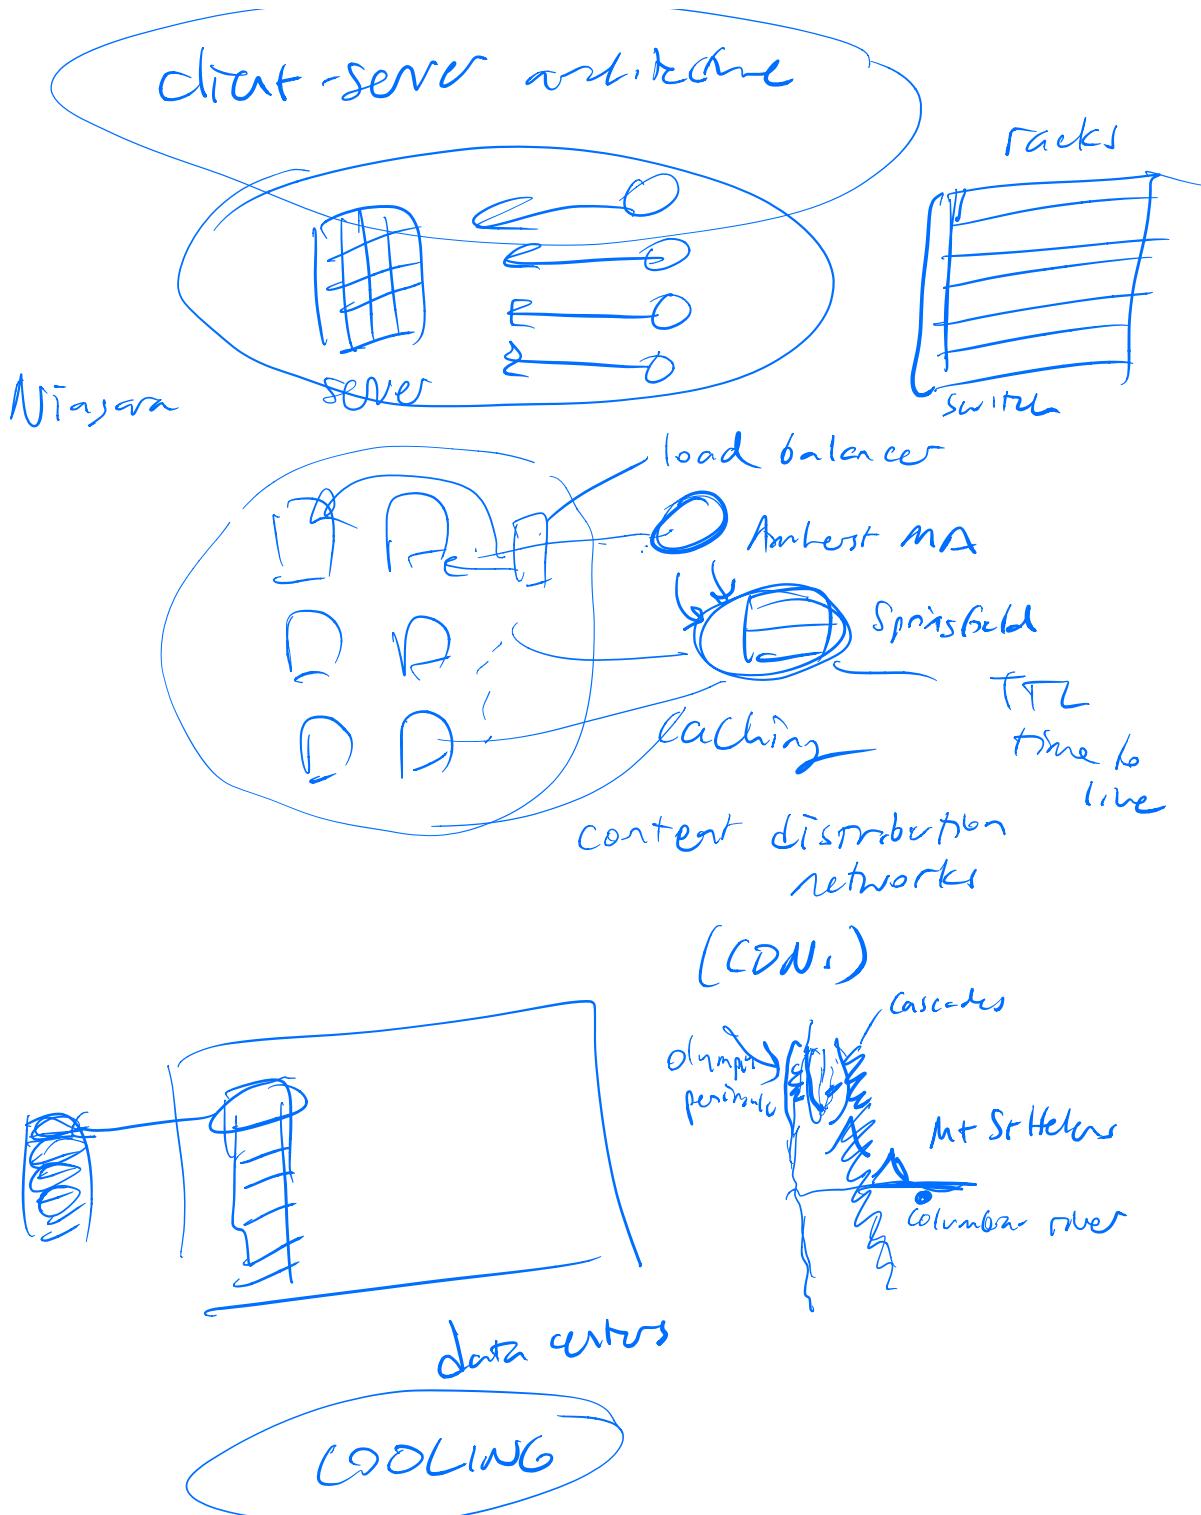
Google  
FB

Gustavson's Law

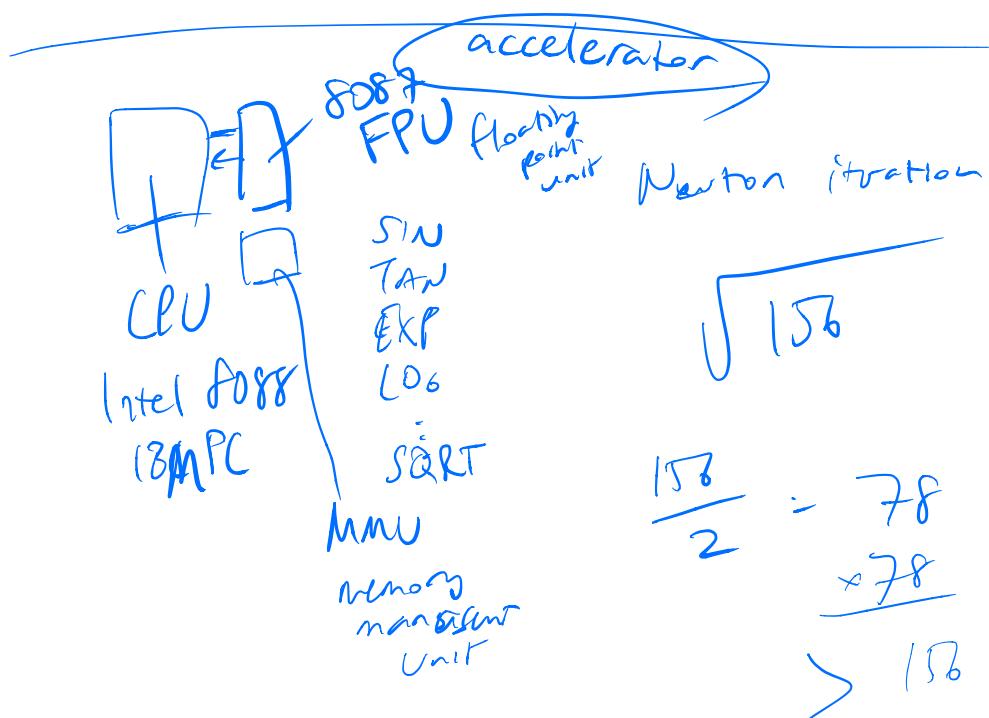
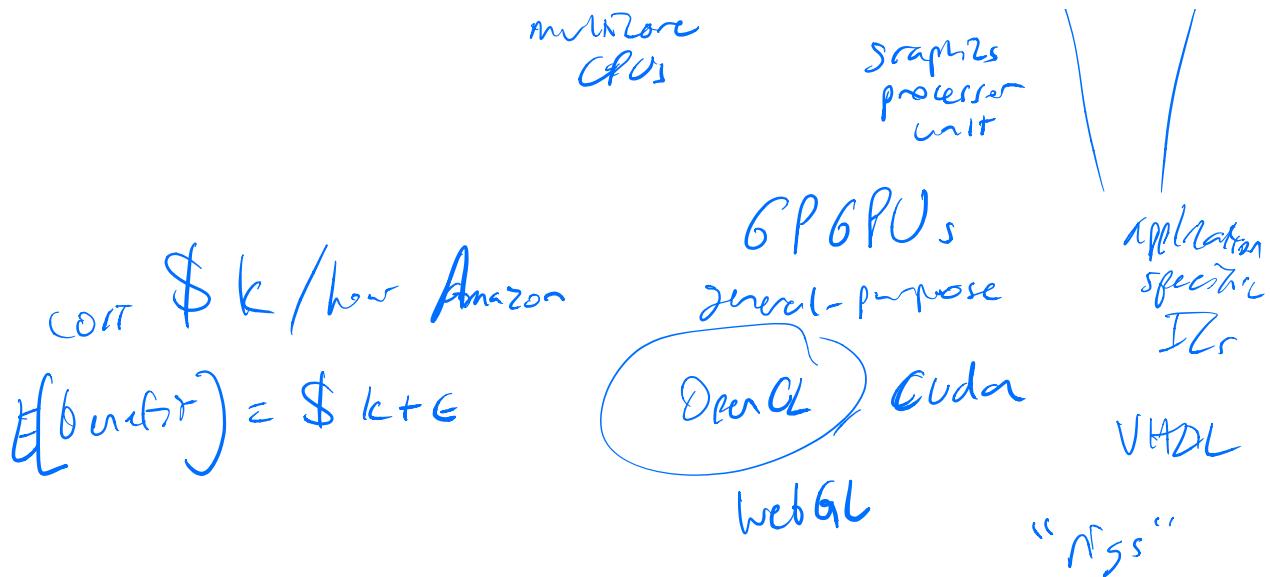
$$\Delta S \gg \Delta P$$

↑  
size of problem      ↑  
# of processors

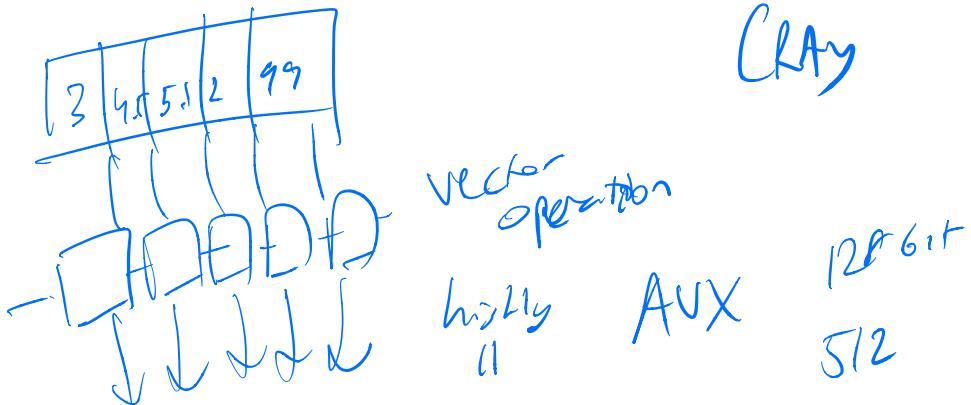
---



CPUs → cluster → GPUs → ASICs



Accelerators = CRYPTO  
Vectors



"Neural" ← HYPE

s/neural/matrix multiply/g

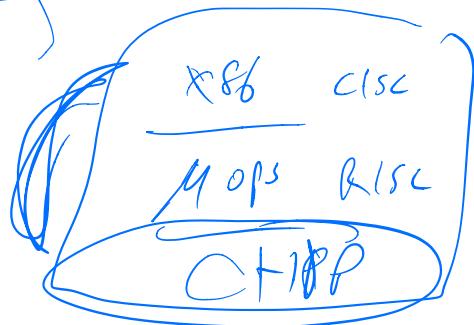
"tensor"

NPU  
TPU → Google  
tensor processing unit

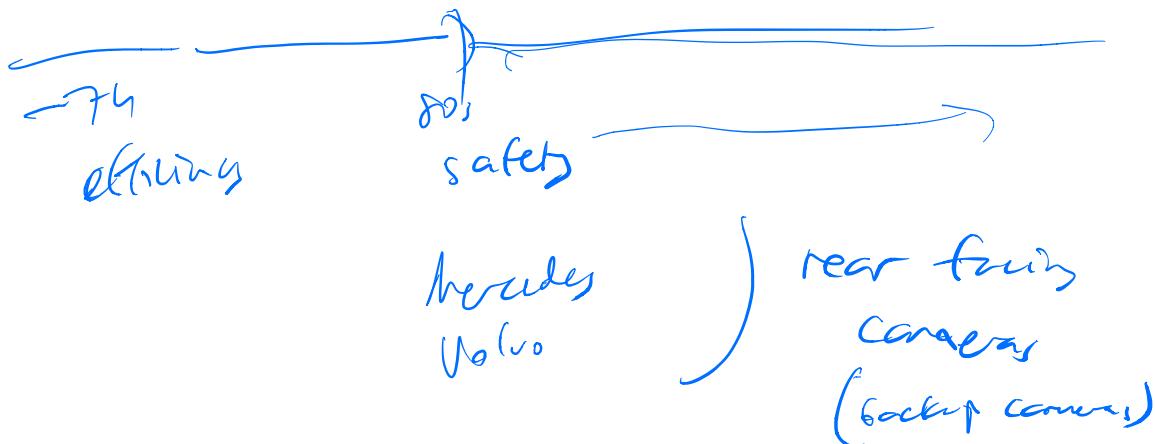
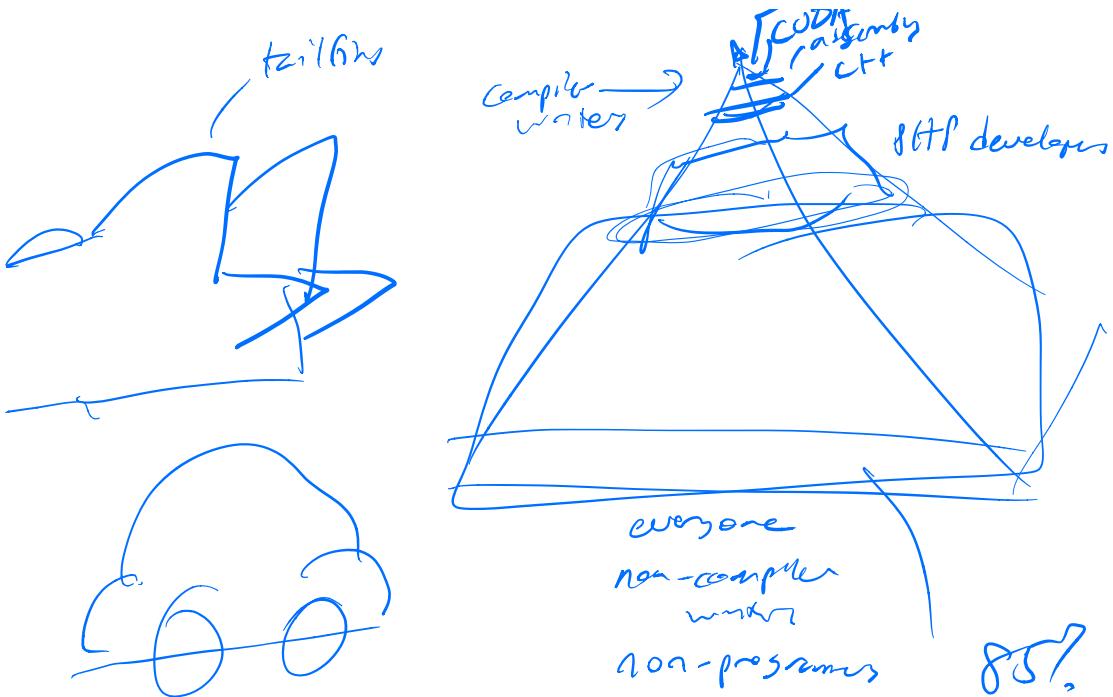
TVM

RISC vs. CISC

every Android Device  
→ ARM



four corelangs



17T H4  $\Rightarrow$  36V H1 - 3 ~~45~~ H8

Demand Scaling

phones

2 core  
4 core  
8 core

cameras  
potassium ≈  
energy / battery life

standard size instructions

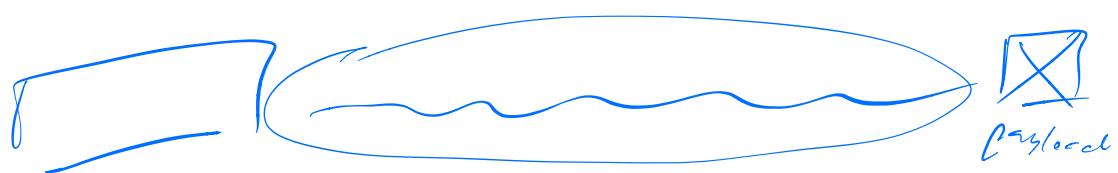
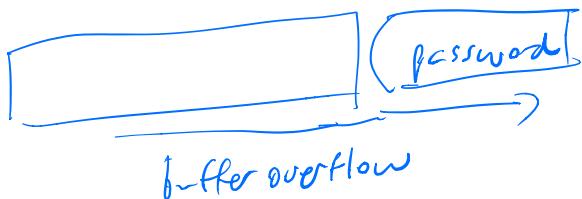
NOP → lock ROP stack — — —  
1 byte



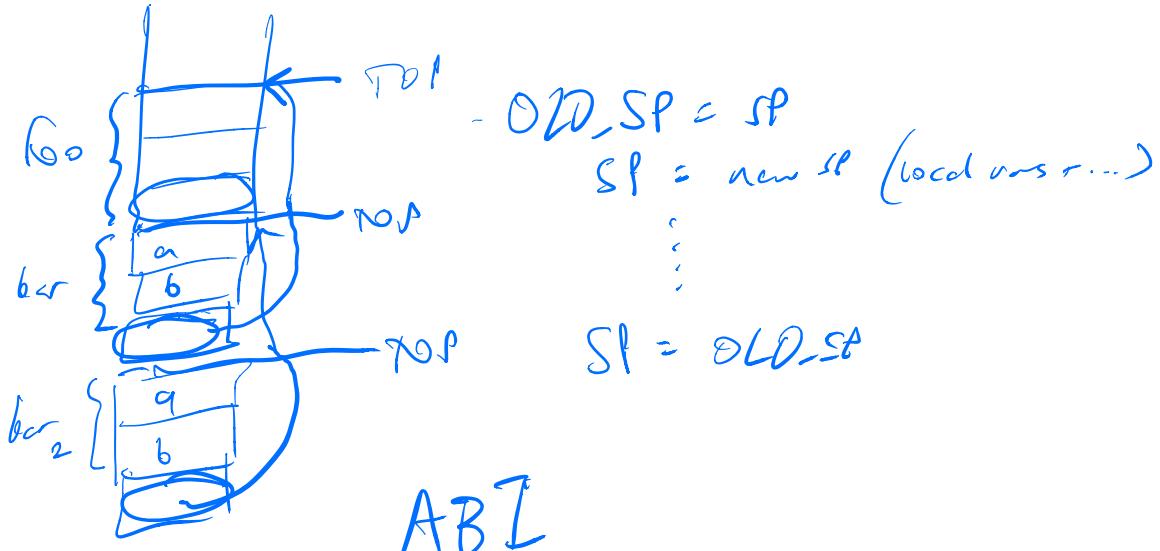
Variable sized instructions

ROP

low level security exploits — memory unsafety



## stack smashing attack



`int a;`  
`int b;`

~~`int c;`~~

`*((b+1))=pc;`

"root"  
"pwn"

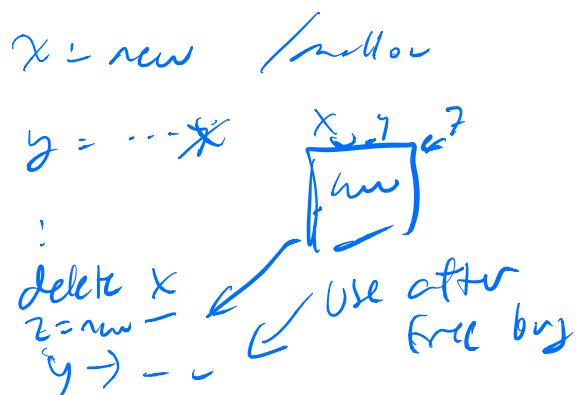
Morris worm 1988

DDoS distributed  
denial-of-service  
attack

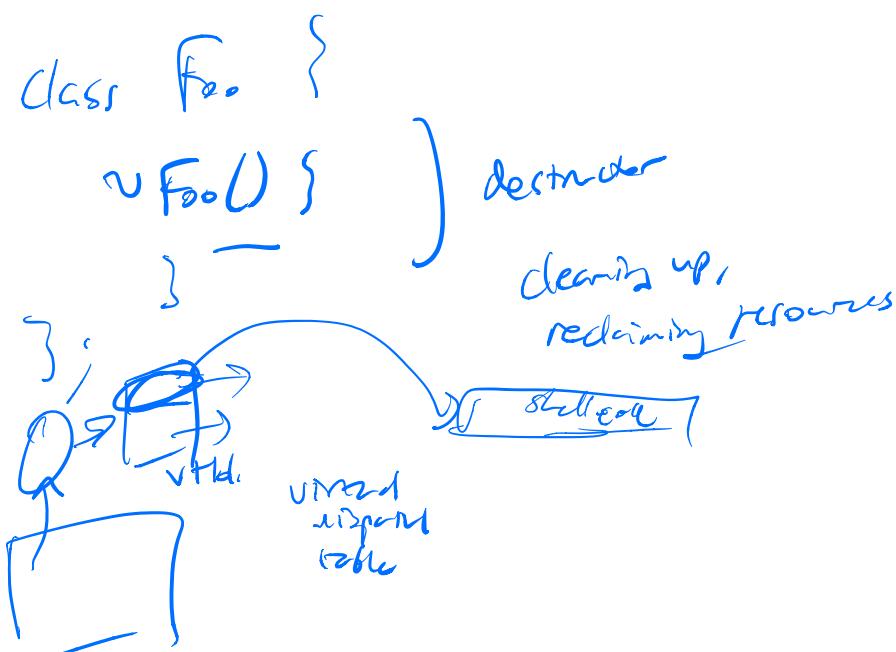
NOP NOP NOP NOP , shellcode

↳ →  
 not std  
 polymorph  
 script kiddies  
 English

dangling ptr / use-after-free

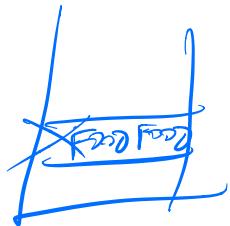


"temporal safety"



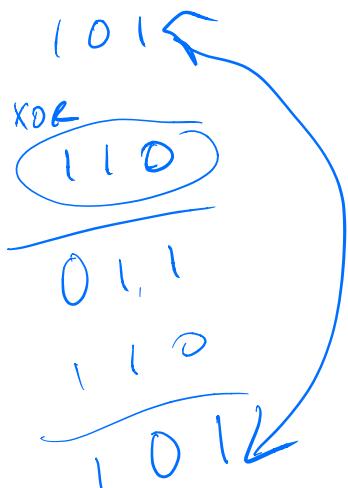
~~delete f;~~

stack smashing: shadow stack  
"Stack canaries"  
XOR ~~"key"~~



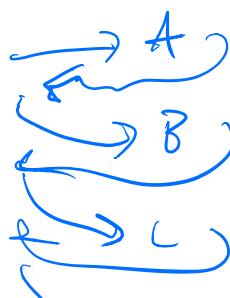
$x \wedge y$

$$\begin{array}{rcl} 1 & \text{xor } \phi & = 0 \\ 0 & \text{xor } 1 & = 1 \\ 1 & \text{xor } 0 & = 1 \\ 0 & \text{xor } 0 & = 0 \end{array}$$

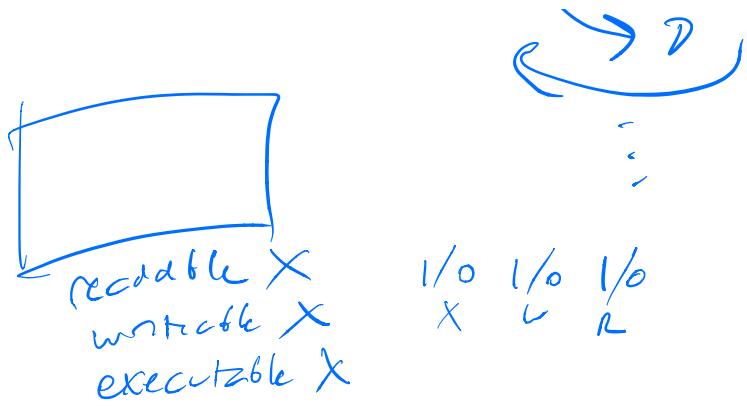


a xor key xor key =  $\alpha$

Dr Harder



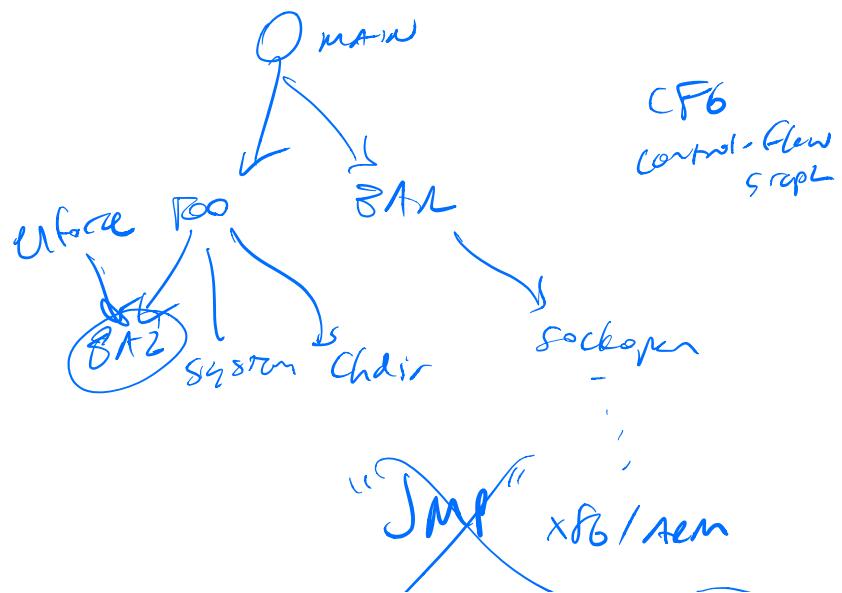
"gadgets"



## NX Protection

return-to-libc attack

"system"  
"rm -rf /"  
CFI      control-flow integrity

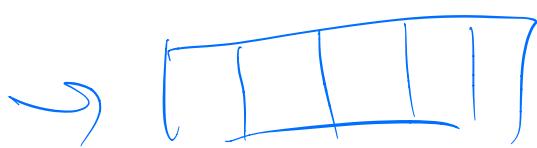




“One level of indirection”  
“all problems in CS”

RISC smaller instruction set  
“fixed-size instructions”

- ① easier to parse “decoder”
- ② harder to exploit
- ③ faster → easier to design how to make fast



PPC  
“TSX”  
HTM

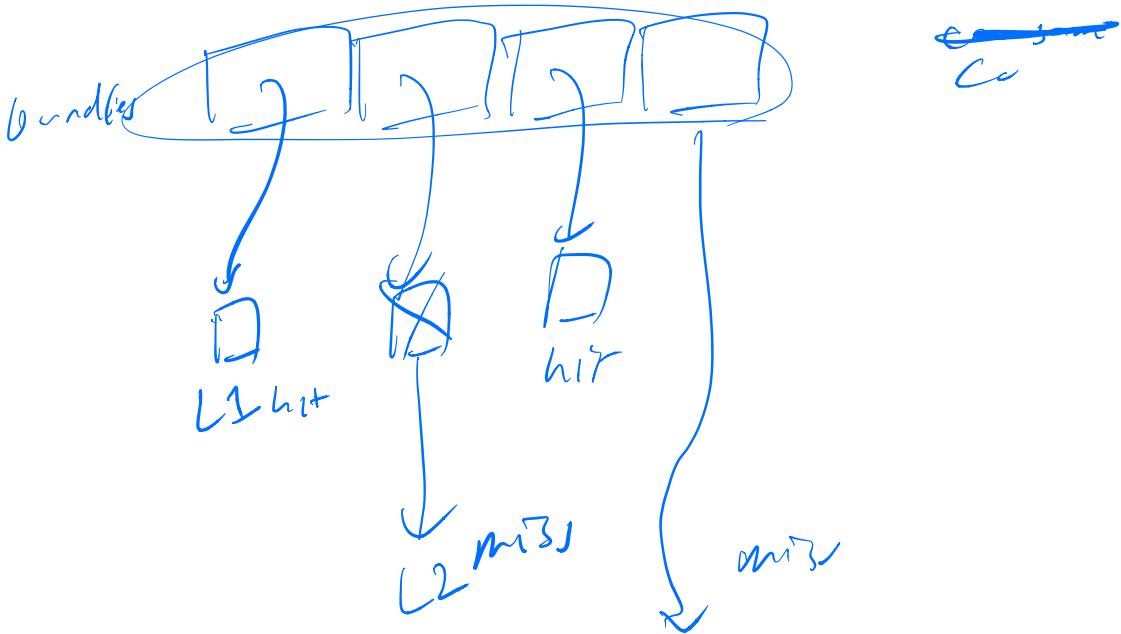
Amd - vector-instruction      Intel MPX  
- crypto      - memory protection

Google "ASan"  
AddressSanitizer

clang -fsanitize=address

VLIW

version instruction wed



EPIC

explicitly  
parallel  
instructions  
Co

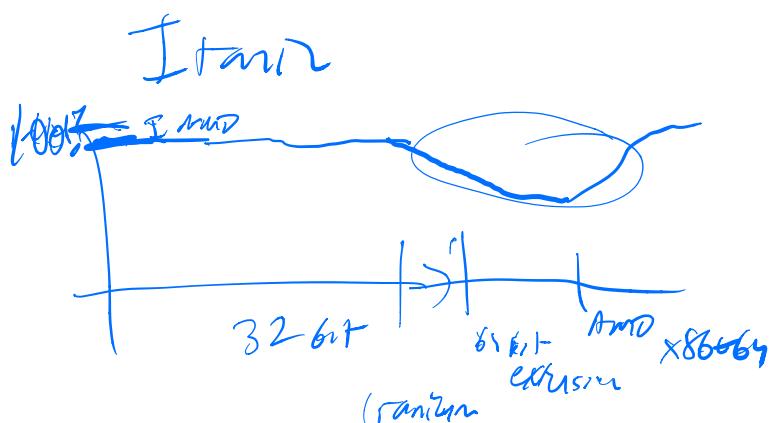
PREDICATION

(IF branch<sub>1</sub>)

DDDD

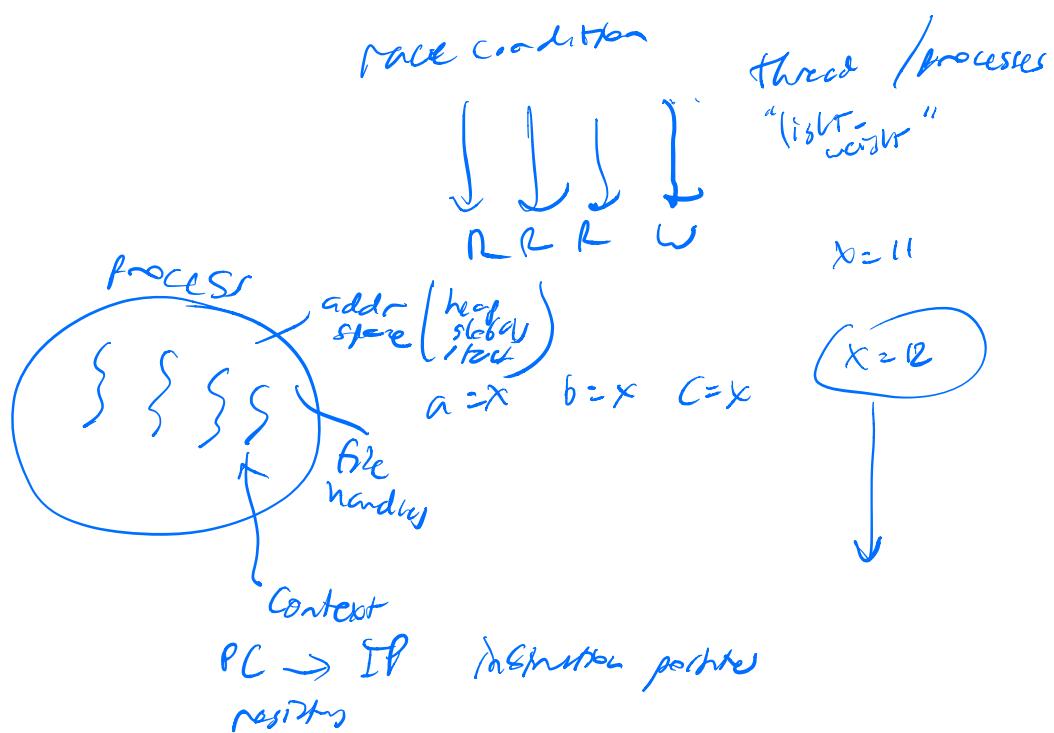
Itanium

IA-64



MOV  $\ddot{Q}$   
ADD  $\ddot{Q}$

## Concurrence



## Concurrence

~~errors~~  
races

atomicity violation

Order violation

deadlock

livelock

) condition variables

— wake-up problem

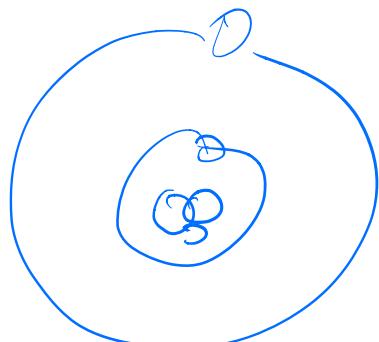
while(<sup>true</sup>  
    lock  
     $\rightarrow$  want()  
     $\wedge$  check state  
     $\wedge$   $T\theta(\cdot)$  break  
    }  
    })

NONDETERMINISTIC

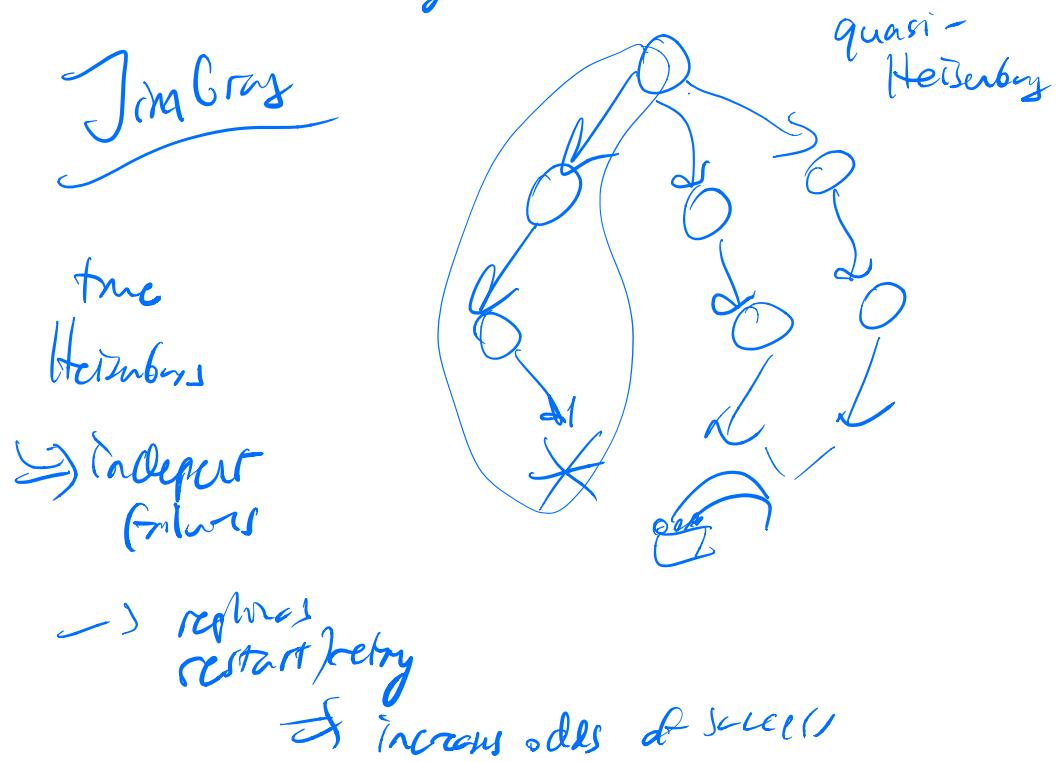
livelock

Heisenboss

DETERMINISTIC boss — Bohr boss



- severity?
- frequency?



Developers

debugging  
determinism!

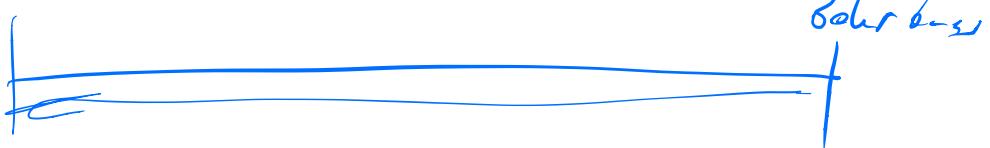
Bohr Bus

Users

Heisenberg

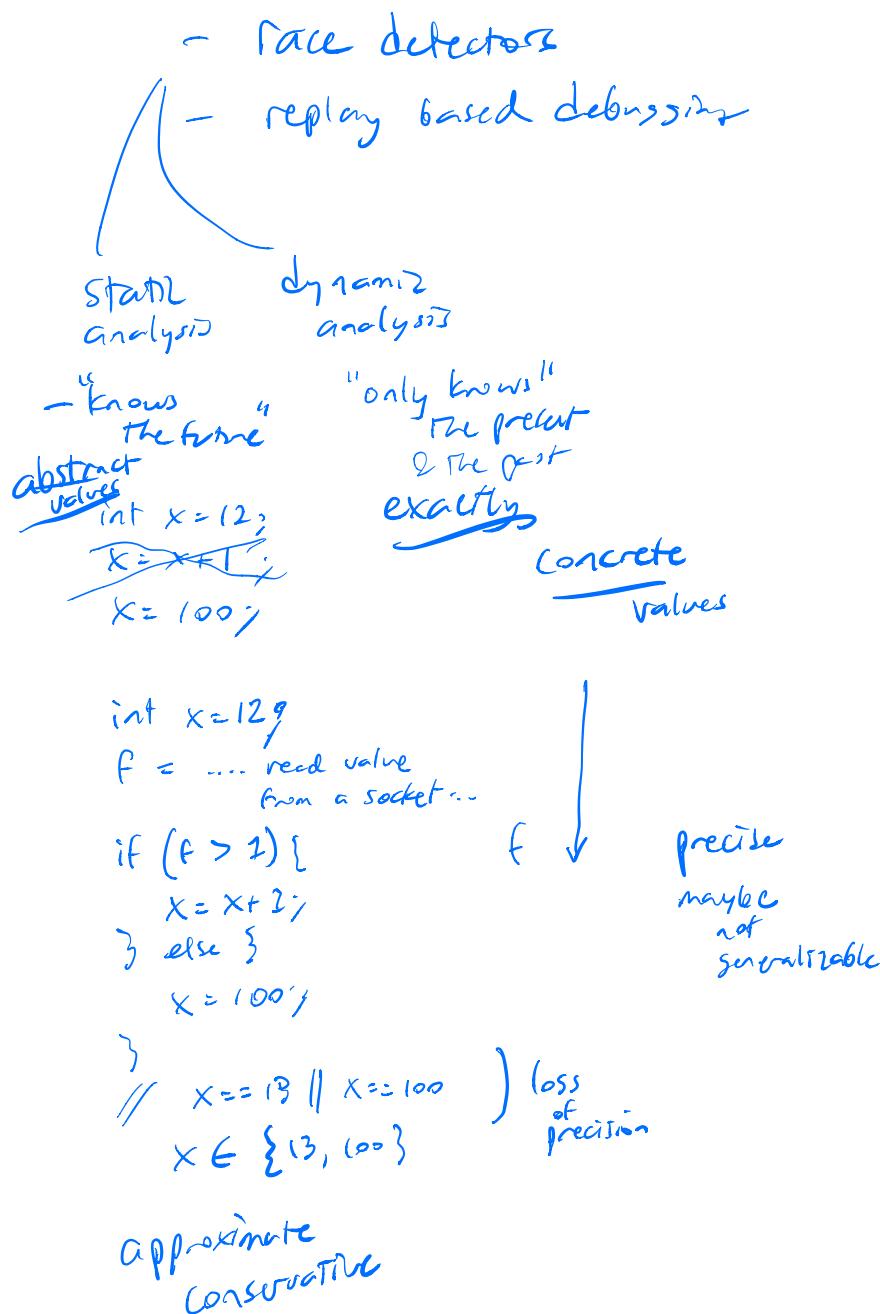
→ retry

→ fault tolerance  
→ indep. failures)



D+E

100%



## bus detection

partial specification

P never derefs NULL

"implicit specification"

segfault buffer overflow  
NPE use-after-free  
races assertion failure

benign races / malignant races

"Hauswald!"

## formal verification

proves program P

implements spec S

code

y

log2

"complete / total  
specification"

abnormally



Word-sized  
read/write

X = 1

001

X = 4

100

101

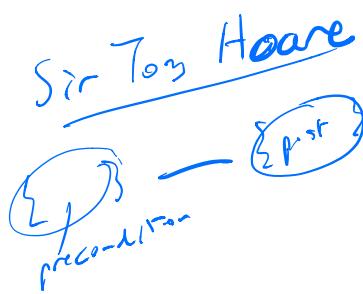
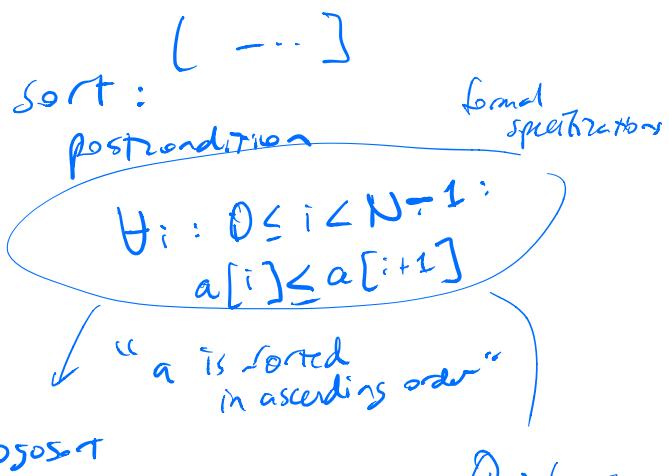


double  
d

d = 0.1      d = 2.0

d = ...

" O-O



Quicksort  
average case  $O(n \log n)$   
worst case  $O(n^2)$

Mergesort  
avg & worst case  $O(n \log n)$

quis custodes custodit

trusted computing base (TCB)

spec + solver  
+ code rev  $\Rightarrow$  selected code

$\approx 1000$  LOC  
just: hundreds!

targets for formal verification }  
Narrow APIs  
well specified  
really important

Project Everest  
"https" TLS F\*

## Heartbleed

static analysis for concurrency errors

~ false positives

precision

recall

print "NO BUG HERE!"

print "BUG!"

dyn analysis — no false positives

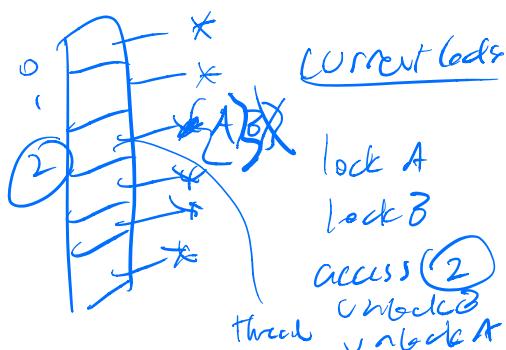
WITHINSS race!

— recall lower —

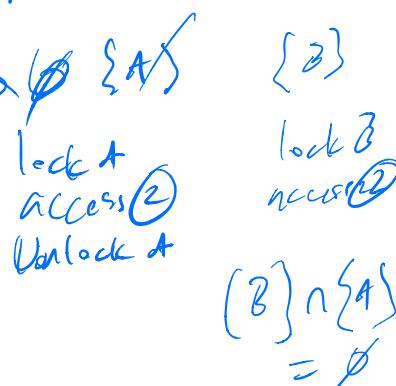
only reports race  
in that execution

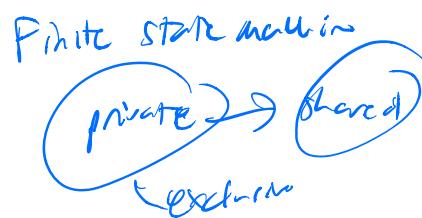
## Thread Sanitizer

### lockset analysis



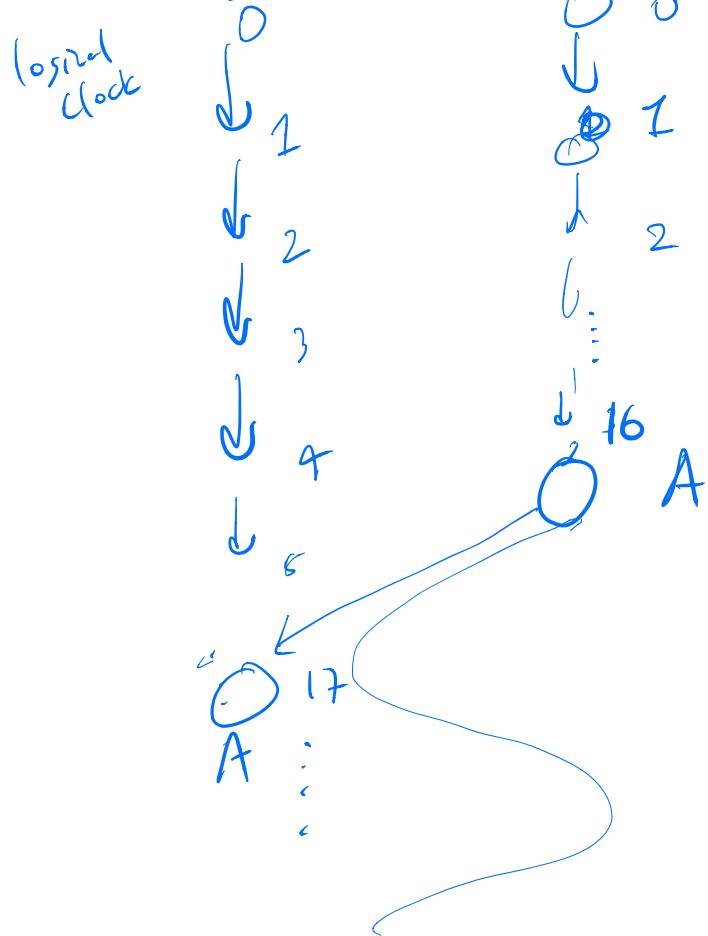
### happens-before analysis





Error

happens-before analysis



Why  
Threads?

- parallelism

- coordination

- hiding latency  $\uparrow$  performance  
I/O  
memory

↳ average the periods  
↳ the delay until "first by"

"throughput" — rate of jobs safely done  
"throughput"

throughput-latency tradeoff

interactive

batch

|  
no context  
switches

CPU util ~ 100%

time  
sharing

/dev/fny

ed

MULTICS "Eunuchs"

1964 - 1967

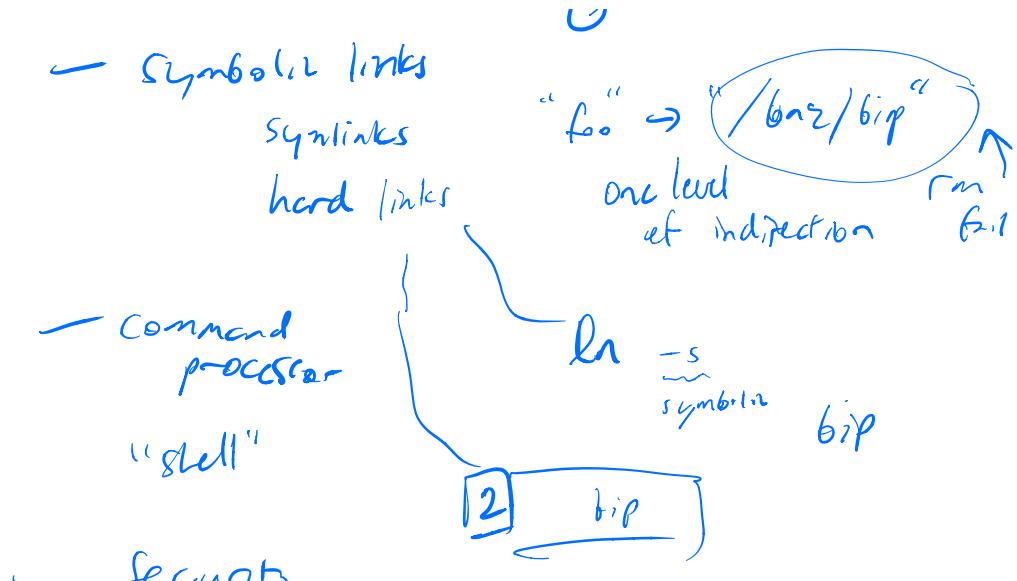
MIT + GE + Bell Labs

- hierarchical FS

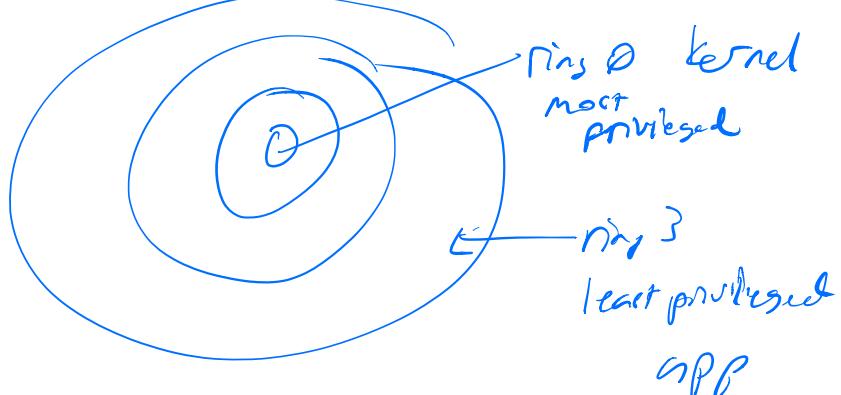
all previous:  
"flat"

relative  
paths





— Security rings




---

Why is choice of quantum size a tradeoff?

---

newObject  
 mark  
 clear  
 sweep  
 gc

```

for (auto it = allocated.begin();  

     it != allocated.end();  

     ++it) {  

  (*it) → clear();
}
  
```

Object \*  
 newObject();

clear();  
 mark(&root);  
 sweep();

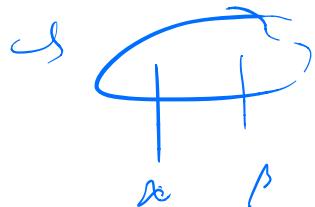
```

\ Objet & obj;
if (freed->empty()) {
    obj = new Object();
} else {
    obj = freed->front();
    freed->pop_front();
}
allocated->push_back(obj);
return obj;
}

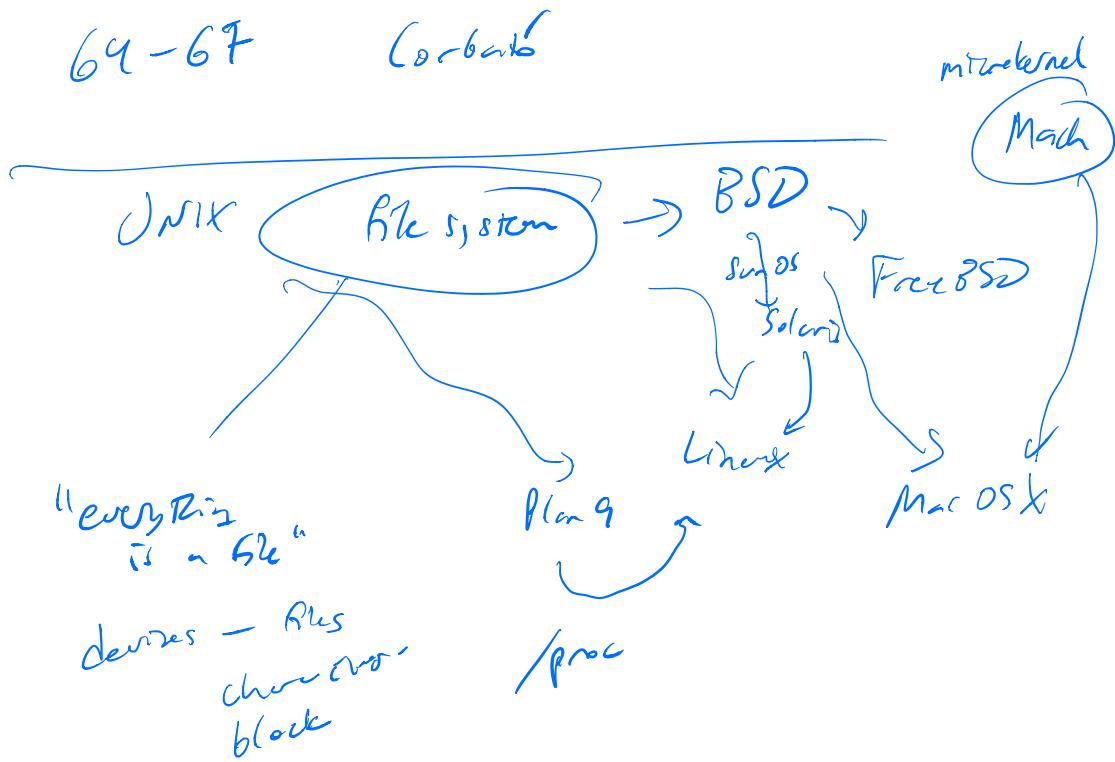
void mark (Object* a) {
    if (!a->isMarked ()) {
        a->mark ();
        if (a->isBinary ()) {
            mark (a->sub1 ());
            mark (a->sub2 ());
        }
    }
}

```

Cons &  $\beta$



## hardware enforced security



stdin stdout  
stderr

"one level of indirection"

echo "hello world" > foo  
 > /dev/l-



0% ( cut -f 1 -d\| ) | sort | uniq | wc -l

pipeline parallelism

"streaming"

data flow graph



awk

Aho Weinberg Kernighan

sed

s/-/ /g

grep

general regular expression parser

Perl - Larry Wall

Python - Guido van Rossum

Ruby - "Matz"

→ killer app

"Rails"  
Ruby on Rails  
Niklaus Wirth

VisiCalc 1979  
Microsoft Word  
Excel  
PowerPoint

Office

Scala - Martin Odersky

EPFL

→ Pascal  
Modula-2  
Oberon

Andreas (Ae)tschba  
Turbo  
Pascal

→ Java  
generics  
Spark

TypeScript

JavaScript      Stockholm syndrome

Barend  
Eijk

David Turner  
≡ KRL  
→ Miranda<sup>TM</sup>

Simon Peyton Jones — MSR  
Phil Wadler → Comb.  
John Hughes → Edinburgh  
Paul Hudak → Yale  
...  
Miranda<sup>TM</sup> ↗  
erden.  
of Res. SFA Ltd.  
Haskell Curry

$f(x, y, z)$

$f x y z$

"f 3"      currying

pure lazy functional prog. language

```

int s = 12;
int void f(x) {
    return x*2 + (s+1);
}

```

1/0

$f(10)$

ones = 1 : ones

$[a \times 2 \mid a \leftarrow [1, 2, 3]]$

print ("foo")

mountable file systems

NFS

/var/fsdb

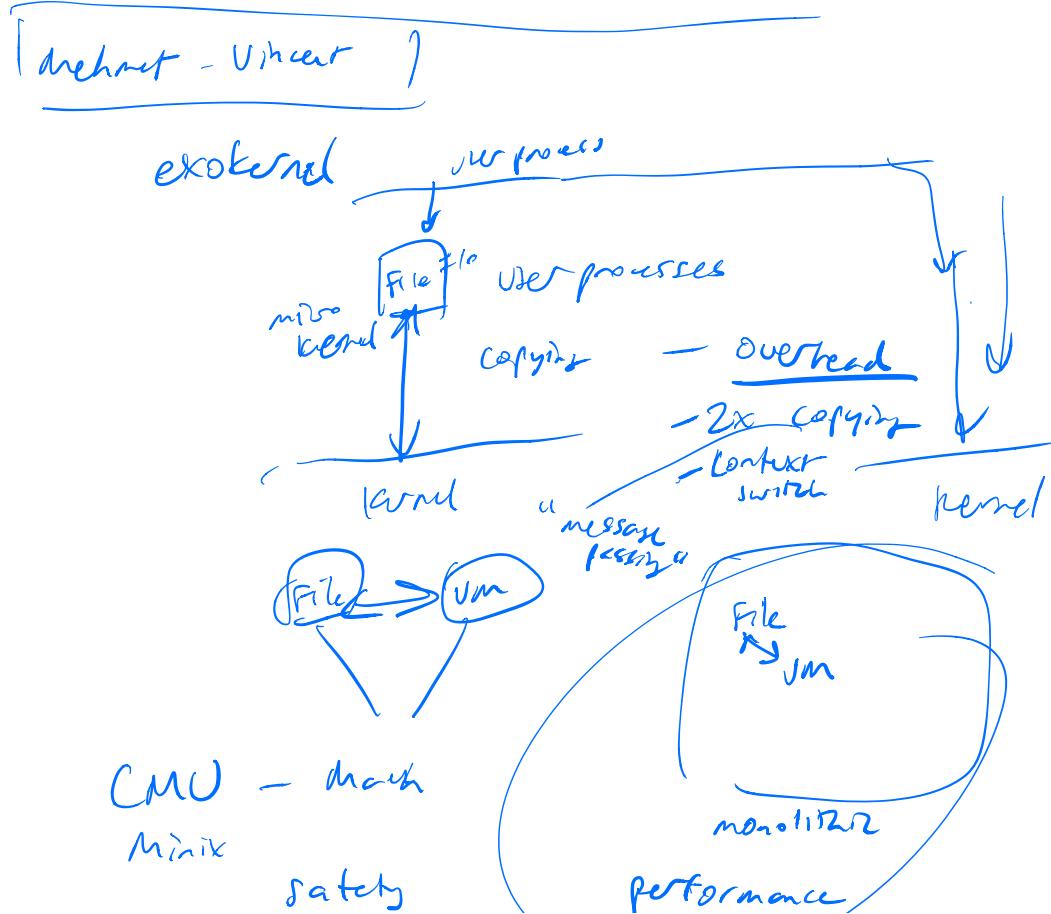
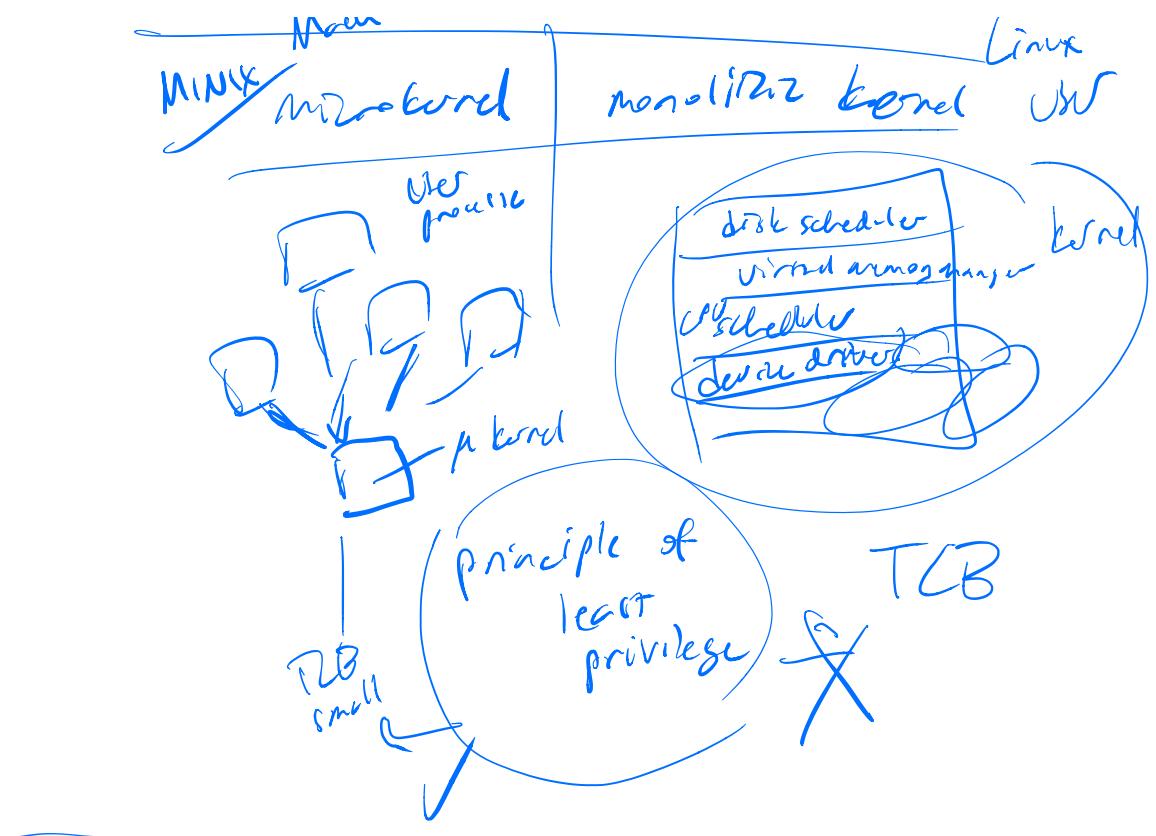
/home

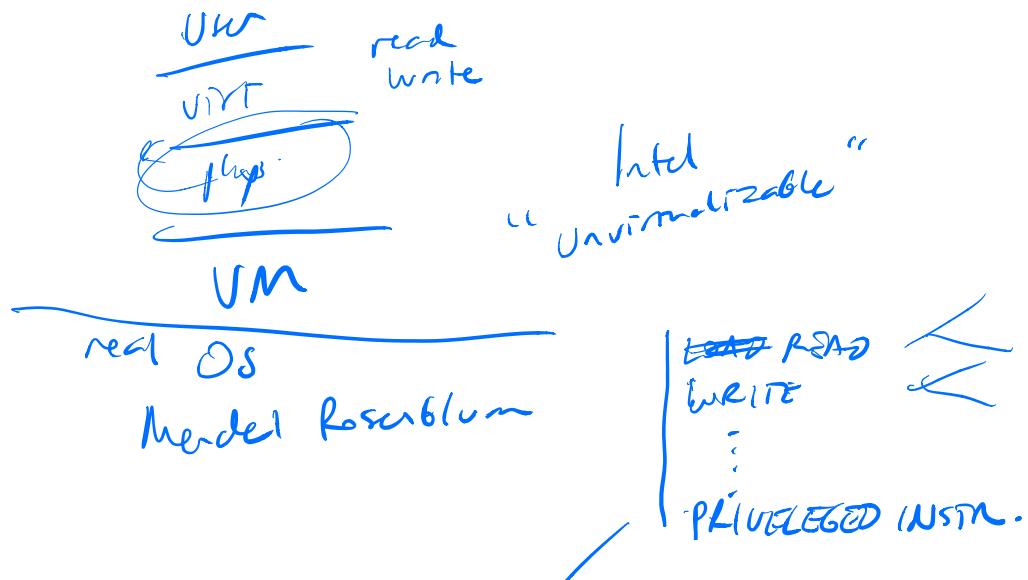
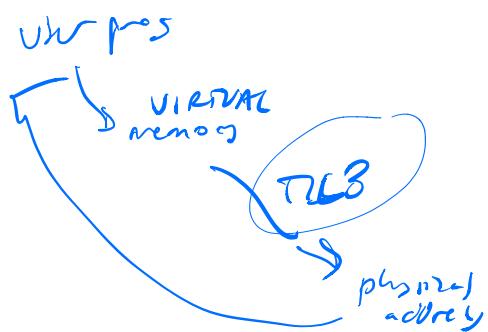
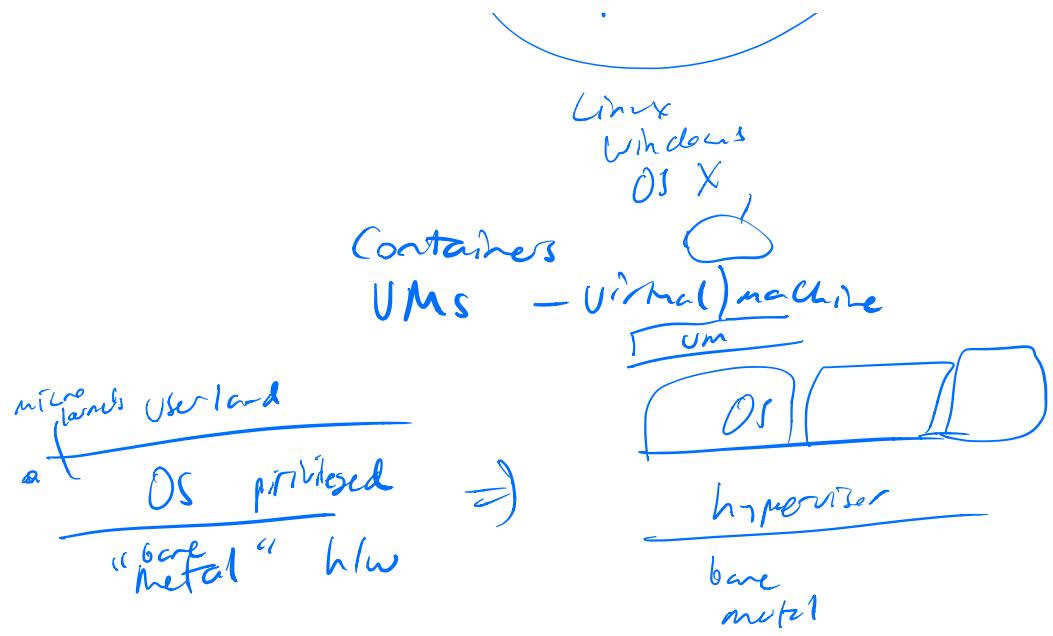
/usr

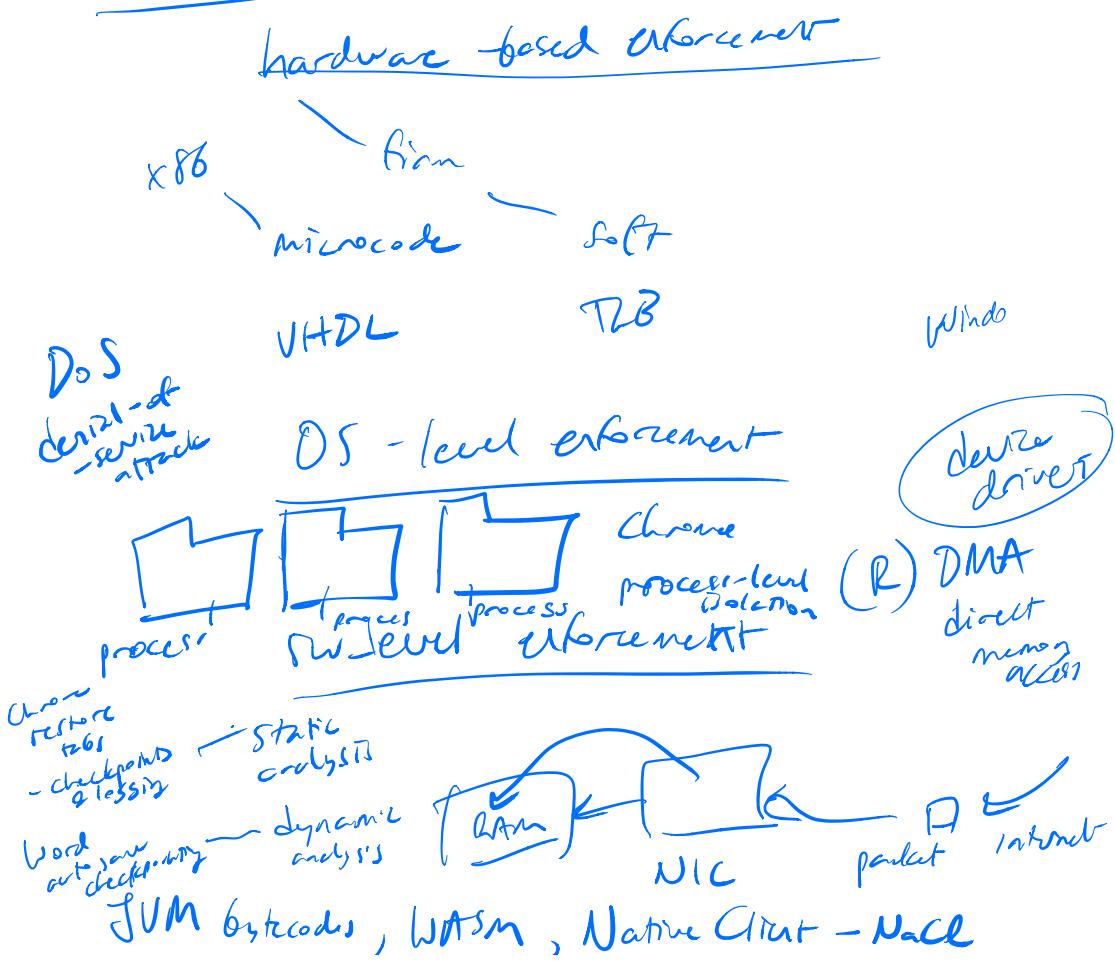
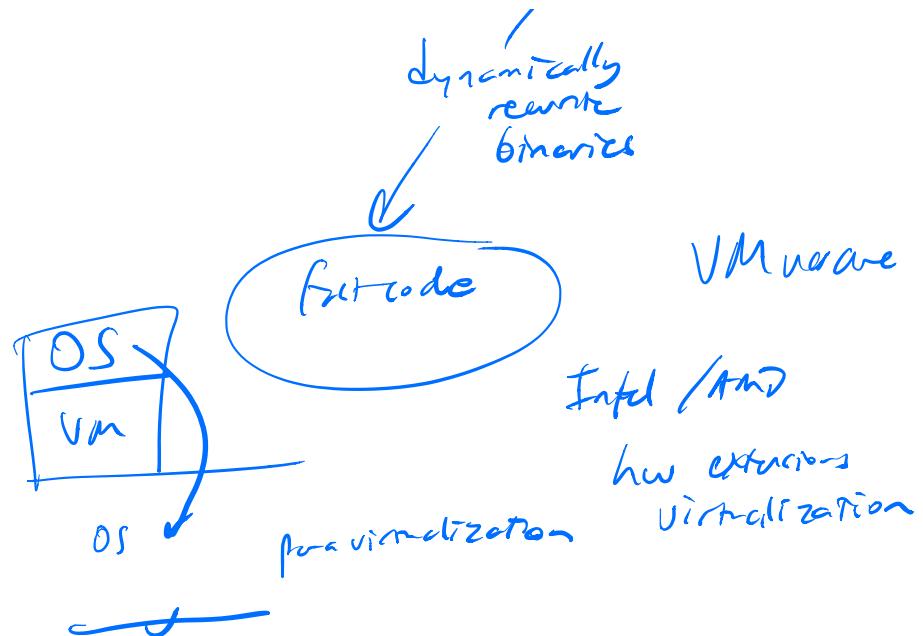
shuttle

Torvalds  
debate  
Flame war

Torvalds  
debate







Windows failure  $\rightarrow$  BSOD pNaCl

kernel panic

DDK

- fails internal integrity check
- hardware error

Wisdom  
of the  
crowd

Galton

VoxPPL

"  
home economics"

rational agent - maximize gain

cost-benefit analysis

$$n_{jk} = P(\text{break } h)$$

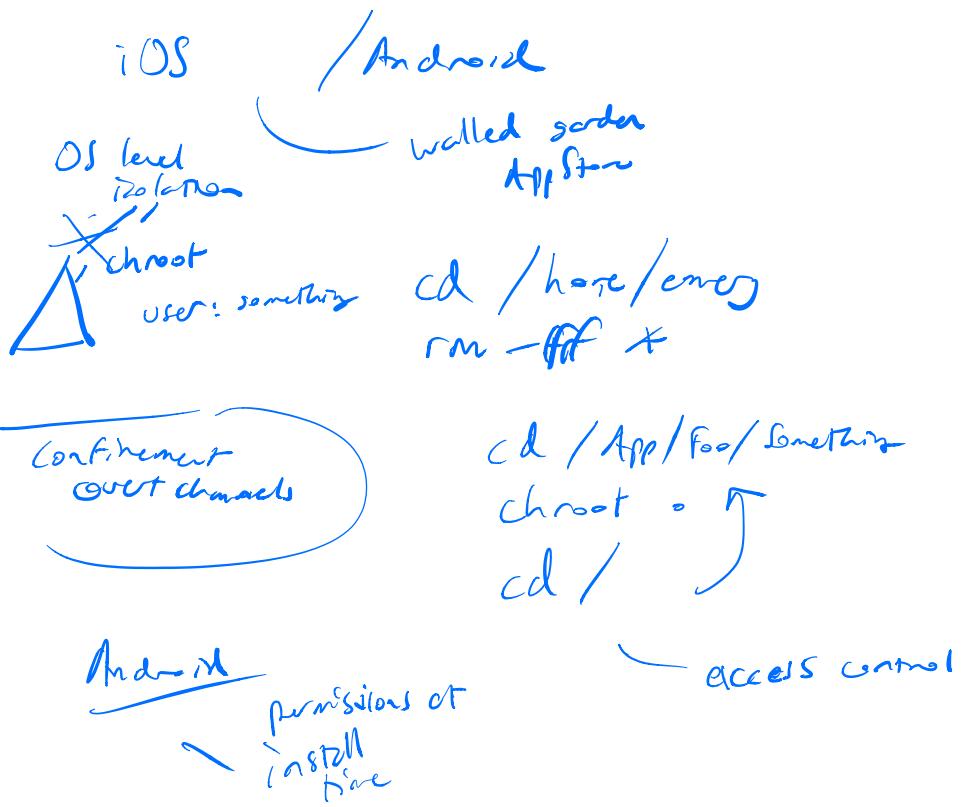
cost =  $\sum \text{cost of replacement}$

$$\begin{array}{rcl} \$30K & & \text{ransomware} \\ \hline 10,000 & + \frac{30,000}{1} & = \\ & \$1 & \$15 \end{array}$$

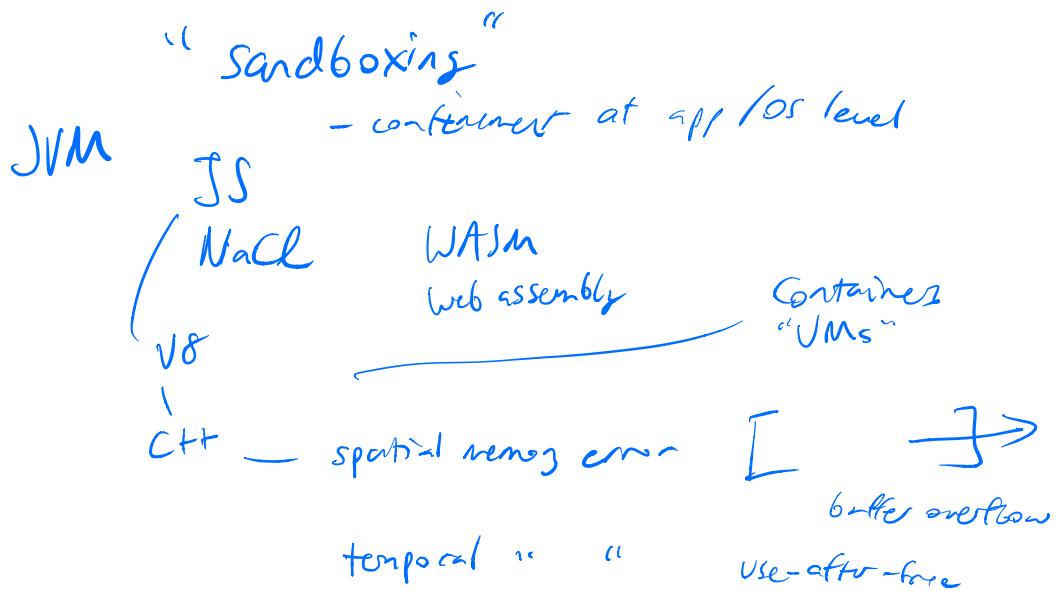
Fort Knox

social engineering  
phishing

asymmetric attack



Scribes: Jun, Kyle



LangSec long-level security

— Systems-level lang  
enforces spatial  
+ temporal safety

- "memory safety"
- "type safety"

contract-kw integr?



- X high-level security  
is mem safety (ctr.) sufficient?
- VMS low-level security  
SQL injection attack
- "safe" PL ?? X

RUST

- memory safety
- concurrency

Ownership types

- no GC

"unsafe"

→ NOT MODULAR



"cat and mouse" game  
"arms race"

# define while if

query = "SELECT \*  
From employees  
WHERE " + cond + "

Cond = "salary > 100000"  
position != 'CEO' OR  
sqlserver.execute(query, results);  
Send : (results),  
SELECT -- ; DROP TABLE

salary > 100000 ;

XKcd Little Bobby Blue Tables

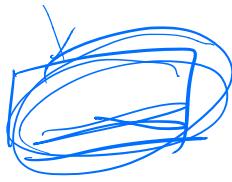
drive-by attacker

↳ zero-day vulnerability

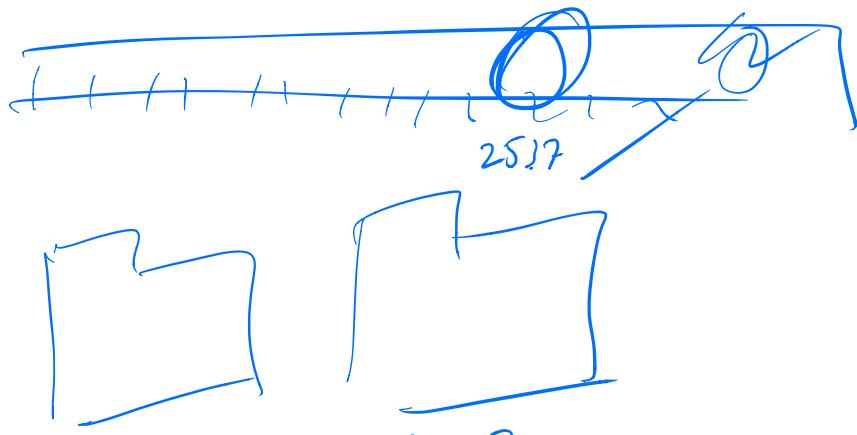
heap spraying  
JIT spraying

s = " shellcode " ;  
;

speculative execution



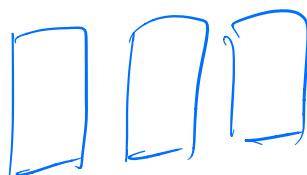
$a[x]$



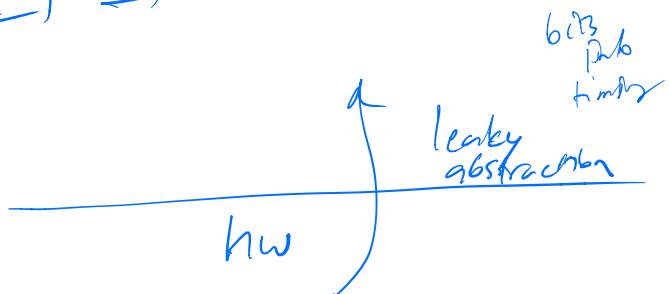
Shared Array Buffers

high res times

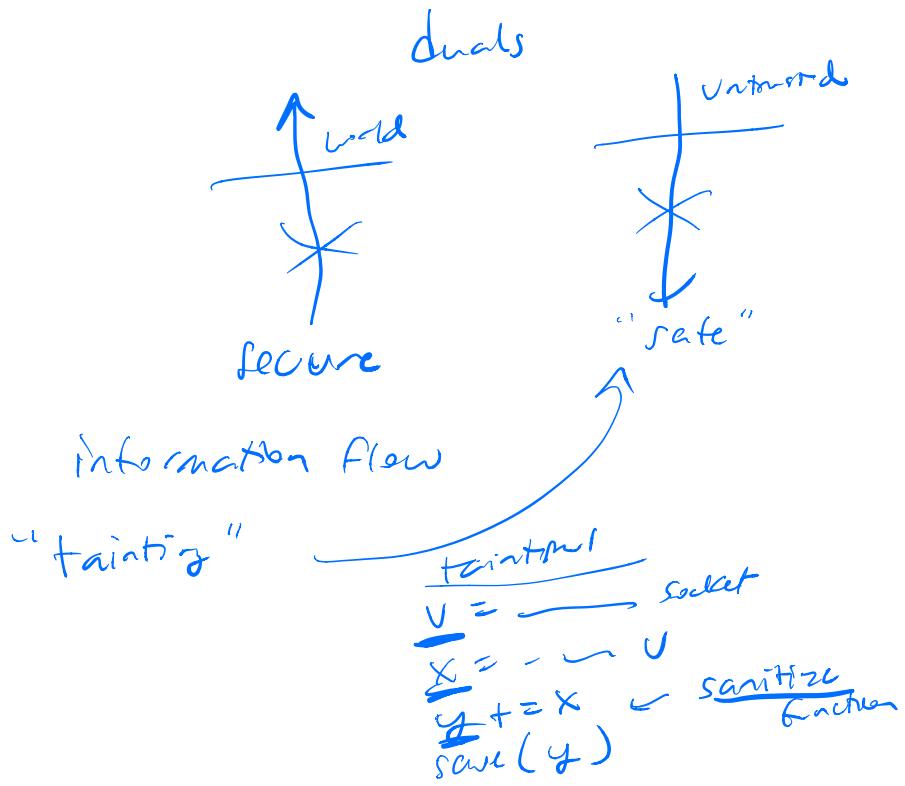
Rowhammer - Our motto



DRAMs

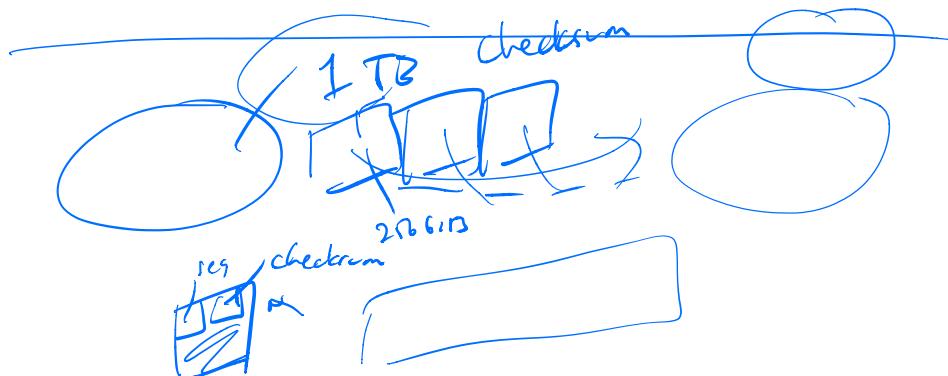


privacy                      integrity



$$\text{sanitize}(\text{untainted}) = \text{untainted}$$

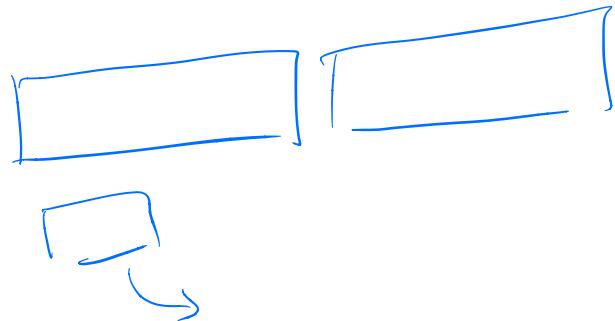
$$(\text{tainted}) = \text{untainted}$$



$$P(\text{covert packet undetected}) = 1,000,000$$

silent film

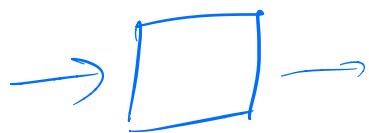
D D D D D  
~~~~~  
≥ 1000 0000

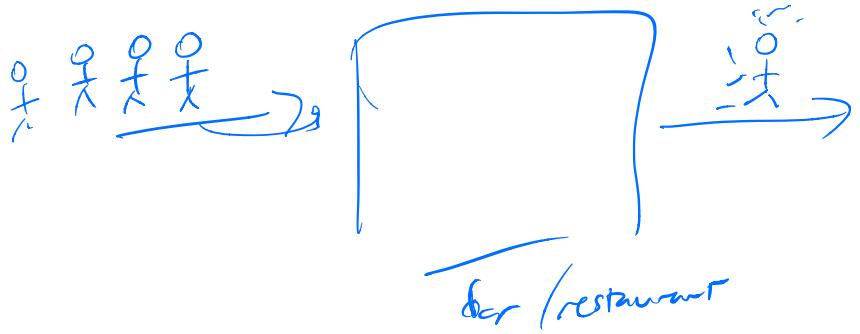


securing integrity

MOAR BITS

shedding load
~ "admission control"





$$L = \lambda W$$

Little's Law

how many people waiting in source

arrival rate service rate

$\frac{100}{hr}$ $1\text{hr}/\text{customer}$

$$E[L] = E[\lambda] \cdot E[W]$$

backlogged

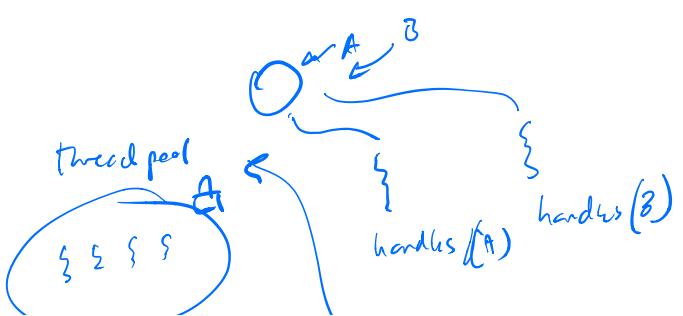
load shedding

admission control — bouncers

QoS

DoS

"thread pool"



$N=30$
 $N=100$

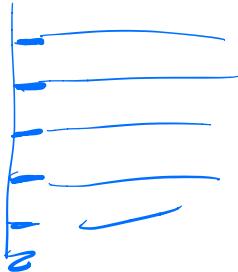
admission control mechanism

Cryptography

"Security
through
obscure"

codewords

"steganography"



"key"

Caesar cipher

key = 5

A B C D E F G ... Symmetric
↓ E F G H I J K crypto
values

A B C D E F ... DEAD

Z Q G A F R ... AFZAA

(frequency analysis
ciphertext)

Q Z D

E
T
A
O
I
N
O

0 --- 0

H
R
D
L
:

One-time pad



XOR

$$\begin{array}{rcl} 0 \text{ XOR } 0 & = & 0 \\ 0 \text{ XOR } 1 & = & 1 \\ 1 \text{ XOR } 0 & = & 1 \\ 1 \text{ XOR } 1 & = & 0 \end{array}$$

rotating keys
(codebooks)

Enigma

known plaintext attack

- decryption!

Goal: encrypted message ≈
random noise

⇒ best you can do
⇒ brute force

make key comb

probabilistic

$$E(\text{decrypt time}) = \frac{1}{2} * \text{key space}$$
$$\left(\text{e.g. } 2^{1024} \right)$$

2²⁰⁴⁸

if you get my key,
you can't : decrypt
encrypt) asymmetric
keys

encrypt w/o decrypt

use RSA

attestation / signing

Diffie Hellman

one way function

$$f(x) \rightarrow \beta \text{ easy!}$$

$$f'(x) \rightarrow \alpha \text{ hard!}$$

prime factoring is hard

$$A \wedge \overline{B} \wedge C \wedge D \wedge \dots$$

A: T

$$\begin{array}{l} C : F \\ B : T \\ D : T \end{array}$$

$$\underline{n} = 12$$

Alice Bob

Mallory Eve "Sybil attacks"

reputation based
systems

Amazon reviews
Gelph
Trip Advisor
Netflix
;



$$E_A(s) \rightarrow \underline{m} \rightarrow D_A(m) \rightarrow s$$

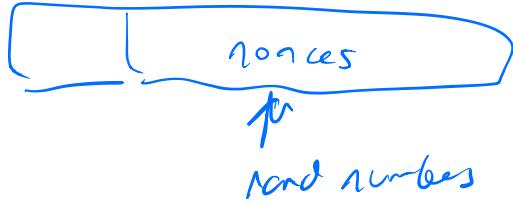
(e, d, n) e constant
d is relatively prime to $(p-1, q-1)$

$$\text{Signature} = D_B(\underline{m}) - s$$

$$\text{or just } E_A(s) \rightarrow \text{Signature}$$

$$\text{Alice } D_A(s_{\text{Alice}}) \rightarrow s_{\text{Alice}}$$

$$\text{now } \underline{m} = E_B(s)$$



PKI
public key infrastructure

not
certificate
authorities
quis custodiet custodes?

$$\text{availability} = \frac{\text{MTBF}}{(\text{MTBF} + \text{MTTR})}$$

mean time between failures
mean time to recover

ROC

recovery oriented computing

fail-safe vs fail-stop

fail-safe
roll-back
(or
"undo")
error

fail-stop
stop or one-

fail-over - replicas
 $\times \rightarrow \square$

~~xx~~ →

State replicated

Replicated state machine

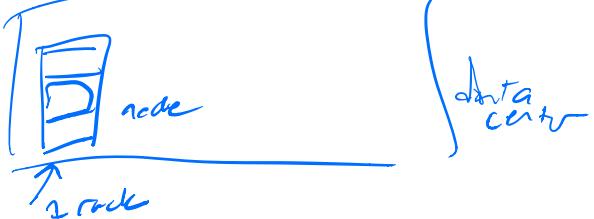
Zossmann & Schneider

replication

Single point of failure

Boeing 737 Max

Geographic distribution



Correlated failure = BAD

GOAL: independence

$$P(\text{failure}_1) \times P(\text{failure}_2)$$

$\frac{1}{1000}$

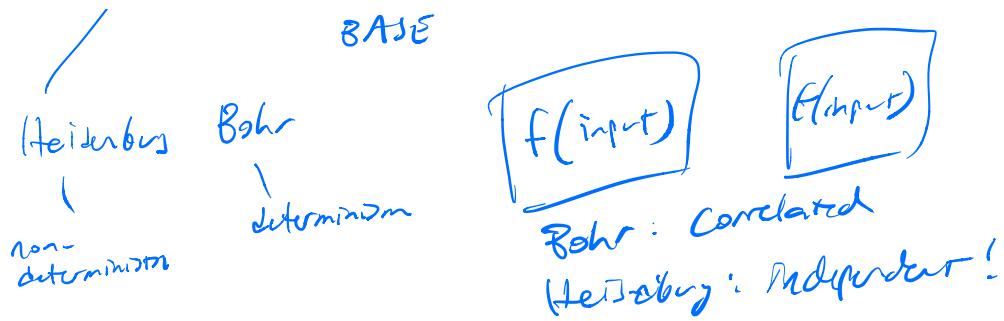
$\frac{1}{1000}$

$$= \frac{1}{1,000,000}$$

fault tolerance

- hardware RAID

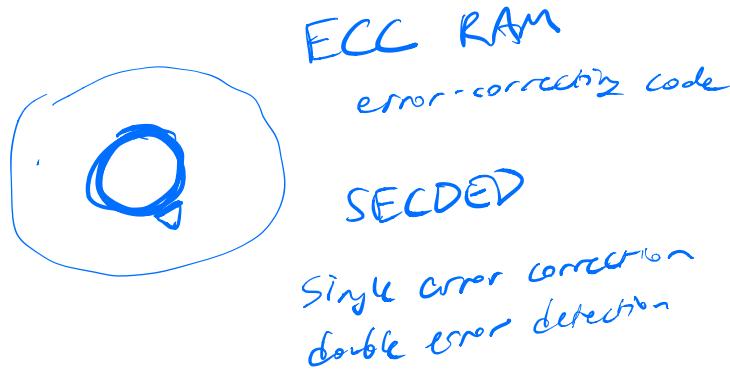
- software transactions) replication
ACID



SAY: $P(\text{failure}) = \frac{1}{100}$
Heideburg

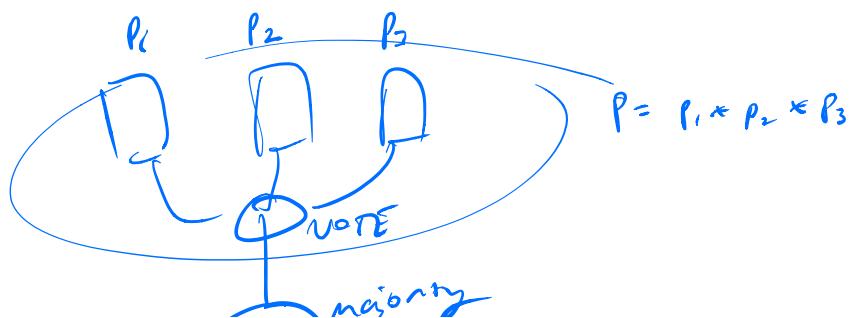
$$E(\# \text{ of people who aware}) = (-P) = 99\%$$

hw non-det failures



trimodular redundancy

3 things



transactions

atomic
all or nothing

consistent
in face of concurrency

log of updates
journal

WRITE ADDED class,

WRITE " "

WRITE SUBTRACT ...
DROPPED

ABOUT to commit

COMMITTED

integrity

durable

commited
"stable storage"
(persistent)
— Hard disk

Checkpoint "snapshot"

Concurrency control..

lock - pessimistic

- optimistic

"conflict" - Abort
roll back & restart
retry

RAID

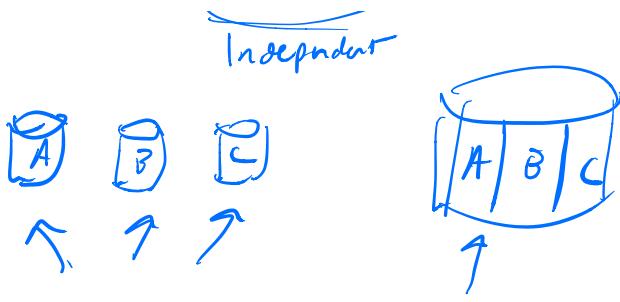
NW

network
of
workstations

L1SC

disks → bigger → more expensive!

Redundant
Array of Inexpensive Disks



$$f = P(\text{failure}) = \frac{1}{1000} = 0.001$$

succes: $1 - f = \frac{999}{1000}$

$$(1 - f)^2 = \frac{998}{1000}$$

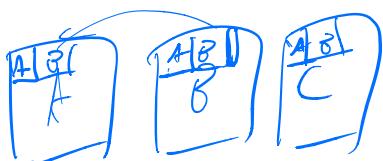
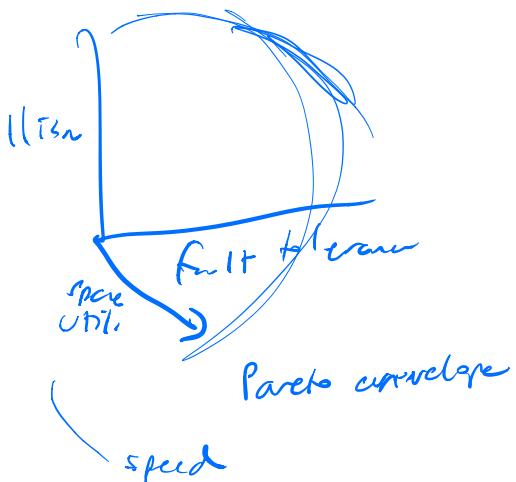
$$\approx 0.997$$

$$10 = 99\%$$

$$- 100 = 90\%$$

$$- 1000 = 37\%$$

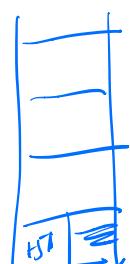
balance || Dm
2 fault tolerance
& space



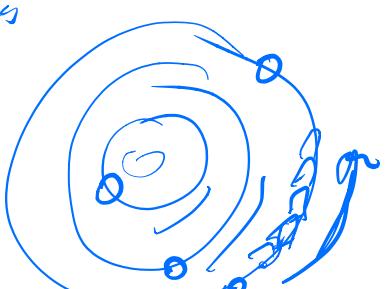
Hamming code

RAID-5

OS file blocks
disk



buffer queue



rotational latency << seek latency

sw fault tolerance

- failure-oblivious
- probabilistic methods

approx computation

loop performance

$i = 8$

$\text{for } (\text{int } i=0; i < N; i++) \{$

/ \}

Bolt



talk tomorrow at noon

Caroline Lemieux

CS 150/1

12:15 · wed

failure-oblivious

general

- assert(0)
- throw exceptions
 - ↑
unhandled

- $x = 2/0$
- $x = * (0)$

} future
programmer
start

- These
are fails

- “always”
↳
 - $x = (\text{null}) \rightarrow \text{foo}$
 - $\text{exit}(-1)$
 - memory safety errors

$x = \text{new char}[9];$ buffer overflow (spewz)
 $x[9] = 'z'$ dangling PTR / race
 $x[\underline{9 \dots 10}]$ after free (temporal)
 $\boxed{\dots}$ “overfitting”

CAVEAT!

- concurrency errors

race

correct execution

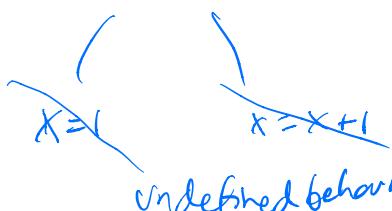
buffer overflow (spewz)

dangling PTR / race
after free (temporal)

“foreign races”

C++
memory model

Santa Adve
Hans Boehm



undefined behavior

~~if ($x > 0$) {~~
~~:~~
~~}~~

depends on undefined behavior

deal

Hogwild!

higher performance

~ undefined behavior!

“int32”

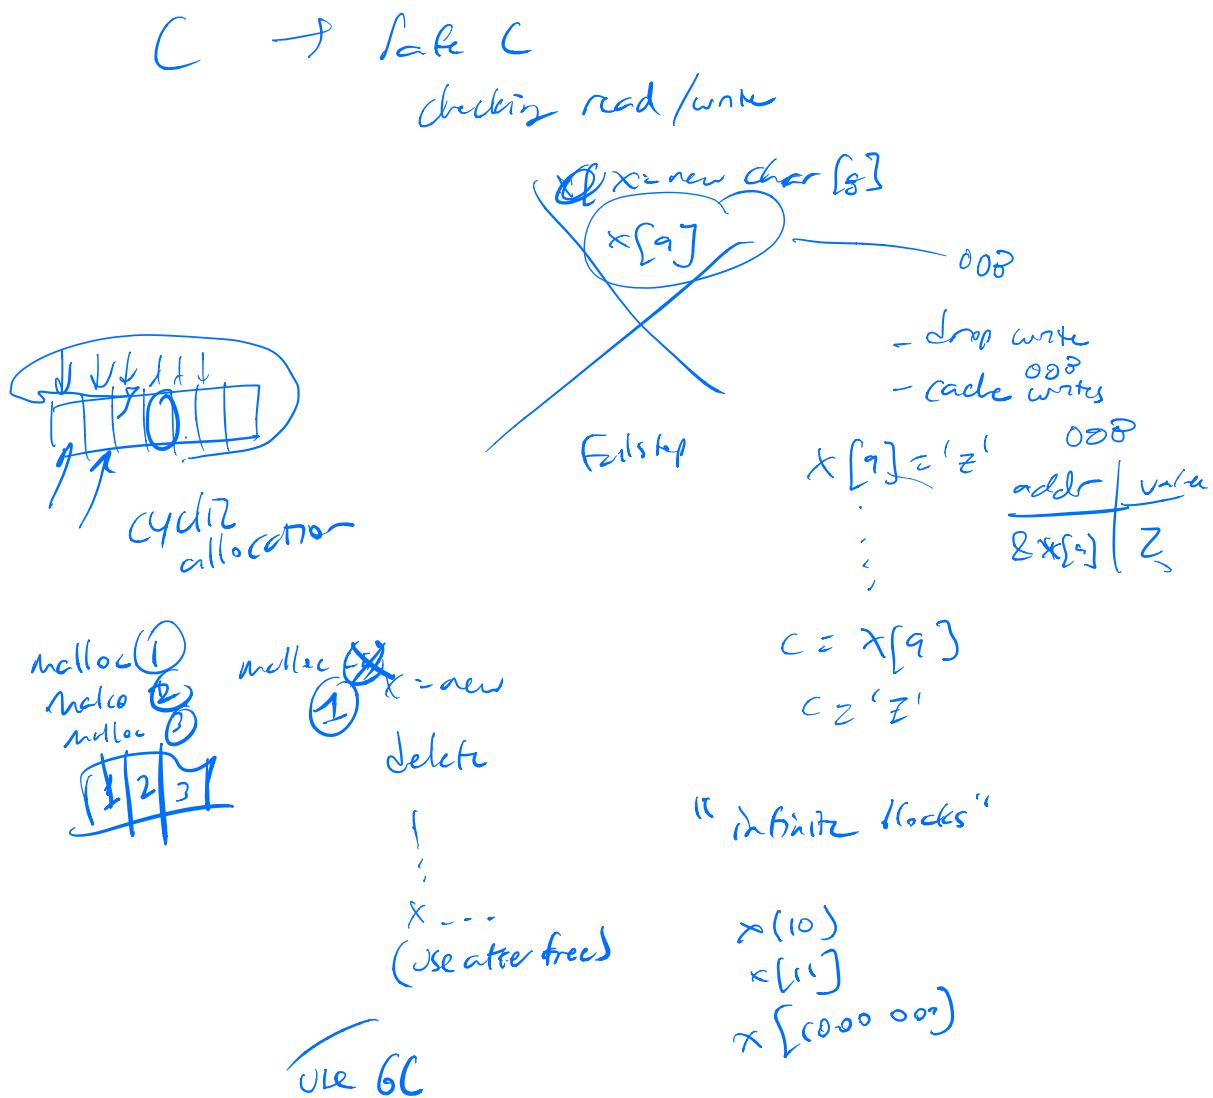
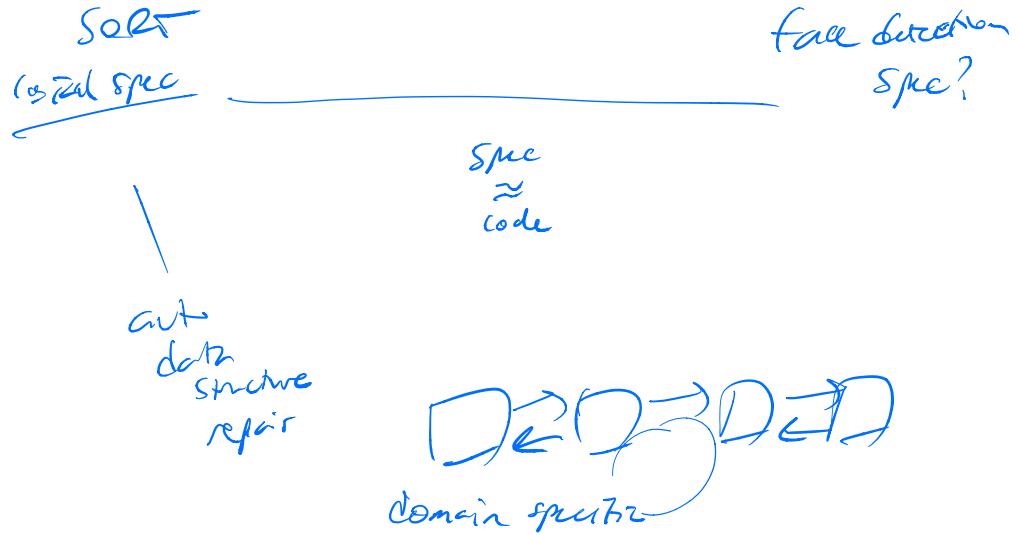
lossy
discrete!

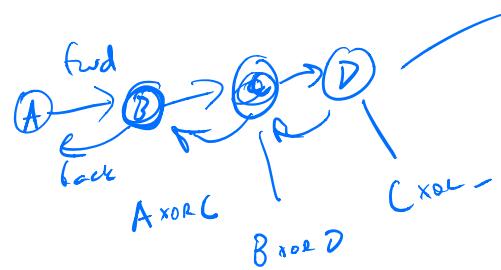
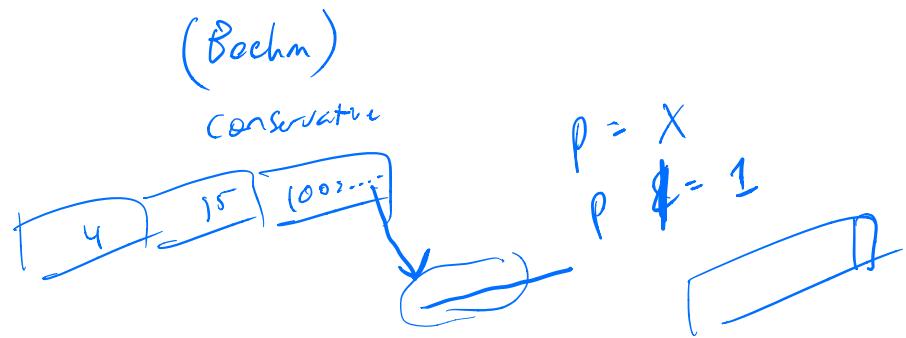
“floating point”

bignum
bignum
exact math

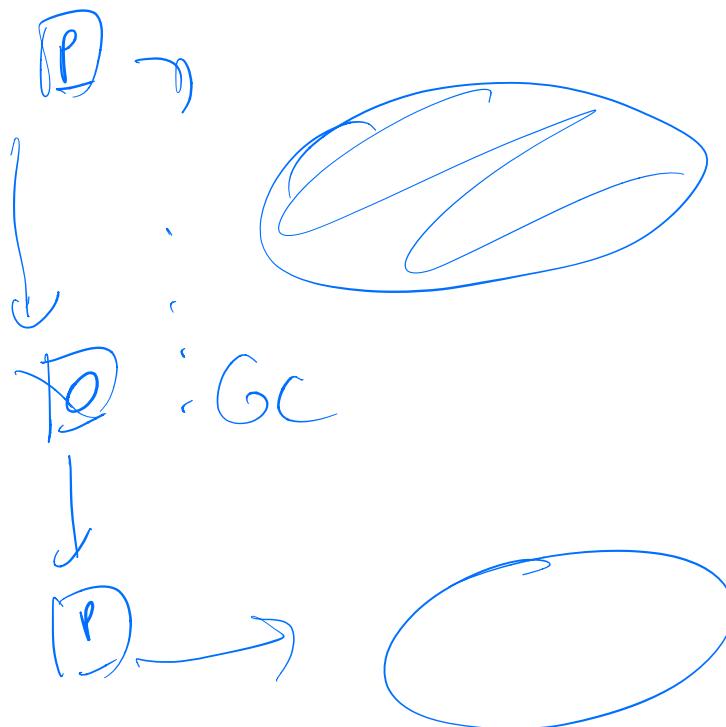
image
processing

NLP
speech
understanding
SGD





$$A \text{ XOR } C \text{ XOR } C \Rightarrow A$$

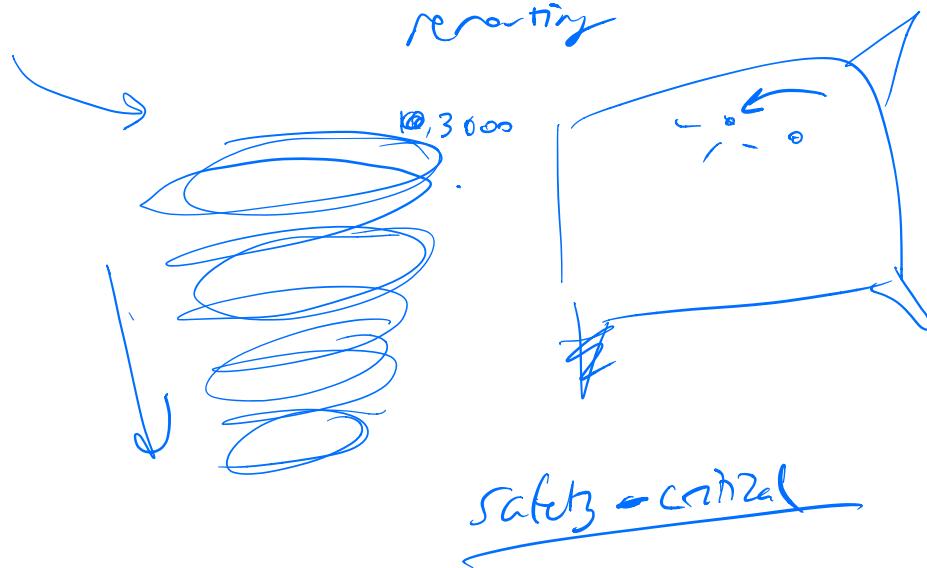


ATC

backoff protocol

reporting

10,300s



checkpointing
backups

phone apps
not safety critical
acting
checkpoints
backups

fsck

filesystem check

OOM killer

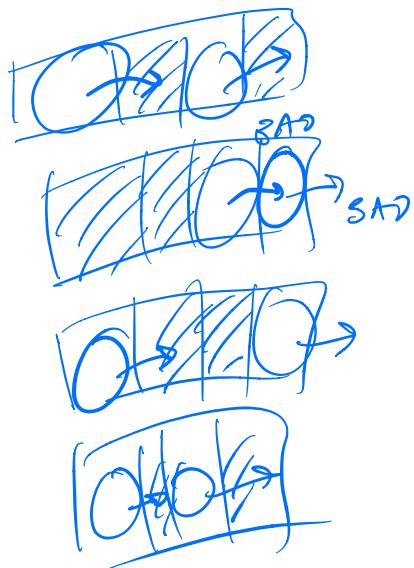
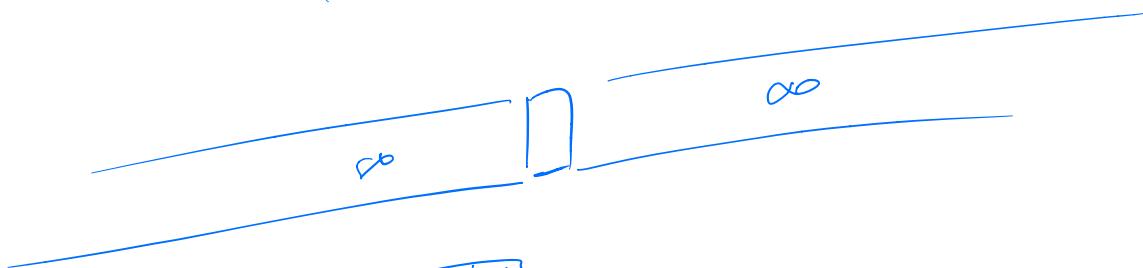
Draft

FO - slow

Archipelago

mem hazard

infinite heap demands



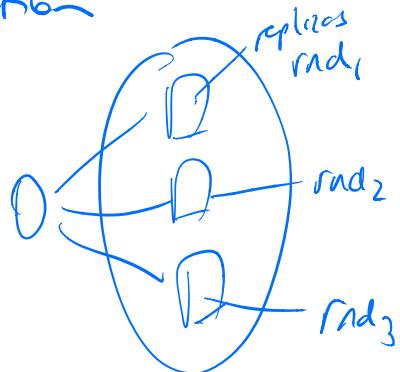
Fahr \rightarrow Heißabzug

$$\begin{aligned} &P \\ &E[\text{# hours failing}] \\ &= p \times N \end{aligned}$$

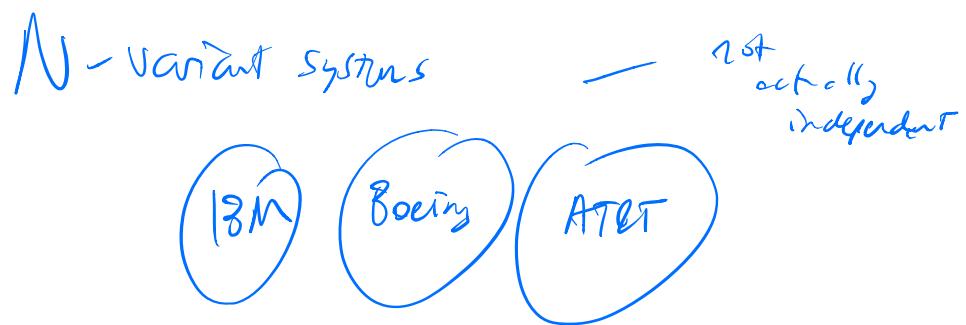
Space for reliability

probably - is not
of probabilistic

replication



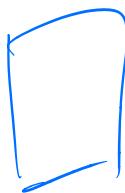
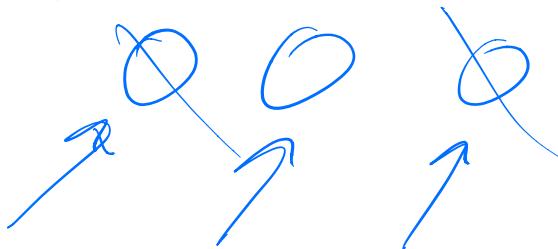
independence
of
failures



replicated state machines "RSM"

Dressendorf & Schneider
(Cornell)

DFA_i



determinizer

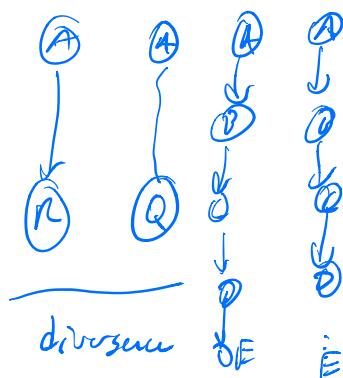


rep^g
implementations
w/ determinization

(threads!
async I/O
(in-order arrival
(of I/O)
req. for det.)

User input

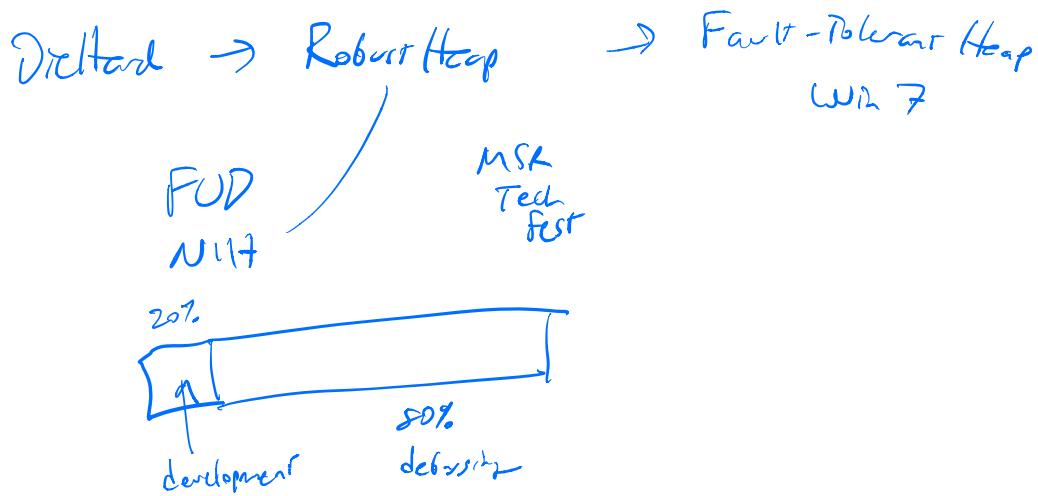
Scheduling



RSMs
 + determinism → Polar
bias
 RSMs
 + Deltard → Herderbrassw/
probabilities

Time
 ↗ RNG
 never roll your own crypto
hash fn
 ↗
 binary search ↗

GPL infection



Security System

how secure is it against attack
 by an attacker
 who can do X

} threat model

fault tolerance
independent
 hardware
 software
 correlated...

Deltard → Win 8





profilers

throughput too low
latency too high

prof

The count

main() {

$\leftarrow l_1++$
 $\leftarrow l_2++$

—
} R. ()
 f-L)
 R-C)
 ——— lin
 }
 S
 }

end-to-end

"kernel" (not os kernel)

4) heavily
Computer -centric
library function

u_m, \dots
def learning

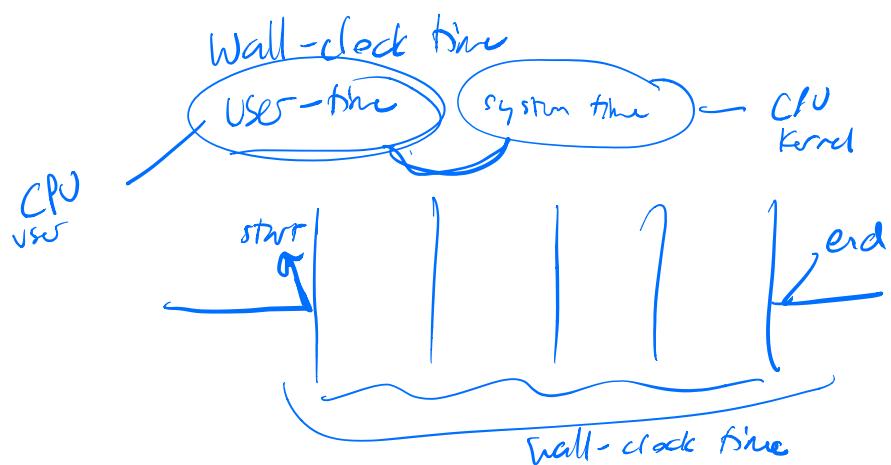
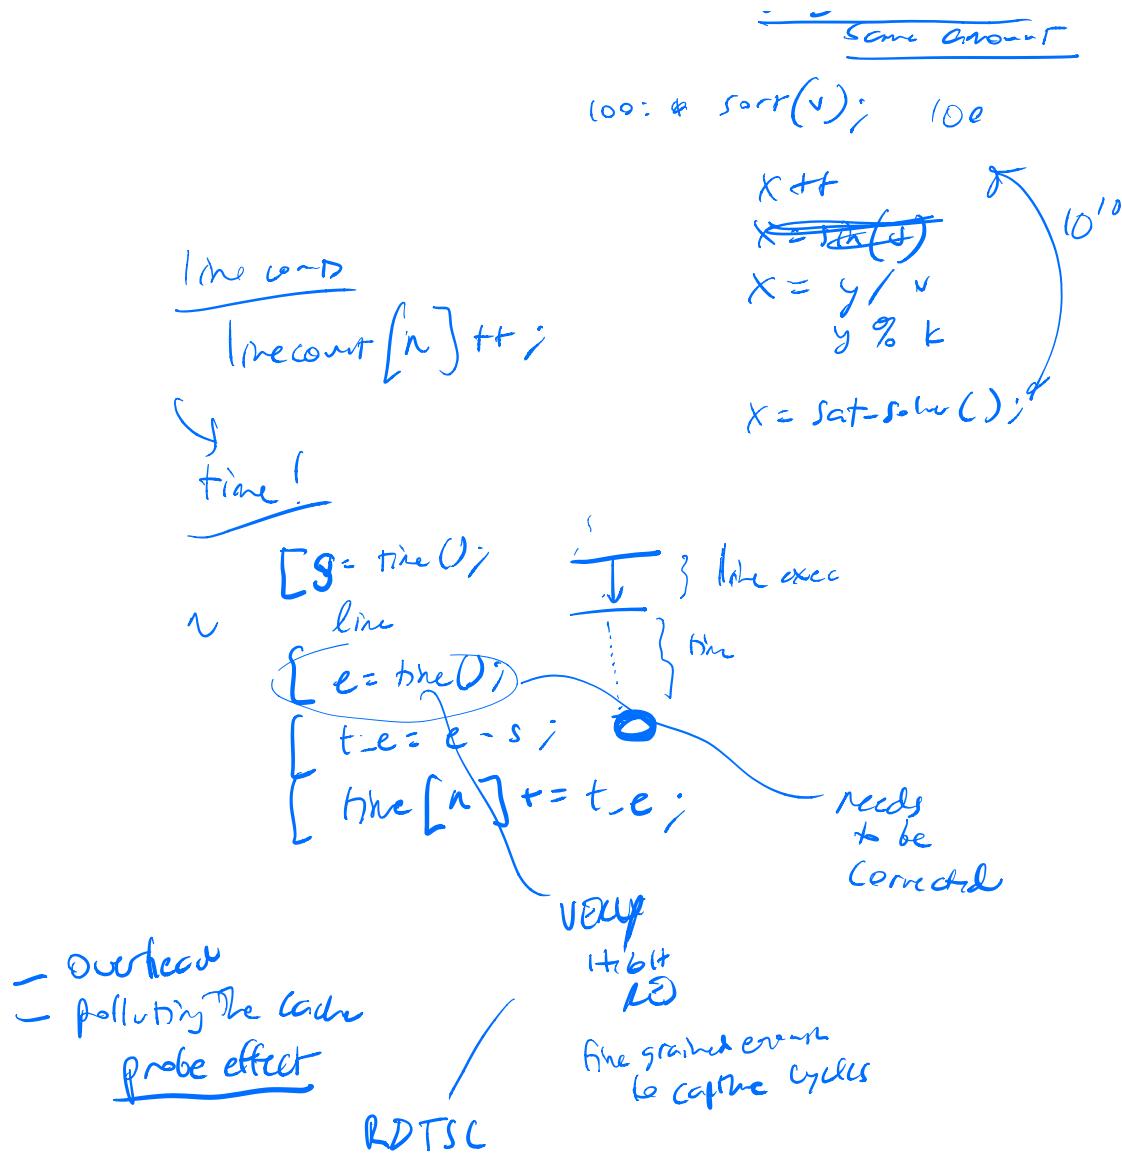
NOTE:
SEMAL
CODE
only!

Advantages (Pros)

If the count (gap count) is 6
→ perfect

Dadurch verhindern (cont.)

- overhead
slow program down
 - assumption —
over time costs



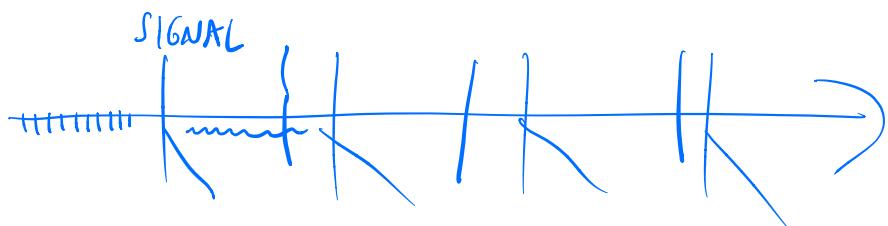
$$U_{\text{let}} + S_{\text{sys}} = 0$$

$$U \xrightarrow{\text{Instrumentation}} U' \xrightarrow{\text{Unbiased}} E[U']$$

- Stat -
- low overhead ✓
 - accuracy ✓
 - unbiased (time not memory) ✓

Statistical probing

BEDWARE of
Sampling bias



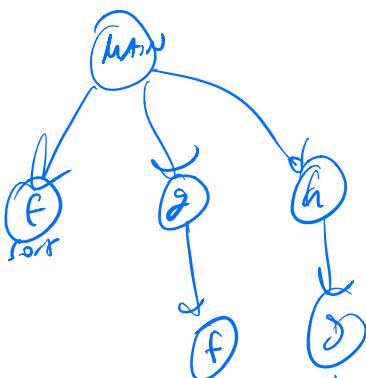
50% line 12

10% line 11

5% line 15

Low & Large Number!

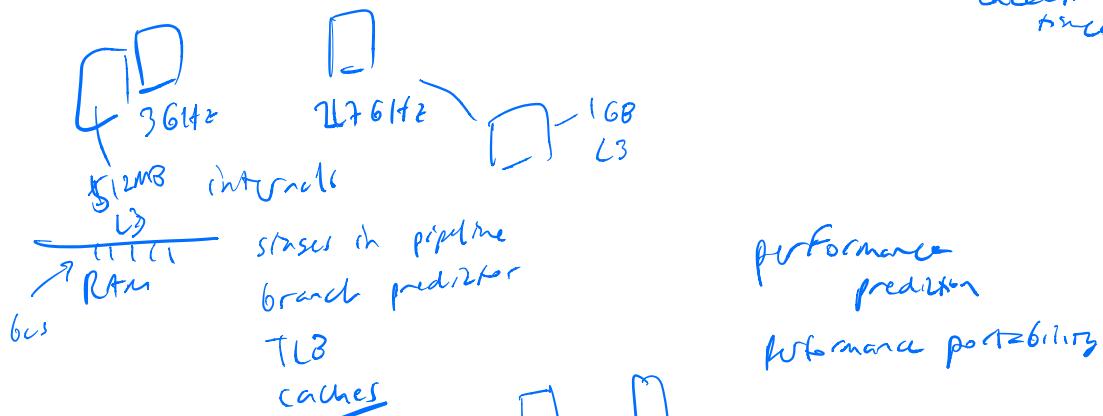
:



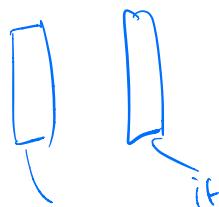
602

Average
total
max quarks = -

Statistical sampling - periodically interrupt
→ Where do you land
collect samples ~ distribution
of observation times



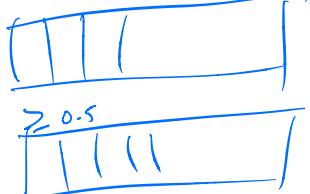
performance
prediction



$\text{RND}(0.0 - 1)$



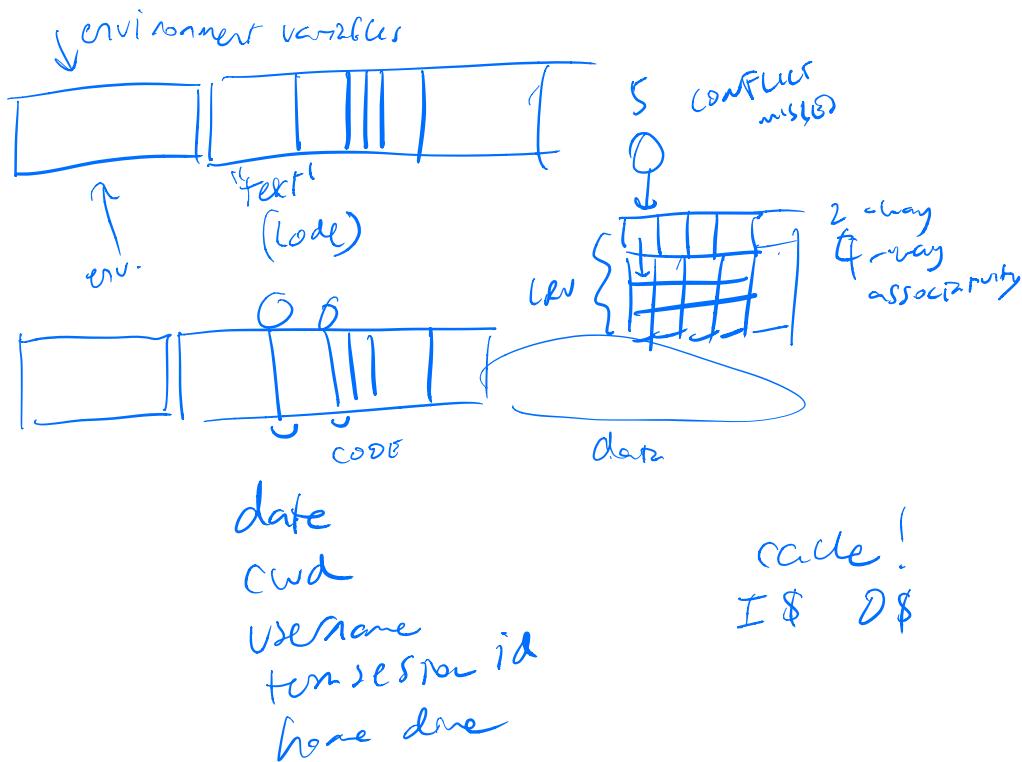
∠ 0.5



if $\dots \cdot = 6$

Patterning Table

Producing wrong data w/o doing anything obviously wrong!
Mytkowicz et al.



On-chip hardware performance counters

perf

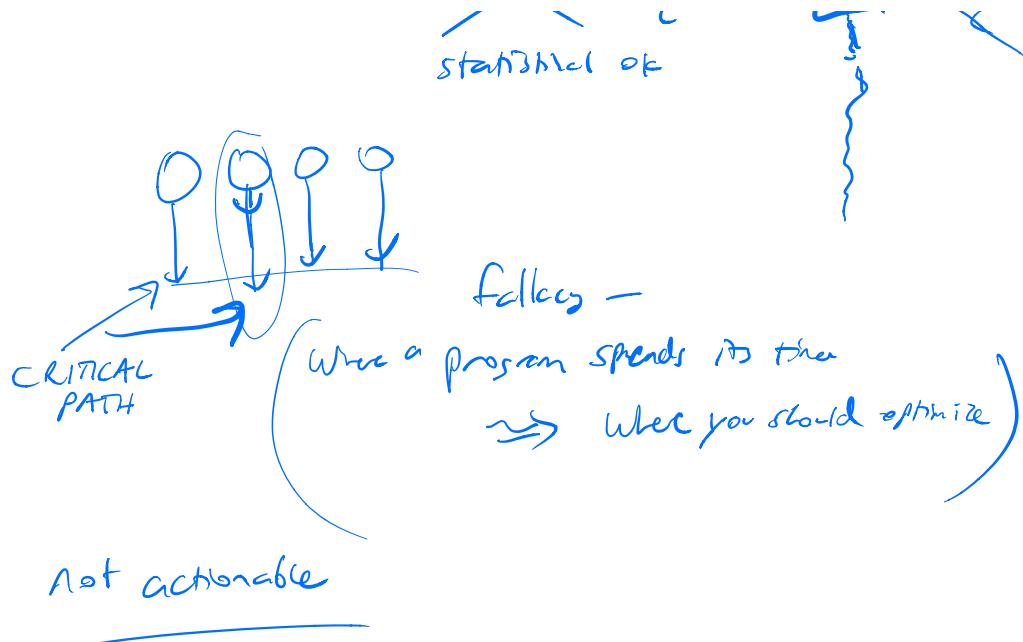
Counts up
countdown timers — signal when perf counter hits zero
 \Rightarrow profiler for real code

the old days

good
no threads
virtually no I/O
non-intrusive
asynchronous

~~prob effect!~~

The diagram shows a tree structure where multiple small circles (representing processes) merge into a single large circle. Below this, a crossed-out note reads "prob effect!".



profiler wrapup -

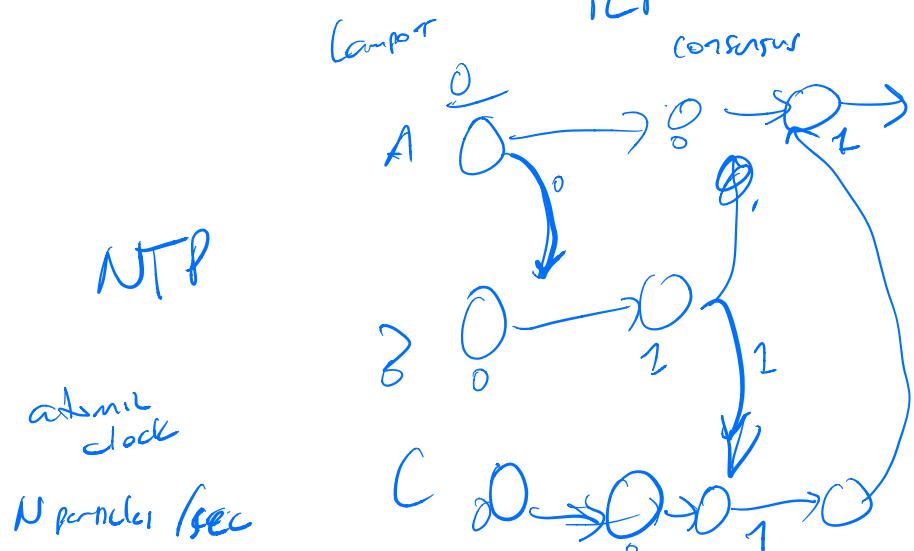
- Instrumentation based
- Sampling based
- perf counter based profiler
- causal profiling
- + concurrency
- + asynchrony
- + parallelism
- distrib. sys?

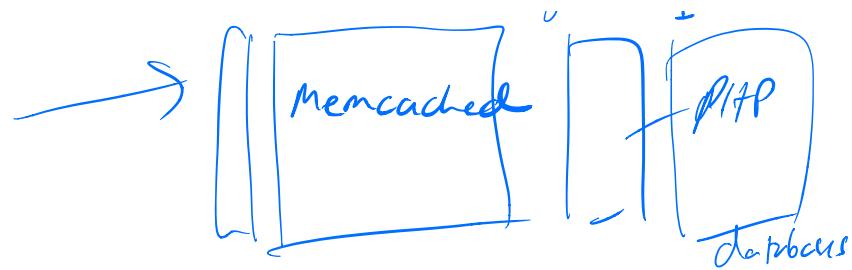
"single node"

distrib sys?

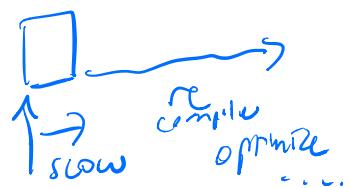
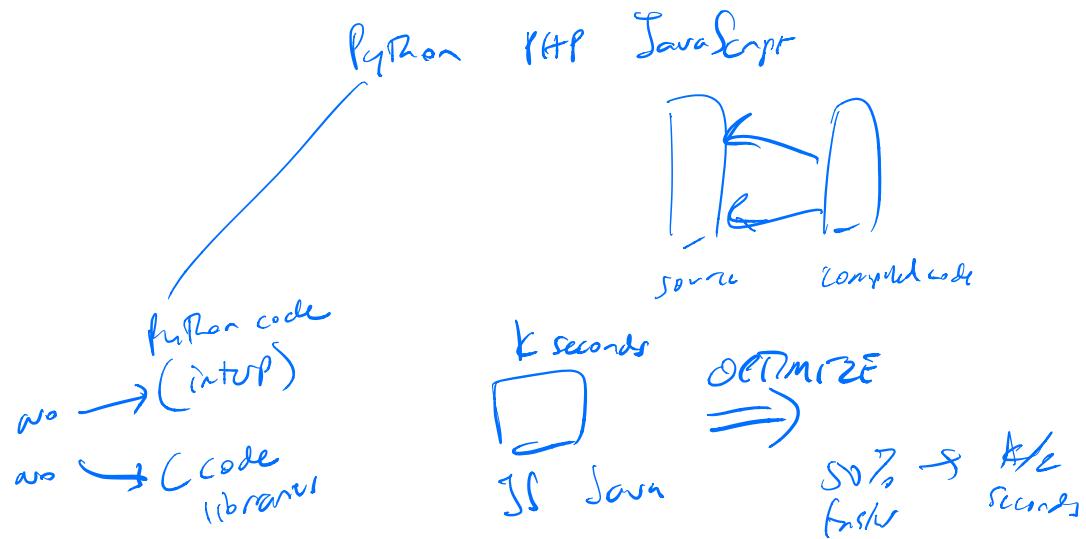
no global sync clock

FLP





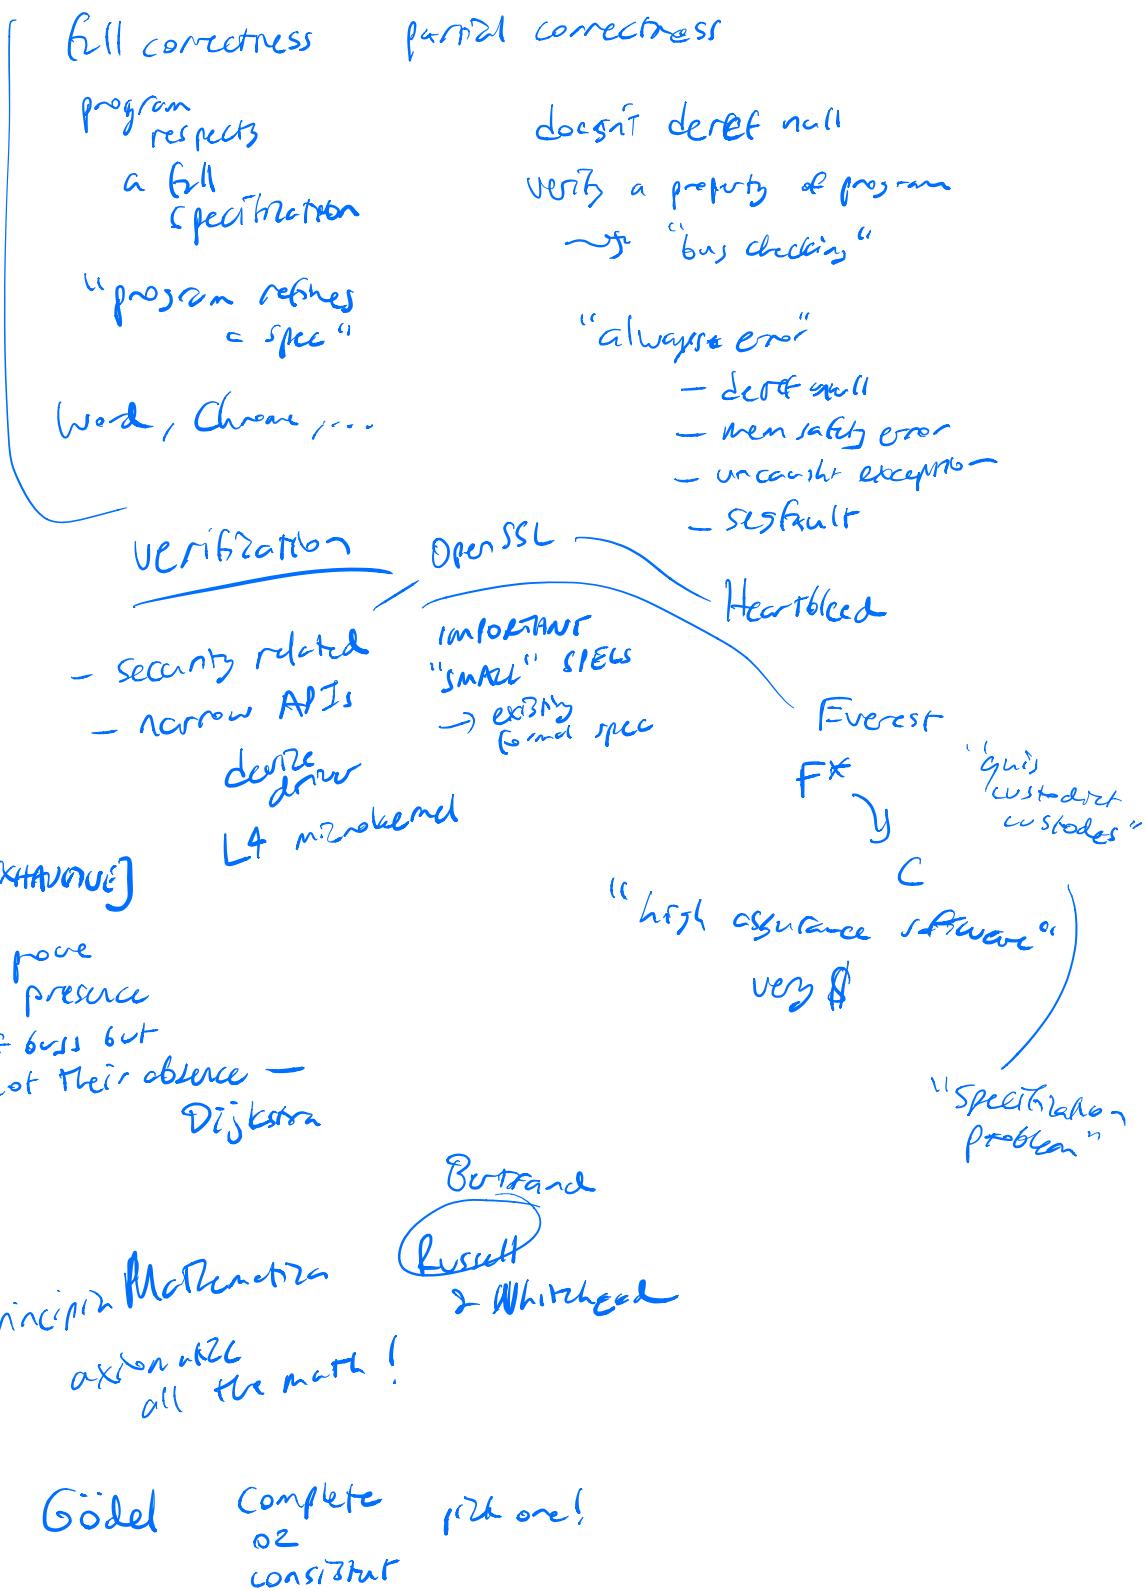
profiling "mixed" languages



perf hard

Correctness hard

bit monotonically increasing
- compositional



dynamic analysis - precise
 static analysis - low false positives
 testing - imprecise
 high false positives

low recall

bias detection

"heuristic static analysis"

"sound"

q: If there's any possibility
of error, report it!

arises due to over
approximation

$x \in \{1..10\}$ $y \in \{-5..0..5\}$

numerator

denominator

Fundamental tension

Soundness - low false
positive ratio

Increasingly - , 0, + x/y sound \Rightarrow
precise error here!

(ptr analysis)
(still sound)

"Scalable"

"Infer" PB

$1/6 \log 0 \dots$

Valgrind
Google's
Dynamic analysis

Sanitizer
Address
Thread
Undefined Behavior

ASan
TSan
UBSan

Clang++
-fsanitize=address

Pin binary instrumentation
- dyn analysis

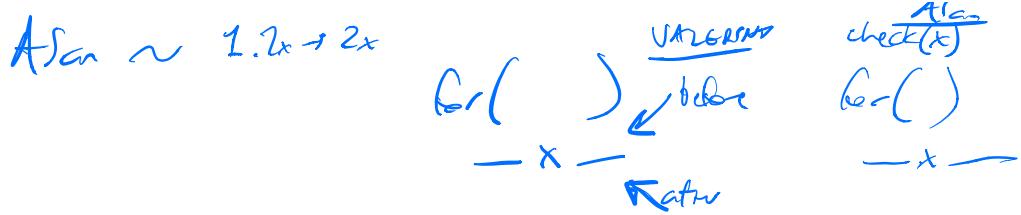
$x = \text{malloc}(32)$

\rightarrow before
read/write

\rightarrow after
read/write

$$\frac{\#_{\text{err}}}{x} = \frac{\#_{\text{err}}}{32}$$

Valgrind ~ 10x-100x



GC — "dynamic analysis"

eliminates owners instead of repeating Exterminator

- lgc



testing

You are the tester
(user)

Correctness - question of
economics

- users are more forgiving

- users love features

- "first mover advantage"

"slip early"

- network effects

second system

How do
we know

software is "ready"?

good enough

- satisfying

1968

"Software Crisis"

- how can we
make
process more
correct?

- software architectures

- process

Waterfall model

"agile"

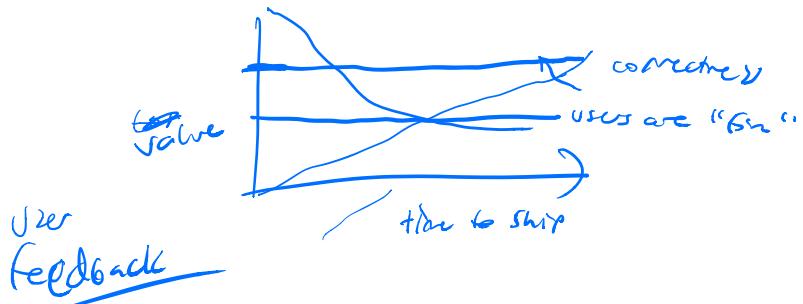
scrum

TDD

test-driven

development

- analysis



bug reports

- sparse
- imperfect

retro case steps required
to trigger error

automatic & high quality bug reports

→ "Watson" Windows Error Reporting
Crash reporting

Why yes?

give up privacy
for efficiency
for "important"
failure for me!
want to help (some people)
easy
default

Why no?

invade on privacy
never looked at
too much work
(not just "no")
follow-up questions

"minidump" - metadata
SOS version
HW
some state stack function calls

PID personally identifiable info) scrub

GDPR retention ~ 120 days
PID

differential privacy

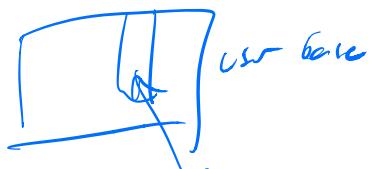
less info → better for privacy
worse for debugging

Cynthia Dwork

Watson — used for triaage

"field experiment"

web apps



gradual

deployment

if bad model

Random

subset has new feature

new version

"rollout"

losses ↗
are lossy
process
log at lot → loses down
log too little → can't reconstruct error

testing

- few
- all inputs
- I control everything
- deterministic

If anything → think harder
end-to-end test

→ no prior on some inputs
observe results

"be smarter"

Unit tests

JUnit

KUnit

class —

@test
void - { }

- lots of tests needed
- rewrite tests
(if they rely on internal state)
- limited -
within module

- tests can be wrong
tests are software

[integration tests]

Unit → CI
- bias!
when are tests written?

- a priori "best practice"
 - TDD - test-driven dev.
 - XP Extreme programming
 - agile
 - pair prg.

- a posteriori - after the fact
add after

tests run - regression testing

- test coverage

- line coverage } proof proxy linear
 - path coverage ↵

- "flaky tests" deterministic
tests
→ MT

fuzzing

- + automatic
- + unbiased

Random \rightarrow unbiased

Likely to explore cases you hadn't considered

explores state space in unbiased manner

\Rightarrow unanticipated corners

- Complex input unlikely to be able to generate it

$\langle \text{HTML} \rangle \subset \langle \text{C/C++} \rangle$

if ($x == 0.0$) { 11,767,033 years

}

STATE SPACE EXPLOSION!

AFL American Fuzzy Lop

Coverage driven
"mutation" of inputs
Random process \rightarrow driven by branch coverage

DART
↓ directed

↓
if $(x == 0)$ {

}

\ systems that do not do much with the heap
constraints to ever variable

Windows
device
drivers

UNPATT

DDK

property-based checking

Quick Check

testing w/o
or oracle

fast w/o spec

+

$$x + 0 = x$$

differential testing

$$x + y = y + x$$

gce don't
FFox Chrome

PROPERTIES

- no spec

- better oracle than oracles / memory

- one impl.

$$\text{rand}_x \Rightarrow +$$

$$\text{rand}_y \Rightarrow +$$

"python property testing"

DISTRIBUTED SYSTEMS

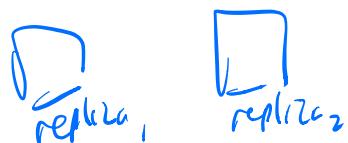
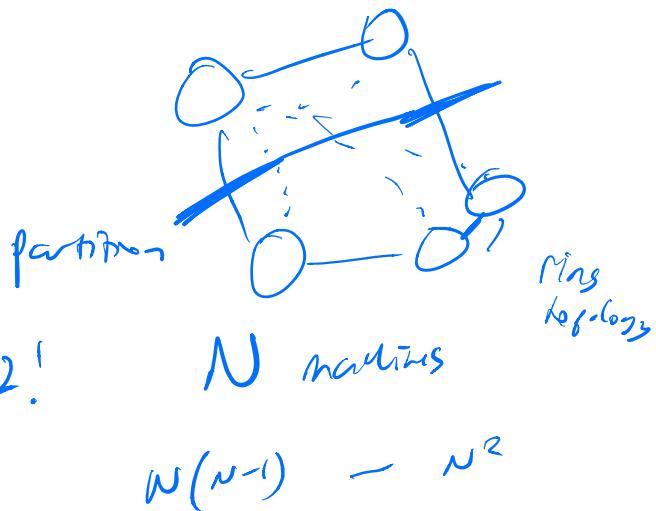
- geo-collocation (locality)
- reduces



CAP theorem

— Eric Brewer

consistency
 availability
 partition-tolerance



in sync
 ⇒ consistent

eventual consistency



Twitter follows
 # of retweets

other domain specific

Consensus



START: initial

binary value (0 or 1)

END:

every non-failed
system has same value

Synchronous

- terminates
- all agree
- value "valid"

some # of
rounds -

Majority vote
you know someone
has crashed

asynchronous

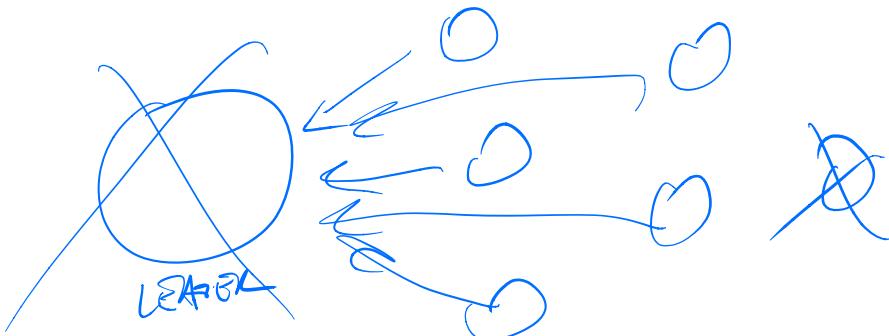
- no upperbound
on delivery time

NO

timed!

majority
vote

outputs
result



Paxos: A Part-Time Parliament

RAFT

- leader election
 - detect crashes by timeouts
- log replication
 - (leader → commands from client applied to log)
 - log → replicated to servers
- only servers w/ up-to-date logs are eligible to become leader
 - "abort" → retry
 - "metadata"

Byzantine faults

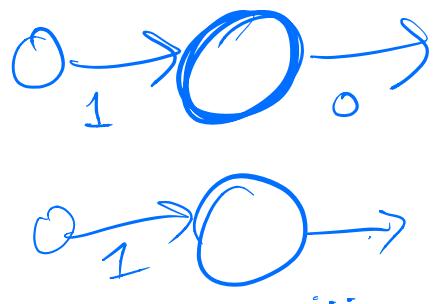
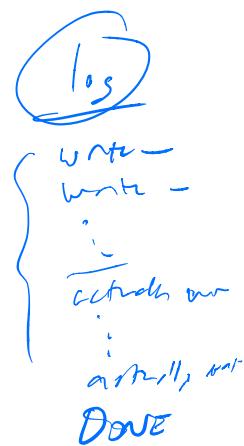
Byz General problem

journalling FS
(transactional)

ACID

atomicity
consistency
isolation
durability

NTFS
etc & q



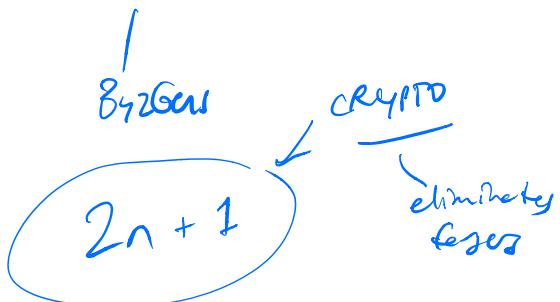
Byz Gen - forge messages
change messages delay messages
drop messages

BFT Byz Fault Tolerance → Libra

↙ Castro

$3n+1$

Practical BFT



↙
Zigzag (s_1)

BChain

not anonymous
very power inefficient
"Proof of Work" = compute

solves no real problem

- decentralized
 - true transactions
- CRYPTO!

microfragment
→ Brave

