

Kongjie Lu (UMN) - me: safe languages aren't enough

examples

- $\log 4 J$

Equifax vulnerability ??
"Apache Struts"

- "untrusted contributors"
(well...
typically have
to submit PRs...)

vulnerability
known
but not patched

11 min

3 pruned approach

4 directions

me: what about
Reinhart,
Spectre
etc.

understanding
"assumptions"?

- fundamental techniques
- analyses, automated patch fixing
"new classes"
- hardening systems,
enforcing least privilege
- "secure compilers"
compiler - introduces

This is so high level!

doesn't
explain CFI

dissim in on indirect calls
& CFG analysis, now

(can collect info at execution
from early, also
ptr analysis, addr taken,
type analysis)

"MTA"

type analysis

or
access to ptr
for invocation?

every layer
multimon

types of
access for disambig.

kernel
stuff

not
clear
how much
this can, no 'polano' yet

compare
to Nook,
but no actual
mechanism, so...

identifying importance
of vulns in Linux
when in Android

"DIFF CUS"

seems pretty meh,
I don't understand
why Android vendors
don't have resources

Future work

ME:

LARGE body
of work on
type analysis
in Java

Copilot
observation
not
new

I think they
are already
looking at this
+ Gd but not bad

Kexin Pei talk

static program analysis - not always required, often problematic
Unusual take on static analysis requiring “rules & heuristics”
OK, use ML to automate program representation analysis, let’s see
ChatGPT, finds issues

Haha Scalene, nice touch (cite me! :))

Google experience?

ML4CODE.github.it

Limitations of code2vec (changing variable names from array to types, model changes)
(Did not actually understand semantics)

“Program semantics does not just manifest in static text”

[me: of course, we can extract more info, even dynamic info for Scalene]

- overfits to spurious textual and task-specific patterns
- distribution shift: program syntax and task requirement changes

Binary analysis (untyped not an issue with WASM, interesting)

- strips variables, arguments, etc. -> lowered to registers

“ML needs to understand semantics, not syntax”

Security requires rigorous understanding of program semantics

Search through binary (binary code similarity), compare to Yahav?

“Data-driven program analysis”

(I really like the program smoothing, which was an idea I tried to get some students interested in)

Oh no, not going into the white box thing for neural nets, too bad

“Programs only make sense when executed”

Hm, ML cannot understand instruction semantics (Carbin work?)

Dynamic analysis - large-scale vulnerability might not scale (maybe because execution is “slow”),
potential false negatives

Key high level idea (22 min)

- Learn execution semantics of program
- Transfer knowledge w/o dynamic analysis

1. How do we collect and represent program behavior?
2. How do we train a model?
3. How do we avoid “expensive” dynamic analysis?

Me: guarantees?

Under-constrained micro-execution

ME: How does this not just lead to SEGVs all the time?

ME: Isn't this also time-consuming? (Running these tests, training the model...)

instructions + data flow states (values updated) + control-flow state (covered) + code addresses

I'd love to hear from Subrhansu how conventional or not this encoding is

transformer self-attention + masked language modeling

Embodies synthesis of instructions, or forward/backward interpretation

Mask more "dataflow states" (actual contents - concrete values of registers)

- Very rough initial response to Brian ("already discussed this" (!))
- Good follow-up though - difficult to train model to extrapolate ("best-effort")

Fine-tuning as static analysis using pretrained model

Binary similarity - maximize agreement between two codes compiled from same source code

Future work:

1. Multi-modal software (heterogeneous "modalities")
 2. Security applications (automate existing security tasks and enable new security apps)
 3. Robustness guarantees (develop & evaluate to ensure strong robustness guarantees)
-
1. ASTs, specs, physical interaction, natural language like comments, documentation, var names, etc.; configurations; execution traces, coarse-grained system logs
 - A. Grounded data-driven program analysis
 - B. Automate existing tasks

Correlate events from physical interactions - cyber-threat investigation or debugging system failures

Document-grounded invariant generation

Interactive debugging / interactive reverse engineering

Principled robustness measurement

- random testing
- "How can we systematically explore the program space?"
 - "Transformation-oriented testing" - focus on well-defined program transformations (?) - semantics-preserving syntax alteration, evaluate if model stays invariant wrt changes
- Develop data-driven program analysis with provable robustness by construction
 - Formalize transformations using "symmetry group"
 - "Open to neuro-symbolic approaches"

Yi Ding talk

A Holistic View on Machine Learning for Systems

Co-designs ML and software systems

UChicago, postdoc MIT

Meta

Heterogeneous hardware (CPU, GPU, FPGA, ASIC)

Software (MySQL, Cassandra, Httpd, Nguni, HDFS, Apache Flink, Transformers (hugging face))

- configuration parameters

Maximizing accuracy of ML is not the goal - instead, “holistic view”

- “understanding cost structure of systems problems” - trade offs

Trade model accuracy for reduced data collection (not all samples are equally useful)

- maximizing efficiency over all samples is not necessarily desirable
- Only configs that represent optimal power-performance are relevant
- asymmetric costs - under prediction may be less or more costly than over prediction

ML for systems

- Neural surrogates
- NURD
- Multi-phase sampling (ISCA)
- NURD (MLSys)

Fundamental ML thrust as well

Minimize energy under performance constraints

SRAD application - all configurations on ARM big.LITTLE system
(# of cores and clock speeds) - also # sockets, memory controllers,
cores/socket, hyper threading

Optimal frontier - lower convex hull

Multi-phase sampling (new sampling technique)


- predict optimal frontier and then sample from it

Phases:

1. uniform sampling (from lowest to highest) - not random [me: could be quasirandom]
2. Predict
3. Selective sampling (predicted energy efficiency = pred perf / pred power)
 - A. Pick from samples from phase 1 with highest efficiency
 - B. Predict perf & power from other unsampled configs

(How do you predict the different ones?)

matrix completion methods & hierarchical Bayesian methods



[ME: is the perf not workload dependent? If not, how do you cope with it? Is average OK? How well does this generalize? **Compare to adaptive approaches?**] — used worst-case inputs
[ME: also isn't perf very closely connected to power? If not, why not?]

Cost structure: not all configs are equally useful (majority is not)

[ME: similar insight to what I said years ago about training for important customers]

Online-straggler prediction (NURD)

- Predict task latency based on task and node features, on-the-fly w/o prior knowledge

Trains model based on which tasks have completed by a time threshold

Use finished tasks, re weight predictions based on a function of dissimilarity between finished and running tasks

Gradient boosted tree for regression predictor

Weighting function:

- adjusted weight \geq threshold \rightarrow straggler

Propensity score - conditional probability that task belongs to a class of finished tasks given its feature at time t

- use logistic regression since there are just two categories

Use this for reweighting

Evaluated on two public traces from Google and Alibaba Cloud

(For me, holistic here really is “end-to-end”)

Meta - “Acela” [arxiv 22]

- over subscribing a server with too many maintenance yields server being withdrawn from production
 - Quantile regression
 - Currently deployed at Meta

AutoML - EuroSys paper - couldn't run experiments

Challenges in ML (or stats, or control theory) for systems?

- must capture causal relationship but ML captures stats
- Rethinking from causal perspectives
 - How to evaluate outcomes w/o bias when controlled experiments are not possible
 - Build causal models that permit observational studies
 - Answer counterfactual questions using offline log data
 - Uncover causality in systems when interference and unobserved confounders exist
 - Build causal models that incorporate domain expertise in computing
 - Develop new intervention strategy for causal discovery

\rightarrow resource scheduling and management

Sustainability & Climate in computing

- Reduce carbon footprints
 - Get high model accuracy w/reduced environmental costs

- Reduce # of samples and environmental cost of each sample
- Explore trade offs between carbon, energy, perf & reliability

Ramesh: big roadblock - explainability of AI is a problem (performance)

GIL+

- interpretability of configurations that lead to low or high performance
- design for interpretability - characterize hierarchical structure (cache misses, cache miss rate, IPC) - new ML perf model