

Energy Berger

energyberger.github.io / CompSCI-630

Systems Lunch @ Fridays - Fall

Random number

TA Abhinav Jangda

Github!

Hot CRP

energy@cs.umass.edu

- Scribes!
  - course load :
    - 1 core / semester for MS
    - ..
    - 1<sup>st</sup> semester MS / PhD

> 1 core = too much!
  - schedule - yes, we have a midterm & final!
  - reviews
- 
- ## FORTRAN (in COBOL)
- 1954
- |  |  |
|--|--|
| military <ul style="list-style-type: none"> <li>- calculating</li> <li>- artillery trajectories</li> <li>- decryption</li> </ul> | first computers<br>— people!                   |
| general-purpose computing machine  | automatic computers                            |
| Turing machine<br>von Neumann machine  | ] data & code together<br>Harvard architecture |
|  | fetch-decode-execute cycle                     |

## assembly language

- lowlevel
- direct access to memory
- operations not abstract

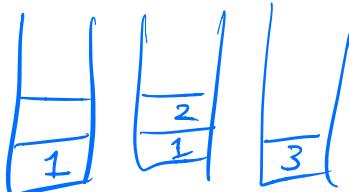
- stack architecture
- register architecture

A, B/C/D  
AX BX CX DX

- portability

ASM ≠  
~~portable~~

MOV 1, A  
MOU 2, B  
ADD A, B, C



PUSH 1  
PUSH 2  
ADD  
POP

"Intel" "ARM" (subset of instr.)  
JVM, .NET, Python bytecodes,

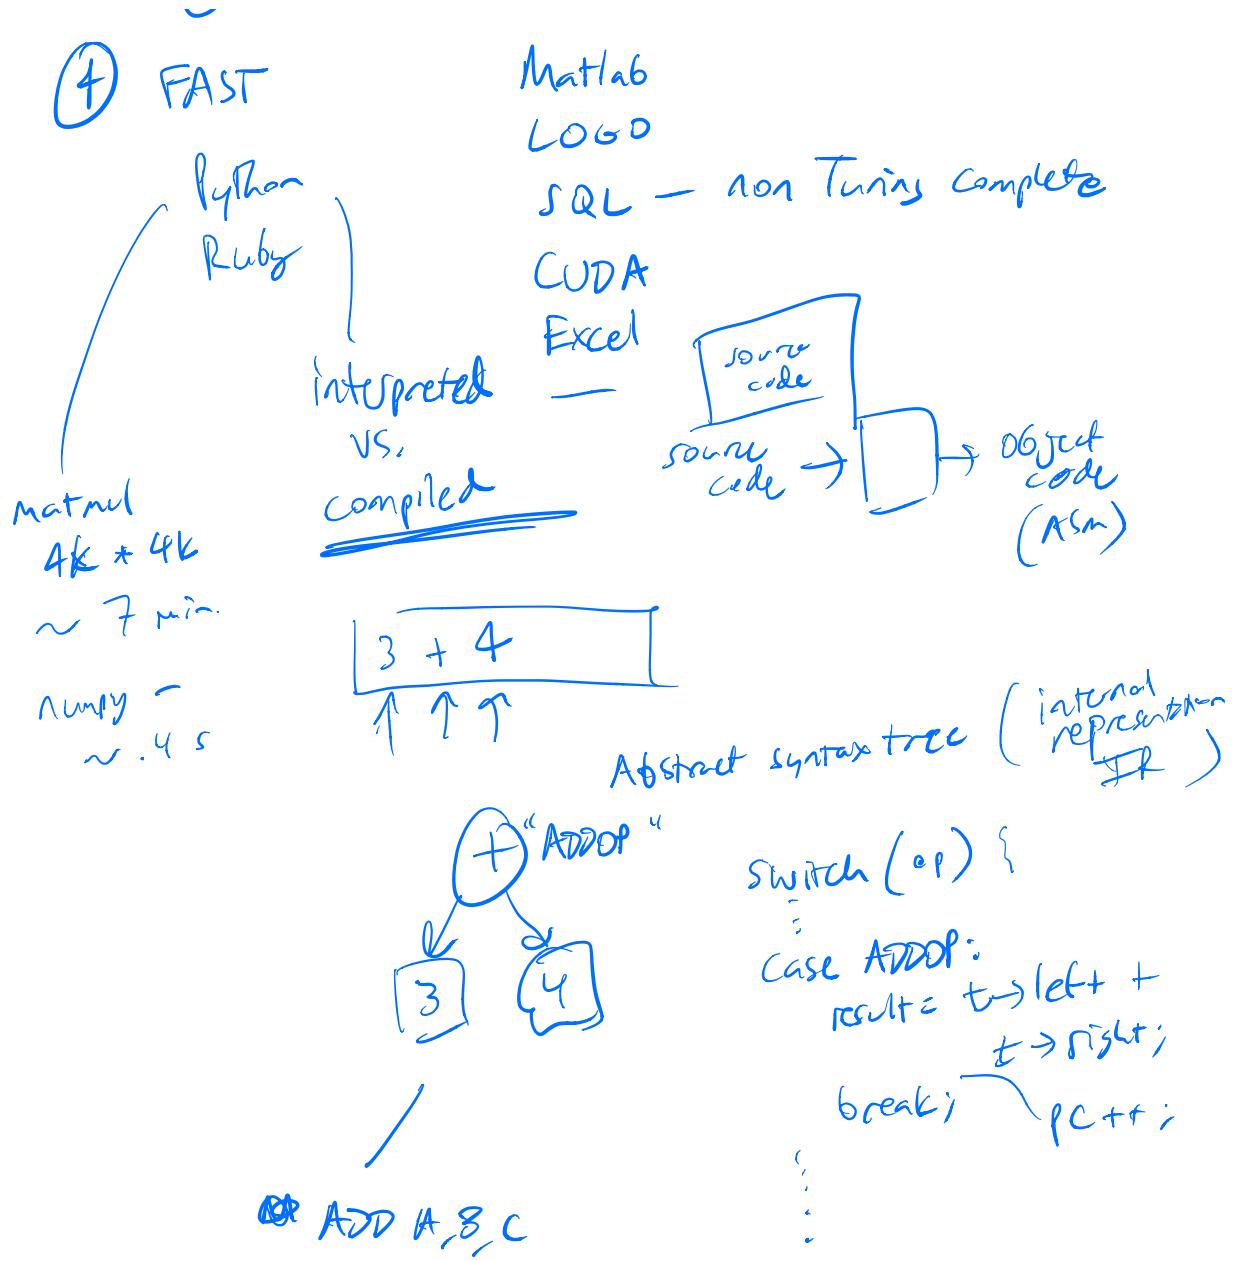
WASM

backward compatibility

- hard to read, hard to debug

- ① PORTABLE
- ② READABLE
- ③ WRITABLE

- COBOL - managers  
FORTRAN - scientists } domain-specific languages



programmer time MORE USEFUL than cycles

— modern POU

- productivity
- debugging
- ~ 100x faster!

interpreter → just-in-time compiler (JIT)

Python      PyPy  
JavaScript    V8  
                Chrome

WASM

Predicates -  
- Parsing  
- Compiling  
- program analysis  
- type coercion  
- scope  
- object orientation  
- modularity  
- structured prog.

FORTRAN  
1st optimizing compiler

I, J, K      integer  
X, Y, Z      float

DO 10 I = 10 10      \*

(was comma) line number

10 CONTINUE

BASIC



$\text{DO } \cup^1 \phi \cup^2 \psi = \cup^1 \phi$

control flow

$\text{DO } \phi \psi = \phi \psi$

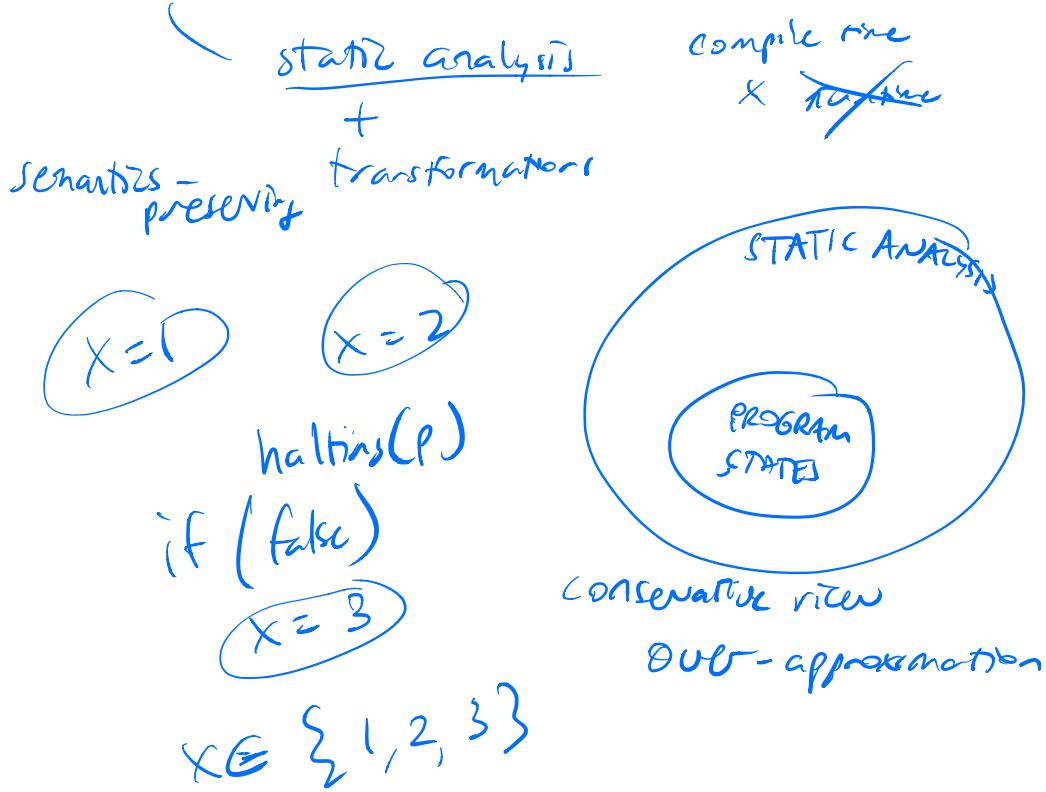
FORTRAN 77, FORTRAN 90 ~

writing

LEGACY  
FAST

The Fastest PC today! 2019 !!

optimization



## Halting problem

halts (program-text)  
→ true if program terminates  
                                  <sup>always</sup>  
→ false otherwise

$p'$       if    halts( $p$ )  
                          run forever  
else                    terminate

### ① Constant folding

$$x = 3 + 4$$

↓

$$x = 7$$

\* PI \* PI

### ② Constant propagation

$$x = 7 \quad z = 12$$

$$y := x + 3 + z$$

↓

$$x = 7 \quad z = 12$$

$$y = 22$$

### ③ ``strength' reduction

$$x = x \times x$$

↓

$$x = x * x$$

/ %

expensive  
cheap

$$x = x \% 8$$

↓

$$x = x \& 7$$

\_\_\_\_\_ 111

### ④ vectorization

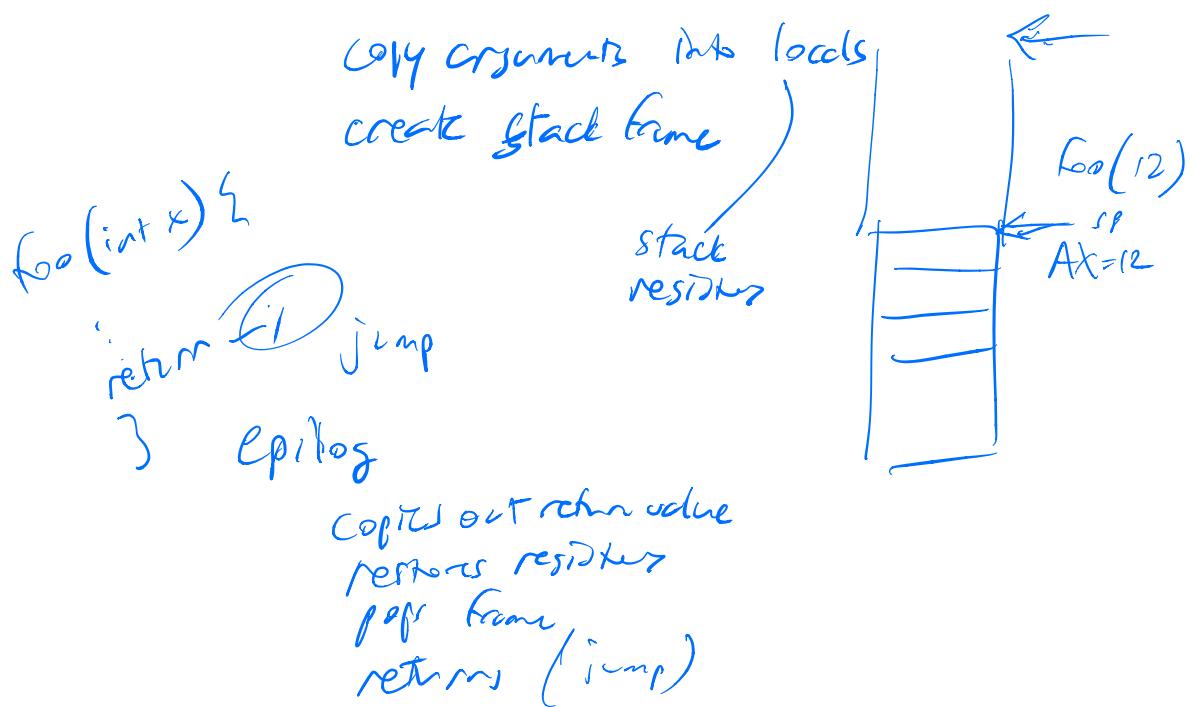
8	12	13	45
↓			
16	24	26	90

```
for (i=0; i<16; i++) {
    a[i] *= 2;
}
|||  

vector ops (4x speedup)
```

### ⑤ inlining

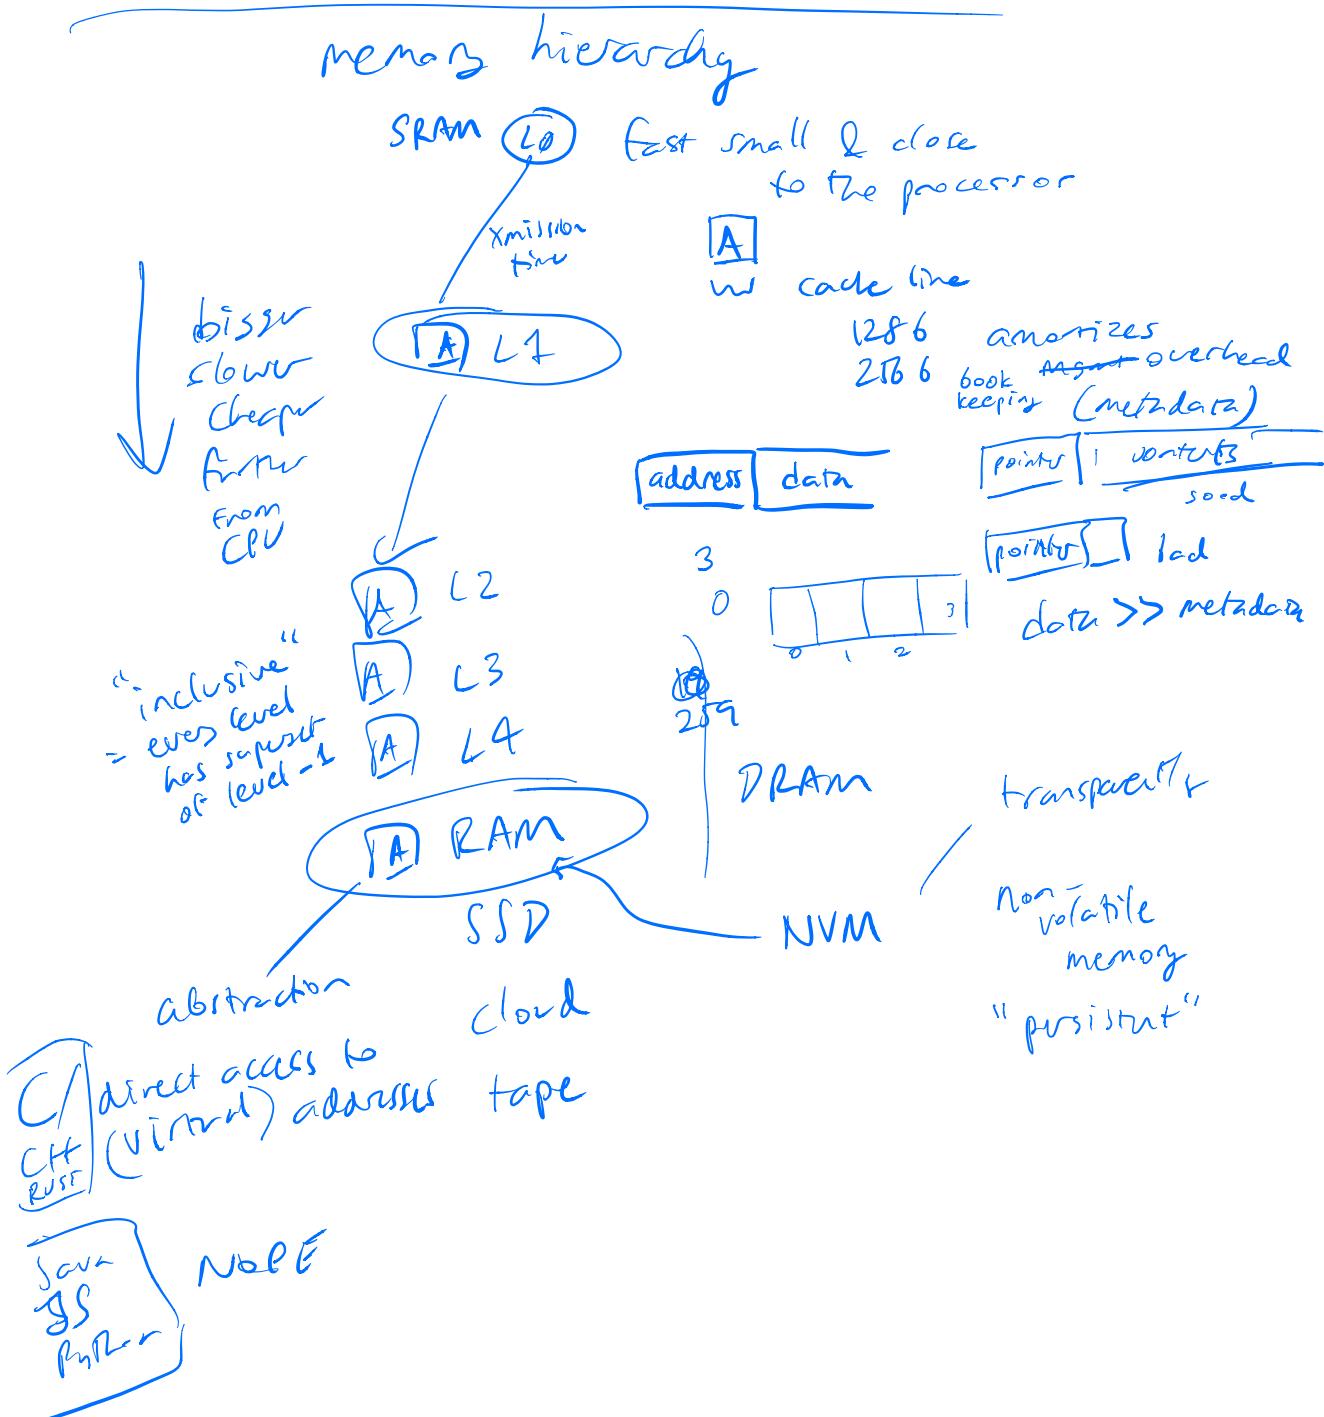
## function prolog

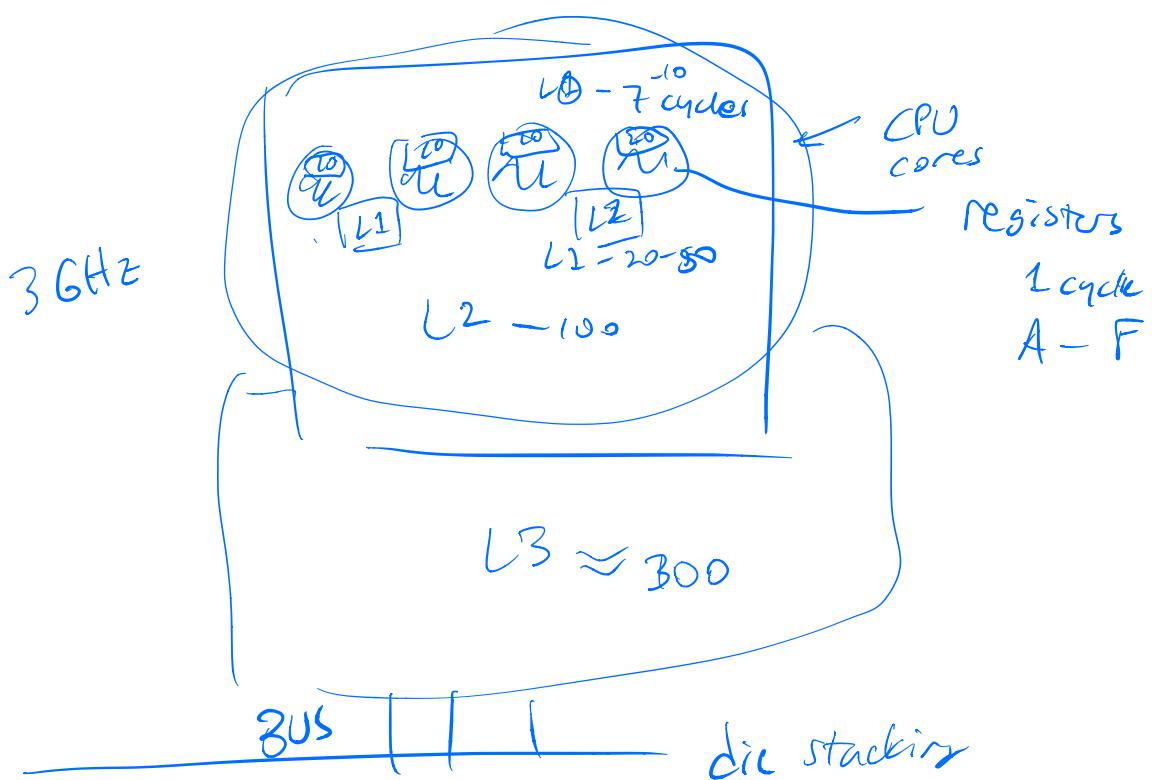


```
int add(int a, int b) {  
    return a + b;  
}
```

foo() {  
 z = add(a, b);  
}  
 ↓  
foo() { z = a + b; }

- code size
- I\$ instruction cache





RAM "1 Gbs"

~~1000~~

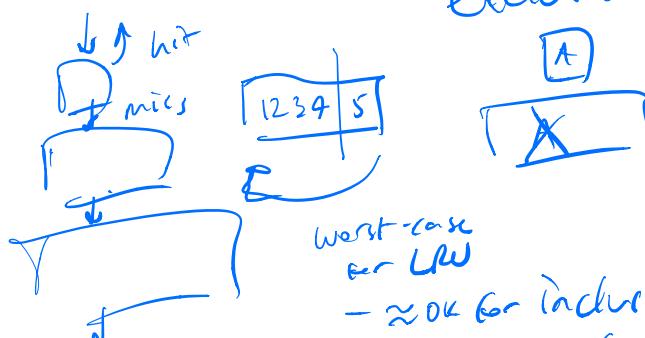
RAM bandwidth

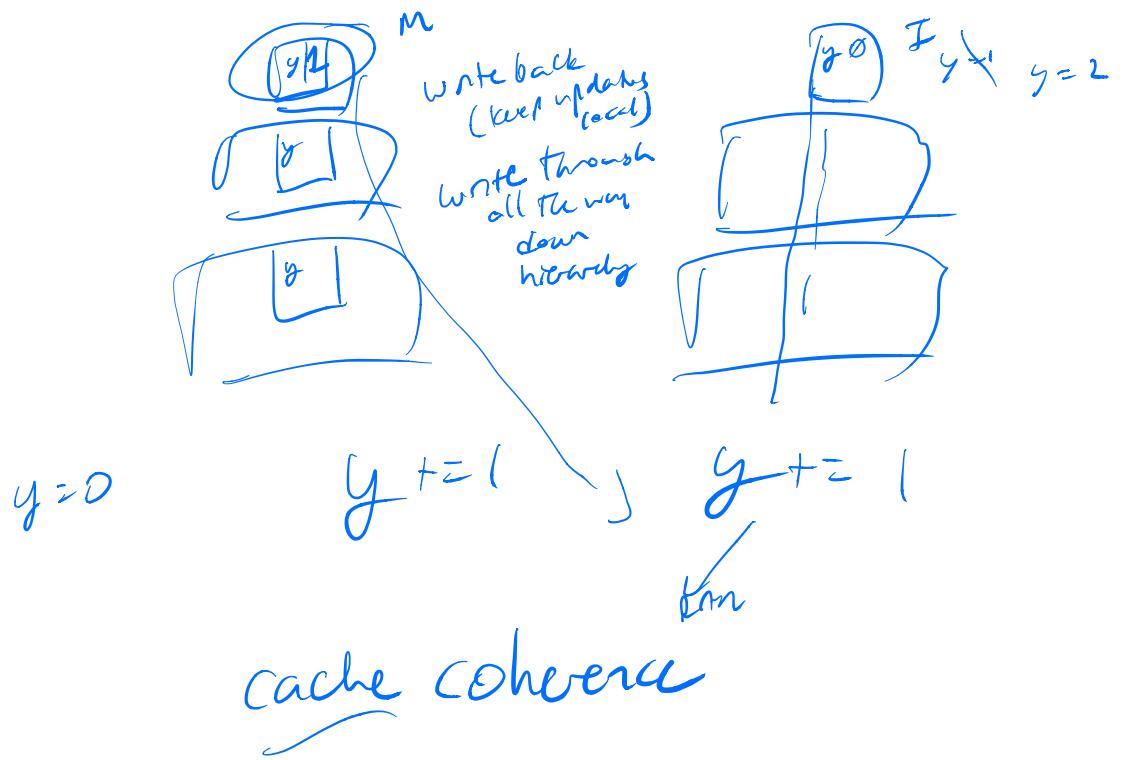
1000  
10000

Caches — "replacement policy"

eviction "LFU"

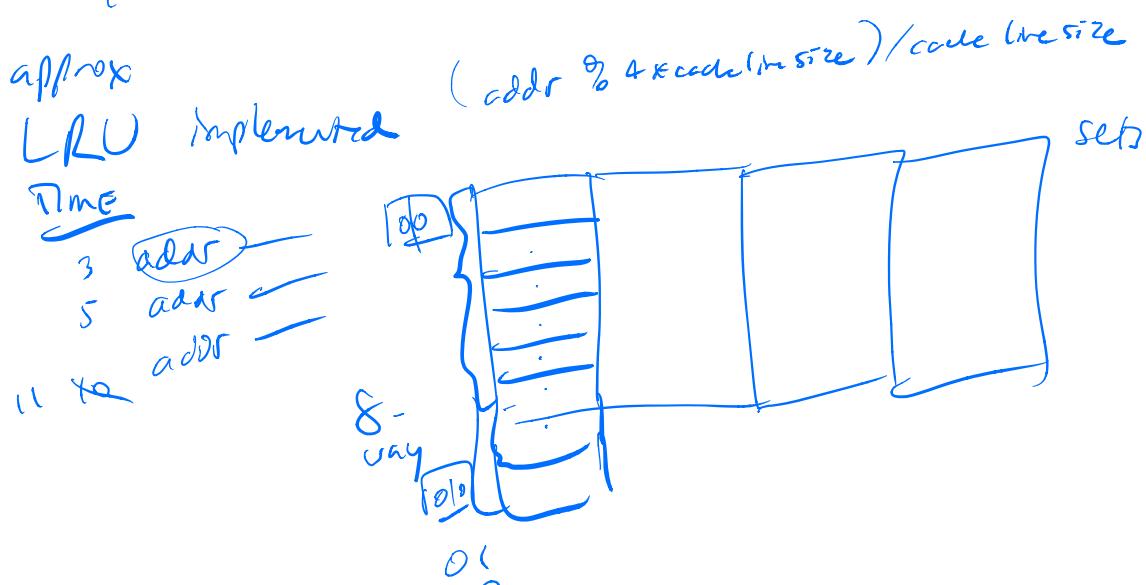
latency  
wait time  
caches wide latency





## MESI protocol

- { Modified
  - { Exclusive
  - { Shared
  - { Invalid
- I changed it  
just me  
several CPUs have it (modification)  
trash



Cache misses

Can be different  
types -

Capacity - "hash collision"  
Conflict -  
Coherence - 3C's model

direct mapped - 1-way

2-way

4-way

8-way

16-way

fully associative -  $\infty$  way

TLB

translation  
lookaside  
buffer

Cache of  
virtual  $\rightarrow$  physical addresses

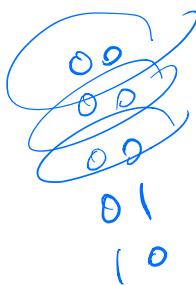
#include < stdio.h >

```
int main() {
    int a;
    printf("a = %p\n", a); 0x10000000
    return 0;
}
```

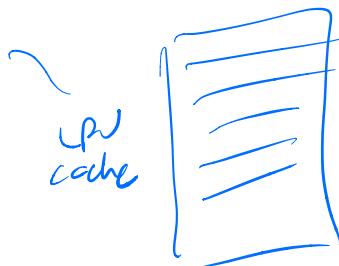
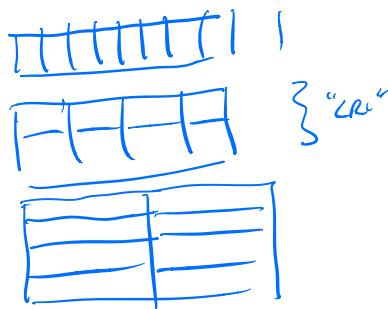
address

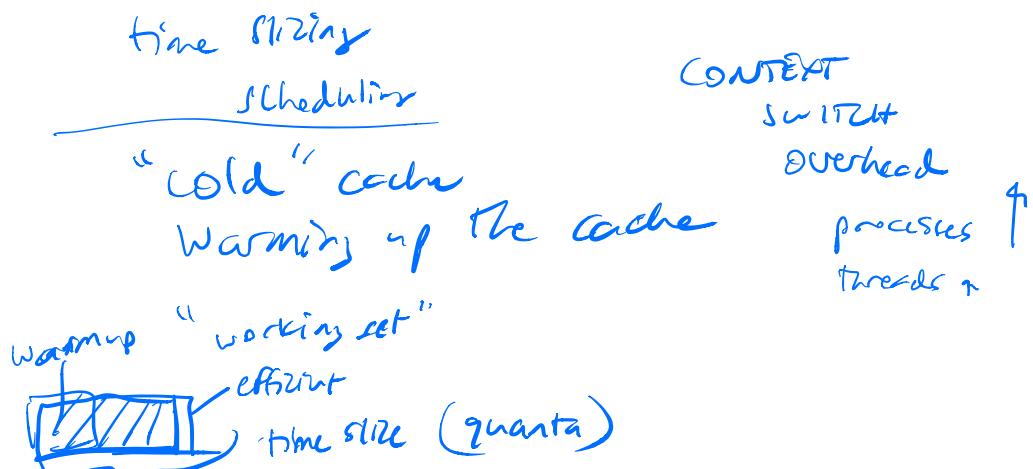
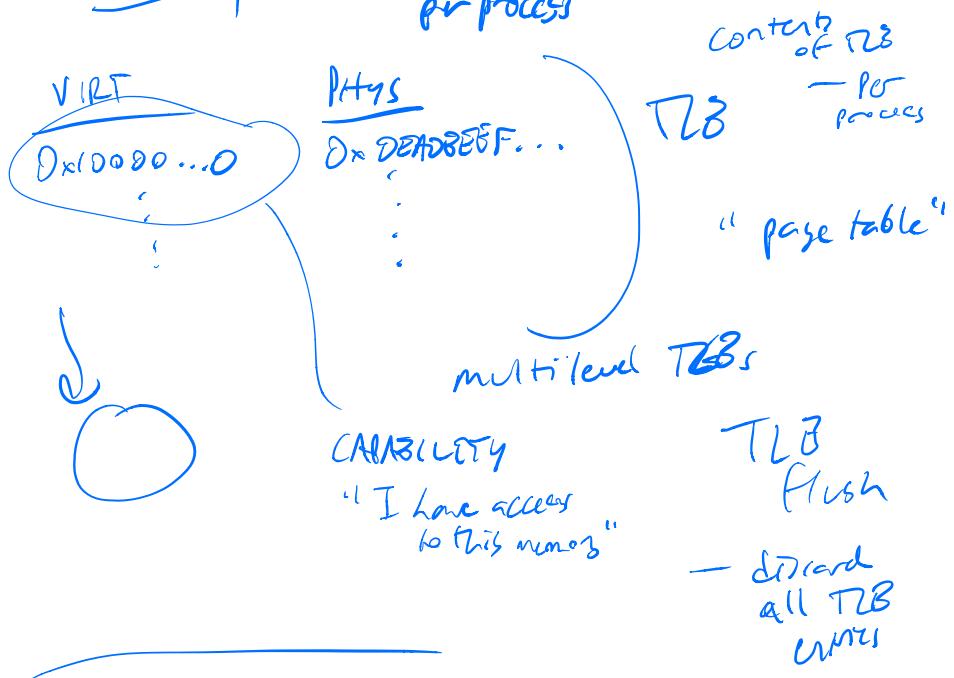
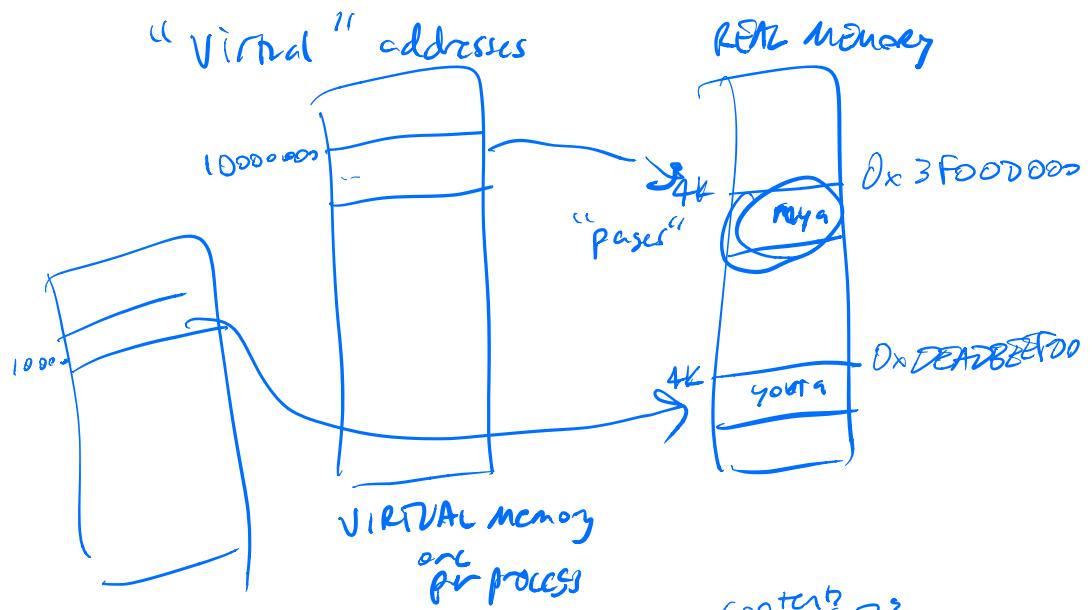
process isolation

10  
11



11  
11  
saturated  
counter







anomalous cleanup  
across the quantum



REGS  $\sqsubseteq$  VARIABLES

register allocation

"spill" - copy  
contents  
of register  
to mem "eviction"!

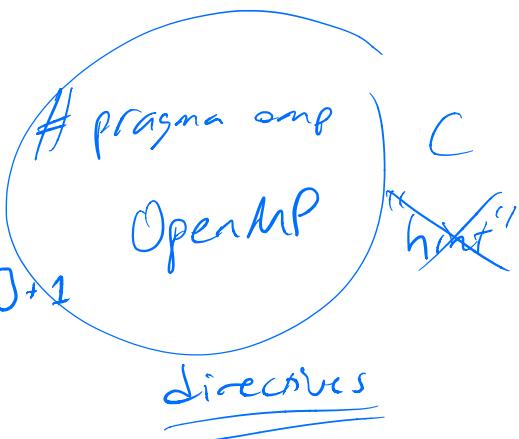
straight-line code  
- predictable!

"graph coloring" NP complete!

Why is FORTRAN fast?

~~OMP PARALLEL~~

```
DO I = 1, 100
  DO J = 1, 100
    A[I][J] = A[I][J-1] + 1
  ENDDO
ENDDO
```



automatic  
parallelization

"holy grail"

main thread  
 S  
 ↓  
 for ( $i=0; i < N; i++$ ) {  
     foo();  
 }  
 ↗  
 are there  
 loop carried  
 dependencies?

P  
 }  
 for ( $i=0; i < N; i++$ ) {  
     t[i] new thread (foo);  
 }  
 ↓  
 foo<sub>0</sub> foo<sub>1</sub> ...  
 [t[i].join()]

$\text{result}(S) \equiv \text{result}(P)$

parallel prog  
semantically equivalent  
to seq prog.

does `foo()` have side effects?

foo()  
 {  
     x++  
     (<sub>global</sub>)
 }

read  $x$  into a register  
 increment  $x$   
 write the result back

$x = 0$

A  
 R<sub>0</sub> ← x  
 R<sub>0</sub> ← R<sub>0</sub> + 1  
 x ← R<sub>0</sub>

$A; B \Rightarrow z$

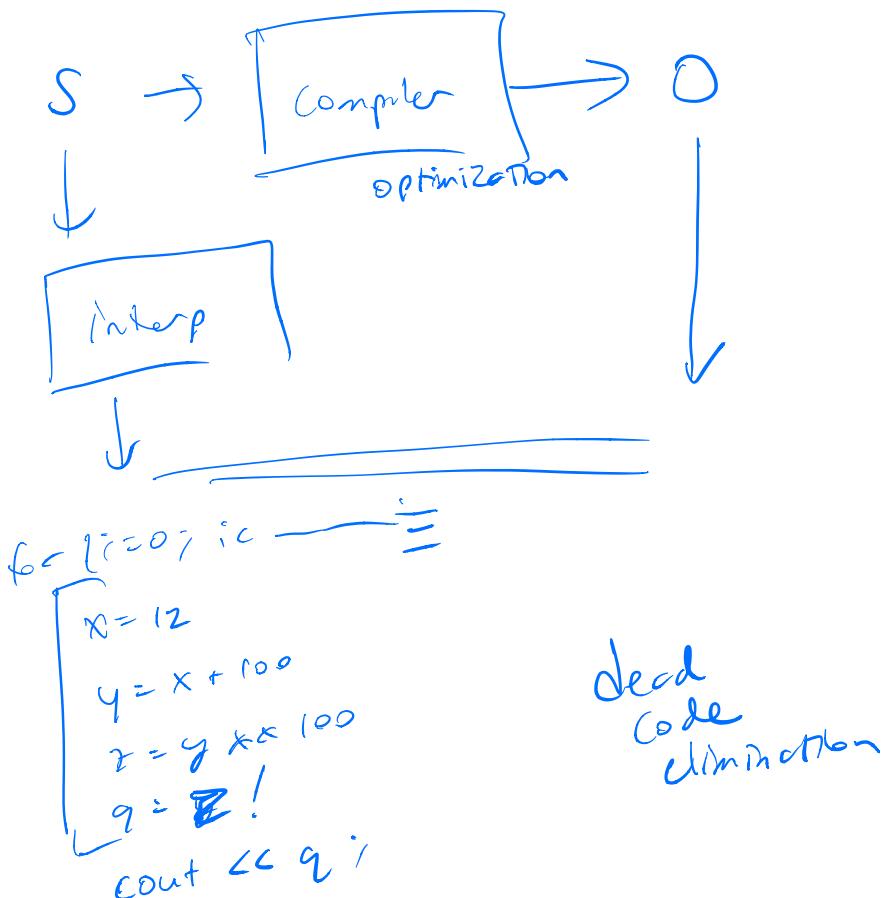
B  
 R<sub>0</sub> ← x  
 R<sub>0</sub> ← R<sub>0</sub> + 1  
 x ← R<sub>0</sub>

$$X \in \{1, 2\}$$

non-deterministic!

Race condition

Concurrency errors



```

foo()
{
    ~x ~
}

```

```

foo (int i)
{
    ~i ~
    return ack(i, i+1);
}

```

Ackermann

```
foo (int *int i a,) {  
    a[i] = ack(i, i+1);  
}
```

```
foo (int * a, int *int i b) {  
    a[i] = -;  
    b[i] = a[i] + 2;  
}
```

```
int x[200]; int *y = &x[0];  
v = foo(x, y, i)
```

### aliasing

$p \{ \quad \}$       alias analysis  
 $q \{ \quad \}$       pointer analysis

$p \cap q = \emptyset$        $\text{int } *p; \quad p \in \{ \}$

if ( $\_\_$ ) {  
  $p = q;$        $p \in \{q\}$

} else {  
  $p = r;$        $p \in \{r\}$

symbolic execution

abstract interpretation  $\rightarrow$   
 $\wp \in \{q, r\}$

$$\begin{aligned} x &= 1^2 \\ x &= x - 100 \\ x &= \dots \end{aligned}$$

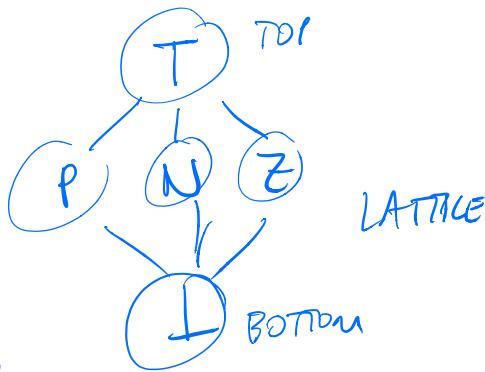
if ( $x > 0$ ) {  
 |  
 }

$x \in O, A, P$

$x = \_ \quad \{O, N, P\}$   
 $x = x - 1 \quad \{ \}$   
 does this run  
 over?  
 $O \Rightarrow O$   
 $N \Rightarrow P$   
 $P \Rightarrow N$

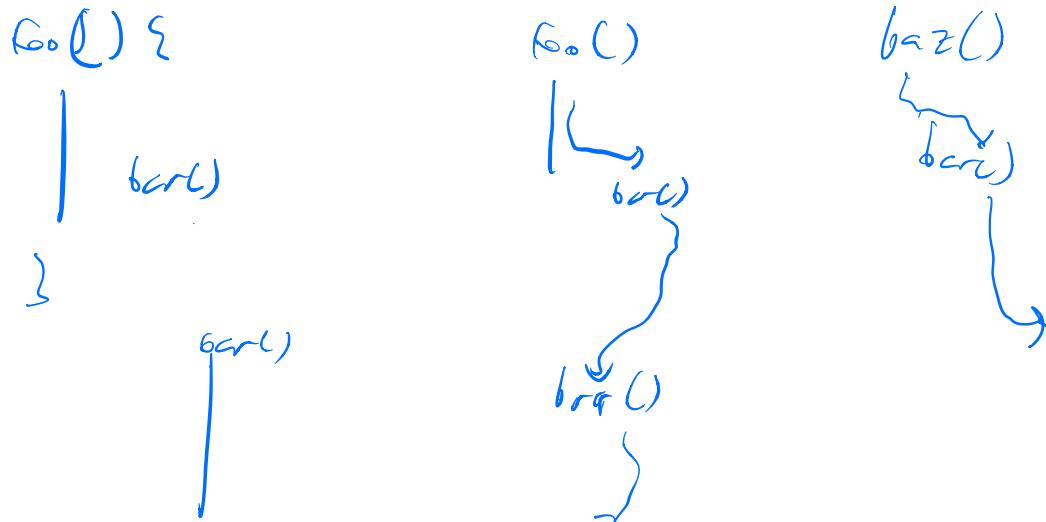
CONSERVATIVE APPROX  
 PROG BEHAVIOR

$\{x : T\}$   
 if ( $\neg$ ) {  
 $x = -1$   
 else  
 $x = 1$   
 }  
 $\{x : \perp\}$



$$\begin{array}{c} \text{if } x : P \\ \text{else } x : P \\ \hline x : P \end{array}$$

## intraprocedural vs. inter procedural



flow sensitivity



Context sensitivity

Steensgaard's algorithm

PCR sensitivity

false positives

precision - recall

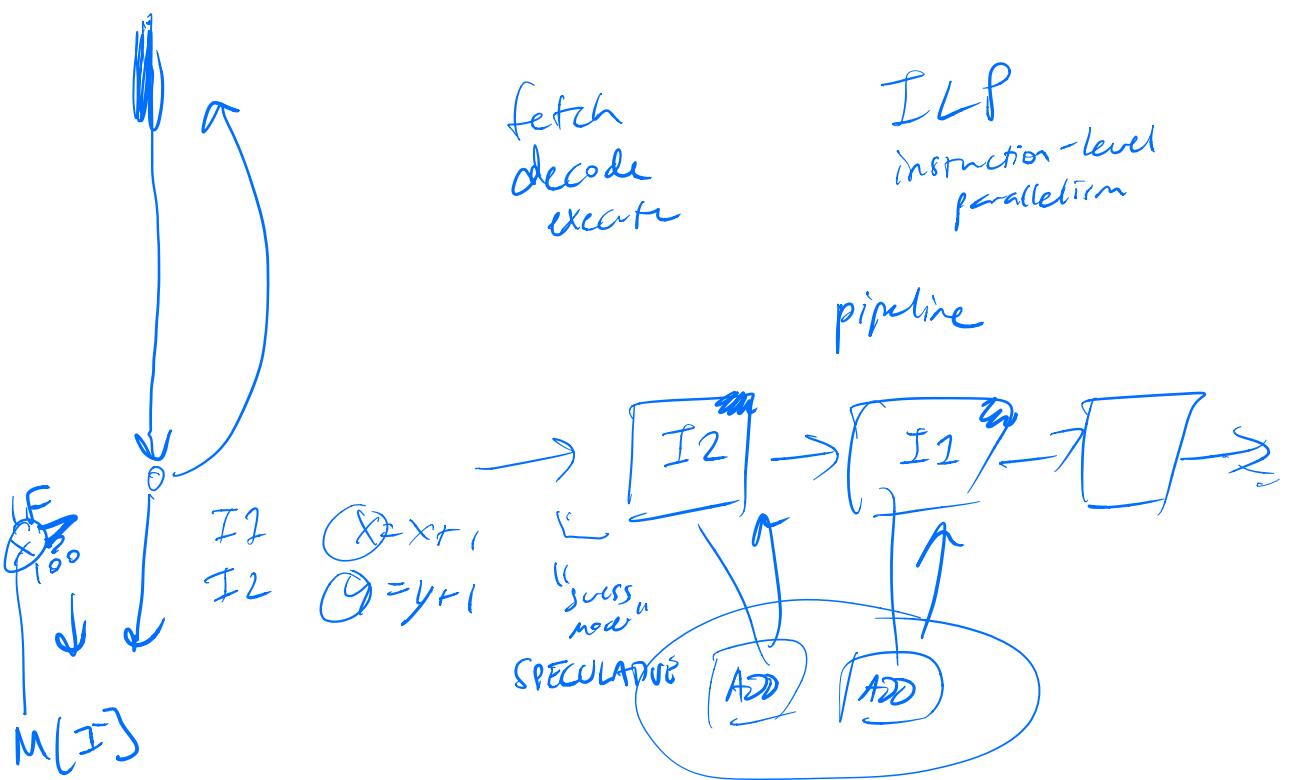
false negatives

tradeoff

$$V \sim \{\emptyset\}$$

-W

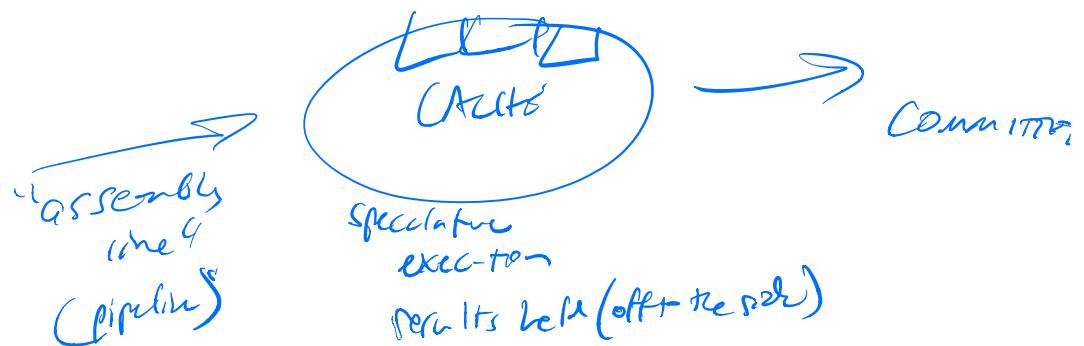
if ( $x == \underline{\hspace{2cm}}$ )  
 $y = x / @ V$



Pipeline stall

"bubble"

"BACK EDGE DATA"



Speculative execution  
hides memory latency

# SPECTRE / Meltdown

timing channel

Covert channel

side channels

last branch taken

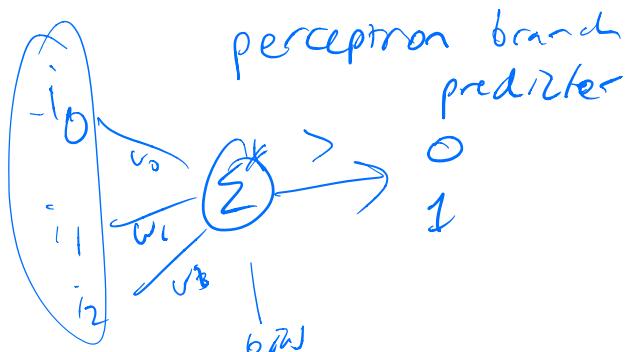
85-90% effective

PHAT

Modern

branch predictors

99.9% effective



---

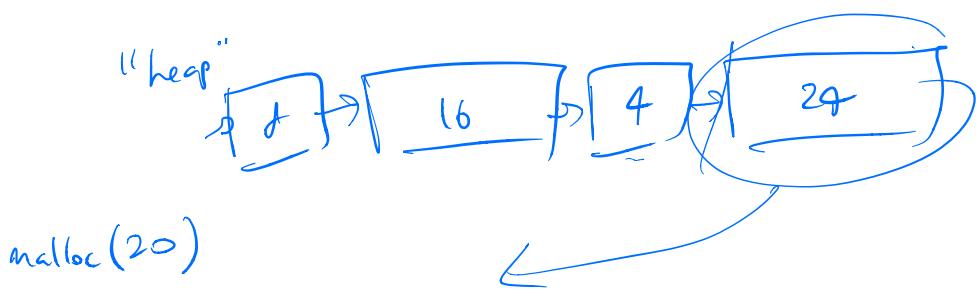
dynam. memory allocation

globally

stack variables

dynamic memory

"the  
heap"



Object 24

First-fit → SPACE INEFF

best-fit

SPACE  
EFFICIENT

slow ...

$$n \cdot O\left(\log \frac{M}{m}\right) * n$$

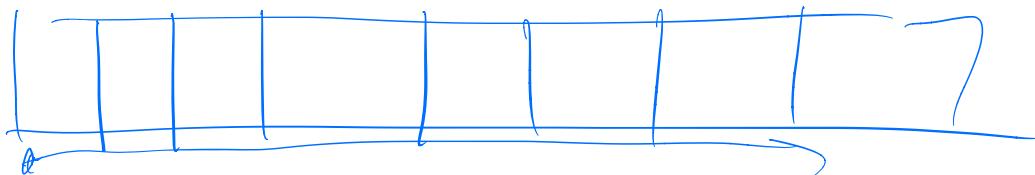
M - biggest  
m - smallest

$$O(M \times m)$$

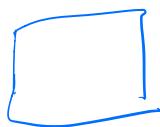
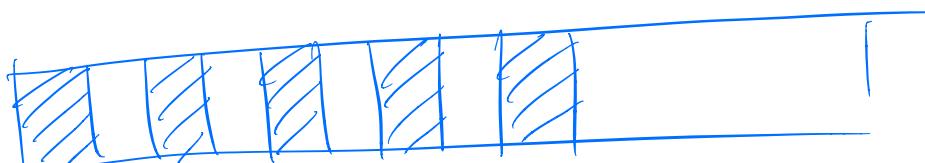
best fit

worst-case  
"fragmentation"

$$= \frac{\text{mem consumed}}{\text{mem requested}}$$



D → D → D



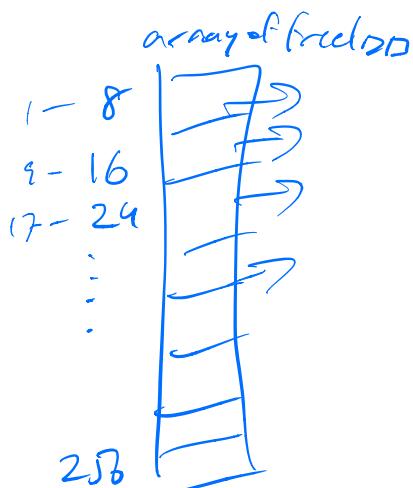
Compaction  
NOT IN C/C++

(L1E')

## Python, Ruby

"custom allocators"  
"ad hoc allocators"

malloc - "too slow"  
Reconsidering Custom Mem. Allocation



Py-object  
- malloc

if ( $sz \leq 28$ ) {

use small obj.

} else

use malloc

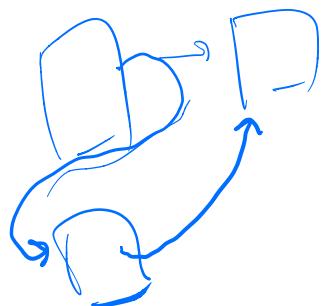
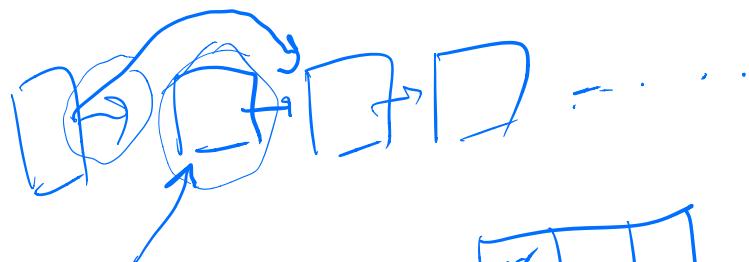
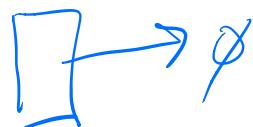
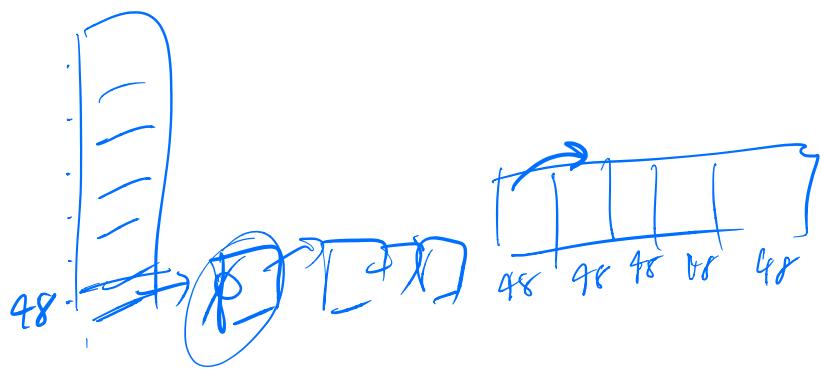
int  $\approx 32$  bits      C      Python  
long  $\approx 64$  bits      48      24/28

"a"        
1/16      50

{ "a": }  
~~2~~

240

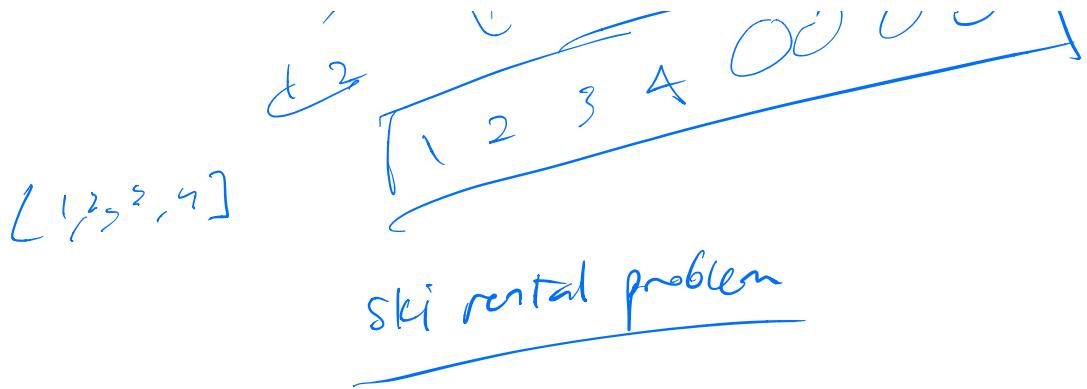
60 20 75 100 140



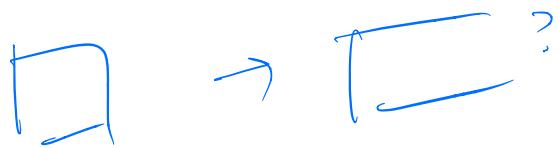
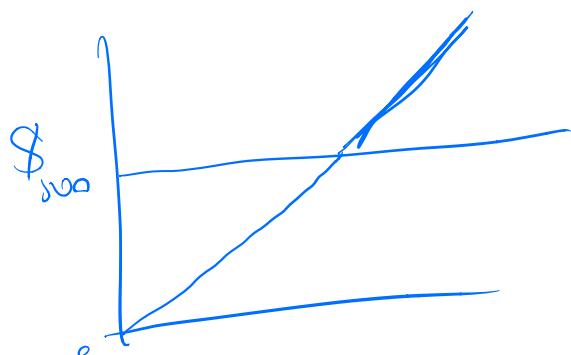
48			
ref count			

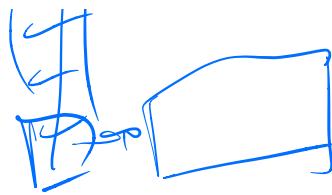
$$x = [1, 2, 3]$$





\$50 rent  
\$500 buy  
rent until  
you reach  
cost of buying

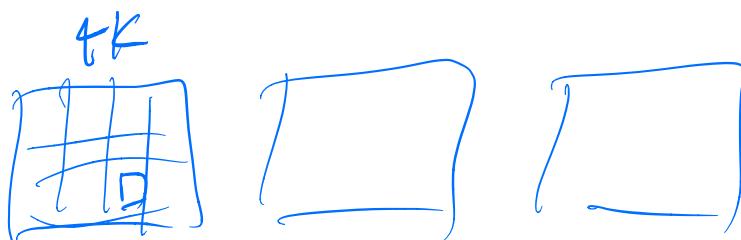
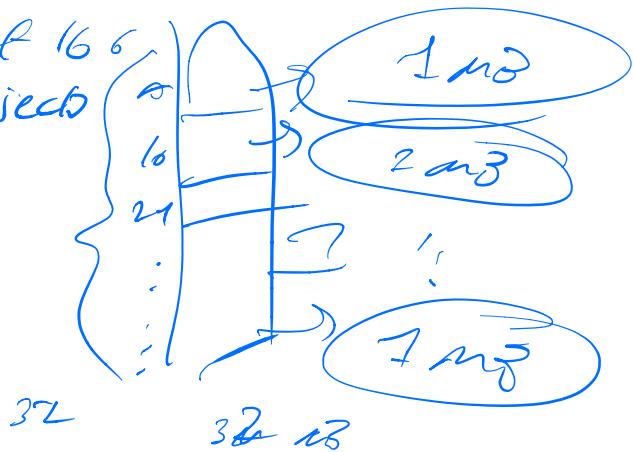




allocate 1 MB of 8 byte objects  
free them all

allocate 1 MB of 16 byte objects  
free them all

actual  
MMT 1 MB  
32x



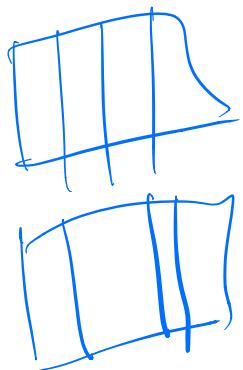
OS passes

Space-time tradeoff

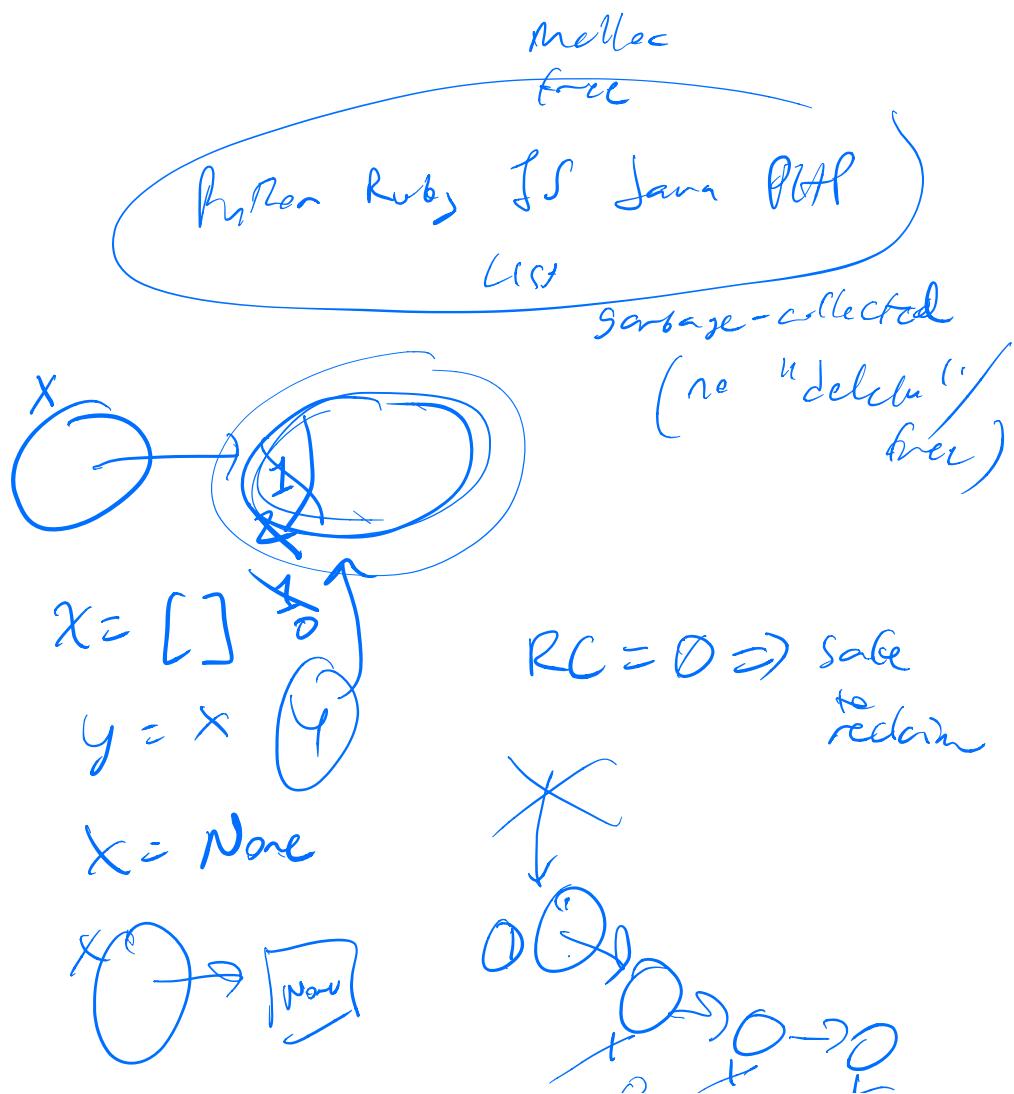
$$f(a) =$$

$$\begin{aligned} f(0) &= a \\ f(1) &= b \\ f(2) &= c \end{aligned}$$

0	a
1	b
2	c

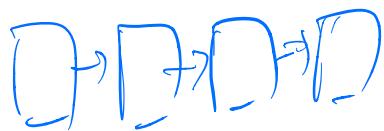


## Reference counting



$y = \text{None}$

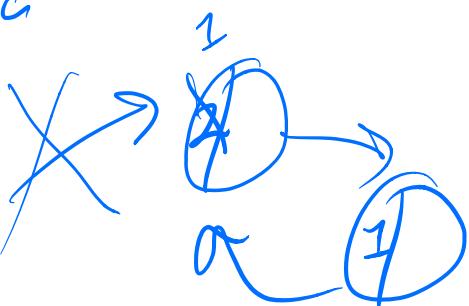
$\leftarrow \circ \rightarrow$



doubly-linked list

$a = \{\text{'key'}: 0\}$  CYCLE

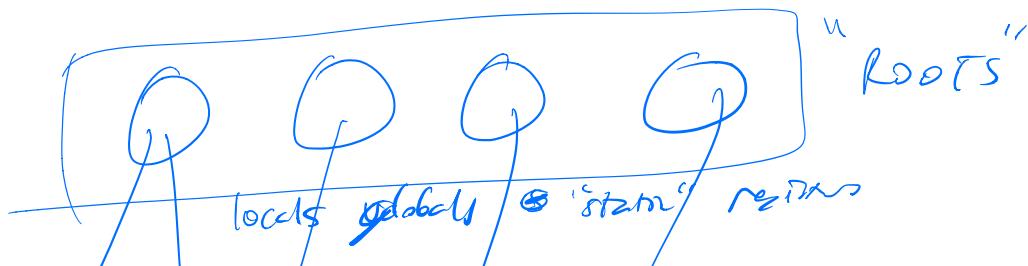
$a[\text{'key'}] = a$

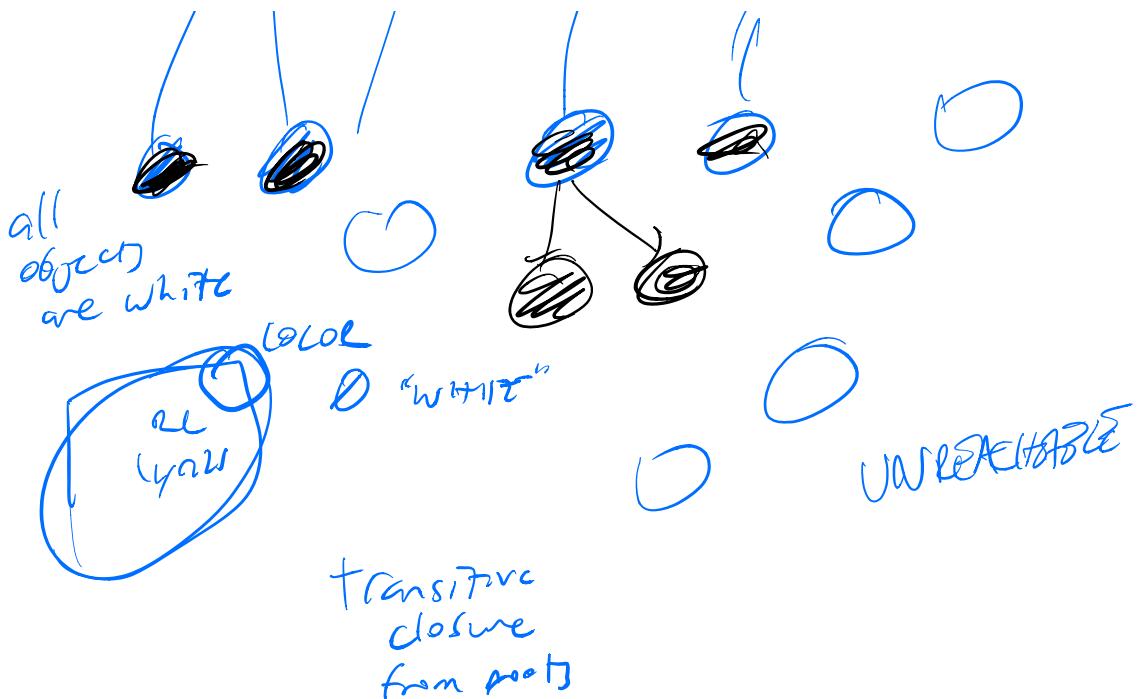


HC  
(1956) Incomplete  
GC algorithm

---

Mark-sweep (1978)  
"generational" GC (1983?)





David Unger

smalltalk self

Jg Holzle  
Gerd Hanberg

Ole Ascan postbyr eratid

dynamiz

Lor  
gak

## GL algorithms

JIT compilation

## dynamic languages

polymoriz

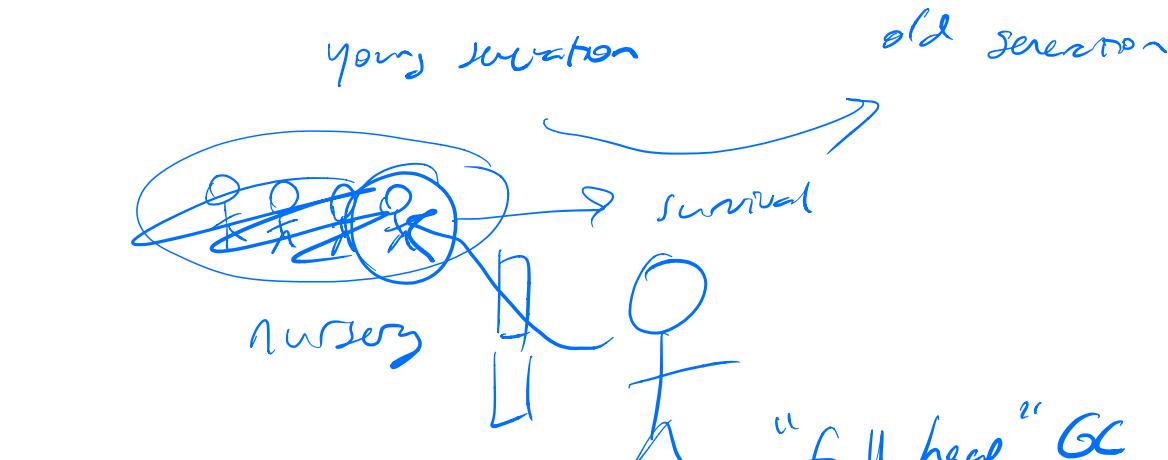
main  
cache

"Pic"

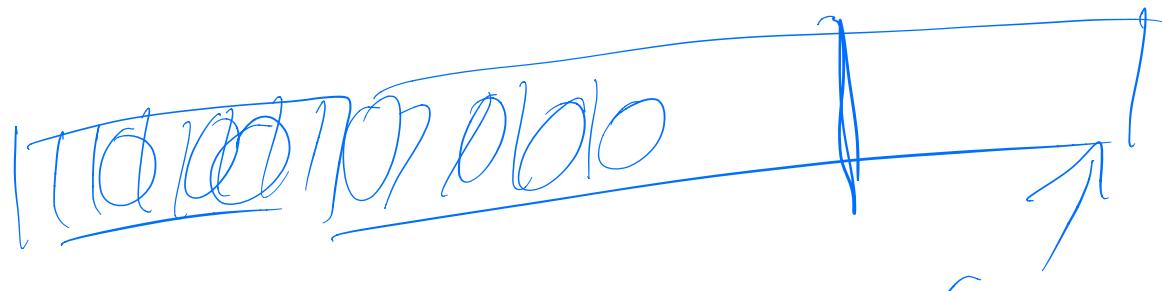
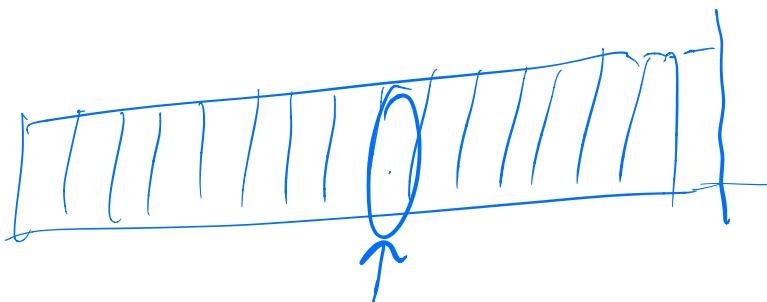
Google

→ U8

most objects die young



gen GC →  
takes pressure  
off the full  
heap GC

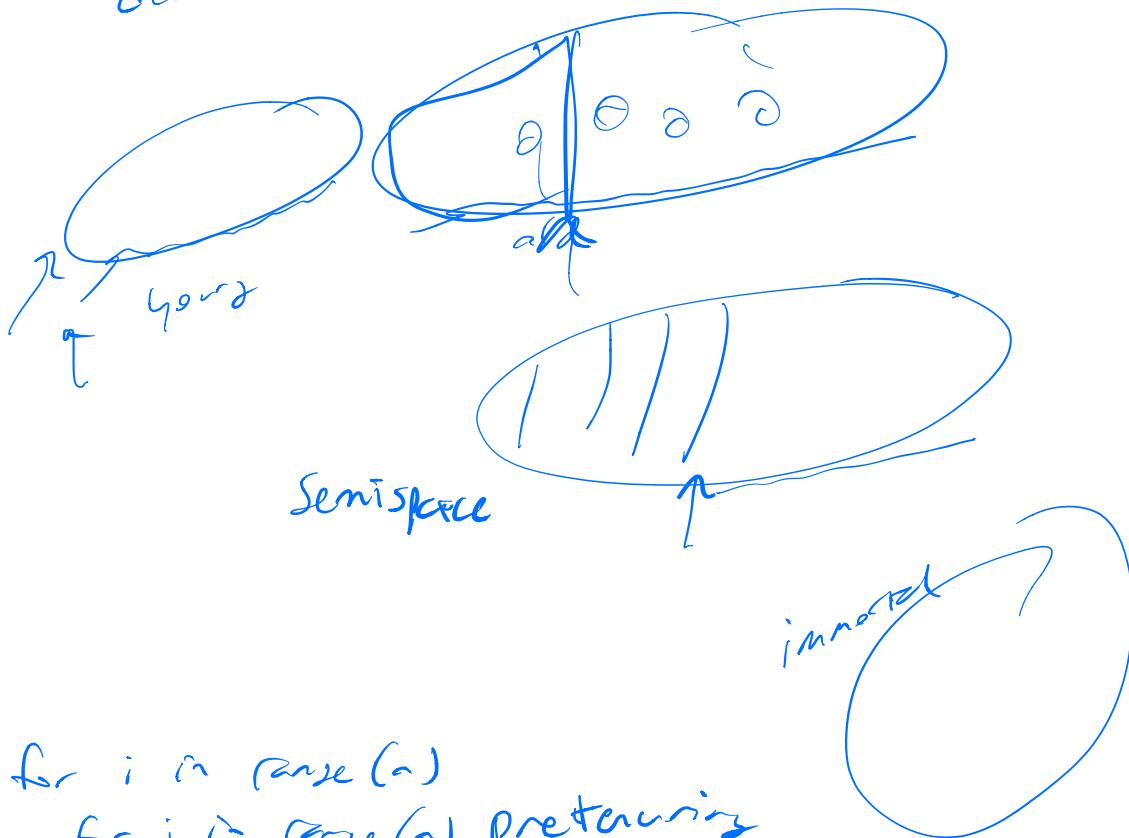


GC

$\sim 2 \rightarrow 5x$

"live  
objects"  
"dead"

Quantizing the cost of  
GC vs. Explicit deallocation



for i in range(a)

    for j in range(n) pretenuring

        for k in range(n)

$$C[i][j] += \delta A[i][k] + \delta[k][j]$$

• pyc

PyPy

# Ontogeny recapitulates phylogeny

SCRIBERS

Lillian  
& Sam

OS: single user, single process, no protection

multiple users, multiple processes

"time-sharing"

time slices

memory protection

file system protection

:

Mac 1984

1 user

1 process

no protection

VM

"MMU"

memory  
management  
unit

PC

Apple  
IBM PC

:= 1981

Motorola 6802  
Intel 8080...  
68000 Macs

) no MMU

MS-DOS

Microsoft

Microsoft BASIC

Altair

Gates & Allen

CP/M

8080

Z-80

8088

Gary Kildall

1972

"security  
through  
obscenity"

AT&T

Unix

BSD

OS X

iPhone

stripped down Unix!

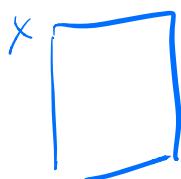
→ multiple processes  
mem protection  
security model!

PL?

→ dynamically typed lang — Values have types  
(statically typed lang) — Variables have types

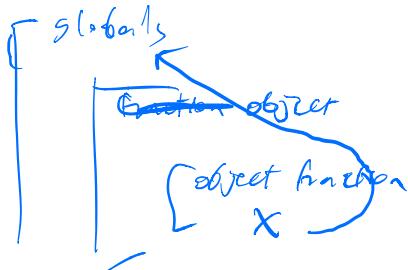
$x = []$   
 $x = {}$   
 $x = ""$

→ dynamically scoped



static / lexically scoped

"weird  
scope  
mics"



Match the innermost scope

foo() {

~x~

} x=x+1

walking call stack

bar() {

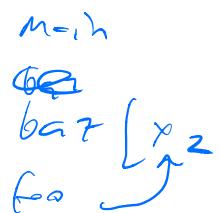
int x;  
foo();

}

bar() {

int x;  
foo();

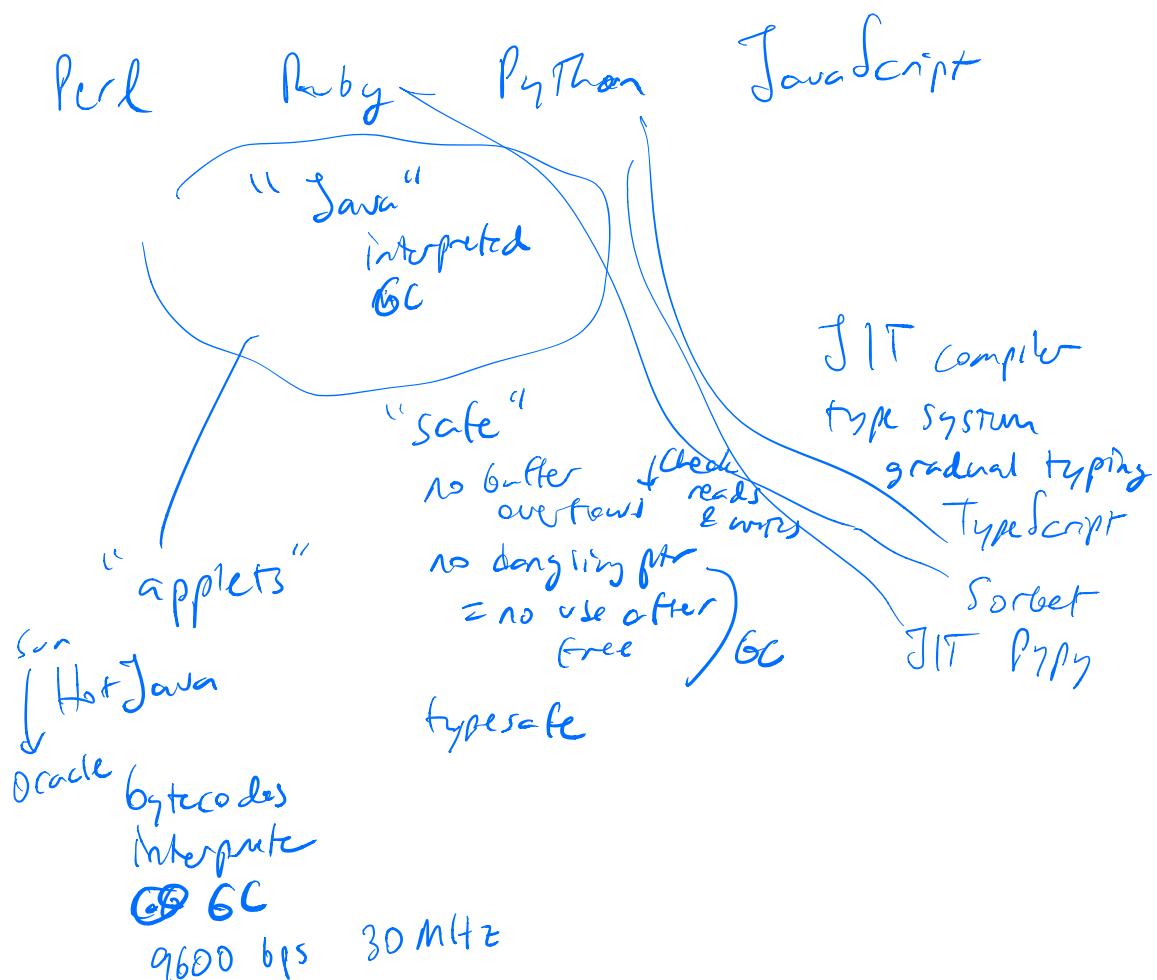
}



→ Interpreter eval

hard to compile  
into efficient code

— reference country GC



Netscape → Mozilla  
↓ Navigator

1995      Brendan Eich

Mozilla      {      }  
"curly-brace languages"

LISP Scheme  
( ( ( ) ) )

LiveScript      10 days

Browsix

Doppio

asm.js

subset that is easy to compile

asm.js  
 $x = (v \ 1 \ \emptyset);$   
 $x = a + b;$

Web Assembly

- safe
- efficient
- fast to parse!!



-- loading form compile connection

- ISA Instruction set architecture
- arbitrary memory (module mem protection)
  - arbitrary structure jmp GOTD
  - hard to analyze

Register-based  
stack-based  
architecture

A8  
CD

NaCl  
p NaCl

Compiler  
Smart  
"safe"  
subset  
of x86  
ARM

||  
PUSH  
POP  
ADD

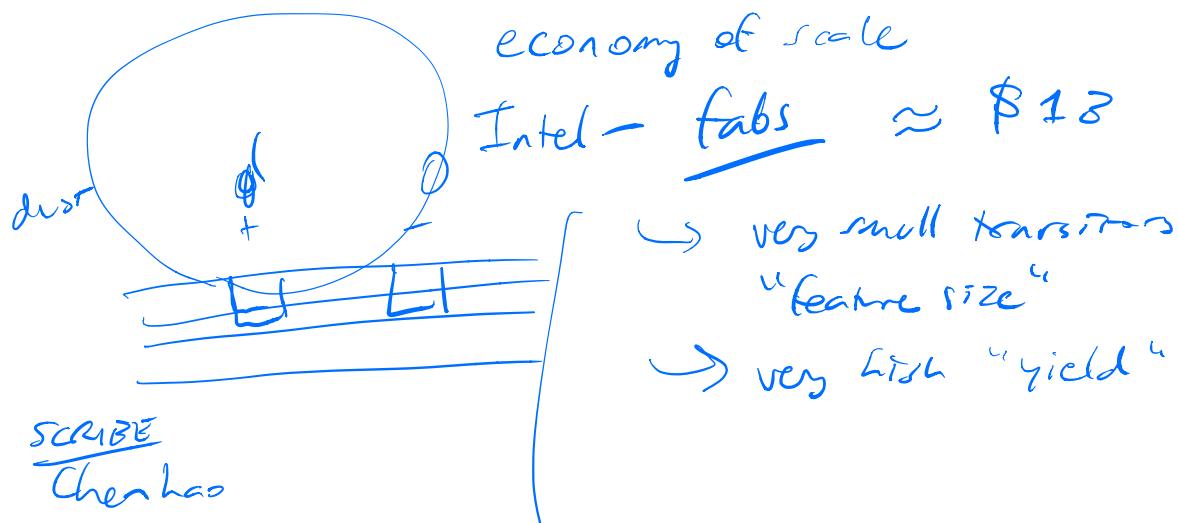
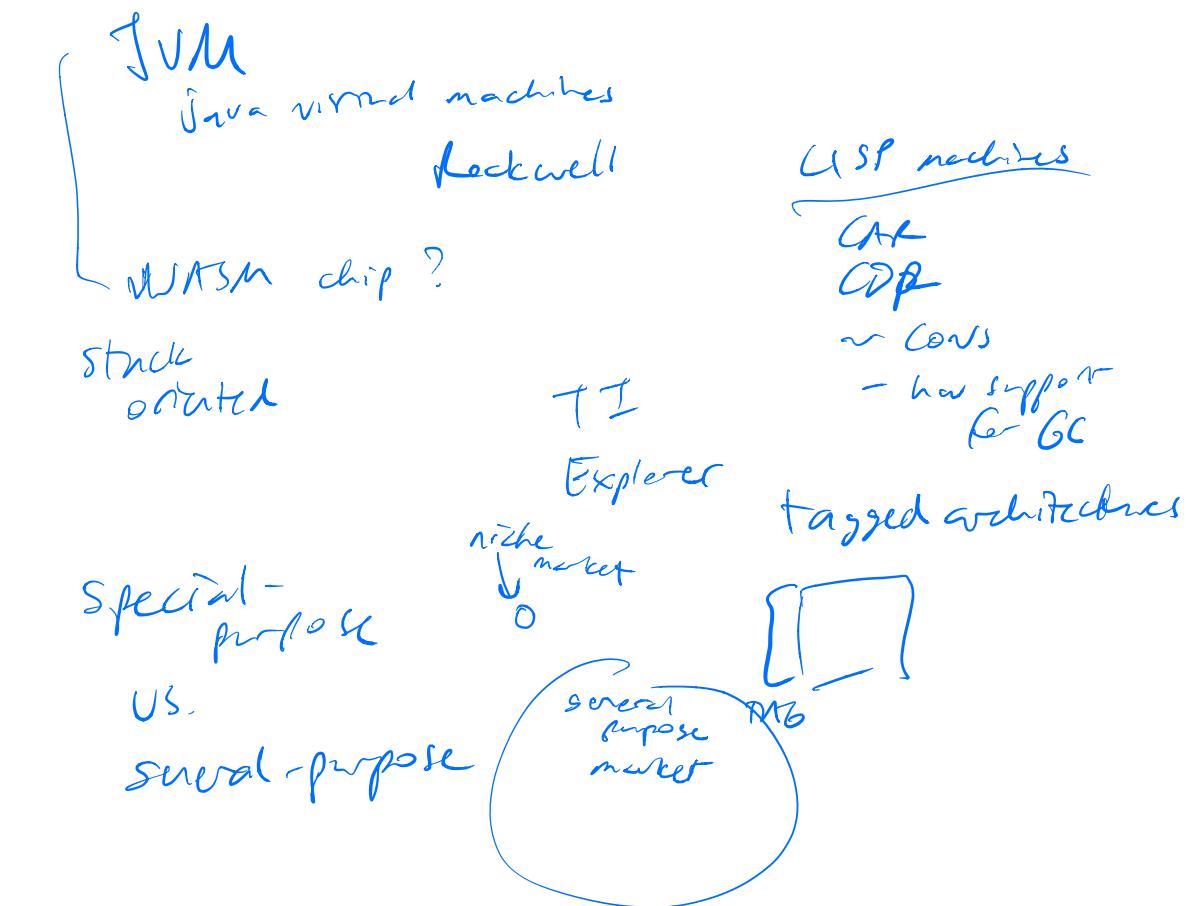
ROP  
Return-oriented  
programming  
Return-to-libc



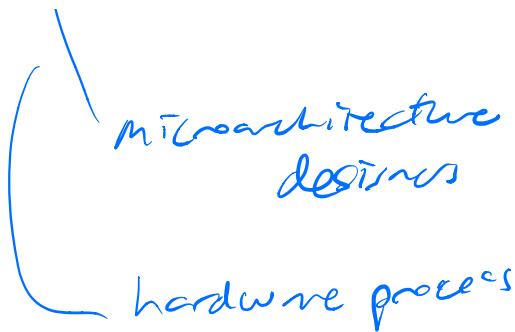
CFI  
Control-  
Flow  
Integrity

# WASM - CISC or RISC?

simple small instruction set



Shahrya

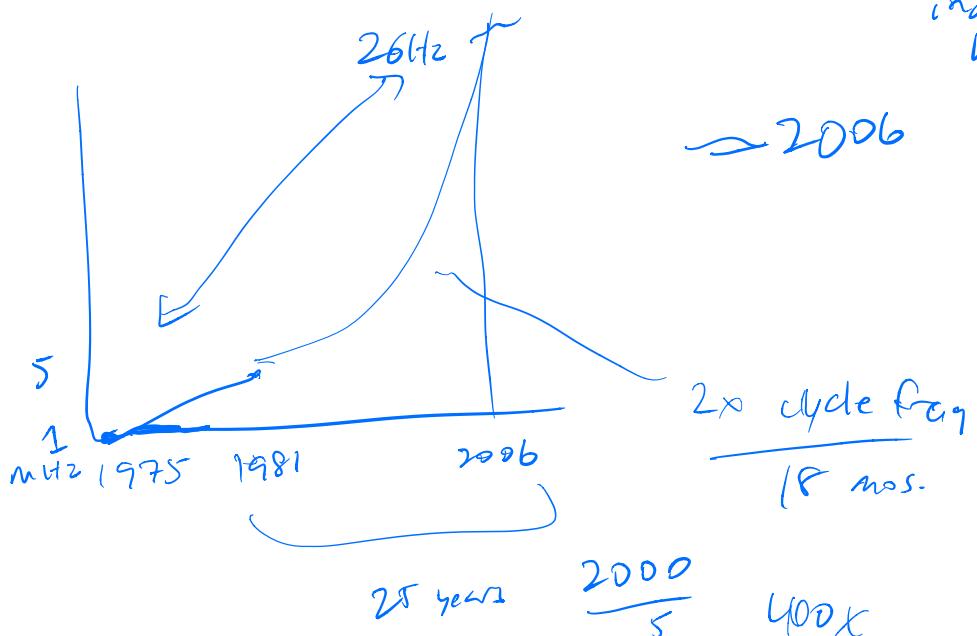


AMD

- fabs - not as good
- ⇒ low yield
- lower profits

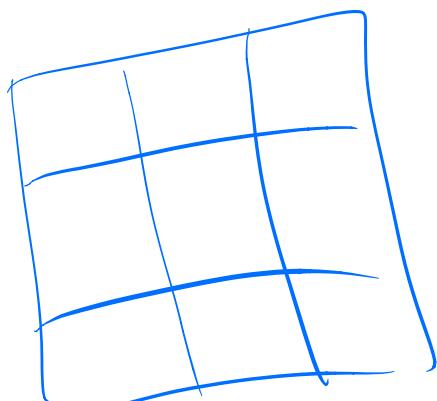
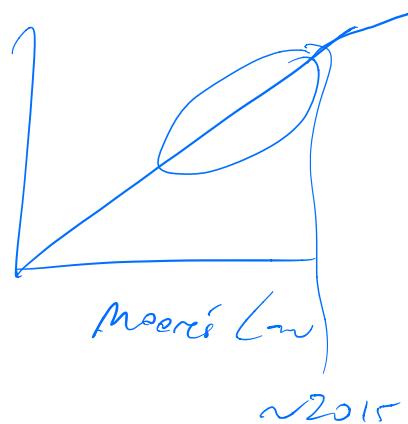
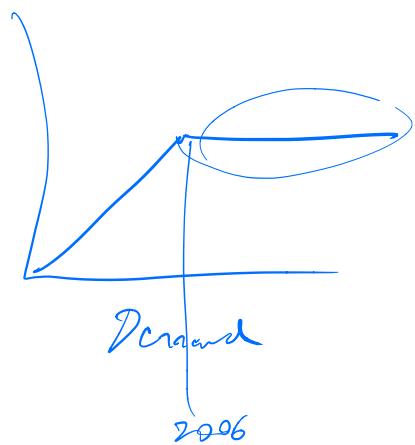
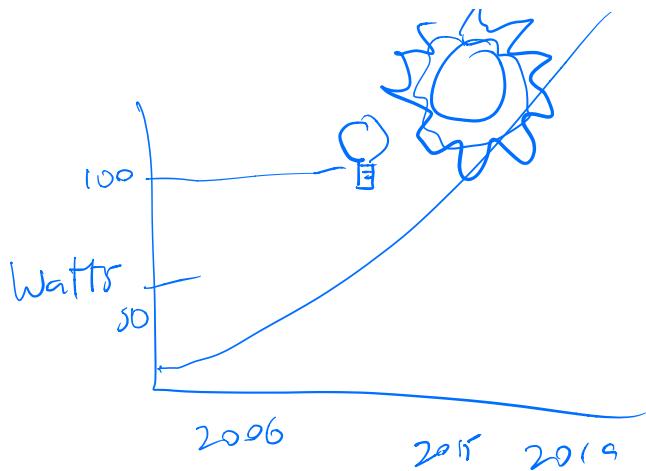
→ 32-bit to 64-bit

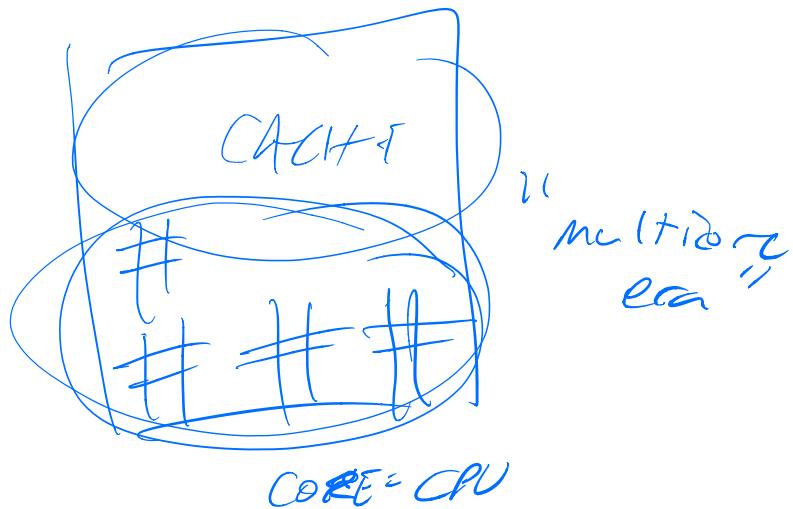
≈ multizone



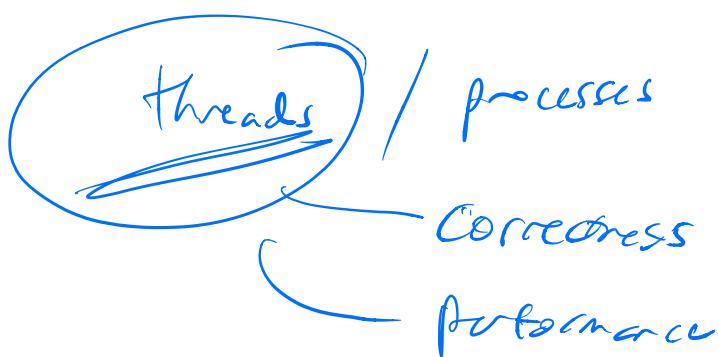
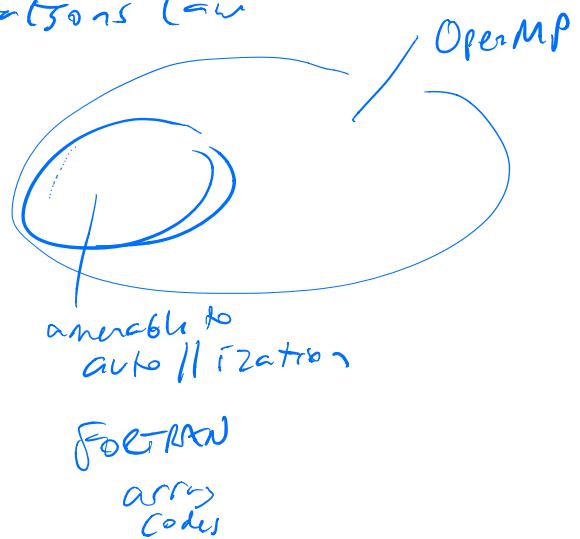
Demand scaling

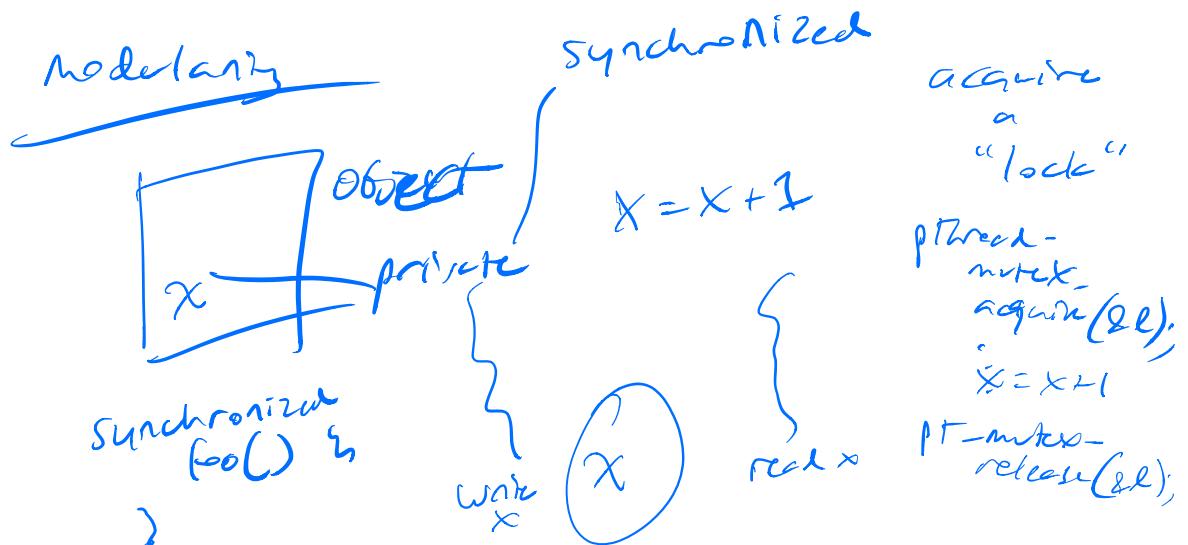
increase clock speed  
w/o increasing voltage  
→ w/o  
increasing  
heat.





Amdahl's Law  
Gustavson's Law





"atomicity"

transactions: all or nothing

atoms - indivisible

atomic

Demostrates

transaction

- commit
- abort
- rollsback

"conflict"

retry

spin lock

lock  
while ( $l \neq 1$ ) // while locked  
  ;              wait — sleep

(atomic) if ( $l == 0$ ) {  $l = 1;$  } // acquire lock atomically

TST  $\wedge$  Lock —

~~~

(unlock)  $l=0;$

"thundering herd"

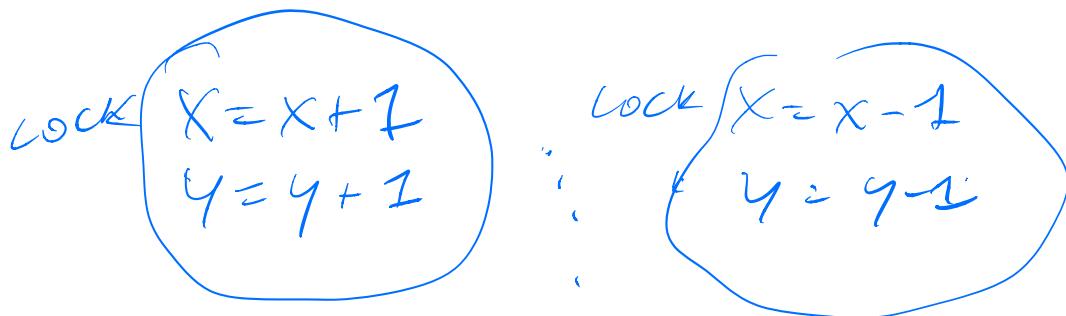
fair locks  
queuing locks

---

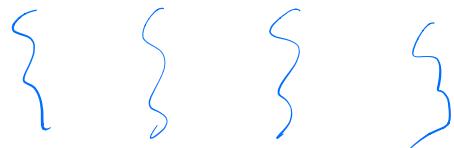
adding locks everywhere

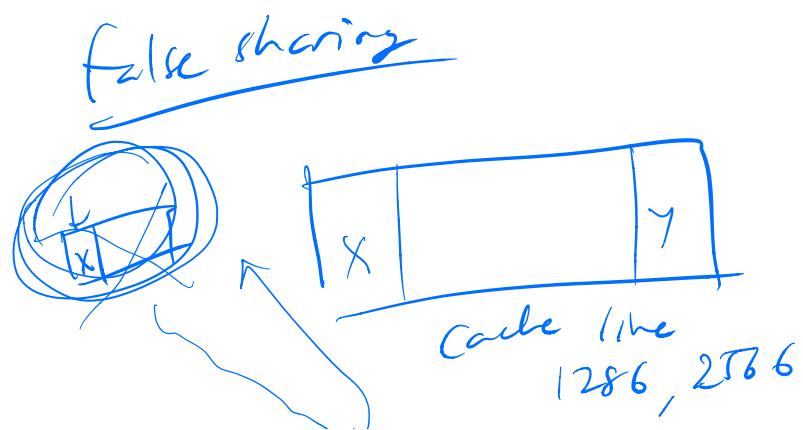
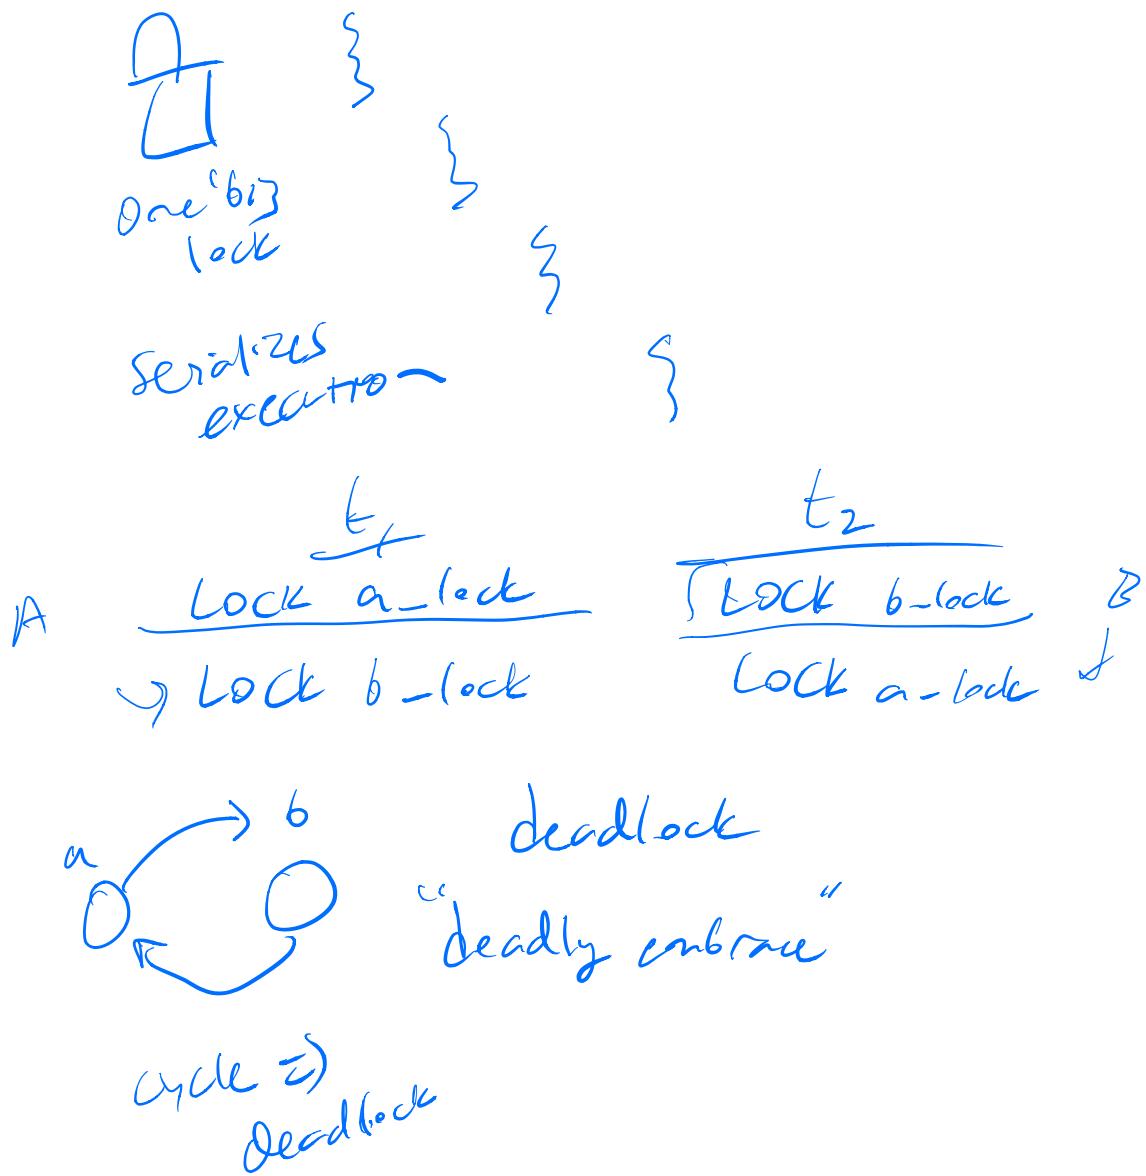
neither necessary (for correctness)  
nor sufficient (for performance)

$$x=0, y=0 \quad \left\{ \begin{array}{l} x=y \\ \end{array} \right.$$



fine-grained locks  $\rightarrow$  coarser  
for atomicity



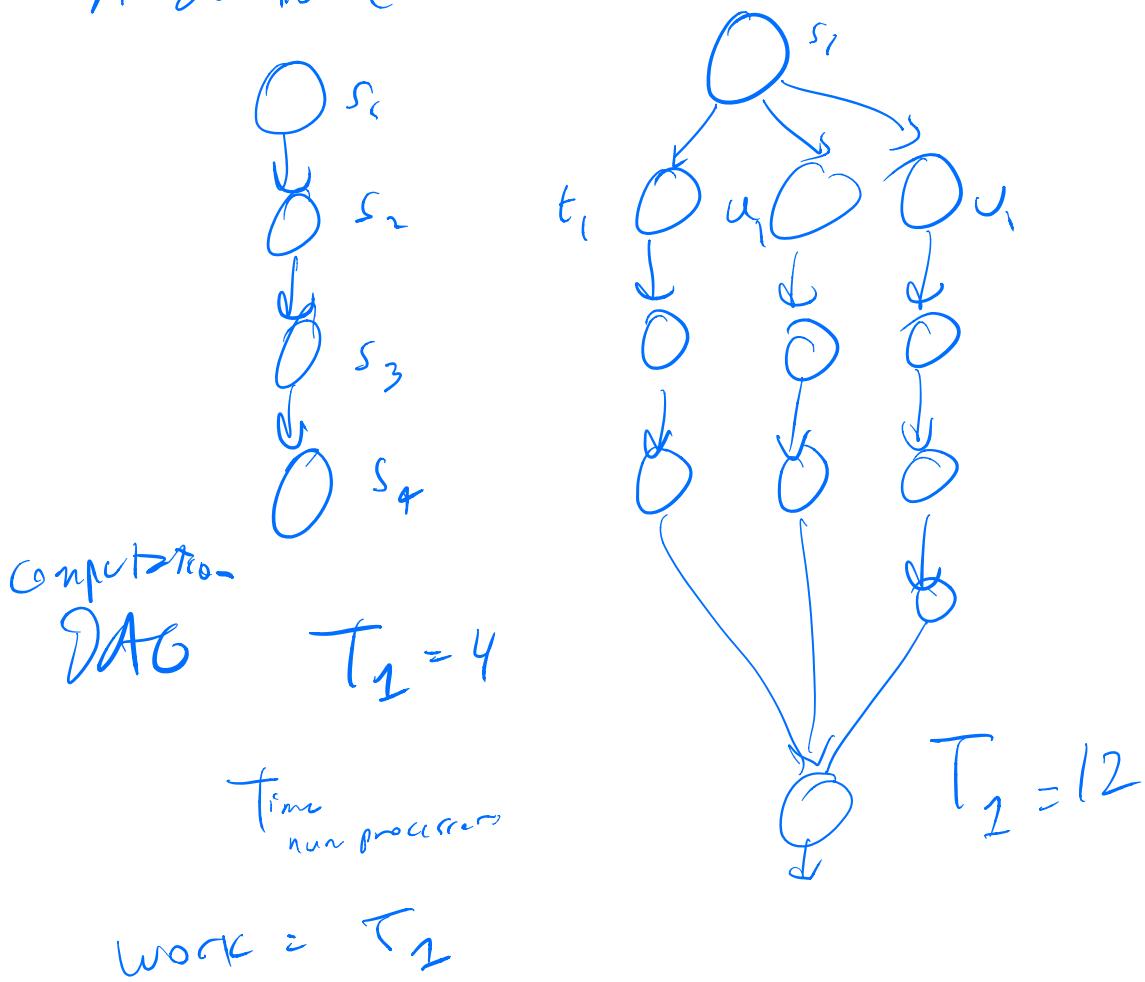


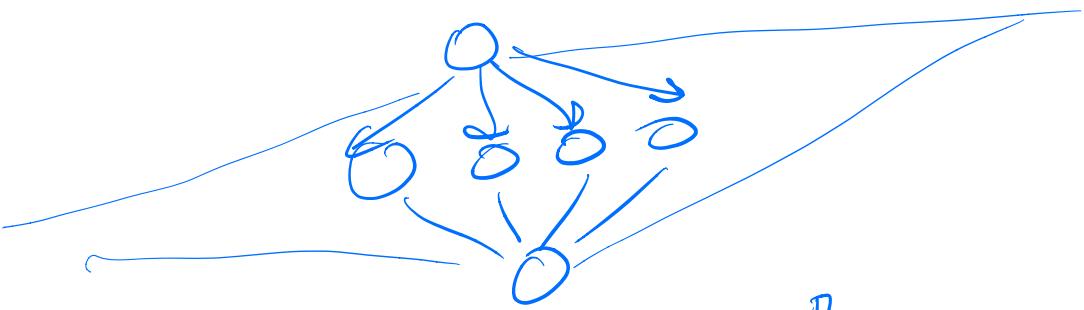
~

Cache line  
pins parsing

$\approx f\text{-}O_x$  slowdowns

Amdahl's Law





$T_{oo}$  = control path  
 max from start to leaf from  
 "header of Data"

"serial part"

$$T_p \approx T_1 + T_{oo}$$

$T_1$

parallel work

P

serial part

parallel work

90% of program — perfectly parallelizable

embarrassingly parallel

$$T_{pp} = \frac{90\%}{PP} + 10\%$$

0

100s → 50s

max speedup  $\rightarrow$  DX

VM hosting  $\approx$  embarrassingly II  
scales w/ # of instances

Amdahl's Law  
→ 1 workload fixed size

today's world  
Netflix → indep. customer moves  
emb.-II  
 $\rightarrow$  # customers

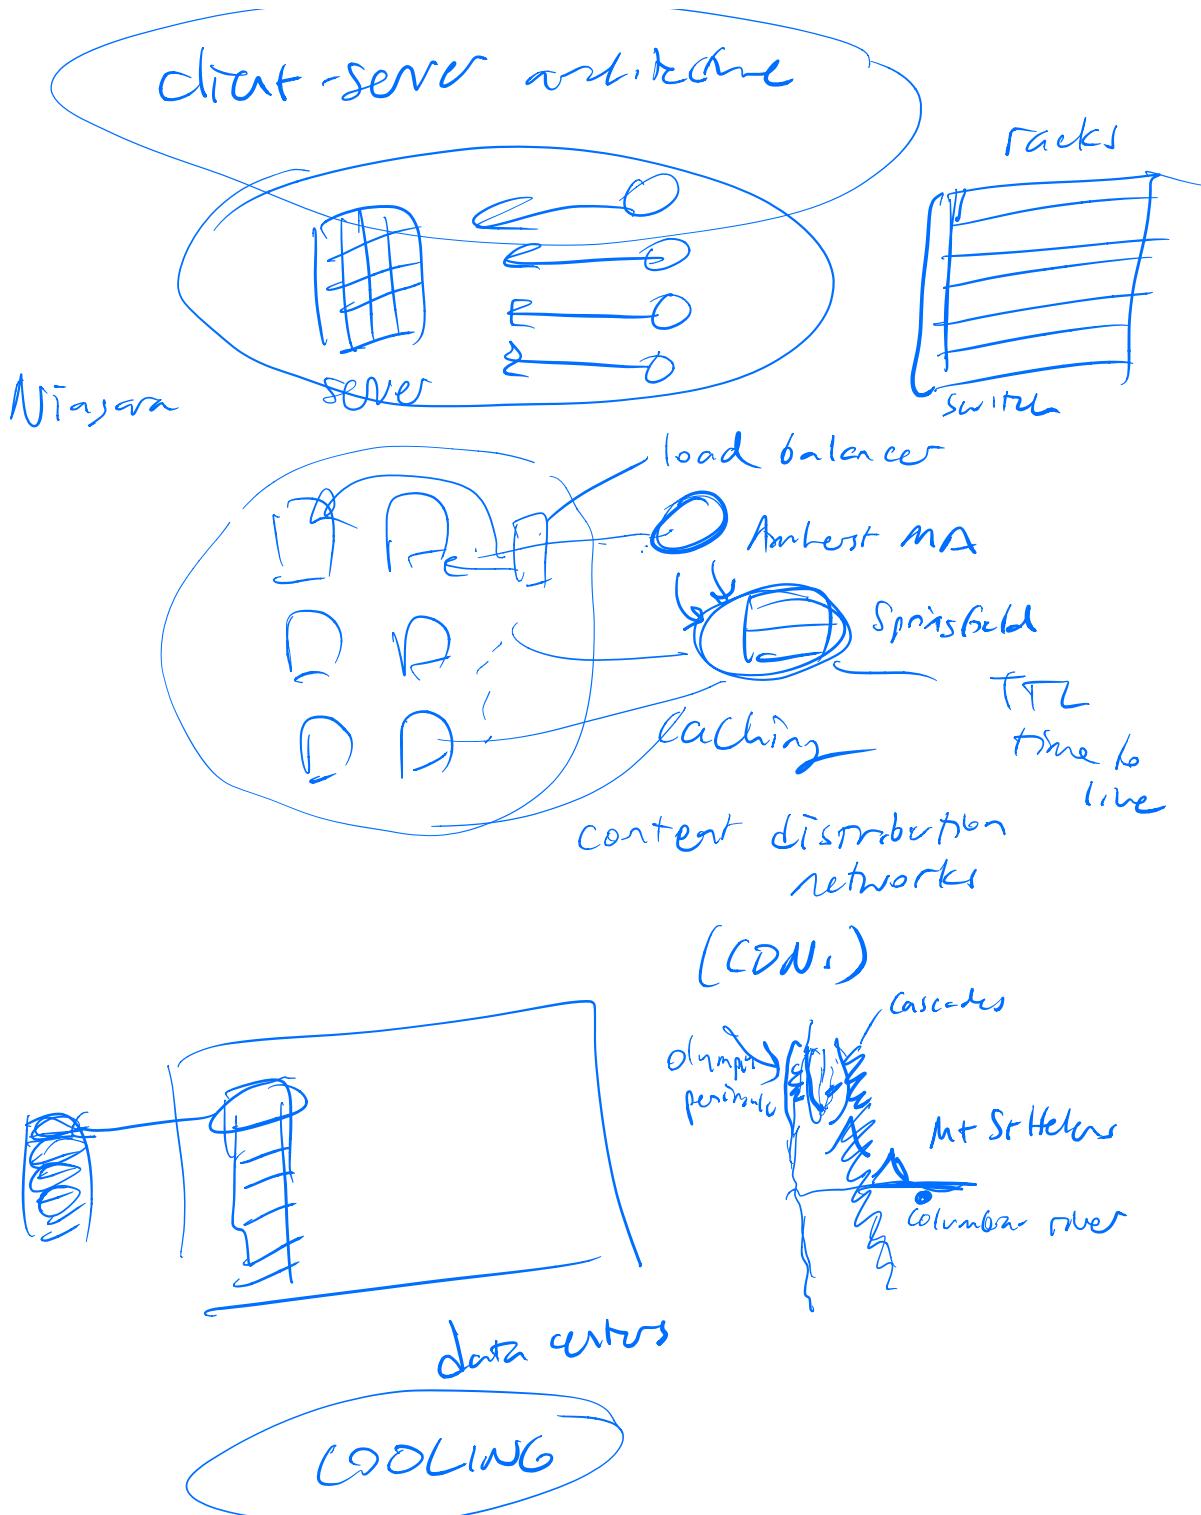
Google  
FB

Gustavson's Law

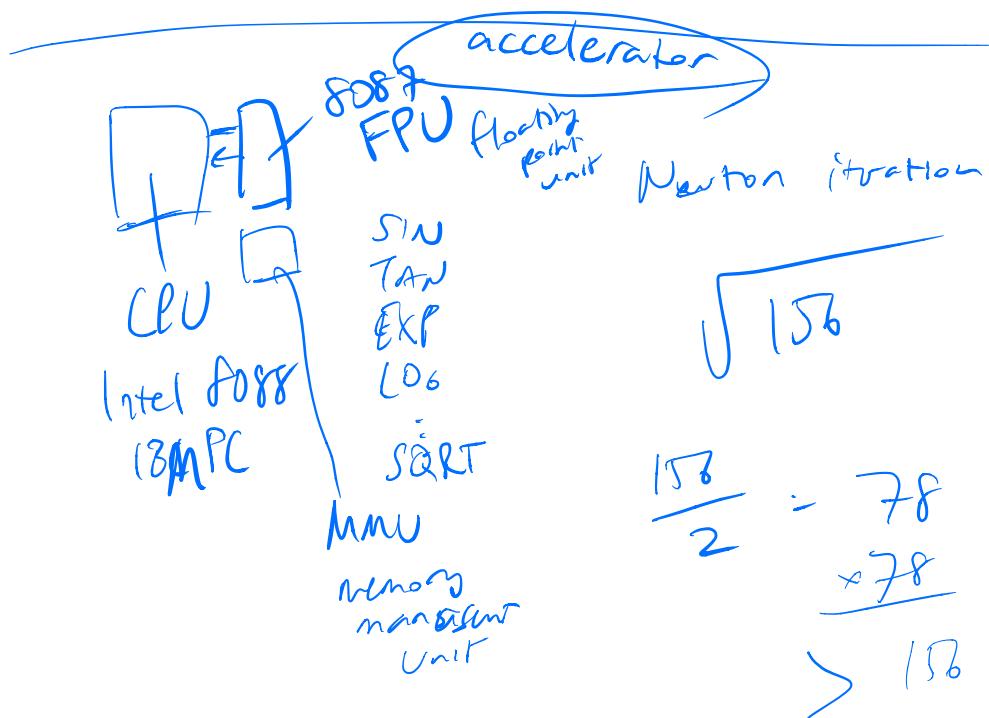
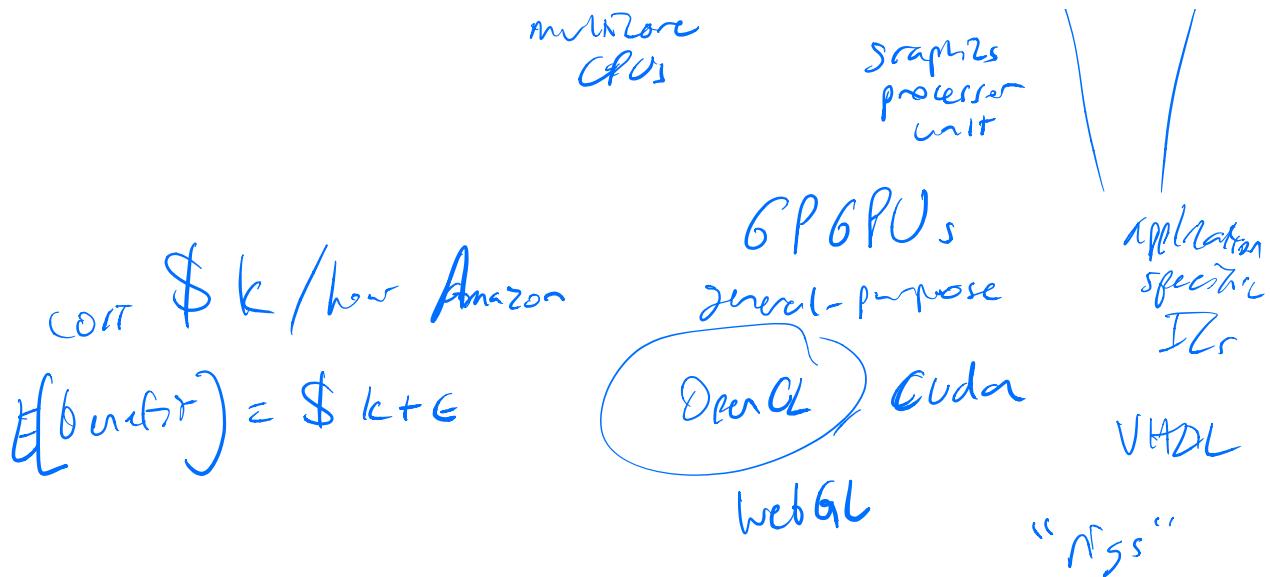
$$\Delta S \gg \Delta P$$

↑  
size of problem      ↑  
# of processors

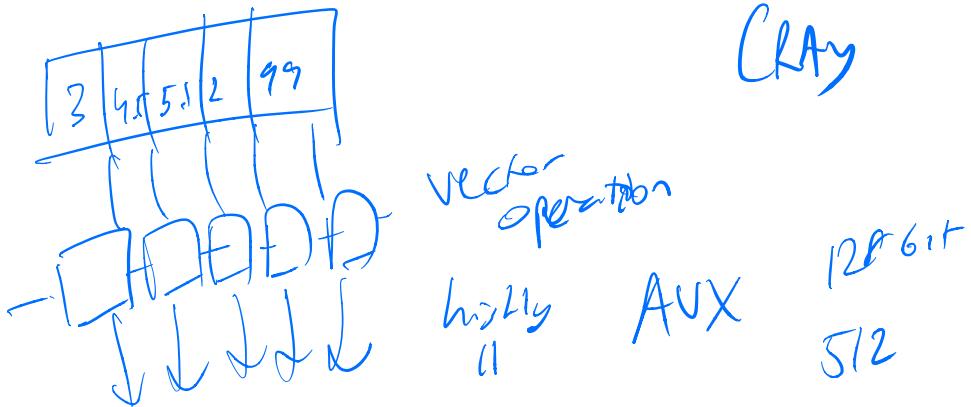
---



CPUs → cluster → GPUs → ASICs



Accelerators = CRYPTO  
Vectors



"Neural" ← HYPE

s/neural/matrix multiply/g

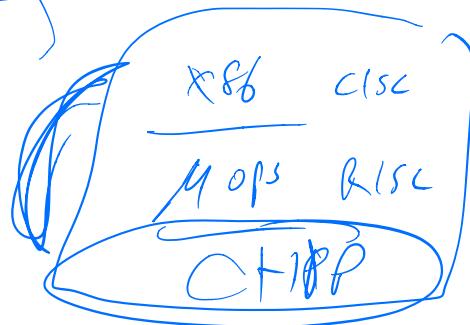
"tensor"

NPU  
TPU → Google  
tensor processing unit

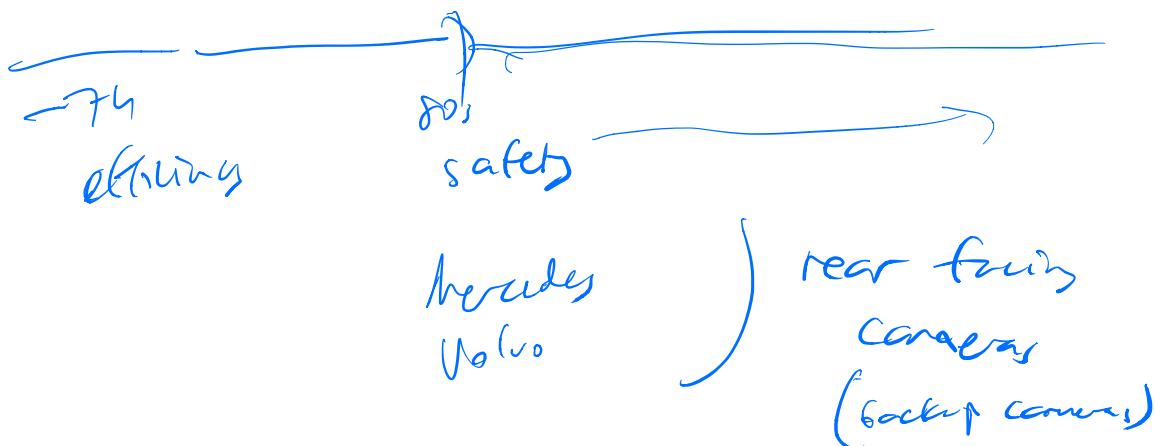
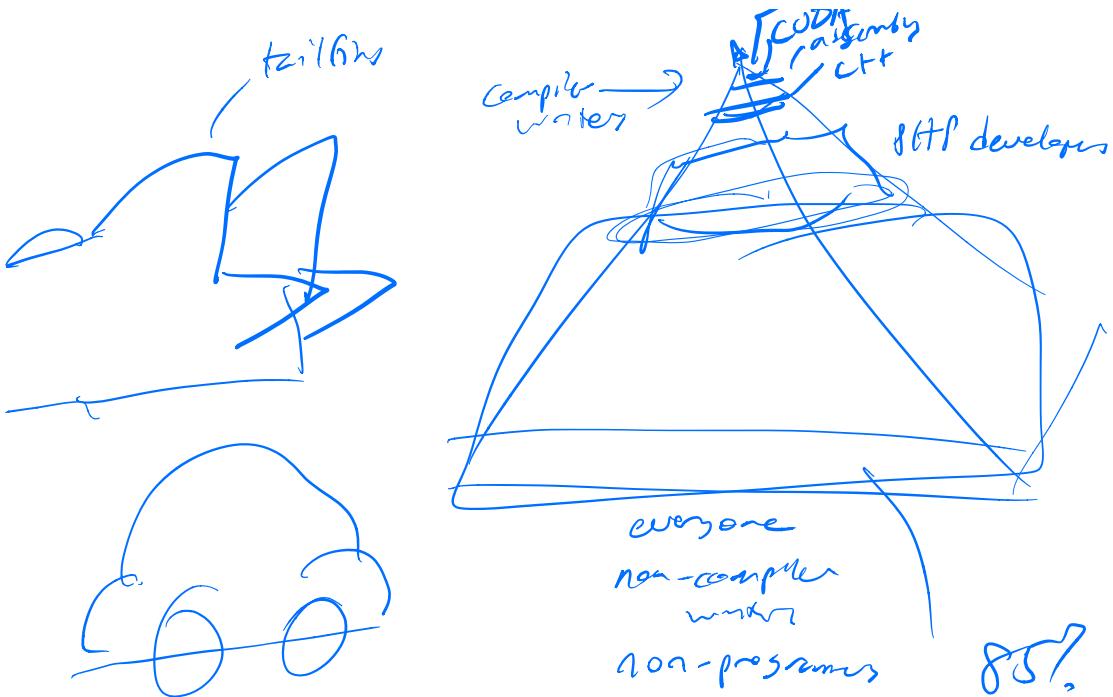
TVM

RISC vs. CISC

every Android Device  
~ ARM



fast  
concurrent access



17T H4  $\Rightarrow$  3GP H4 - 3 45 8GP

Demand Scaling

phones

2 core  
4 core  
8 core

cameras  
potassium ≈  
energy / battery life

standard size instructions

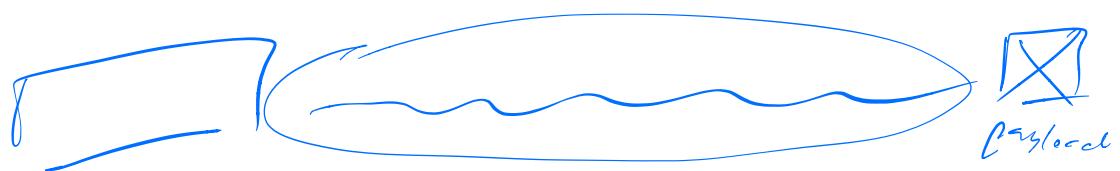
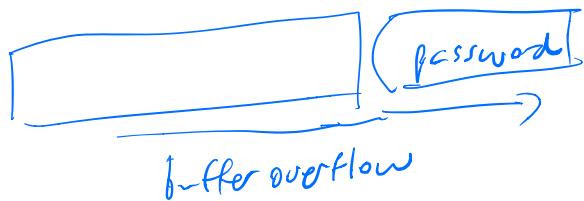
NOP → lock ROP stack — — —  
1 byte



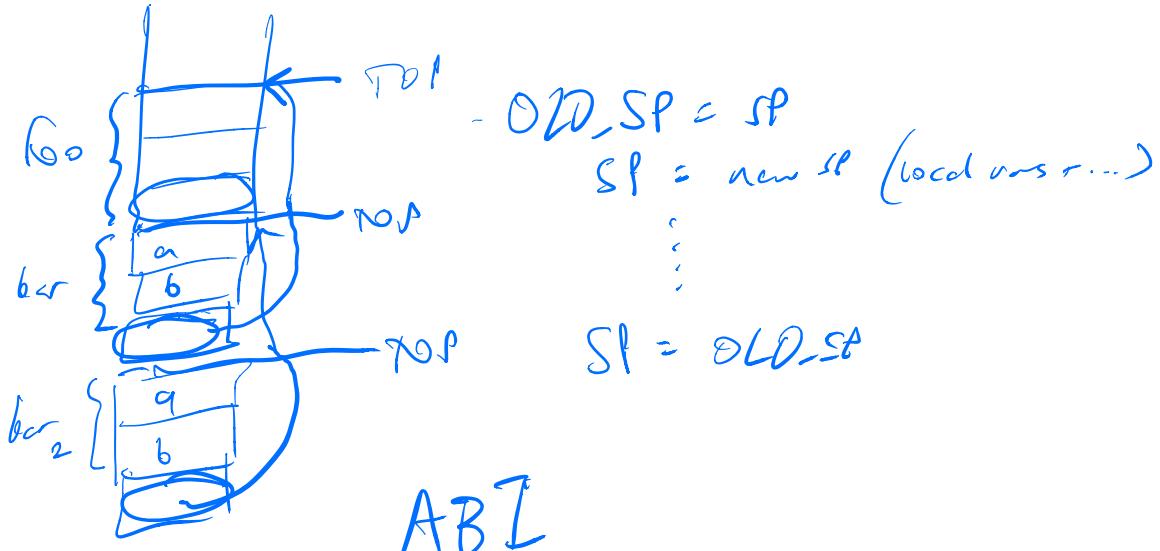
Variable sized instructions

ROP

low level security exploits — memory unsafety



## stack smashing attack



ABI

application binary interface

int a;  
int b;



`*((b + 1)) = pc;`

"root"  
"pwn"

Morris worm 1988

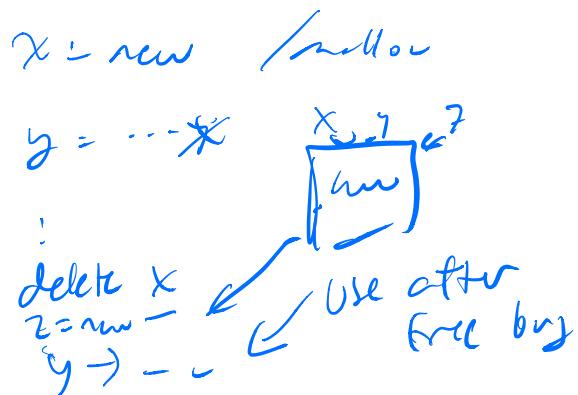
DDoS distributed

denial-of-service  
attack

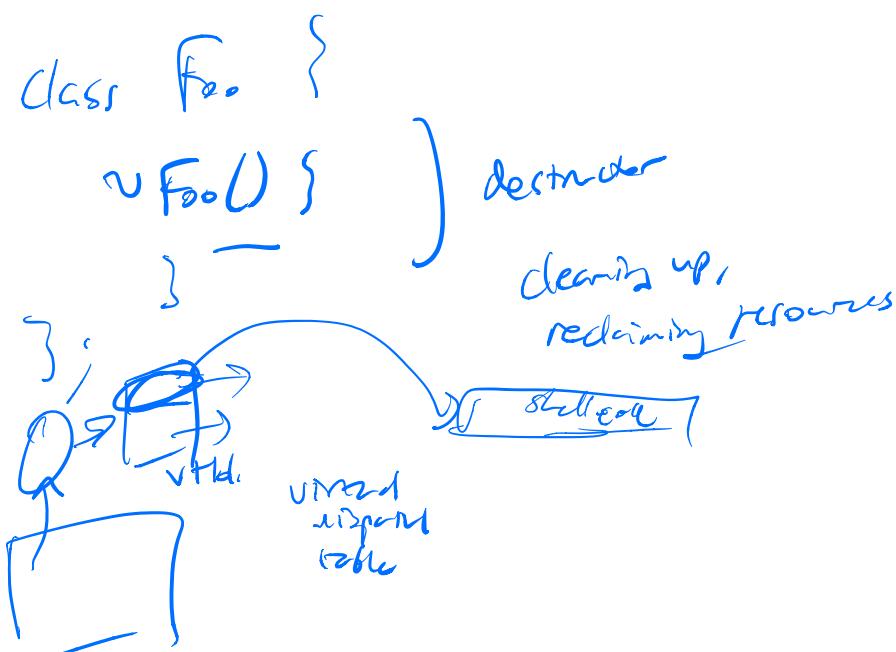
NOP NOP NOP NOP , shellcode

↳ →  
 not std  
 polymorph  
 script kiddies  
 English

dangling ptr / use-after-free

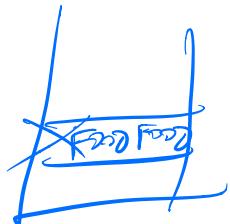


"temporal safety"



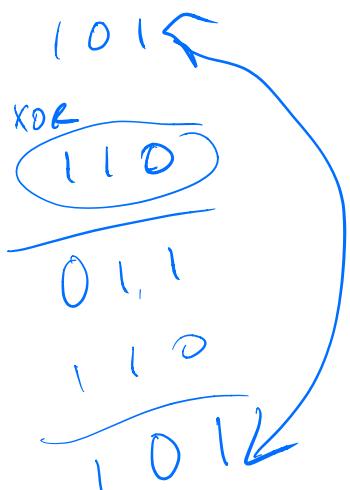
~~delete f;~~

stack smashing: shadow stack  
"Stack canaries"  
XOR ~~"key"~~



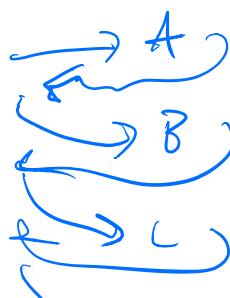
$x \wedge y$

$$\begin{array}{rcl} 1 & \text{xor } \phi & = 0 \\ 0 & \text{xor } 1 & = 1 \\ 1 & \text{xor } 0 & = 1 \\ 0 & \text{xor } 0 & = 0 \end{array}$$

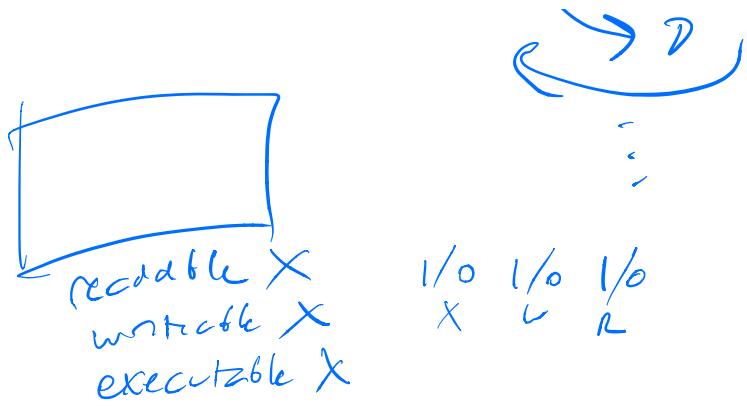


a xor key xor key =  $\alpha$

Dr Harder



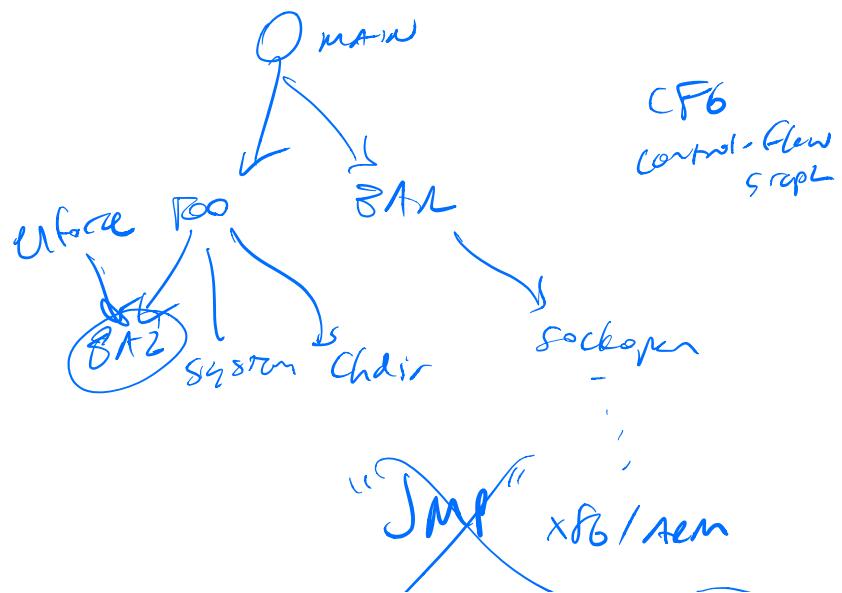
"gadgets"



## NX Protection

return-to-libc attack

"system"  
"rm -rf /"  
CFI      control-flow integrity

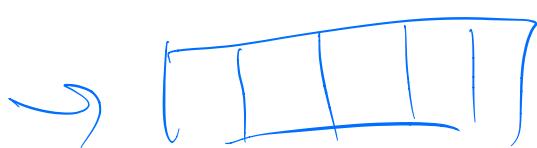




“One level of indirection”  
“all problems in CS”

RISC smaller instruction set  
“fixed-size instructions”

- ① easier to parse “decoder”
- ② harder to exploit
- ③ faster → easier to design how to make fast



PPC  
“TSX”  
HTM

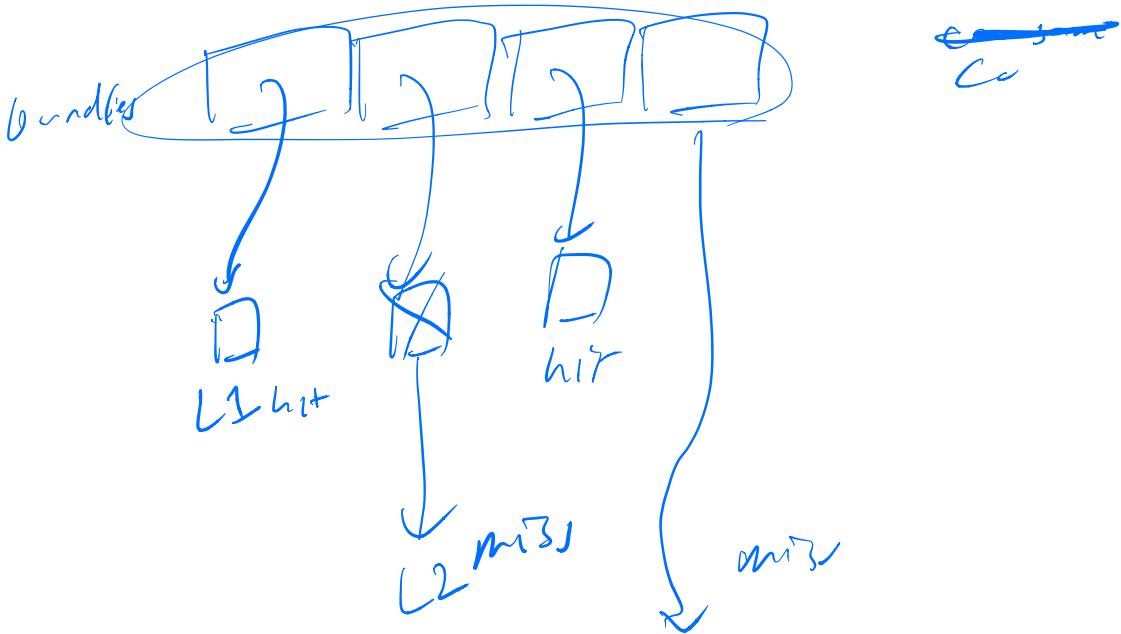
Amd - vector-instruction      Intel MPX  
- crypto      - memory protection

Google "ASan"  
AddressSanitizer

clang -fsanitize=address

VLIW

version instruction wed



EPIC

explicitly  
parallel  
instructions  
Co

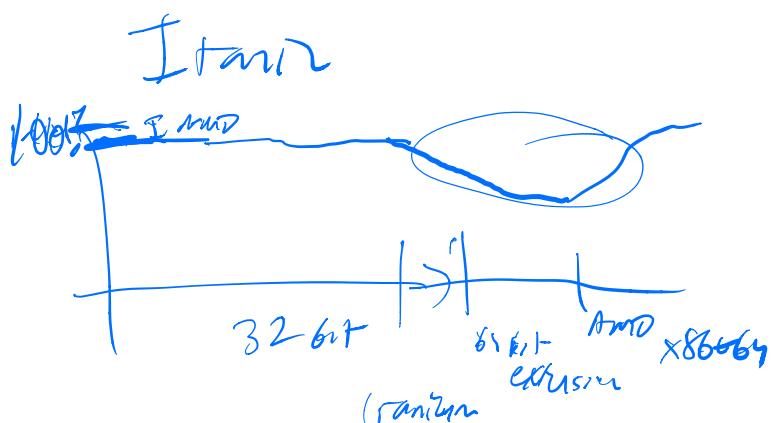
PREDICATION

(IF branch<sub>1</sub>)

DDDD

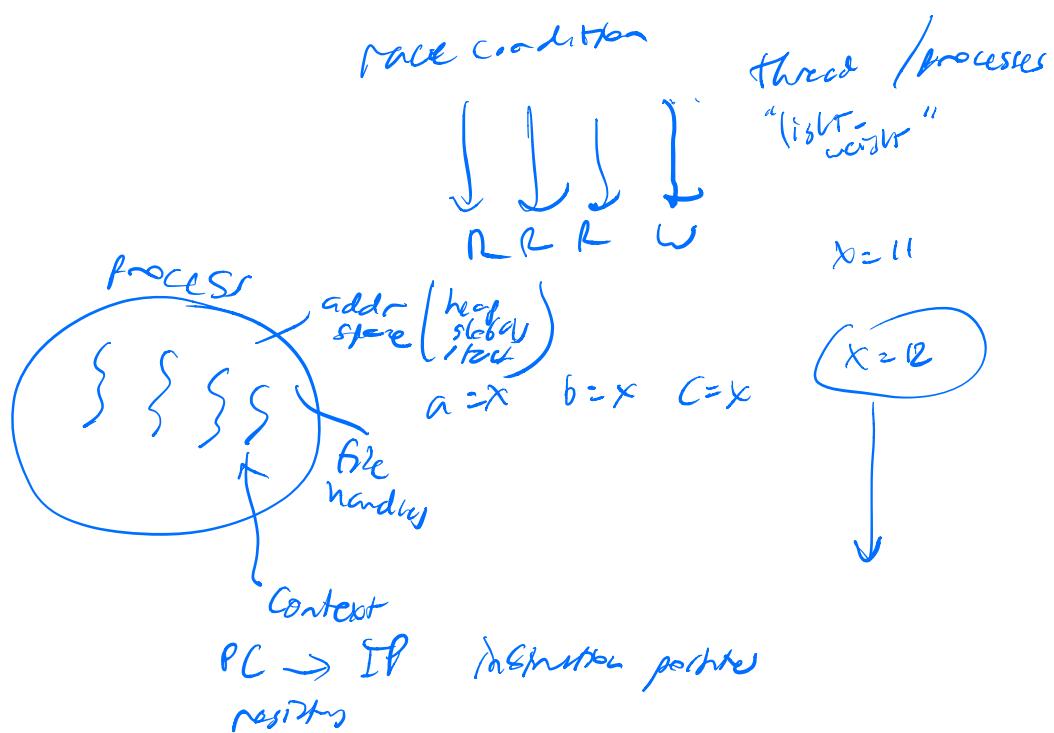
Itanium

IA-64



MOV  $\ddot{Q}$   
ADD  $\ddot{Q}$

## Concurrence



## Concurrence

~~errors~~  
races

atomicity violation

Order violation

deadlock

livelock

) condition variables

— wake-up problem

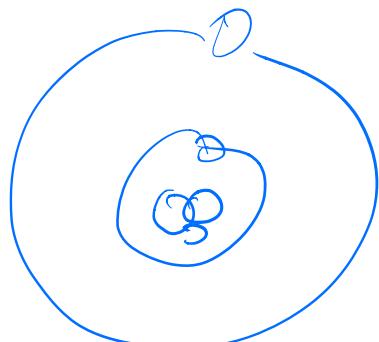
while(<sup>true</sup>  
    lock  
     $\rightarrow$  want()  
     $\wedge$  check state  
     $\wedge$   $T\theta(\cdot)$  break  
    }  
    })

NONDETERMINISTIC

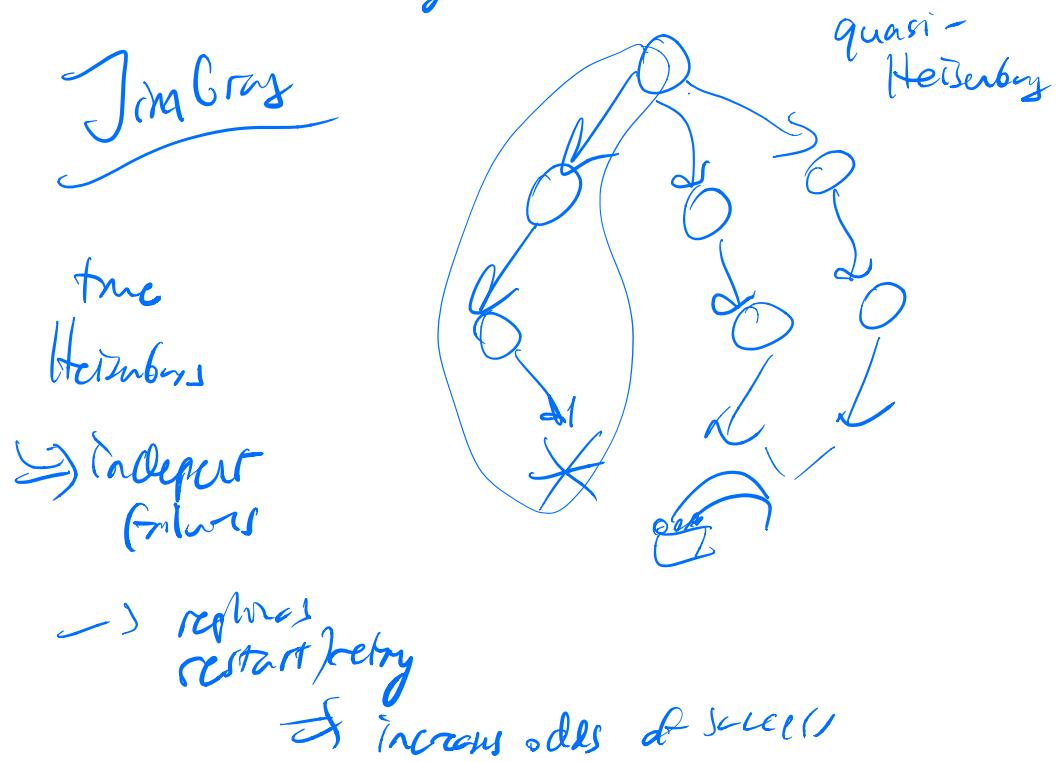
livelock

Heisenboss

DETERMINISTIC boss — Bohr boss



- severity?
- frequency?



Developers

debugging  
determinism!

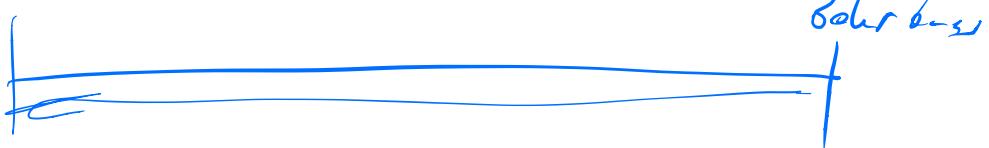
Bohr Bus

Users

Heisenberg

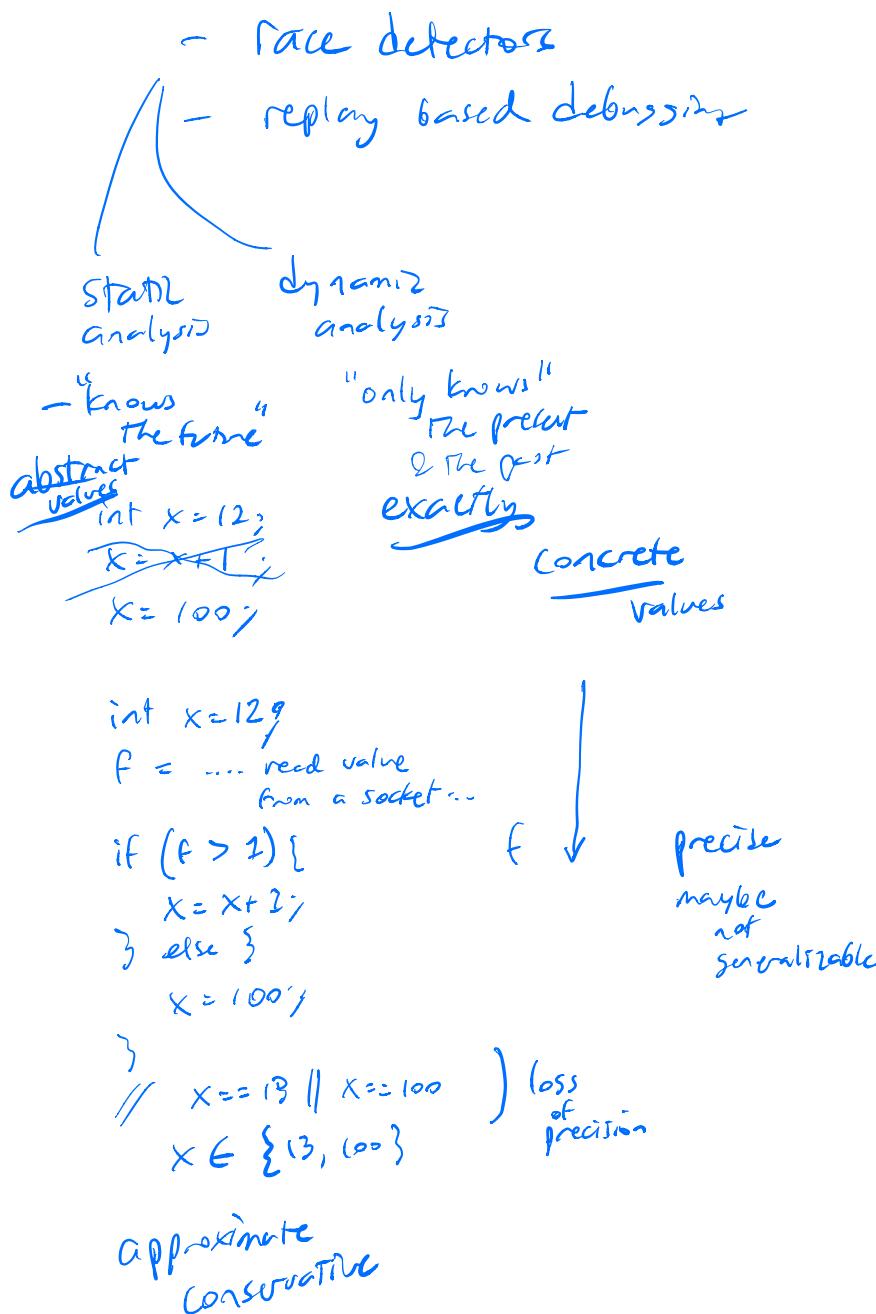
→ retry

→ fault tolerance  
→ Indep. failures)



D+E

100%



## bus detection

partial specification

P never derefs NULL

"implicit specification"

segfault      buffer overflow  
NPE            use-after-free  
races          assertion failure

benign races / malignant races

"Hauswald!"

## formal verification

proving program P

implements spec S

code

y

logiz

"complete / total specification"

abnormally



Word-sized  
read/write

X = 1

001

X = 4

100

101



double  
d

d = 0.1

d = 2.0

d = ...

" O-O

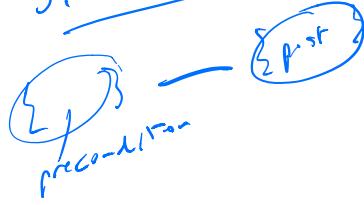
( ... )  
sort :  
postcondition      formal specifications

$H_i : 0 \leq i \leq N-1 :$   
 $a[i] \leq a[i+1]$

↓ " a is sorted  
in ascending order "

Postsort

Sir Tony Hoare



Quicksort  
average case  $O(n \log n)$   
worst case  $O(n^2)$

Mergesort

avg & worst case  $O(n \log n)$

quis custodes custodit

trusted computing base (TCB)

$\underbrace{\text{spec} + \text{solver} + \text{code sv}}_{\text{TCB}} \Rightarrow \text{selected code}$

$\approx 1000 \text{ LOC}$   
just : hundreds !

targets for formal verification }  
Narrow APIs  
well specified  
really important

Project Everest  
"https" TLS F\*

## Heartbleed

static analysis for concurrency errors

~ false positives

precision

recall

print "NO BUG HERE!"

print "BUG!"

dyn analysis — no false positives

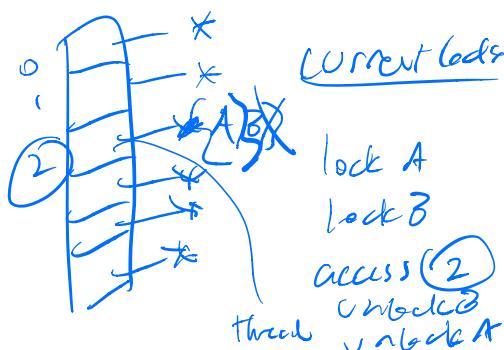
WITHACSS race!

— recall lower —

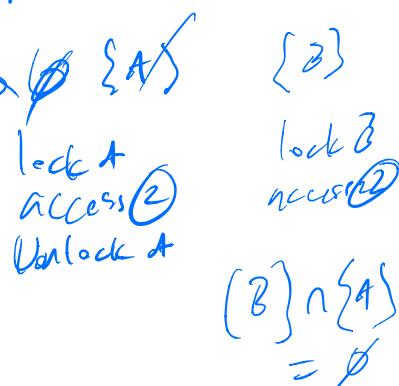
only reports race  
in that execution

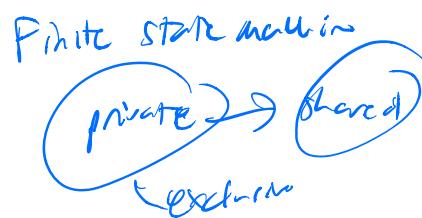
## Thread Sanitizer

### lockset analysis



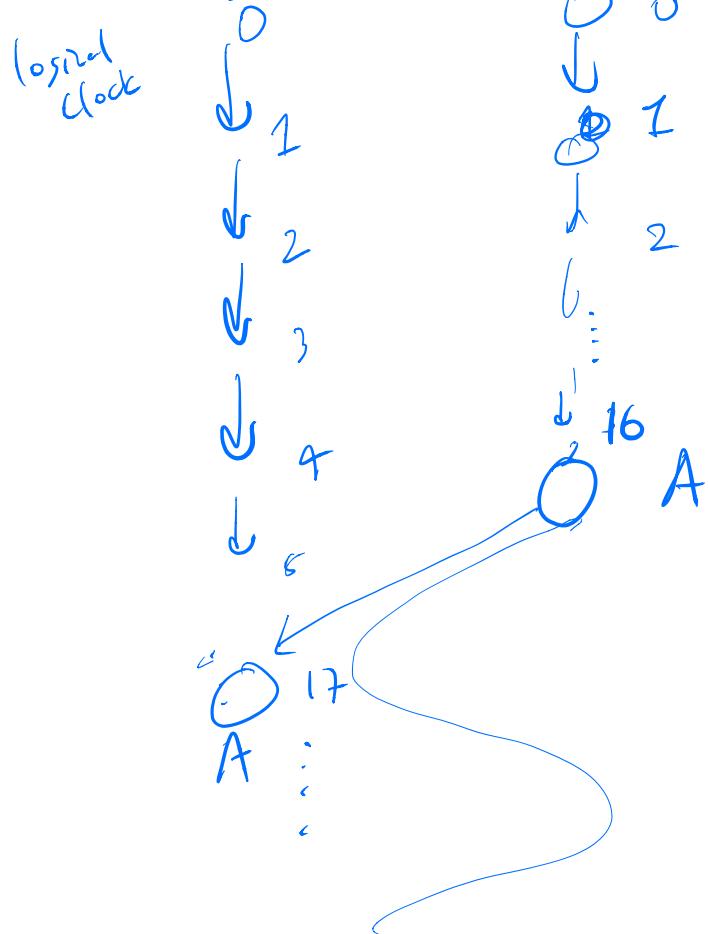
### happens-before analysis





Error

happens-before analysis



Why  
Threads?

- parallelism

- coordination

- hiding latency
  - I/O
  - memory

↑ performance

↳ average the periods  
↳ the delay until "first by"

"throughput" — rate of jobs safely done  
"throughput"

throughput-latency tradeoff

interactive

batch

|  
no context  
switches

CPU util ~ 100%

time  
sharing

/dev/fny

ed

MULTICS "Eunuchs"

1964 - 1967

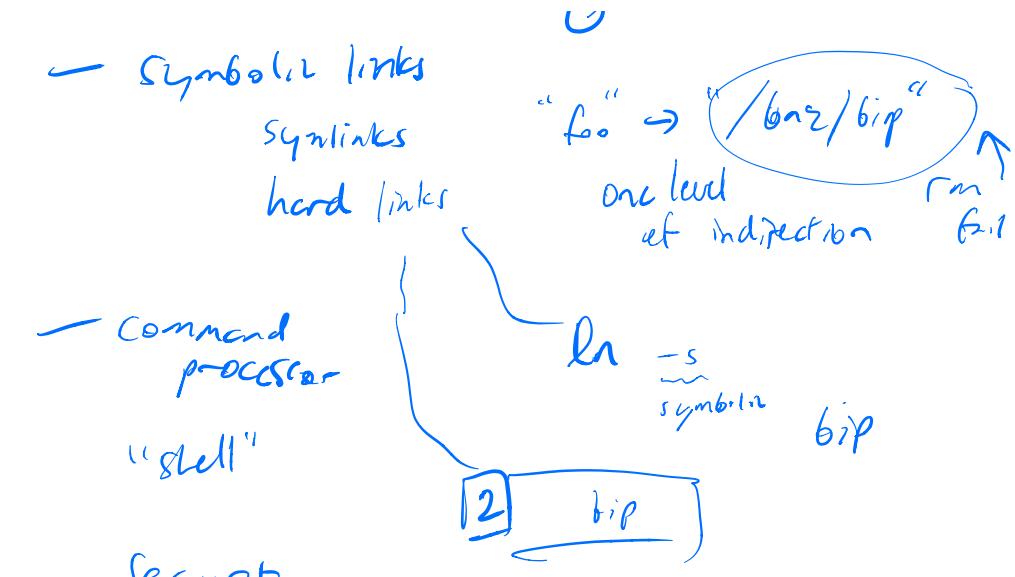
MIT + GE + Bell Labs

- hierarchical FS

all previous:  
"flat"

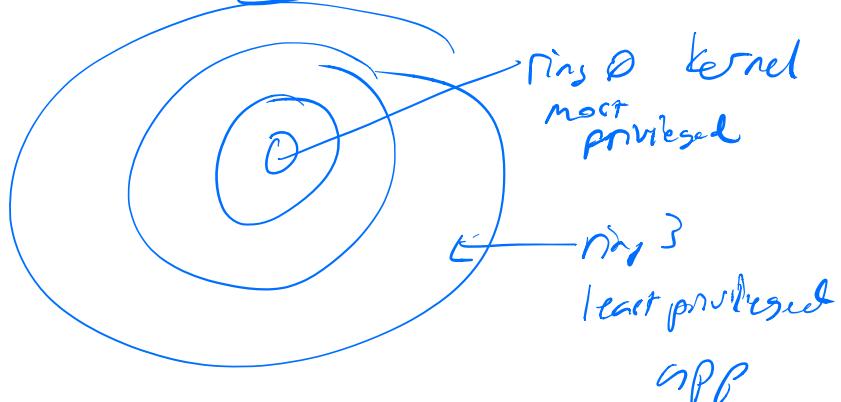
relative  
paths





— Command processor  
"shell"

— Security rings




---

Why is choice of quantum size a tradeoff?

---

```

newObject
mark
clear
sweep
gc
Object * newObject();
    
```

for (auto it = allocated.begin();  
 it != allocated.end();  
 it++) {  
 (\*it) → clear();
 }

clear();
 mark(&root);
 sweep();

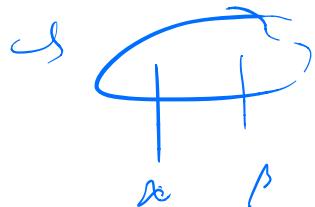
```

\ Objet & obj;
if (freed->empty()) {
    obj = new Object();
} else {
    obj = freed->front();
    freed->pop_front();
}
allocated->push_back(obj);
return obj;
}

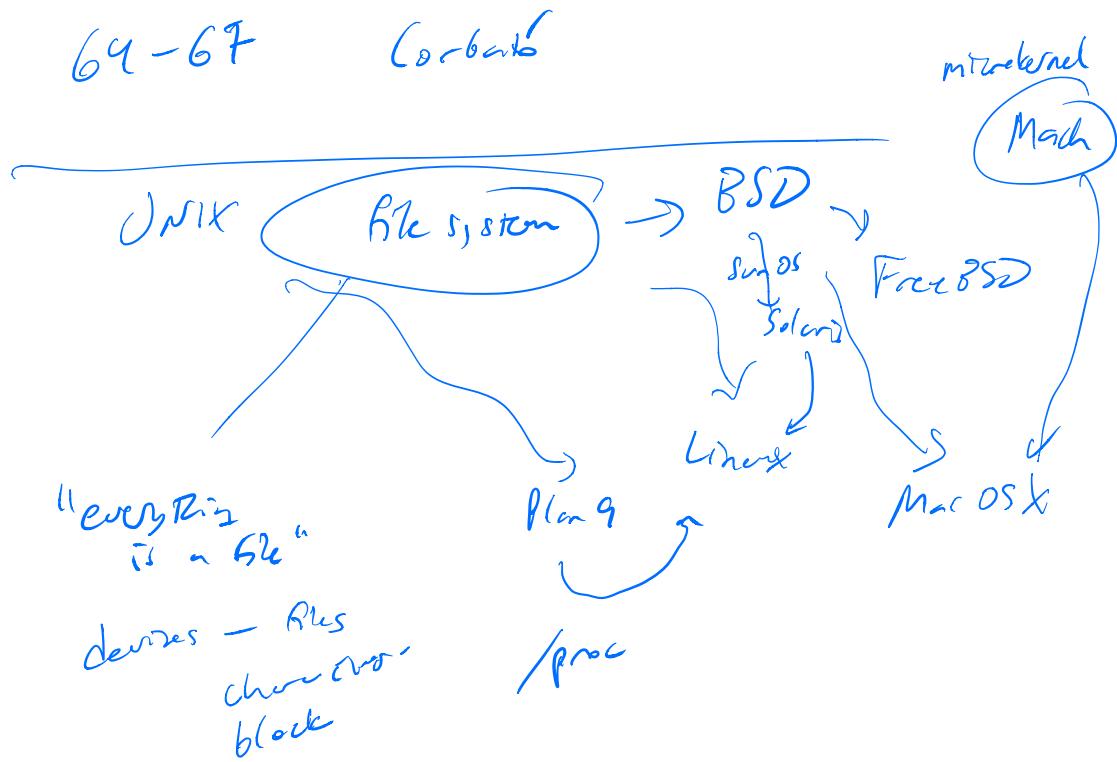
void mark (Object* a) {
    if (!a->isMarked ()) {
        a->mark ();
        if (a->isBinary ()) {
            mark (a->sub1 ());
            mark (a->sub2 ());
        }
    }
}

```

Cons &  $\beta$



## hardware enforced security



stdin stdout  
stderr

"one level of indirection"

echo "hello world" > foo  
 > /dev/l-

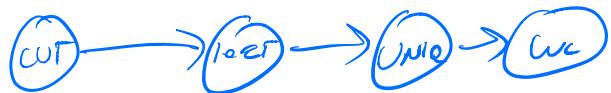


0% ( cut -f 1 -d\n ) | sort | uniq | wc -l

pipeline parallelism

"streaming"

data flow graph



awk

Aho Weinberg Kernighan

sed

s/-/ /g

grep

general regular expression parser

Perl - Larry Wall

Python - Guido van Rossum

Ruby - "Matz"

→ killer app

"Rails"  
Ruby on Rails  
Niklaus Wirth

VisiCalc 1979  
Microsoft Word  
Excel  
PowerPoint

Office

Scala - Martin Odersky

EPFL

→ Pascal  
Modula-2  
Oberon

Andreas (Ae)tschba  
Turbo  
Pascal

→ Java  
generics  
Spark

TypeScript

JavaScript      Stockholm syndrome

Barend  
Eijk

David Turner  
≡ KRL  
→ Miranda<sup>TM</sup>

Simon Peyton Jones — MSR  
Phil Wadler → Comb.  
John Hughes → Edinburgh  
Paul Hudak → Yale  
...  
Miranda<sup>TM</sup> ↗  
erden.  
of Res. SFA Ltd.  
Haskell Curry

$f(x, y, z)$

$f x y z$

"f 3"      currying

pure lazy functional prog. language

```

int s = 12;
int void f(x) {
    return x * 2 + (s + t);
}

```

I/O

$f(10)$

ones = 1 : ones

$[a \times 2 \mid a \leftarrow [1, 2, 3]]$

print ("foo")

mountable file systems

NFS

/var/fsdb

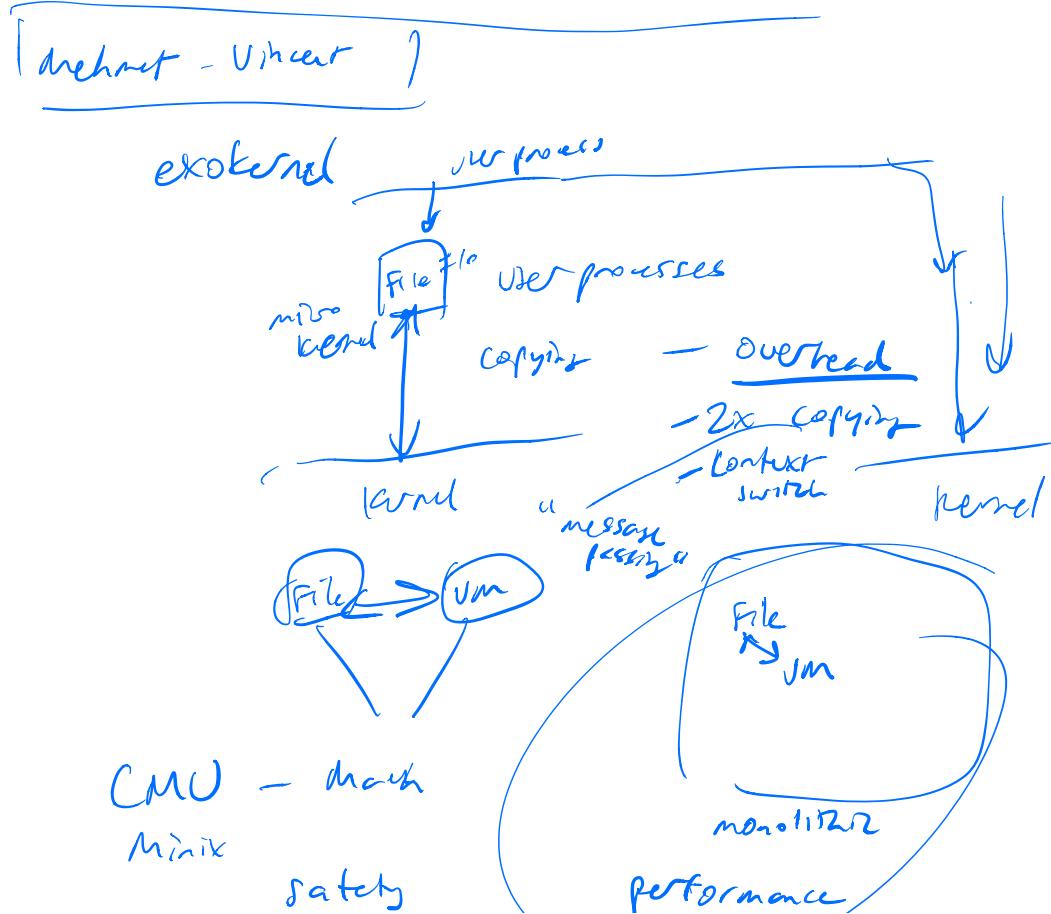
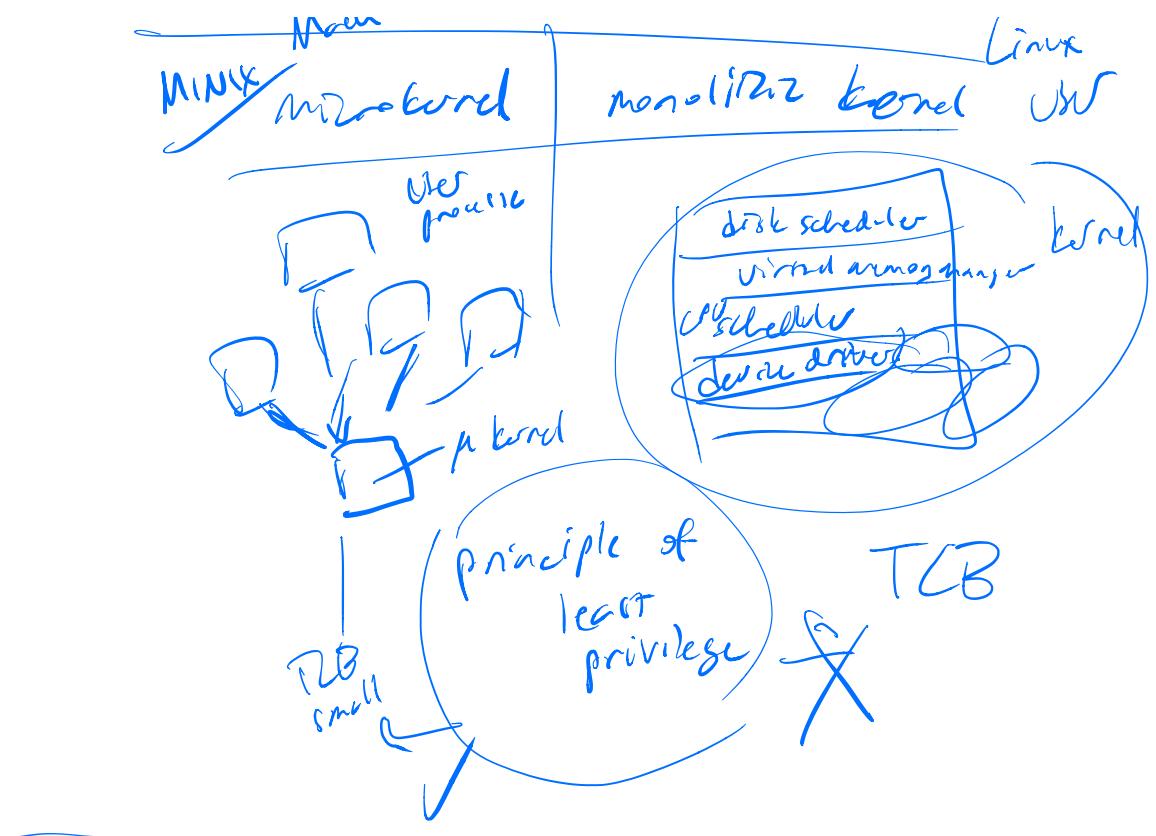
/home

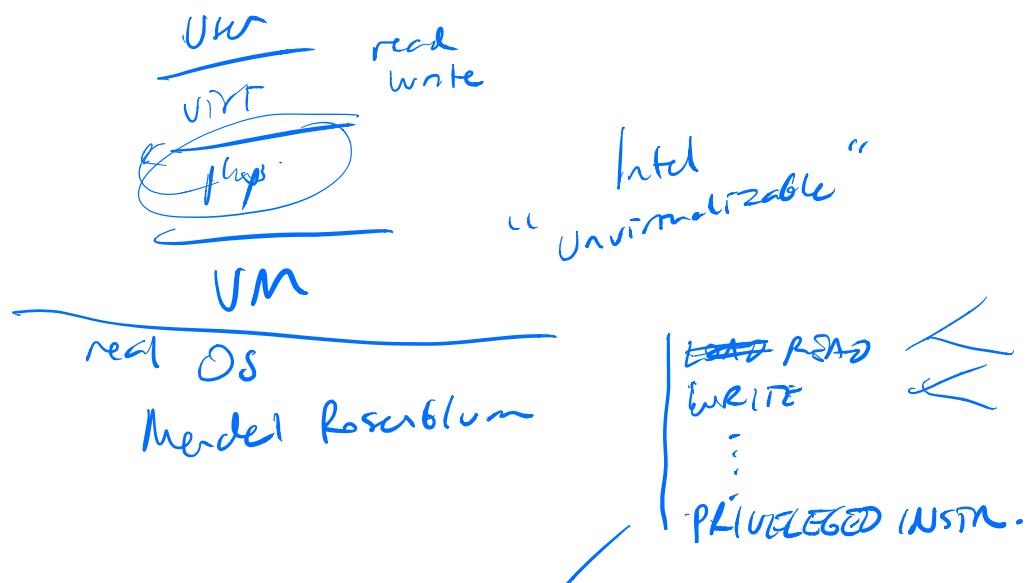
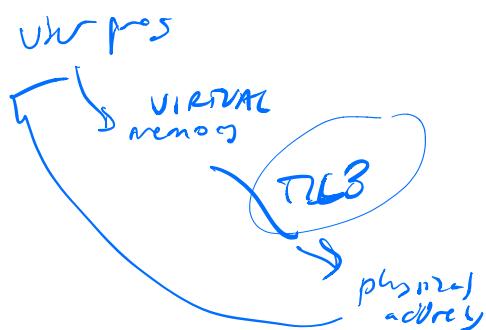
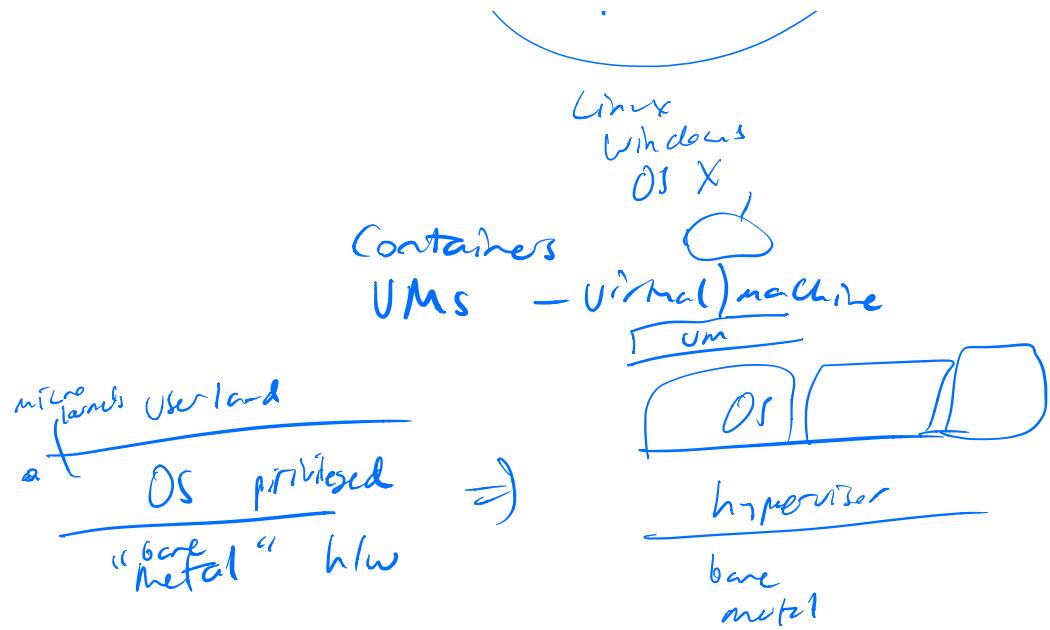
/usr

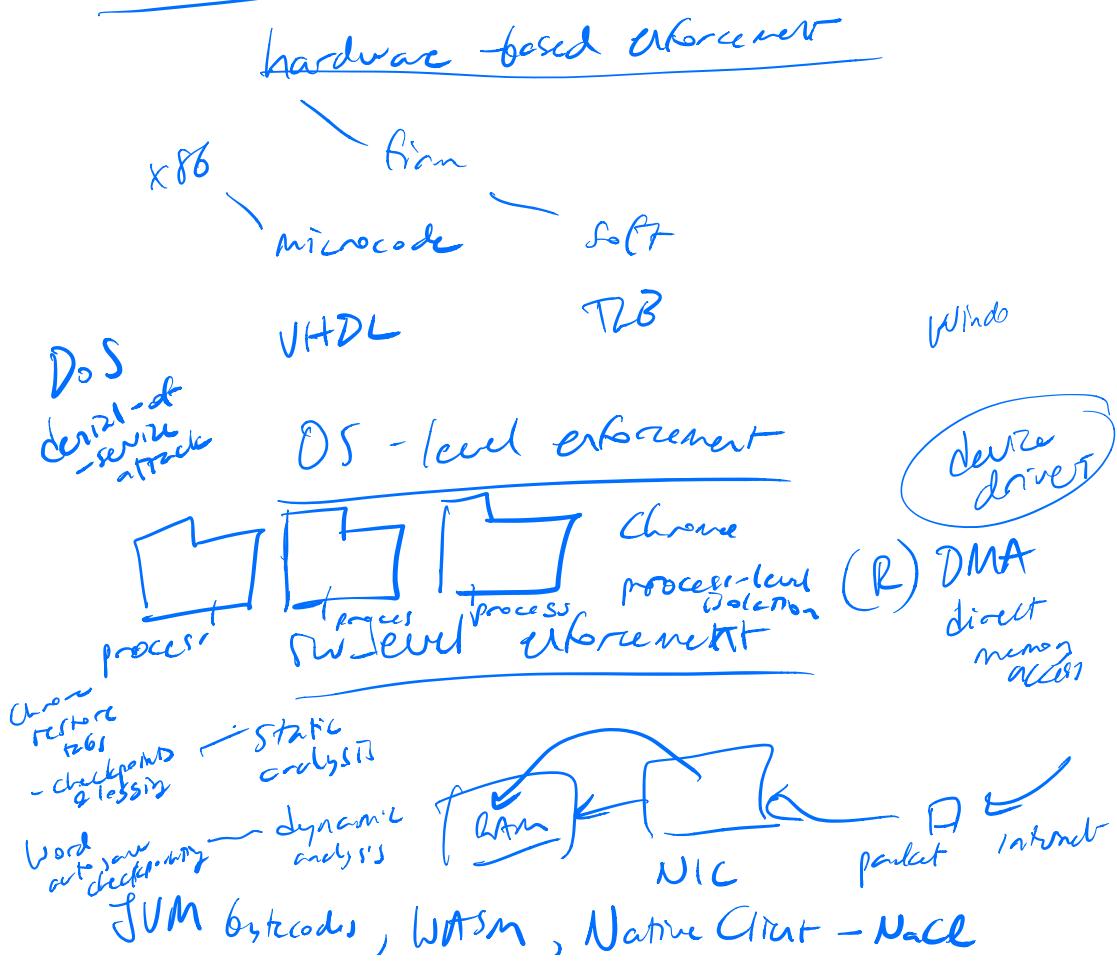
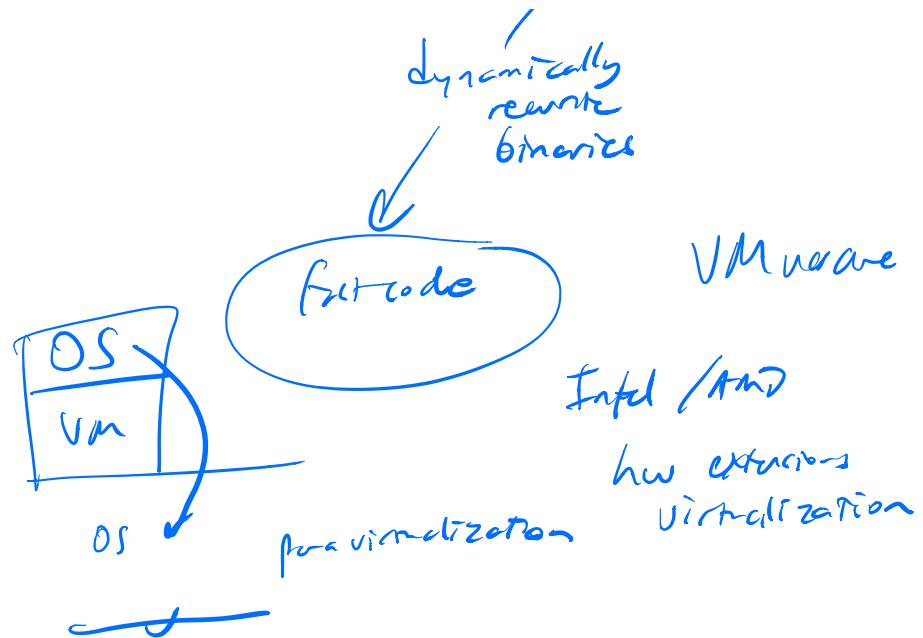
shuttle

Torvalds  
debate  
Flame war

... .







Windows failure  $\rightarrow$  BSOD pNaCl

kernel panic

DDK

- fails internal integrity check
- hardware error

Wisdom  
of the  
crowd

Galton

VoxPPL

"  
home economics"

rational agent - maximize gain

cost-benefit analysis

$$n_{jk} = P(\text{break } h)$$

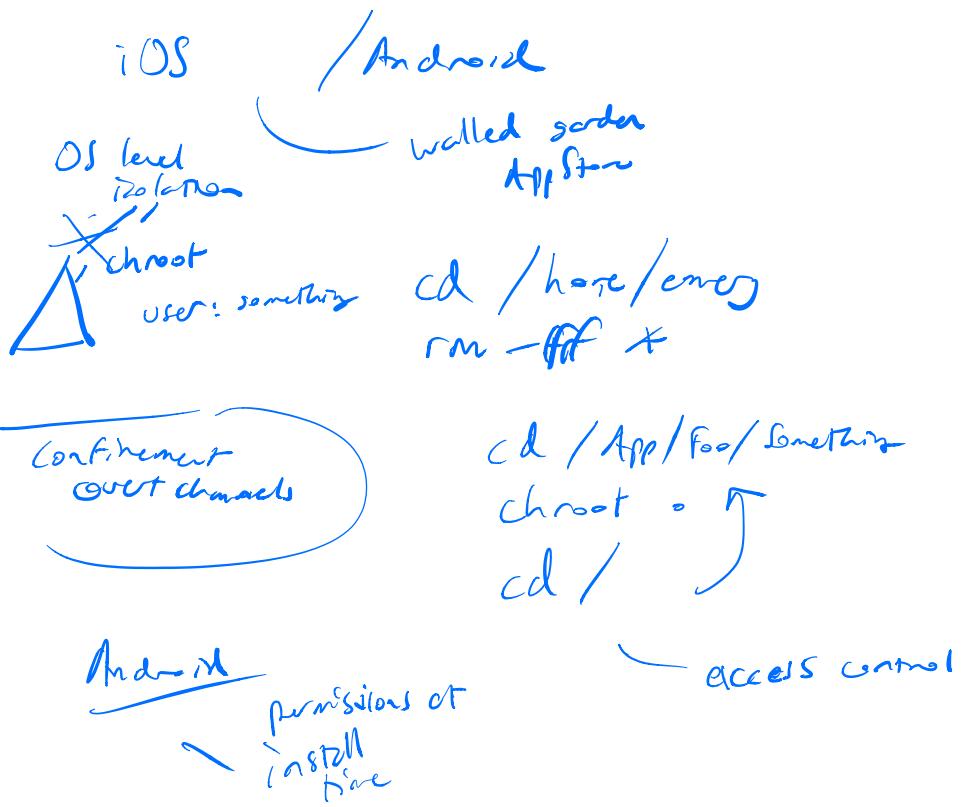
cost =  $\sum \text{cost of replacement}$

$$\begin{array}{rcl} \$30K & & \text{ransomware} \\ \hline 10,000 & + \frac{30,000}{1} & = \\ & \$1 & \$15 \end{array}$$

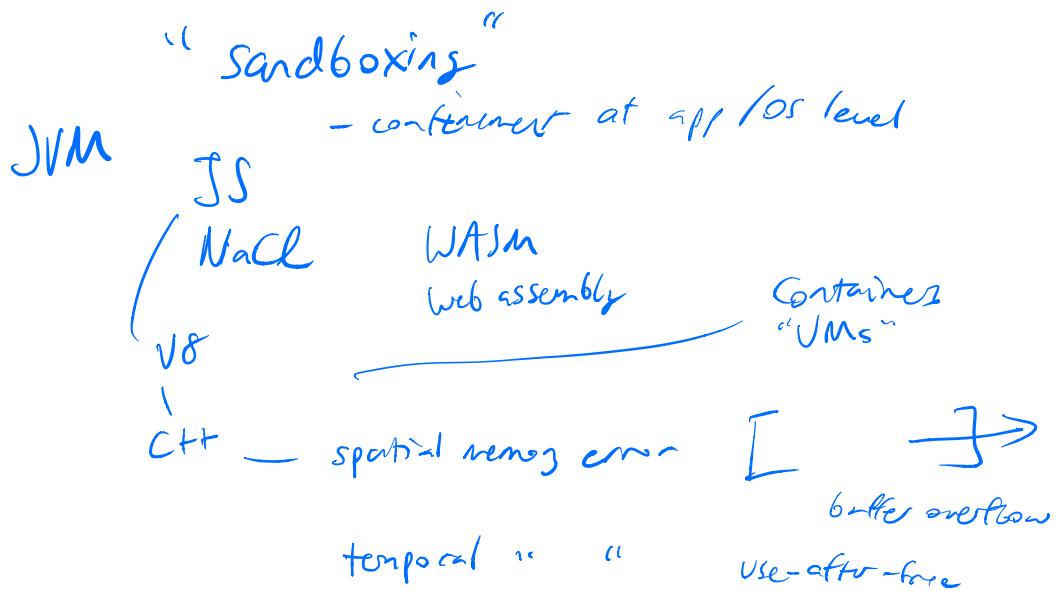
Fort Knox

social engineering  
phishing

asymmetric attack



Scribes: Jun, Kyle



**LangSec** lang-level security

— Systems-level lang  
enforces spatial  
+ temporal safety

- "memory safety"
- "type safety"

contract-kw integr?



- X high-level security  
is mem safety (ctr.) sufficient?
- VMS low-level security  
SQL injection attack
- "safe" PL ?? X

RUST

- memory safety
- concurrency

Ownership types

- no GC

"unsafe"

→ NOT MODULAR



"cat and mouse" game  
"arms race"

# define while if

query = "SELECT \*  
From employees  
WHERE " + cond + "

cond = "salary > 100000"  
position != 'CEO' OR

sqlserver.execute(query, results);  
Send : (results),  
SELECT -- ; DROP TABLES  
[ salary > 100000 ; ]

XKcd Little Bobby Tables

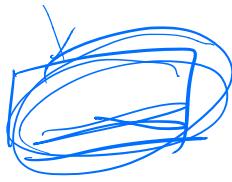
drive-by attacker

↳ zero-day vulnerability

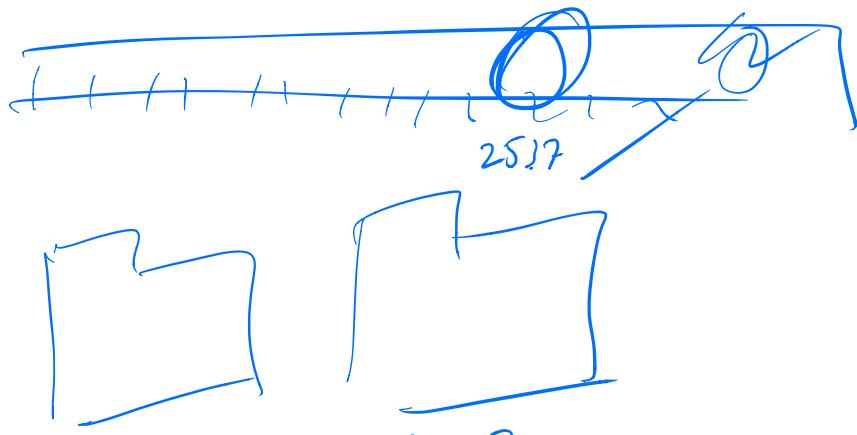
heap spraying  
JIT spraying

s = " shellcode " ;  
;

speculative execution



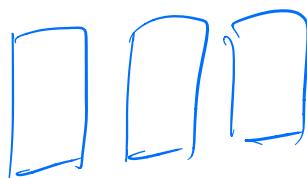
$a[x]$



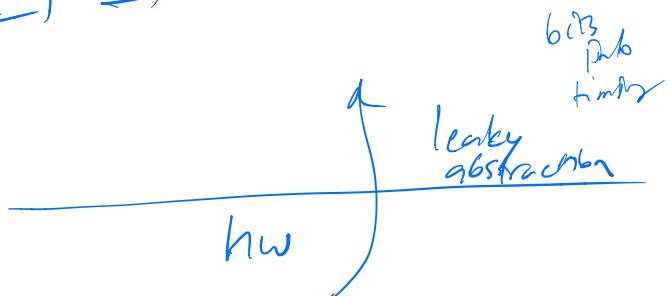
Shared Array Buffers

high res times

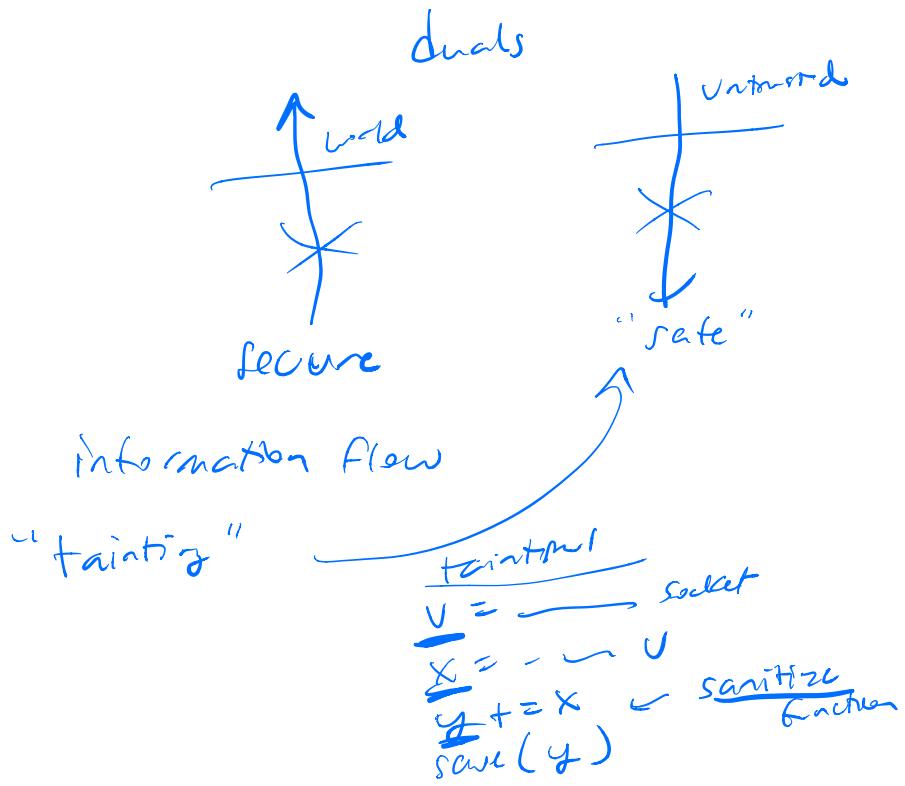
Rowhammer - Our motto



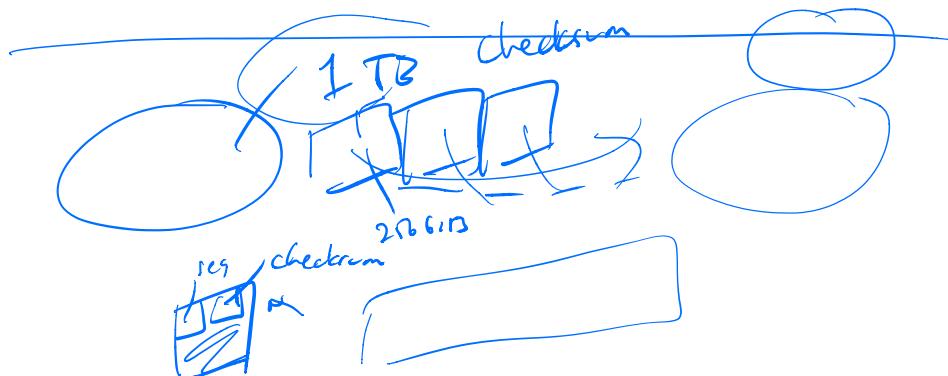
DIMMs



privacy                      integrity

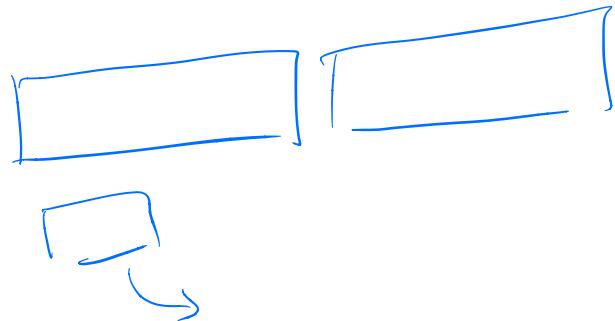


$$\text{sanitize}(\text{untainted}) = \text{untainted}$$
$$(\text{tainted}) = \text{untainted}$$



$P(\text{covert packet undetected}) = 1,000,000$   
silent film

D D D D D  
~~~~~  
≥ 1000 0000

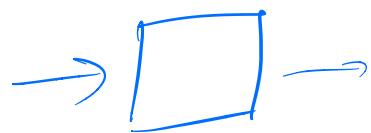


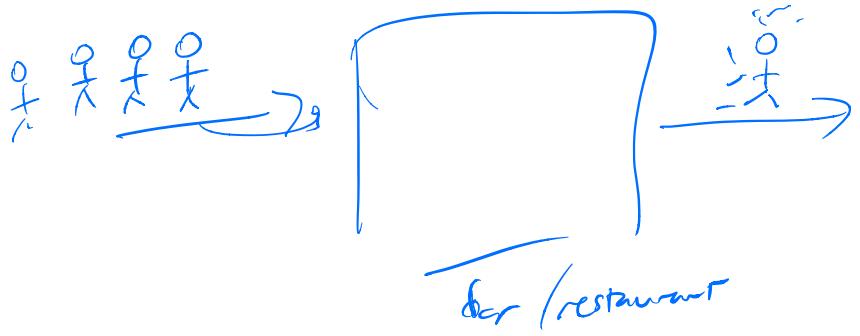
securing integrity

MOAR BITS

---

shedding load  
~ "admission control"





$$L = \lambda W$$

Little's Law

how many people waiting in source

arrival rate      service rate

$\frac{100}{hr}$        $1\text{hr}/\text{customer}$

$$E[L] = E[\lambda] \cdot E[W]$$

backlogged

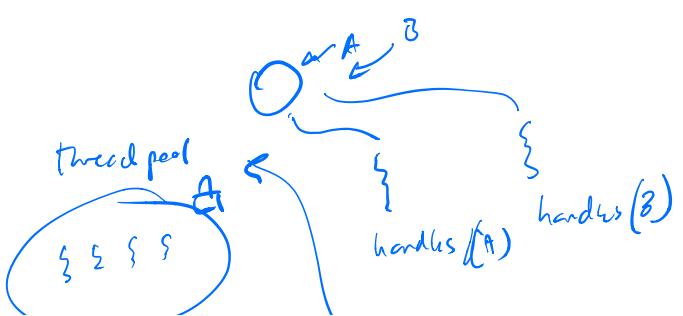
load shedding

admission control — bouncers

QoS

DoS

"thread pool"



$N=30$   
 $N=100$

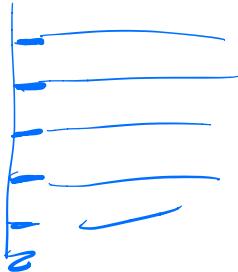
admission control mechanism

Cryptography

"Security  
through  
obscure"

codewords

"steganography"



"key"

Caesar cipher

key = 5

A B C D E F G ...      Symmetric  
↓ E F G H I J K      crypto  
values

A B C D E F ...      DEAD

Z Q G A F R ...      AFZAA

(frequency analysis  
ciphertext)

Q Z D

E  
T  
A  
O  
I  
N  
O

0 --- 0

H  
R  
D  
L  
:

One-time pad



XOR

$$\begin{array}{rcl} 0 \text{ XOR } 0 & = & 0 \\ 0 \text{ XOR } 1 & = & 1 \\ 1 \text{ XOR } 0 & = & 1 \\ 1 \text{ XOR } 1 & = & 0 \end{array}$$

rotating keys  
(codebooks)

Enigma

known plaintext attack

- decryption!

Goal: encrypted message ≈  
random noise

⇒ best you can do  
⇒ brute force

make key comb

probabilistic

$$E(\text{decrypt time}) = \frac{1}{2} * \text{key space}$$
$$\left( \text{e.g. } 2^{1024} \right)$$

2<sup>2048</sup>

if you get my key,  
you can't : decrypt  
encrypt ) asymmetric  
keys

encrypt w/o decrypt

use RSA

attestation / signing

Diffie Hellman

one way function

$$f(x) \rightarrow \beta \text{ easy!}$$

$$f'(x) \rightarrow \alpha \text{ hard!}$$

prime factoring is hard

$$A \wedge \overline{B} \wedge C \wedge D \wedge \dots$$

A : T

$$\begin{array}{l} C : F \\ B : T \\ D : T \end{array}$$

$$\underline{n} = 12$$

Alice Bob

Mallory Eve "Sybil attacks"

reputation based  
systems

Amazon reviews  
Gelph  
Trip Advisor  
Netflix  
;



$$E_A(s) \rightarrow \underline{m} \rightarrow D_A(m) \rightarrow s$$

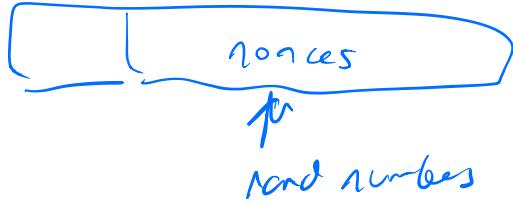
$(e, d, n)$  e constant  
d is relatively prime to  $(p-1, q-1)$

$$\text{Signature} = D_B(\underline{m}) - s$$

$$\text{or just } E_A(s) \rightarrow \text{Signature}$$

$$\text{Alice } D_A(s_{\text{Alice}}) \rightarrow s_{\text{Alice}}$$

$$\text{now } \underline{m} = E_B(s)$$



**PKI**  
public key infrastructure

not  
certificate  
authorities  
quis custodiet custodes?

$$\text{availability} = \frac{\text{MTBF}}{(\text{MTBF} + \text{MTTR})}$$

mean time between failures  
mean time to recover

**ROC**

recovery oriented computing

fail-safe vs fail-stop

rollback  
(or  
"undo")  
error

stop or one-

fail-over - replicas



~~xx~~ →

State replicated

Replicated state machine

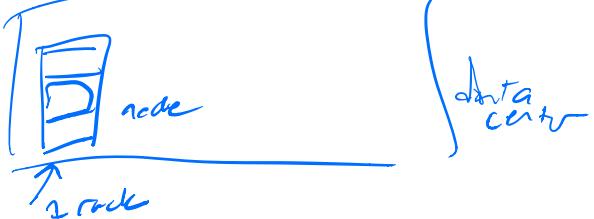
Zossmann & Schneider

replication

Single point of failure

Boeing 737 Max

Geographic distribution



Correlated failure = BAD

GOAL: independence

$$P(\text{failure}_1) \times P(\text{failure}_2)$$

$\frac{1}{1000}$

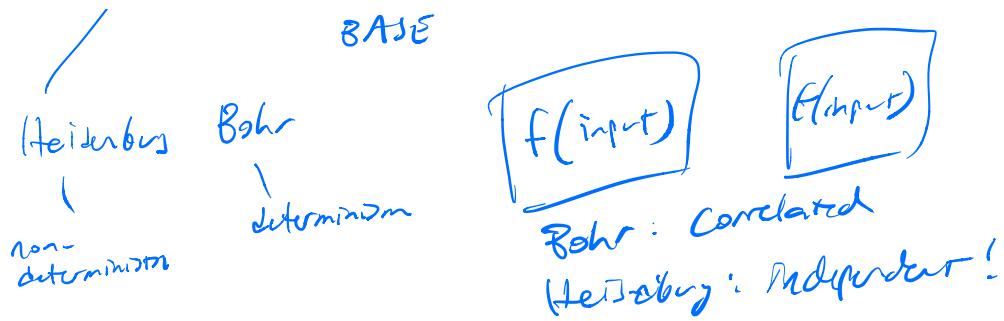
$\frac{1}{1000}$

$$= \frac{1}{1,000,000}$$

fault tolerance

- hardware RAID

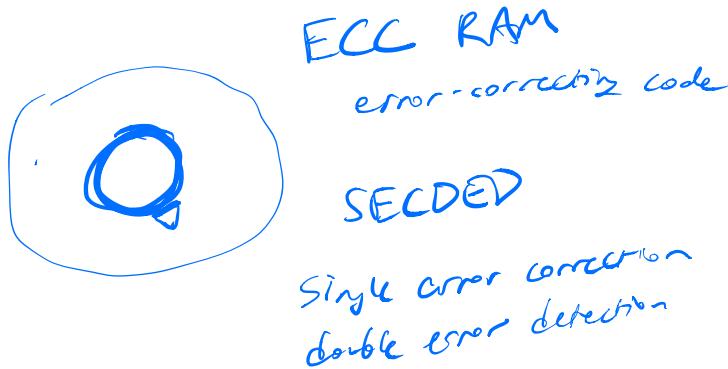
- software transactions ) replication  
ACID



SAY:  $P(\text{failure}) = \frac{1}{100}$   
Heideburg

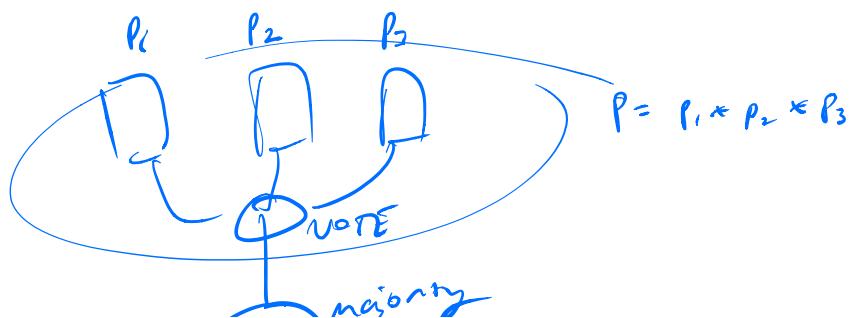
$$E(\# \text{ of people who aware}) = (-P) = 99\%$$

hw non-det failures



trimodular redundancy

3 things



transactions

atomic  
all or nothing

consistent  
in face of concurrency

log of updates  
journal

WRITE ADDED class,

WRITE " "

WRITE SUBTRACT ...  
DROPPED

ABOUT to commit

COMMITTED

integrity

durable

commited  
"stable storage"  
(persistent)  
— Hard disk

Checkpoint "snapshot"

Concurrency control..

lock - pessimistic

- optimistic

"conflict" - Abort  
roll back & restart  
retry

RAID

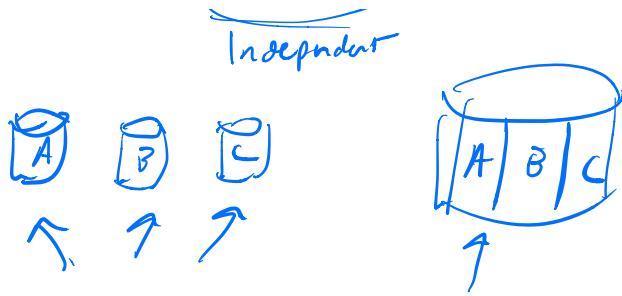
NW

network  
of  
workstations

L1SC

disks → bigger → more expensive!

Redundant  
Array of Inexpensive Disks



$$f = P(\text{failure}) = \frac{1}{1000}$$

success:  $1 - f = \frac{999}{1000}$

$$(1 - f)^2 = \frac{998}{1000}$$

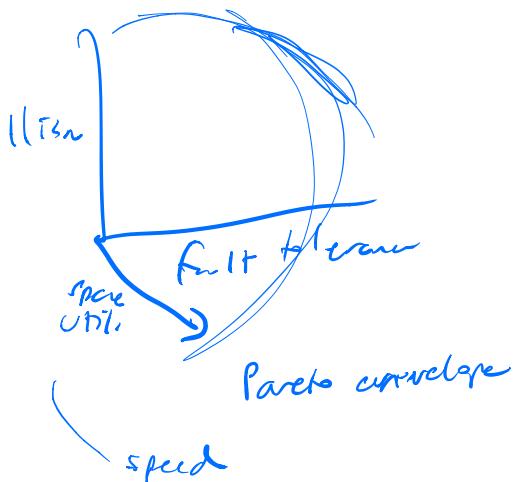
$$\approx .997$$

$$10 = 99\%$$

$$- 100 = 90\%$$

$$- 1000 = 37\%$$

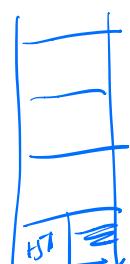
balance || Dm  
2 fault tolerance  
& space



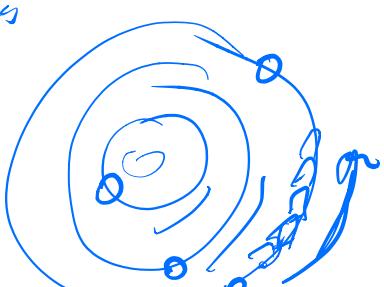
Hamming code

RAID-5

OS file blocks  
disk



buffer queue



rotational latency << seek latency

sw fault tolerance

- failure-oblivious
- probabilistic methods

approx computation

loop performance

$i = 8$

$\text{for } (\text{int } i=0; i < N; i++) \{$

/ \}

Bolt



---

talk tomorrow at noon

Caroline Lemieux

CS 150/1

12:15 · wed

---

failure-oblivious

general

- assert(0)
- throw exceptions
  - ↑  
unhandled

- $x = 2/0$
- $x = * (0)$

} later  
programmer  
start

- These  
are fails

- “always”  
↳
  - $x = (\text{null}) \rightarrow \text{foo}$
  - $\text{exit}(-1)$
  - memory safety errors

$x = \text{new char}[9];$       buffer overflow (spewz)  
 $x[9] = 'z'$       dangling PTR / race  
 $x[\underline{9 \dots 10}]$       after free (temporal)  
 $\boxed{\dots}$       “overfitting”

**CAVEAT!**

- Concurrency errors
- race

correct execution

buffer overflow (spewz)

dangling PTR / race  
after free (temporal)

“foreign races”

C++  
memory model

Santa Adve  
Hans Boehm

$x > 1$        $x = x + 1$

undefined behavior

~~if ( $x > 0$ ) {~~  
~~:~~  
~~}~~

depends on undefined behavior

deal

Hogwild!

higher performance

~ undefined behavior!

“int32”

lossy  
discrete!

“floating point”

bisint  
bisint  
exact math

image  
processing

NLP  
speech  
understanding  
SGD

Sort  
total spec

face detection  
spec?

Spec  
≈  
code

data structure



C → late C

Checking read / write

(Q) x = new char [8]

3

1

down with

= cache <sup>օօջ</sup> լույս

$x[9] = 'z'$       008  
 addr | value  
~~8 \* [9] | z~~

$$c = x[9]$$

C>‘F’

malloc ①  
malloc ②  
malloc ③

~~metac~~ ~~(1)~~ = new

Deleter

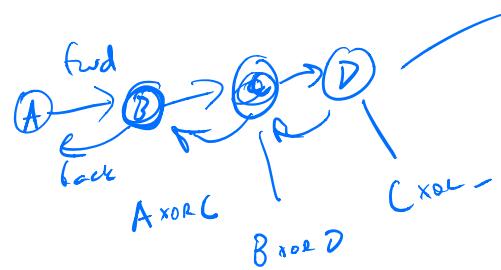
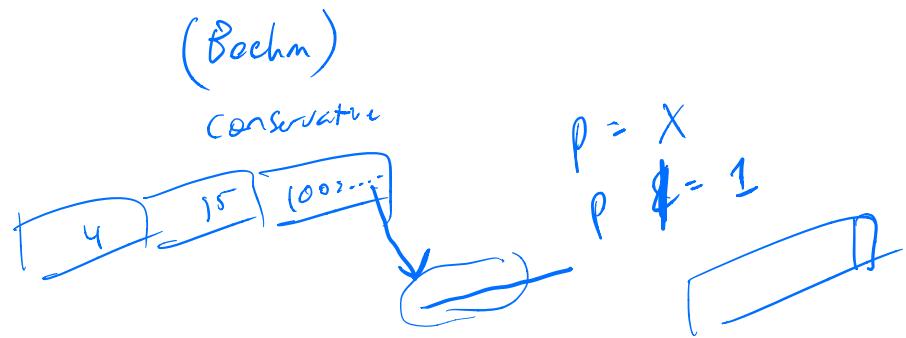
1

X ---  
(use after free)

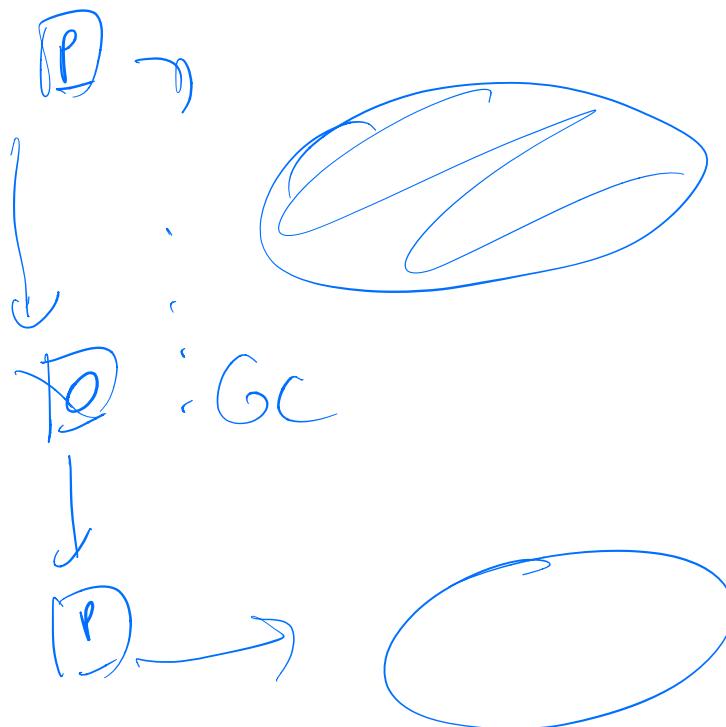
ULE 6C

"infinite blocks"

$$\begin{array}{r} x(10) \\ \times (1) \\ \hline x \{ (0000\ 00) \end{array}$$



$$A \text{ XOR } C \text{ XOR } C \Rightarrow A$$

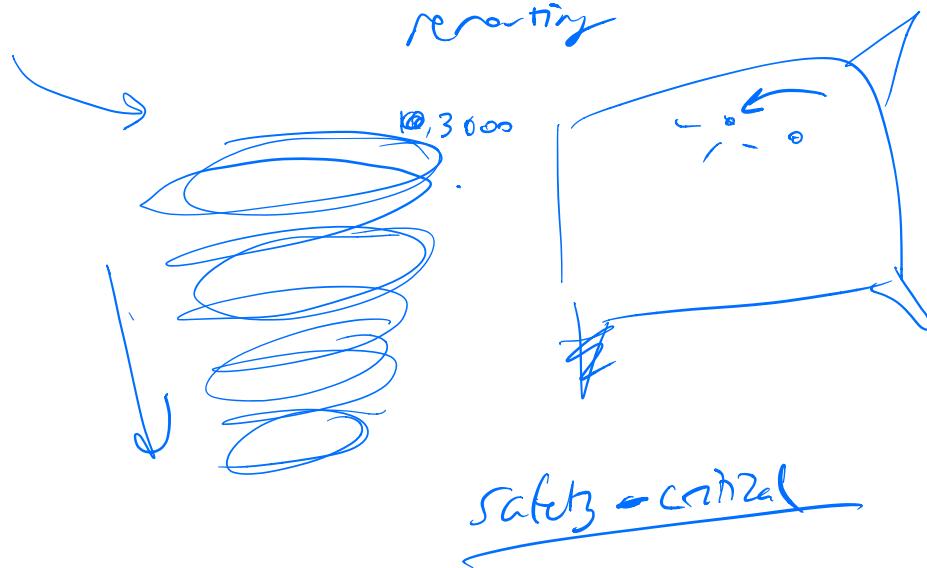


ATC

backoff protocol

reporting

10,300s



safety-critical

checkpointing  
backups

phone apps

not safety critical

retry

checkpointing  
backups

fsck

filesystem check

OOM killer

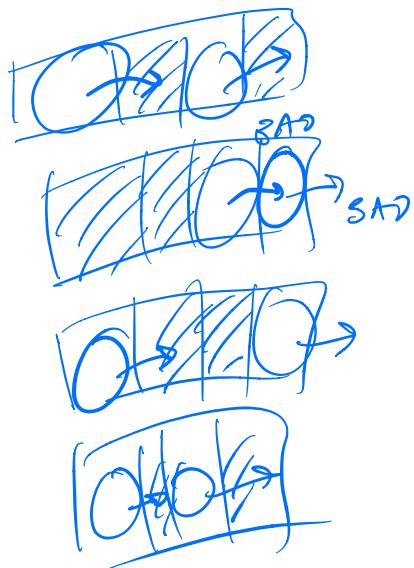
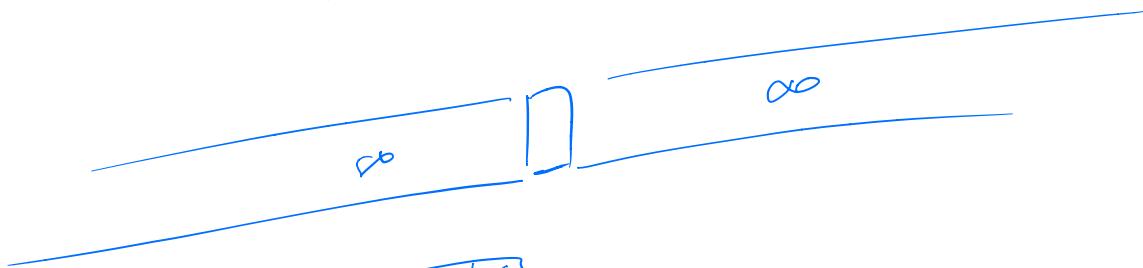
Draft

FO - slow

Archipelago

memor hazard

infinite heap demands



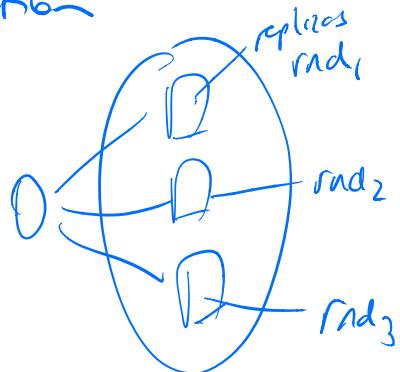
Fahr  $\rightarrow$  Heißabzug

$$\begin{aligned} &P \\ &E[\text{hours failing}] \\ &= p \times N \end{aligned}$$

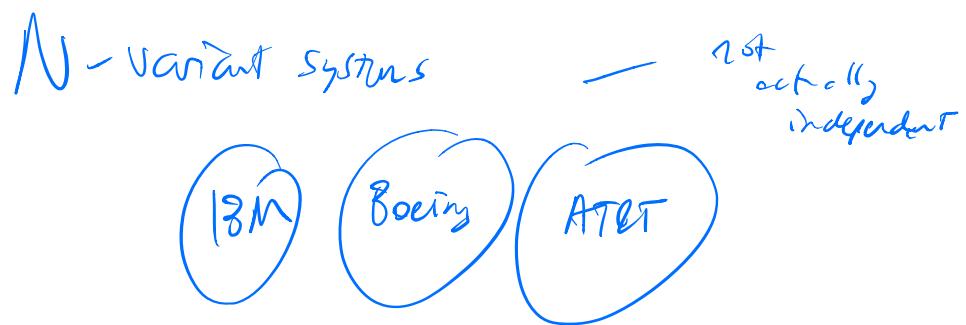
Space for reliability

probably - is not  
of probabilistic

replication



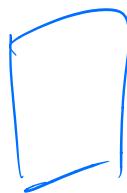
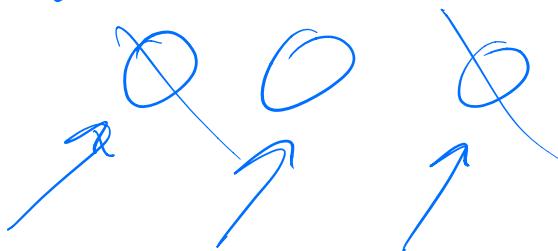
independence  
of  
failures



replicated state machines "RSM"

Dressendorf & Schneider  
(Cornell)

DFA<sub>i</sub>



determinizer

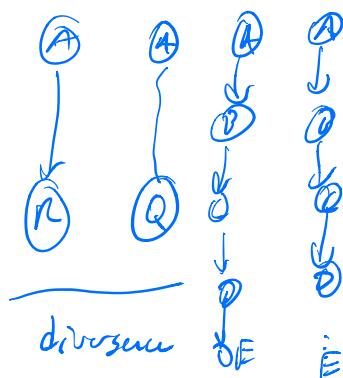


rep<sup>g</sup>  
implementations  
w/ determinization

(threads!  
async I/O  
(in-order arrival  
(of I/O)  
req. for det.)

User input

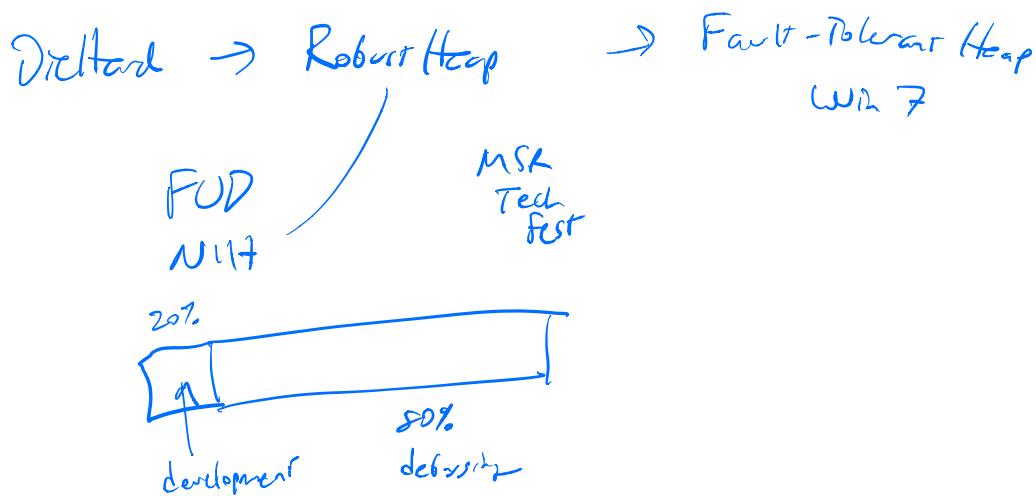
Scheduling



Time

RSMs	+ determinism → Polar beats	never roll your own crypto hash fn <u>MJ</u> binary search :-)	→ RNG
RSMs	+ Deltard → Herderbrassw/ probabilities		

## GPL infection



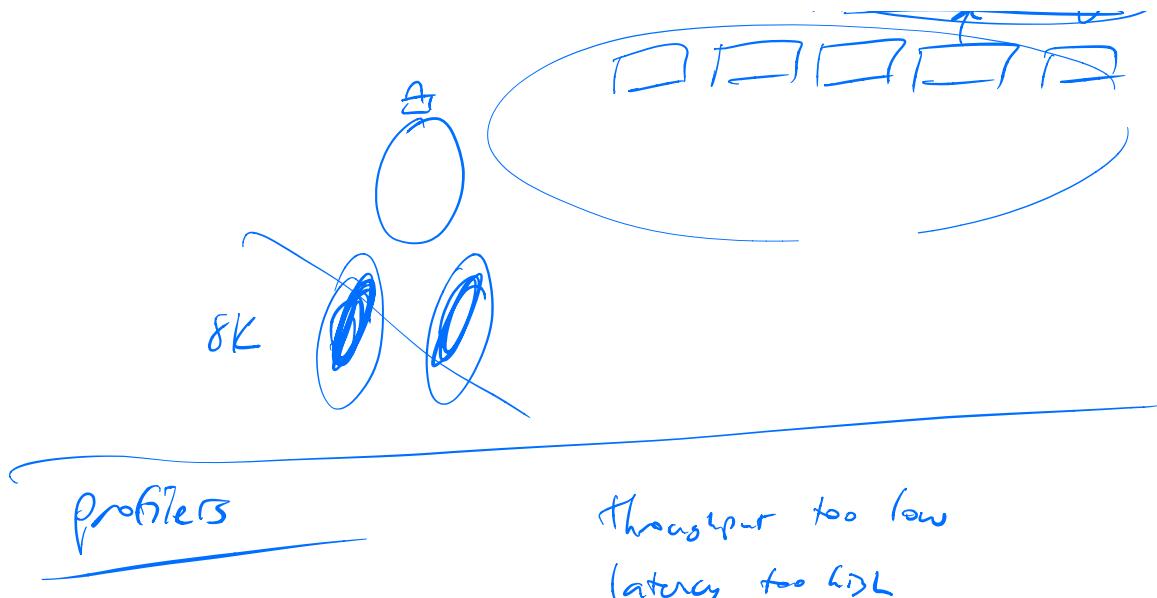
## Security System

how secure is it against attack  
by an attacker  
who can do X } threat model

Fault tolerance  
independent  
hardware  
software  
correlated ...

Deltard → Win 8





prof

line count

main() {

$\downarrow$  l1++

$\downarrow$  l2++

$\downarrow$  <sup>too</sup> l3++

$\downarrow$  l4++

}

f( )

f( )

f( )

{ }

$\downarrow$  l5++

$\downarrow$  l6++

$\downarrow$  l7++

$\downarrow$  l8++

$\downarrow$  l9++

$\downarrow$  l10++

$\downarrow$  l11++

$\downarrow$  l12++

$\downarrow$  l13++

$\downarrow$  l14++

$\downarrow$  l15++

$\downarrow$  l16++

$\downarrow$  l17++

$\downarrow$  l18++

$\downarrow$  l19++

$\downarrow$  l20++

$\downarrow$  l21++

$\downarrow$  l22++

$\downarrow$  l23++

$\downarrow$  l24++

$\downarrow$  l25++

$\downarrow$  l26++

$\downarrow$  l27++

$\downarrow$  l28++

$\downarrow$  l29++

$\downarrow$  l30++

$\downarrow$  l31++

$\downarrow$  l32++

$\downarrow$  l33++

$\downarrow$  l34++

$\downarrow$  l35++

$\downarrow$  l36++

$\downarrow$  l37++

$\downarrow$  l38++

$\downarrow$  l39++

$\downarrow$  l40++

$\downarrow$  l41++

$\downarrow$  l42++

$\downarrow$  l43++

$\downarrow$  l44++

$\downarrow$  l45++

$\downarrow$  l46++

$\downarrow$  l47++

$\downarrow$  l48++

$\downarrow$  l49++

$\downarrow$  l50++

$\downarrow$  l51++

$\downarrow$  l52++

$\downarrow$  l53++

$\downarrow$  l54++

$\downarrow$  l55++

$\downarrow$  l56++

$\downarrow$  l57++

$\downarrow$  l58++

$\downarrow$  l59++

$\downarrow$  l60++

$\downarrow$  l61++

$\downarrow$  l62++

$\downarrow$  l63++

$\downarrow$  l64++

$\downarrow$  l65++

$\downarrow$  l66++

$\downarrow$  l67++

$\downarrow$  l68++

$\downarrow$  l69++

$\downarrow$  l70++

$\downarrow$  l71++

$\downarrow$  l72++

$\downarrow$  l73++

$\downarrow$  l74++

$\downarrow$  l75++

$\downarrow$  l76++

$\downarrow$  l77++

$\downarrow$  l78++

$\downarrow$  l79++

$\downarrow$  l80++

$\downarrow$  l81++

$\downarrow$  l82++

$\downarrow$  l83++

$\downarrow$  l84++

$\downarrow$  l85++

$\downarrow$  l86++

$\downarrow$  l87++

$\downarrow$  l88++

$\downarrow$  l89++

$\downarrow$  l90++

$\downarrow$  l91++

$\downarrow$  l92++

$\downarrow$  l93++

$\downarrow$  l94++

$\downarrow$  l95++

$\downarrow$  l96++

$\downarrow$  l97++

$\downarrow$  l98++

$\downarrow$  l99++

$\downarrow$  l100++

throughput too low  
latency too high

end-to-end

"kernel" (not OS kernel)

↳ heavily compute-centric library function

Math, ...  
deep learning

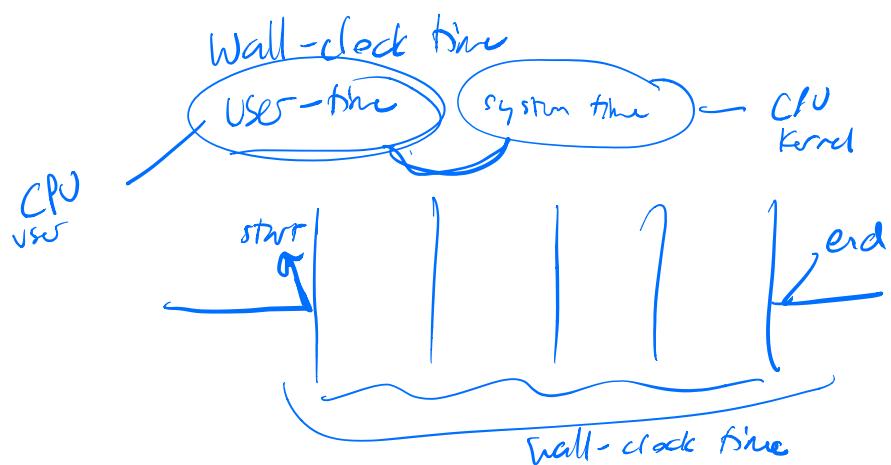
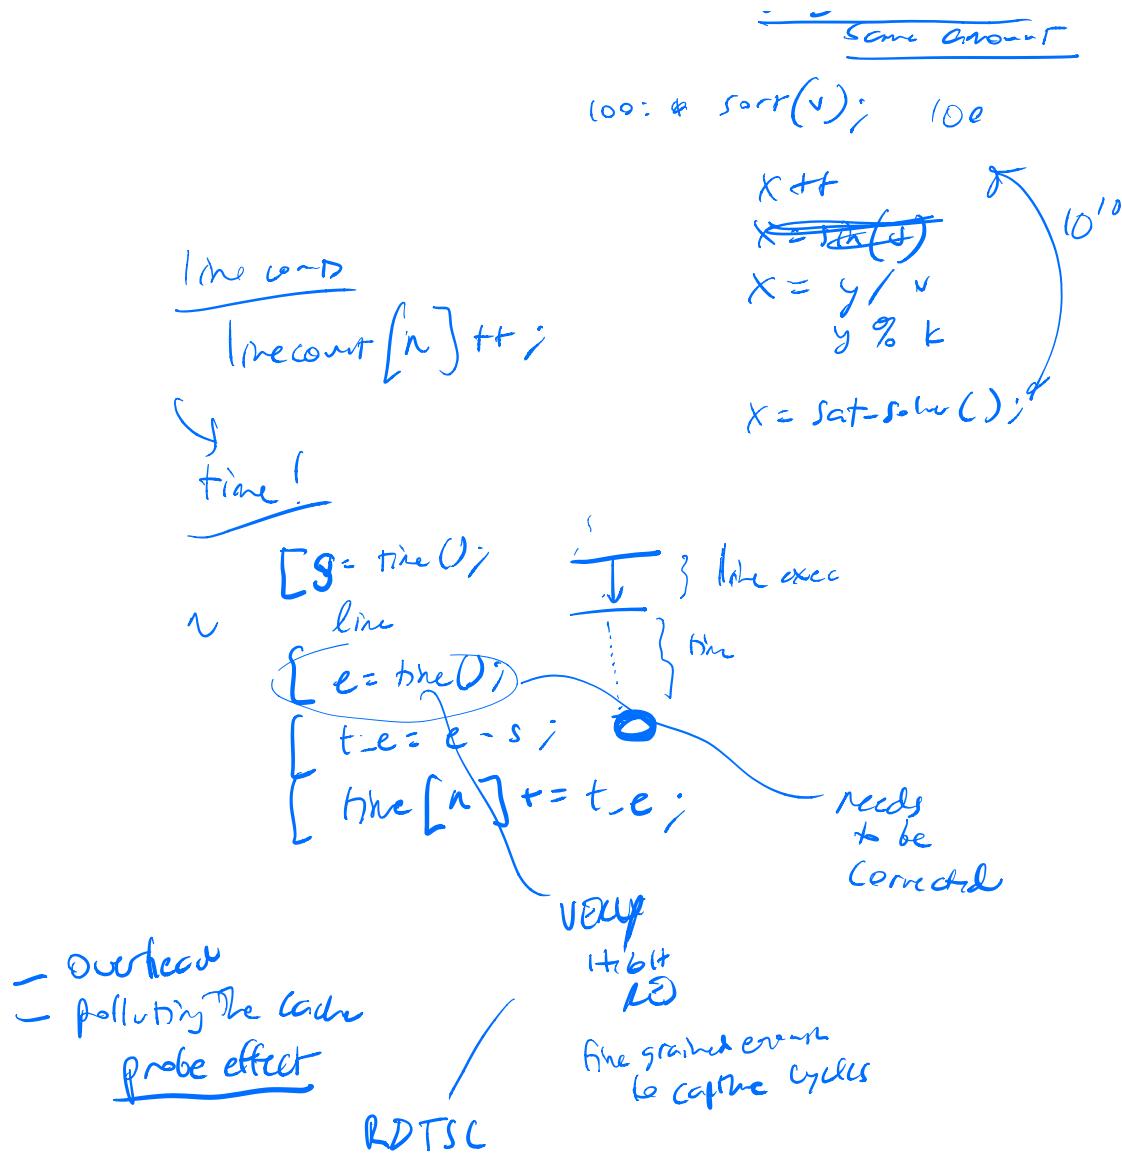
NOTE:  
SERIAL  
CODE  
only!

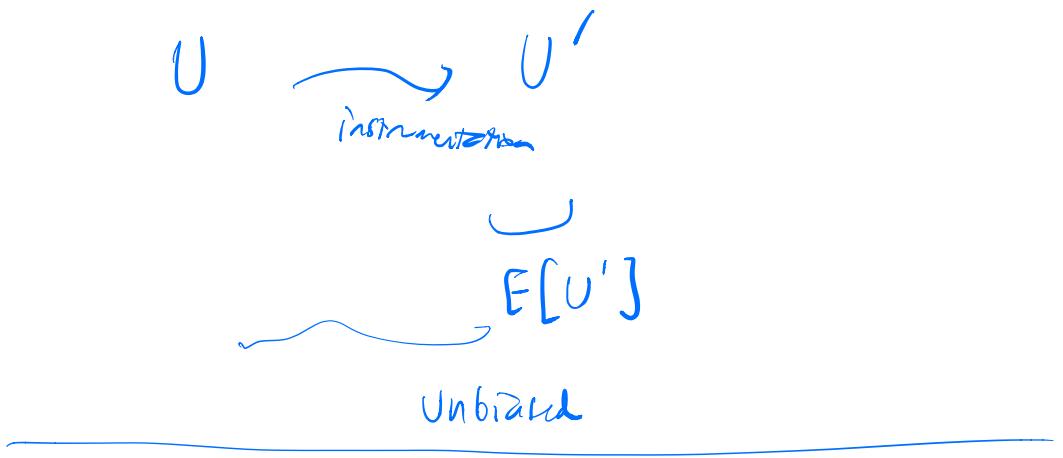
Advantages (pros)

If line count (trap count) high  
→ perf bottleneck

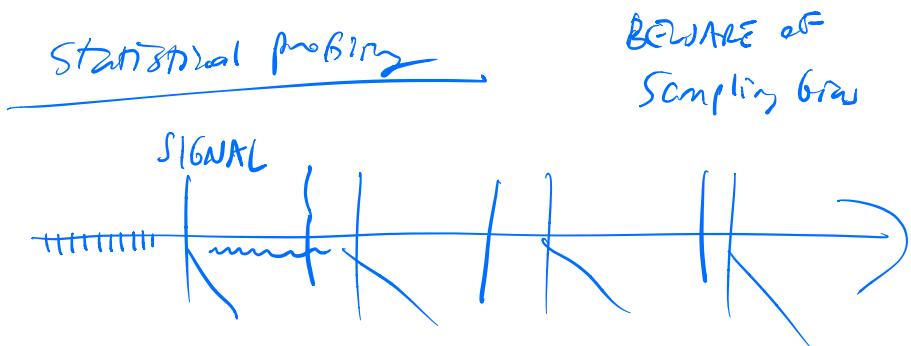
Disadvantages (cons)

- overhead  
slow program down
- assumption -  
over time costs





- 5% off - low overhead ✓
- accuracy ✓
- unbiased (true not known) ✓



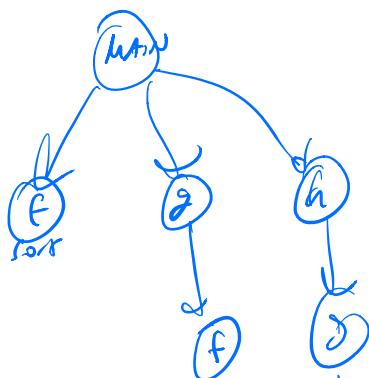
50% line 12

10% line 11

5% line 15

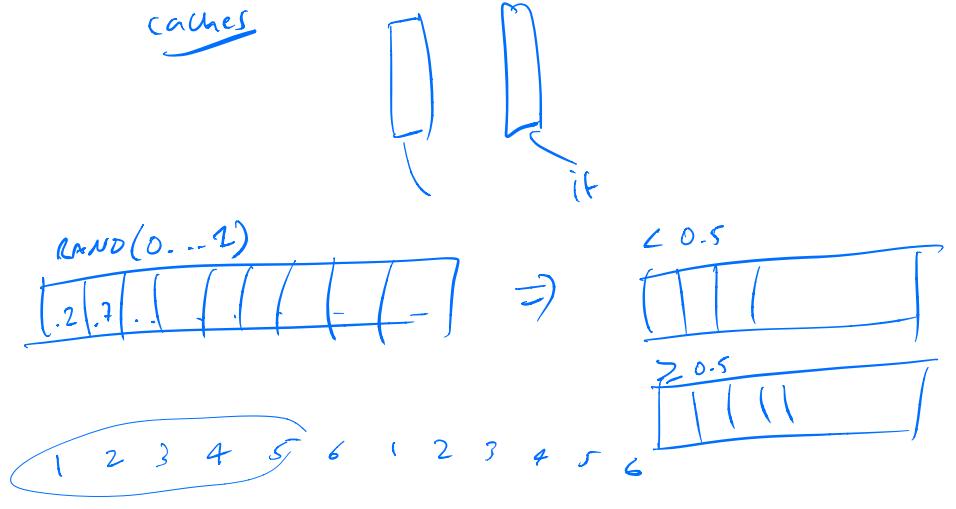
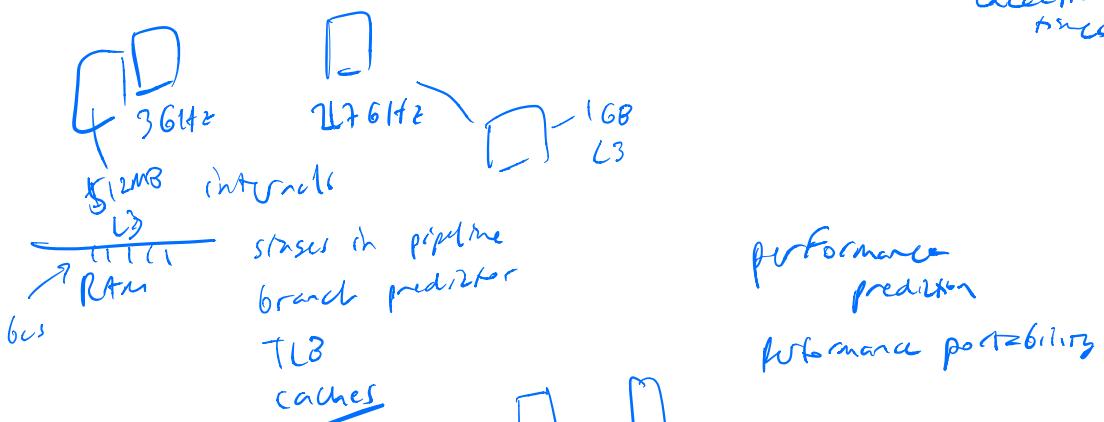
Low & Large Number!

:





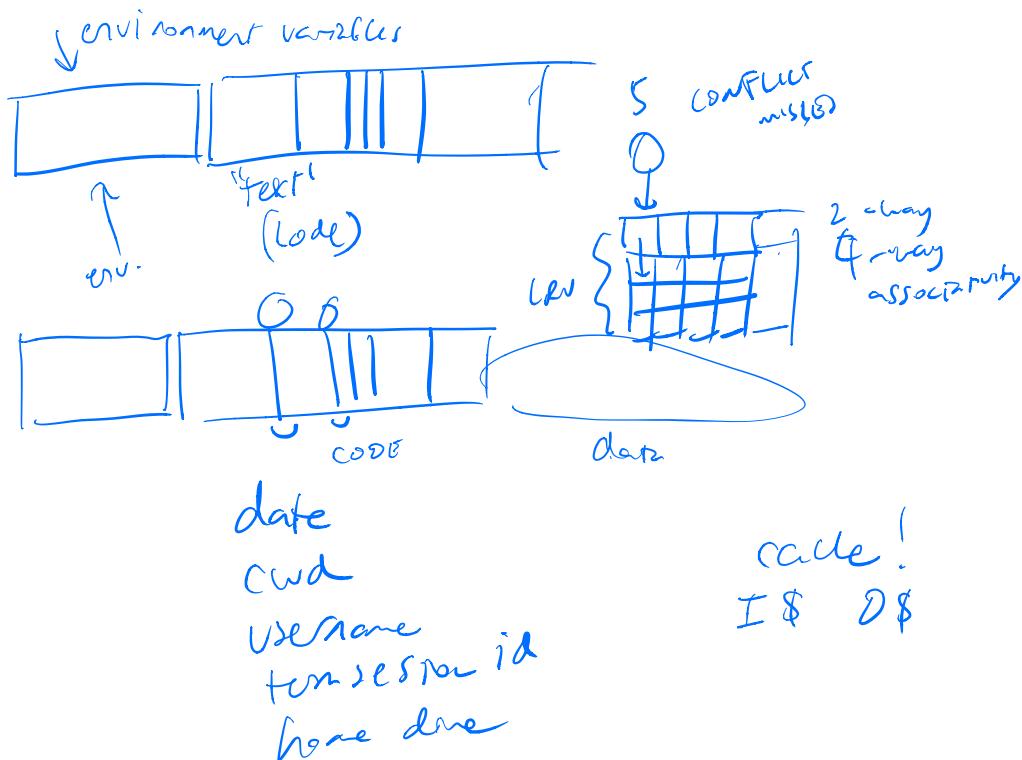
Statistical Sampling - Periodically interrupt  
 → Where do you land  
 collected samples ~ distribution  
 of  
 execution times



if ... == 6

Pattern History Table

Producing wrong data w/o doing anything obviously wrong!  
Mytkowicz et al.



### On-chip hardware performance counters

perf

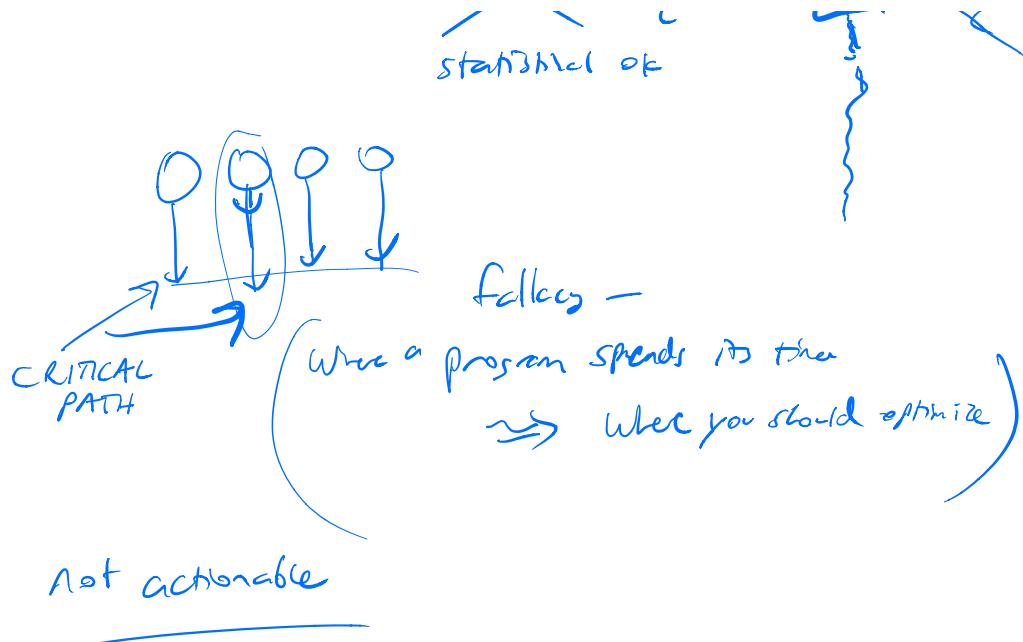
Counts up  
countdown timers — signal when perf counter hits zero  
 $\Rightarrow$  profiler for real code

The old days

good  
no threads  
virtually no I/O  
non-intrusive  
asynchronous

~~prob effect!~~

The diagram shows two sets of four circles representing threads. Arrows indicate the interleaved execution of threads between two cores. Below the cores, a large bracket groups the threads, with a double-headed arrow indicating they are being processed by both cores simultaneously. This visualizes how multiple threads can be active at once, contributing to the complexity of performance counter management.



profiler wrapup -

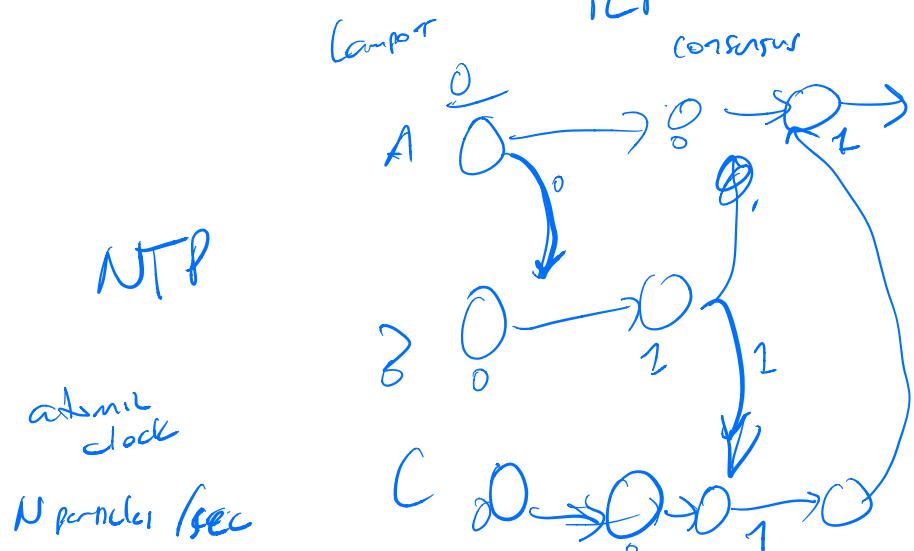
- Instrumentation based
- Sampling based
- perf counter based profiler
- causal profiling
- + concurrency
- + asynchrony
- + parallelism
- distrib. sys?

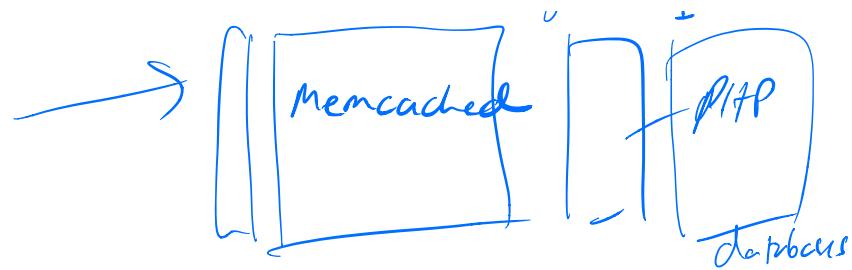
"single node"

distrib sys?

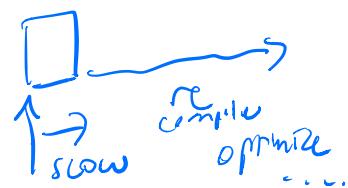
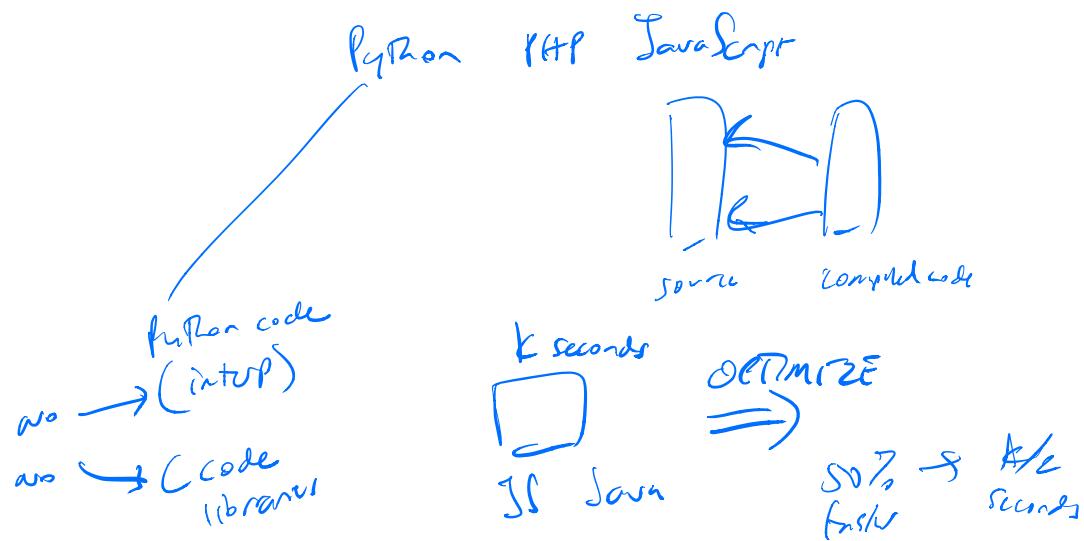
no global sync clock

FLP





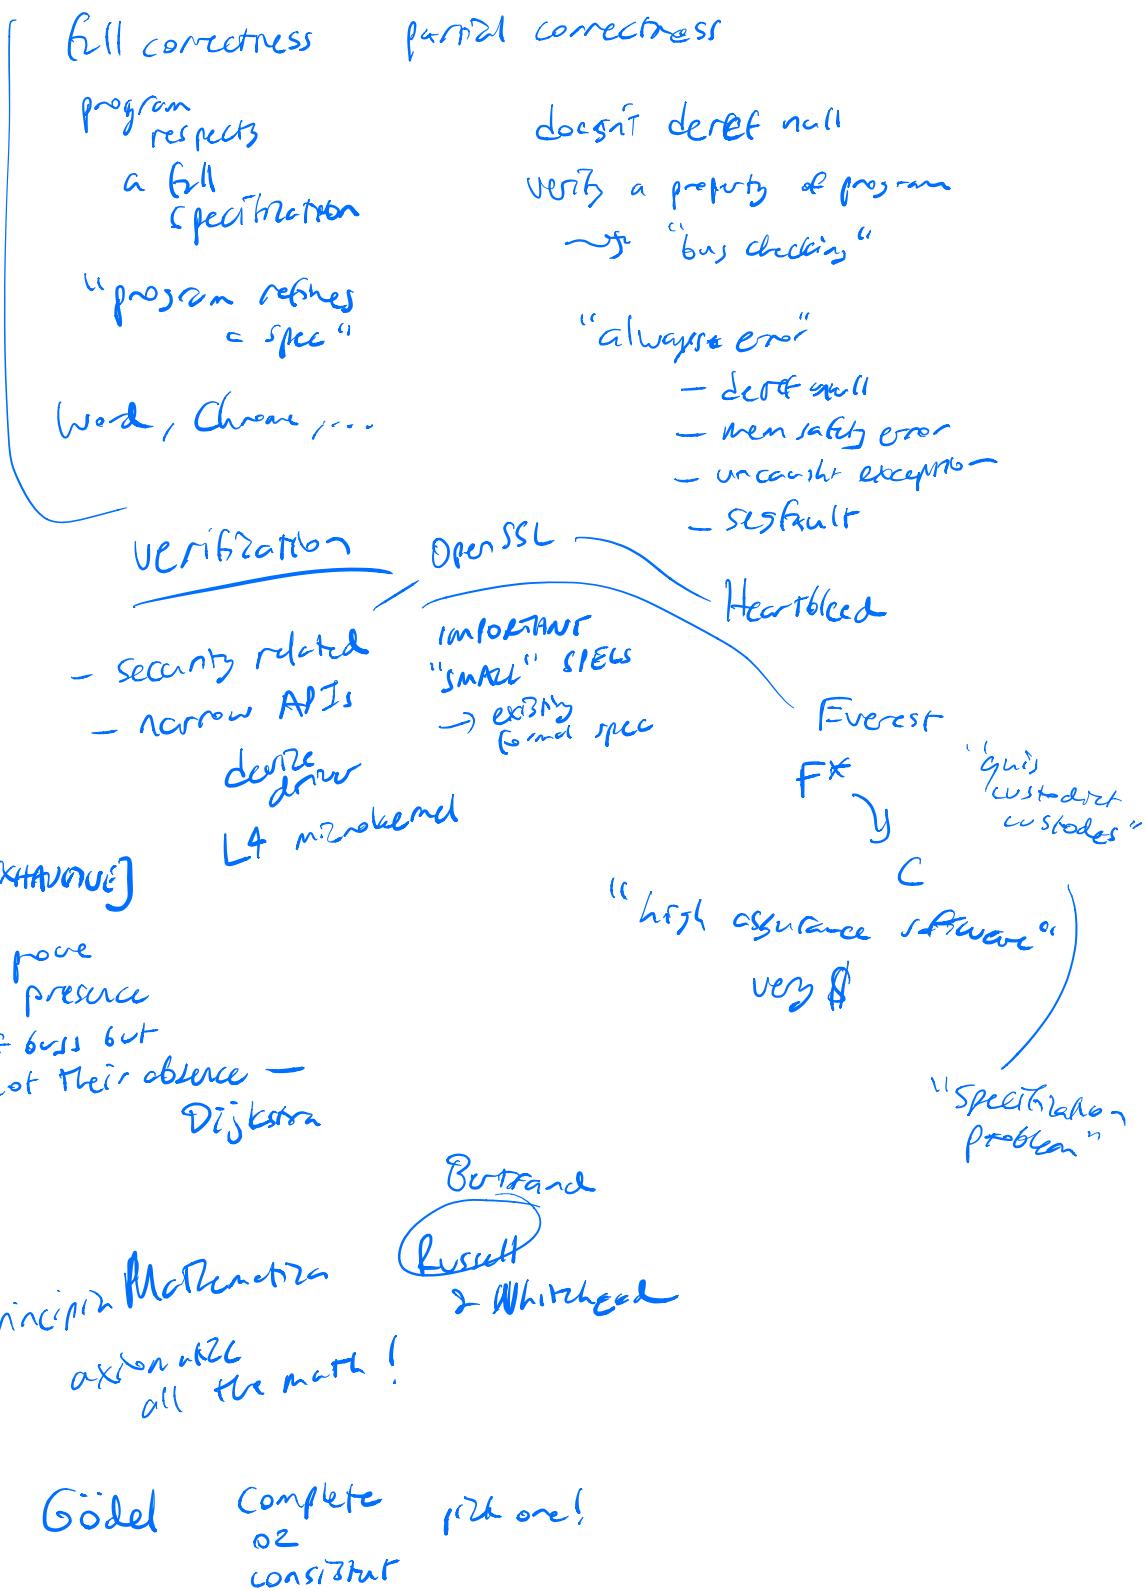
profiling "mixed" languages



perf hard

Correctness hard

bit monotonically increasing  
- compositional



dynamic analysis — precise  
 static analysis — imprecise  
 testing — high false positives

low recall } bugs detection  
 "heuristic static analysis"

"sound"

q If there's any possibility  
of error, report it!

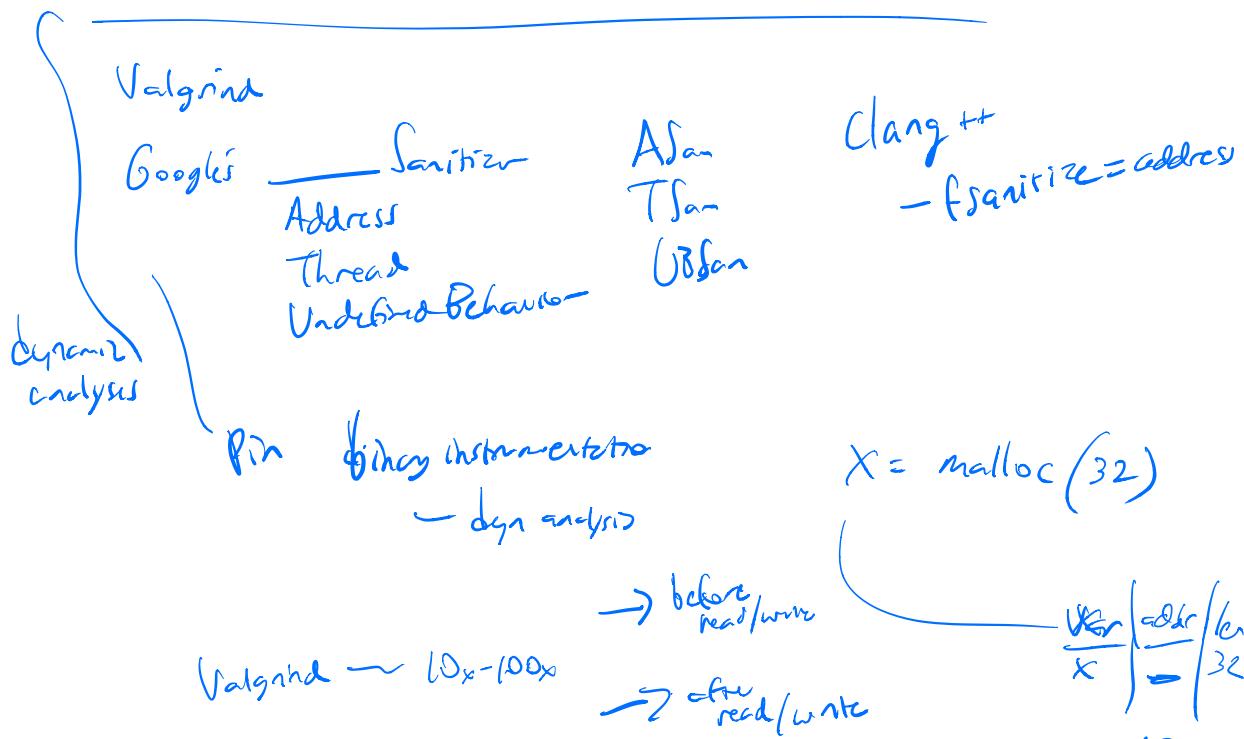
## Coverity

fundamental tension  
 soundness — low false positive ratio

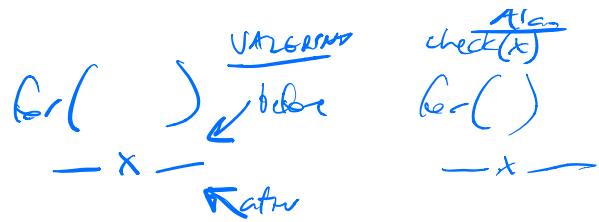
causes due to over approximation

$x \in \{1..10\}$        $y \in \{-5..0..5\}$   
 numerator      denominator

increasing /, - , 0, +       $x/y$       sound  $\Rightarrow$  corner here!  
 precise  
 (ptr analysis)  
 (still sound)  
 "scalable"      "Infer" PB



Alloc ~  $1.2x + 2x$



GC — "dynamic analysis"

eliminates owners instead of repeating

Extinction

- lgc



testing